



# ALGORITHMS OF BIOINFORMATICS

4

## Assembling Genomes

*20 November 2025*

Prof. Dr. Sebastian Wild


## Outline

# 4 Assembling Genomes


- 4.1 Exploding Newspapers
- 4.2 The De-Novo Sequencing Problem
- 4.3 Assembly with Overlap Graphs
- 4.4 De Bruijn Graphs
- 4.5 Practical Assemblers

## 4.1 Exploding Newspapers


# Exploding Newspapers





stack of NY Times, June 27, 2000



stack of NY Times, June 27, 2000  
on a pile of dynamite



this is just hypothetical



so, what did the June 27, 2000 NY  
Times say?

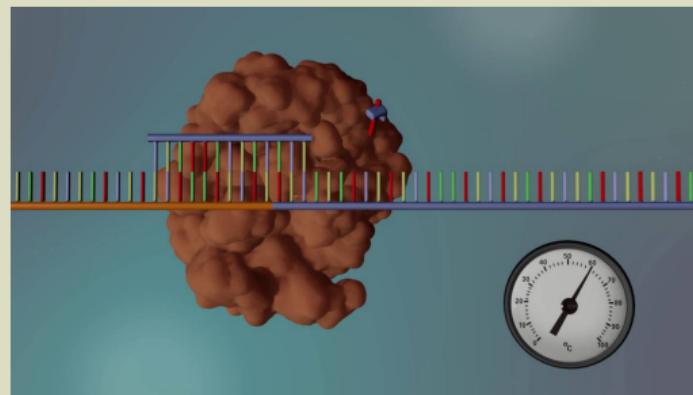
Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.1  
<https://cogniterra.org/lesson/29884/step/2?unit=21982>

# DNA Sequencing

How can we possibly “read out” DNA sequences?

Through clever combinations of chemistry and engineering ...

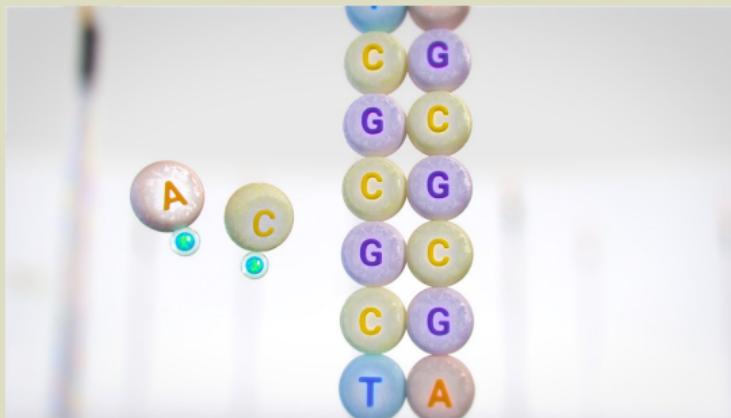
- ▶ Historically first method: *Sanger sequencing*
  - ▶ developed 1977 by Frederick Sanger
  - ▶ earning him his *second* Nobel prize in chemistry  
(the first one was for determining the amino acid sequence of insulin)



► DNA Sequencing - 3D  
<https://youtu.be/ONGdehkB8jU>

# “Next Generation Sequencing”

- ▶ group of methods that are massively parallel
- ~~ easier to automate, much cheaper per read
- ▶ commercially available since 2005



■ Overview of Illumina Sequencing by Synthesis Workflow  
<https://youtu.be/EDVKxSNdSic>

# Limitations fo Sequencing Techniques

- ▶ Sequencing machines can produce somewhat reliable *reads* of a few hundred to a few thousand bases.
- ~~ NGS will produce many reads, covering much of sequence but we need to *assemble* entire DNA sequence

# Overview


## Overview of Sequencing

Multiple identical copies of a genome

Shatter the genome into reads

Sequence the reads

Assemble the genome using overlapping reads



AGAATATCA  
GAGAATATC  
TGAGAATAT  
... TGAGAATATCA ...

Compeau & Pevzner, *Bioinformatics Algorithms*, Fig. 4.13  
<https://cogniterra.org/lesson/29918/step/2?unit=22015>

In the following: Consider an idealized situation.

## 4.2 The De-Novo Sequencing Problem

# Formalization of Sequence Assembly

- ▶  $k$ -mer = length- $k$ -(sub)string
  - ▶ Given text (genome)  $T[0..N]$  and  $k$ ,  
the  $k$ -mer composition of  $T$  is the multiset of all  $k$ -mers in  $T$ :  
$$\text{Composition}_k(T) = \{\{T[0..k], T[1..1+k], T[2..2+k], \dots, T[n-k, n]\}\}$$
  - ▶ Example:  $\text{Composition}_3(\text{TATGGGGTGC}) = \{\{\text{ATG}, \text{GGG}, \text{GGG}, \text{GGT}, \text{GTG}, \text{TAT}, \text{TGC}, \text{TGG}\}\}$
  - ▶ (De-Novo) Sequence Assembly Problem:
    - ▶ Given multiset of  $k$ -mers  $\mathcal{R}$
    - ▶ Find a string  $T$  such that  $\text{Composition}_k(T) = \mathcal{R}$
  - ▶ Example:  $\mathcal{R} = \{\{\text{AAT}, \text{ATG}, \text{GTT}, \text{TAA}, \text{TGT}\}\}$ 
    - ▶ form overlapping pairs
    - ▶  $T$  must start with **TTA** as no 3-mer ends with **TA** and end with **GTT**
- TAA  
AAT  
ATG  
TGT  
GTT  
TAATGTT

## ... not always as easy

Now consider  $\mathcal{T} = \{\text{AAT, ATG, ATG, ATG, CAT, CCA, GAT, GCC, GGA, GGG, GTT, TAA, TGC, TGG, TGT}\}$

Let's start again with **TTA**.

The first few steps are forced.

TAA  
AAT  
ATG  
TAATG

Then we have there are 3 options **TG[CGT]**.

Say we pick **TGT**

TAA  
AAT  
ATG  
TGT  
GTT  
TAATGTT 

We've reached a dead end!

If we instead choose more wisely,  
we can continue:

TAA  
AAT  
ATG  
TGC  
GCC  
CCA  
CAT  
ATG  
TGG  
GGG  
GGA  
GAT  
ATG  
TGT  
GTT  
TAATGCCATGGGATGTT

# Escaping the Dead Ends

*How can we systematically avoid having to try all options?*

*Repeats make life challenging.*




## 4.3 Assembly with Overlap Graphs

# Read Overlap Graphs

Given a  $k$ -mer composition  $\mathcal{R} = \{R_0, \dots, R_{n-k}\}$ ,  
construct the *overlap graph* as directed graph  $G_O(\mathcal{R}) = (V, E)$

- ▶  $V = \{r_0, \dots, r_{n-k}\}$  with  $r_i$  labeled with  $R_i$   
can have duplicate  $k$ -mers
- ▶  $r_i r_j \in E$  if  $R_i[1..k] = R_j[0..k-1]$


Example Overlap Graph (Alphabetical)



Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.8  
<https://cogniterra.org/lesson/29886/step/7?unit=21984>

# Hamiltonian Paths

Given the overlap graph  $G(\mathcal{R})$  of a  $k$ -mer composition  $\mathcal{R}$   
a sequence assembly solution is a *Hamiltonian path* in  $G(\mathcal{R})$



Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.9  
<https://cogniterra.org/lesson/29886/step/8?unit=21984>

- ▶ **Problem 1:** Result may not unique (as in this example)
- ▶ **Problem 2:** DIRECTED HAMILTONIAN PATH is NP-complete.

## Bad News Again

*It seems we once again ran into a hard problem.*

(And we haven't even accounted for computing the overlap graph . . . )

. . . but sometimes, a different way to look at the problem helps.

## 4.4 De Bruijn Graphs

# Origins

Nicolaas Govert de Bruijn's  $k$ -universal string puzzle:

How long does a binary string need to be to contain all  $2^k$  binary  $k$ -mers as substrings?  
How to construct a shortest such string?

Examples:

- ▶  $k = 2$ 
  - ▶ must contain 00, 01, 10, 11
    - ~~ a 2-universal string needs at least 5 letters
      - 4 starting positions, each for a 2-letter substring
  - ▶ Also achievable: 00110
    - ~~ Answer: 2-universal requires 5 characters.
- ▶  $k = 3$  ~~ a shortest 3-universal string is 0001110100

Same question makes sense for strings over alphabet  $\Sigma = [0..\sigma)$ .


~~  $T$  is  $k$ -universal  $\iff$  contains all  $\sigma^k$   $k$ -mers as substring.

# de Bruijn Graphs

Given a  $k$ -mer composition  $\mathcal{R} = \{R_0, \dots, R_{n-k}\}$ ,  
construct the *de Bruijn graph* as directed multigraph  
 $G_B(\mathcal{R}) = (V, E)$

- ▶  $V = \{R[0..k-1] : R \in \mathcal{R}\}$
  - ▶ For each read  $R \in \mathcal{R}$ , add the edge  $vw$  (with label  $R[0..k]$ ) to  $E$   
for  $v = R[0..k-1]$  and  $w = R[1..k]$
- ↔ Sequence using each  $k$ -mer once = Euler path in  $G_B$
- ▶ Euler path efficiently computable!

DEBRUIJN<sub>3</sub>(TAATGCCATGGGATGTT)



Compeau & Pevzner, *Bioinformatics Algorithms*, Fig. 4.1  
<https://cogniterra.org/lesson/29910/step/2?unit=22007>


## *k*-universal strings

The de Bruijn graph “solves” the universal string problem for  $\mathcal{R}$  all  $\sigma^k$   $k$ -mers

- ▶  $\sigma^{k-1}$  vertices,  $\sigma^k$  edges
  - ▶ each vertex has out- and in-degree =  $\sigma$
- ~~  $G_B$  is Eulerian
- ~~ Compute Euler path (see below)

# Euler Cycles

**Euler Walk:** Walk using every edge in  $G = (V, E)$  exactly once.



## Euler's Theorem:

Euler walk exists iff  $G$  connected and 0 or 2 vertices have odd degree.

' $\Rightarrow$ ' trivial (need to enter and exit intermediate vertices equally often)

' $\Leftarrow$ ' Following algorithm *constructs* Euler walk under this assumption



# Euler Cycles – Hierholzer’s Algorithm

- ▶ use an *edge-centric DFS*

- ▶ We mark *edges* (not vertices)
- ~~ stack = **edge-simple walk**
- ▶ We remember iterator *i* globally per *v* to resume traversal

---

```
1 procedure eulerWalk(G):
2     // Assume G = (V, E) is connected (multi)graph
3     Vodd := {v ∈ V : d(v) odd}
4     if |Vodd| ≠ {0, 2} return NOT_EULERIAN
5     if Vodd = {x, y} then s := x else s := 0
6     euler[0..m) := NONE; j := m - 1
7     visited[0..n, 0..n) := false // mark edges as visited
8     for v := 0, ..., n - 1
9         // globally remember next unexplored edge
10        nextEdge[v] := G.adj[w].iterator()
11        edgeDFS(s)
12        return euler
```

---

```
1 procedure edgeDFS(s):
2     frontier := new Stack;
3     frontier.push(s)
4     while ¬frontier.isEmpty()
5         v := frontier.top(); i := nextEdge[v]
6         if ¬i.hasNext() // v has no unused edge
7             frontier.pop()
8             if ¬frontier.isEmpty()
9                 // assign edge leading here largest free index
10                euler[j] := (frontier.top(), v); j := j - 1
11            end if
12        else
13            w := i.next()
14            if ¬visited[v, w]
15                visited[v, w] := true
16                visited[w, v] := true
17                frontier.push(w)
18            end if
19        end if
20    end while
```

---

# Directed Euler Cycles

- ▶ de Bruijn graphs are *directed* ... what changes?
- ▶ **Euler's Theorem, directed version:**
  - Euler walk exists iff  $G$  is (a) strongly connected and
    - (b)  $d_{\text{out}}(v) - d_{\text{in}}(v) = 0$  for all vertices  $v$  or
      - 0 for all but 2 vertices where it is +1 and -1, respectively.
- ' $\Rightarrow$ ' need to enter and exit intermediate vertices equally often
- ' $\Leftarrow$ ' Hierholzer's algorithm also works on directed graphs, with changes below.
- ▶ Changes to Hierholzer's algorithm:
  - ▶ Check balance ( $d_{\text{out}}(v) - d_{\text{in}}(v)$ ) of vertices
  - ▶ start at the unique vertex with  $d_{\text{out}}(s) - d_{\text{in}}(s) = +1$  if  $s$  exists;  
otherwise arbitrary
  - ▶ in edgeDFS, do not mark the reverse edge as also visited

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space<sup>\*</sup>  
Ziad Ismaili Alaoui<sup>†</sup>   Detlef Plump<sup>‡</sup>   Sebastian Wild<sup>§</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n\lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicities.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the existence of Eulerian cycles and their enumeration is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classic algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both employ a backtracking search that explores all possible ways to extend the current cycle. Being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over the standard algorithm. Our algorithm does not require any auxiliary data structures; we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space occupied by auxiliary data structures during the computation. Our algorithm relies on an assumption that the input graph is connected, which is satisfied in most adjacency matrices. In the case of directed graphs, we require that we iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and that finds an Eulerian cycle. Our algorithm also works in reduced models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

<sup>\*</sup>Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V000947/1.

<sup>†</sup>University of Liverpool, United Kingdom (ziad.ismaili-alaeui@liverpool.ac.uk).

<sup>‡</sup>University of York, United Kingdom (detlef.plump@york.ac.uk).

<sup>§</sup>Friedrich-Alexander-Universität Erlangen-Nürnberg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

**SOSA 2026**

Copyright © 2026  
by authors.

# Space-Efficient Hierholzer – Implementing colors

**Assumption:** Can iterate over outgoing **and** incoming edges

- ▶ Every vertex has at most one outgoing **RED** edge
  - ~~ store explicitly in  $O(n)$  space
- ▶ edges are visited (**RED** or **GREEN**) iff before current iterator in adjacency list
  - ~~ stored implicitly!

## Back to Genomes

- ▶ Eulerian walk usually not unique or doesn't exist!
  - ~~ often stop with assembling *contigs*, long contiguous substrings of the genome to be further processed
- ▶ contigs can be found from de Bruijn graph using vertex balances

## 4.5 Practical Assemblers

# Practical Complications

Challenges for DNA sequencing in practice

- ▶ DNA is *not* a single string; one per chromosome!
- ▶ we have (most) chromosomes twice (mom's and dad's), probably similar in large parts!
- ▶ imperfect coverage: some regions might not be covered by any/enough reads
- ▶ sequencing isn't exact; small reading errors need to be tolerated
- ▶ DNA is double stranded ↗ get reads from both strands mixed up

# Some solutions

## Imperfect coverage

- ▶ Rather unlikely to see a read for every possible starting position.
- ▶ But this is harmless; we can break reads of length  $\ell$  into artificial  $k$ -mers for some  $k < \ell$  and obtain perfect coverage!

## Chromosome pairs

- ▶ problematic because of high similarity
- ▶ possible in lab to avoid sequencing both

Inexact Overlaps seem to be the most problematic.

# There's life in the old dog yet!

- ▶ Inexact reads actually much easier to handle in **overlap graph**  
weighted edges for length / similarity of overlap
- ▶ highly effective heuristic: preprocessing overlap graph to transitive reduction
- ▶ Hamiltonian path / TSP tour not the bottleneck      "hard in theory, easy in practice"
- ▶ challenge there: how to compute overlap graph (edge weights) efficiently  
~~ we will come back to this