



ALGORITHMS OF BIOINFORMATICS

5 String Matching

4 December 2025

Prof. Dr. Sebastian Wild

Outline

5 String Matching


- 5.1 Read Mapping
- 5.2 String Matching with Finite Automata
- 5.3 Constructing String Matching Automata
- 5.4 The Knuth-Morris-Pratt algorithm
- 5.5 The Aho-Corasick Algorithm
- 5.6 Alignment-Based Matching
- 5.7 Longest Common Extensions
- 5.8 SNPs via filtering

5.1 Read Mapping

Shotgun Sequencing

Recall:

- *Shotgun sequencing* approach to determine a (human) genome:



Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.1
<https://cogniterra.org/lesson/29884/step/2?unit=21982>

- For a single human genome need 300M reads of 200bp (30x coverage)
 - ~~~ 60 GB of raw data
 - ~~~ genome assembly from those is expensive and error prone

We now have carefully assembled *reference genomes* to compare with!

Shotgun Sequencing in Medicine

Predominant medical use of whole genome sequencing:
detecting known markers (mutations / gene combinations) for diseases

Example: (more details in Compeau & Pevzner 2015)

- ▶ *Ohdo syndrome* (form of mental retardation, “mask-like” face)
- ▶ known to be indicated by single protein-truncating mutation

Often even a **single base replacement** w.r.t. a human reference genome,
a **SNP (single nucleotide polymorphism)**

~~> *To find SNPs, we don't need a new patient's genome fully assembled!*

Read Mapping

- ▶ We thus work towards solving the *read mapping problem*
 - ▶ **Given:** genome/text $T[0..n]$, reads/patterns $P[0..p]$, $P[r] = P_r[0..m_r]$
 - ▶ **Goal:** locations i_r of “best match” for P_r in T for $r \in [0..p]$.
- ▶ “best match” can be interpreted in several ways, leading to different problems:
 - (a) best **semilocal alignment** of P_r to T (gold standard, usually too expensive)
 - (b) match with **fewest mismatches**
 - (c) match with $\leq d$ **mismatches** or **NO_MATCH** if no such exists
for SNPs can even set $d = 1$
 - (d) **exact match** or **NO_MATCH** if no such exists

We will first focus on *exact matches*.

- ▶ simplifies the problem (to get started)
- ▶ the SNPs variants can be reduced to it (using postprocessing)

Part I

Exact matches

Notation

- ▶ *alphabet* Σ : finite set of allowed **characters**; $\sigma = |\Sigma|$ “*a string over alphabet Σ* ”
 - ▶ focus on nucleotides $\{A, C, G, T\}$ and amino acids
 - ▶ but try to keep methods generic
 - ▶ letters (Latin, Greek, Arabic, Cyrillic, Asian scripts, . . .) Unicode characters
 - comprehensive standard character set including emoji and all known symbols
- ▶ $\Sigma^n = \Sigma \times \cdots \times \Sigma$: strings of **length n** $\in \mathbb{N}_0$ (n -tuples)
 - ▶ $\Sigma^* = \bigcup_{n \geq 0} \Sigma^n$: set of **all** (finite) strings over Σ , $\Sigma^+ = \bigcup_{n \geq 1} \Sigma^n$
 - ▶ $\varepsilon \in \Sigma^0$: the **empty** string (same for all alphabets)
- ▶ for $S \in \Sigma^n$, write $S = S[0..n]$, so $S[i]$ (other sources: S_i) for **i th character** ($0 \leq i < n$)
 - zero-based (like arrays)!
- ▶ for $S, T \in \Sigma^*$, write $ST = S \cdot T$ for **concatenation** of S and T
- ▶ for $S \in \Sigma^n$, write $S[i..j]$ for the **substring** $S[i] \cdot S[i + 1] \cdots S[j - 1]$ ($0 \leq i \leq j \leq n$)
 - ▶ $S[i..i] = \varepsilon$
 - ▶ $S[0..j]$ is a **prefix** of S ; $S[i..n]$ is a **suffix** of S

String matching – Definition

Search for a string (pattern) in a large body of text

► **Input:**

- $T \in \Sigma^n$: The *text* being searched within
- $P \in \Sigma^m$: The *pattern* being searched for; typically $n \gg m$

► **Output:**

- the *first occurrence (match)* of P in T : $\min\{i \in [0..n-m] : T[i..i+m] = P\}$
- or `NO_MATCH` if there is no such i (“ P does not occur in T ”)
- sometimes also: find **all** occurrences of P in T .

try all starting positions

- trivially solvable with $(n - m + 1) \cdot \cancel{m} \sim nm$ character comparisons
- ~~ too slow for real mapping!
- string matching available, e.g., Java in `String.indexOf`, Python in `str.find`
- not always robust enough for bioinformatics data
(small alphabet, long repetitions)


5.2 String Matching with Finite Automata

Theoretical Computer Science to the rescue!

- ▶ string matching = deciding whether $T \in \Sigma^* \cdot P \cdot \Sigma^*$
- ▶ $\Sigma^* \cdot P \cdot \Sigma^*$ is *regular* formal language
 - ~~ \exists deterministic finite automaton (DFA) to recognize $\Sigma^* \cdot P \cdot \Sigma^*$
 - ~~ can check for occurrence of P in $|T| = n$ steps!



Job done!



WTF!?

We are not quite done yet.

- ▶ (Problem 0: programmer might not know automata and formal languages ...)
- ▶ Problem 1: existence alone does not give an algorithm!
- ▶ Problem 2: automaton could be very big!


String matching with DFA

- ▶ Assume first, we already have a deterministic automaton
- ▶ How does string matching work?

Example:

$T = \text{aabacaababacaa}$

$P = \text{ababaca}$



text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3	0	1	1	2	3	4	5	6	7	7

String matching DFA – Intuition

Why does this work?


- ▶ Main insight:

State q means:

*"we have seen $P[0..q]$ until here
(but not any longer prefix of P)"*

$$T = aabacaababacaa$$

$$P = ababaca$$



text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3	0	1	1	2	3	4	5	6	7	7

- ▶ If the next text character c does not match, we know:

- text seen so far ends with $P[0..q] \cdot c$
- $P[0..q] \cdot c$ is not a prefix of P
- without reading c , $P[0..q]$ was the *longest* prefix of P that ends here.

$$T = \dots [P[0..q]] [c] [P[0..q']]$$

with $q' < q$


~~ New longest matched prefix will be (weakly) shorter than q

~~ All information about the text needed to determine it is contained in $P[0..q] \cdot c$!

5.3 Constructing String Matching Automata

NFA instead of DFA?

It remains to *construct* the DFA.

- trivial part: 
- that actually is a *nondeterministic finite automaton* (NFA) for $\Sigma^* P \Sigma^*$

~~ We *could* use the NFA directly for string matching:

- at any point in time, we are in a *set of states*
- accept when one of them is final state

Example:

text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	0,1	0,1	0,2	0,1,3	0	0,1	0,1	0,2	0,1,3	0,2,4	0,1,3,5	0,6	0,1,7	0,1,7

But maintaining a whole set makes this slow ...

Computing DFA directly




You have an NFA and want a DFA?

Simply apply the power-set construction
(and maybe DFA minimization)!

The powerset method has exponential state blow up!

I guess I might as well use brute force string matching ...



Ingenious algorithm by Knuth, Morris, and Pratt: construct DFA *inductively*:

Suppose we add character $P[j]$ to automaton A_j for $P[0..j]$ to construct A_{j+1}

- ▶ add new state and matching transition \rightsquigarrow easy
- ▶ for each $c \neq P[j]$, we need $\delta(j, c)$ (transition from j when reading c)
- ▶ $\delta(j, c) =$ length of the longest prefix of $P[0..j]c$ that is a suffix of $P[1..j)c$
= state of automaton after reading $P[1..j)c$
 $\leq j \rightsquigarrow$ can use known automaton A_j for that!

\rightsquigarrow can directly compute A_{j+1} from A_j !

seems to require simulating automata $m \cdot \sigma$ times

State q means:
“we have seen $P[0..q]$ until here
(but not any longer prefix of P)”

Computing DFA efficiently

- ▶ KMP's second insight: simulations in one step differ only in last symbol
- ~~ simply maintain state x , the state after reading $P[1..j]$.
 - ▶ copy its transitions
 - ▶ update x by following transitions for $P[j]$

```
1 procedure constructDFA( $P[0..m]$ ):  
2     //  $\delta[q][c] = \text{target state when reading } c \text{ in state } q$   
3     for  $c \in \Sigma$  do  
4          $\delta[0][c] := 0$   
5          $\delta[0][P[0]] := 1$   
6          $x := 0$   
7         for  $j = 1, \dots, m - 1$  do  
8             for  $c \in \Sigma$  do // copy transitions  
9                  $\delta[j][c] := \delta[x][c]$   
10                 $\delta[j][P[j]] := j + 1$  // match edge  
11                 $x := \delta[x][P[j]]$  // update  $x$ 
```

Example: $P[0..7] = \text{ababaca}$

$\delta(c, q)$	0	1	2	3	4	5	6
a	1	1	3	1	5	1	7
b	0	2	0	4	0	4	0
c	0	0	0	0	0	6	0

String matching with DFA – Discussion

► Time:

- ▶ Matching: n table lookups for DFA transitions
- ▶ building DFA: $\Theta(m\sigma)$ time (constant time per transition edge).
 $\leadsto \Theta(m\sigma + n)$ time for string matching.

► Space:

- ▶ $\Theta(m\sigma)$ space for transition matrix.



fast matching time actually: hard to beat!



total time asymptotically optimal for small alphabet (for $\sigma = O(n/m)$)




substantial space overhead, in particular for large alphabets

5.4 The Knuth-Morris-Pratt algorithm

Failure Links


- ▶ Recall: String matching with DFA is fast,
but needs table of $m \times \sigma$ transitions.
- ▶ in fast DFA construction, we used that all simulations differ only by *last symbol*
- ~~ **KMP's third insight:** do this last step of simulation from state x during *matching!*
... but how?
- ▶ **Answer:** Use a new type of transition: \times , the *failure links*
 - ▶ Use this transition (only) if no other one fits.
 - ▶ \times does not consume a character. ~~ might follow several failure links



~~ Computations are deterministic (but automaton is not a classic DFA.)

Failure link automaton – Example

Example: $T = \text{abababaaaaca}$, $P = \text{ababaca}$



$T : \quad \begin{matrix} a & b & a & b & a & b & a & a & b & a & b \\ P : \quad \begin{matrix} a & b & a & b & a & \times & & & & & & \end{matrix} \end{matrix}$

	a	b	a	b	a	\times					
P :	a	b	a	b	a	\times					
		(a)	(b)	(a)	b	a	\times				
								a	b	a	b

to state 3
to state 1

$q : \quad \begin{matrix} 1 & 2 & 3 & 4 & 5 & 3,4 & 5 & 3,1,0,1 & 2 & 3 & 4 \end{matrix}$

(after reading this character)

The Knuth-Morris-Pratt Algorithm

```
1 procedure KMP( $T[0..n]$ ,  $P[0..m]$ ):  
2      $fail[0..m]$  := failureLinks( $P$ )  
3      $i := 0$  // current position in  $T$   
4      $q := 0$  // current state of KMP automaton  
5     while  $i < n$  do  
6         if  $T[i] == P[q]$  then  
7              $i := i + 1$ ;  $q := q + 1$   
8             if  $q == m$  then  
9                 return  $i - q$  // occurrence found  
10            else // i.e.  $T[i] \neq P[q]$   
11                if  $q \geq 1$  then  
12                     $q := fail[q]$  // follow one ×  
13                else  
14                     $i := i + 1$   
15    end while  
16    return NO_MATCH
```

- ▶ only need single array $fail$ for failure links
- ▶ (failureLinks on next slide)

Analysis: (matching part)

- ▶ always have $fail[j] < j$ for $j \geq 1$
- ~~ in each iteration
 - ▶ either advance position in text ($i := i + 1$)
 - ▶ or shift pattern forward (guess $i - q$)
- ▶ each can happen at most n times
- ~~ $\leq 2n$ symbol comparisons!

Computing failure links

- ▶ failure links point to error state x (from DFA construction)
 - ~~ run same algorithm, but store $\text{fail}[j] := x$ instead of copying all transitions

```
1 procedure failureLinks( $P[0..m]$ ):  
2      $\text{fail}[0] := 0$   
3      $x := 0$   
4     for  $j := 1, \dots, m - 1$  do  
5          $\text{fail}[j] := x$   
6         // update failure state using failure links:  
7         while  $P[x] \neq P[j]$   
8             if  $x == 0$  then  
9                  $x := -1$ ; break  
10            else  
11                 $x := \text{fail}[x]$   
12            end while  
13             $x := x + 1$   
14        end for
```

Analysis:

- ▶ m iterations of for loop
- ▶ while loop always decrements x
- ▶ x is incremented only once per iteration of for loop
 - ~~ $\leq m$ iterations of while loop *in total*
 - ~~ $\leq 2m$ symbol comparisons

Knuth-Morris-Pratt – Discussion

► Time:

- ▶ $\leq 2n + 2m = O(n + m)$ character comparisons
- ▶ clearly must at least *read* both T and P in the worst case
 - ~~ KMP has optimal worst-case complexity

► Space:

- ▶ $\Theta(m)$ space for failure links

 total time asymptotically optimal (for any alphabet size)

 reasonable extra space

The KMP prefix function

- ▶ It turns out that the failure links are useful beyond KMP
 - ▶ a slight variation is (more?) widely used: (for historic reasons)
the (KMP) *prefix function* $F : [1..m - 1] \rightarrow [0..m - 1]$:
- $F[j]$ is the length of the longest prefix of $P[0..j]$
that is a suffix of $P[1..j]$.*
- ▶ Can show: $fail[j] = F[j - 1]$ for $j \geq 1$, and hence

$fail[q] = \text{length of the}$
 $\text{longest prefix of } P[0..q)$
 $\text{that is a suffix of } P[1..q).$

← memorize this!

- ▶ EAA Buch: String indices are 1-based, but definition of failure links matches! $\Pi_P(q) = fail[q]$
- $\Pi_P : [1..m] \rightarrow [0..m - 1]$ with $\Pi_P(q) = \max\{k \in \mathbb{N}_0 : k < q \wedge P[0..k] \sqsupseteq P[0..q]\} = fail[q]$

5.5 The Aho-Corasick Algorithm

Multiple-Pattern Matching

- ▶ So far: process a single pattern P in one pass over T .
 - ~~> Cannot beat time $\Omega(p \cdot n)$
- ▶ The essence of KMP can be generalized to deal with several patterns!
 - ~~> The *Aho-Corasick Algorithm*

The Multiple-Pattern Matching Problem

- ▶ Given: text $T[0..n]$, patterns $P[0..p], P[r] = P_r[0..m_r]$
 - ▶ all over $\Sigma = [0..\sigma]$ for constant σ
 - ▶ total length of patterns: $m := \sum_{0 \leq r < p} m_r$
- ▶ Goal: all matches, i. e., all pairs (i, r) such that $P_r = T[i..i + m_r]$

Aho-Corasick can do this with $\mathcal{O}(m)$ preprocessing and $\mathcal{O}(n + \text{output})$ matching time!
Here *output* is the number of match pairs (i, r) .

single pass!

Aho-Corasick Automaton – Overview

Aho-Corasick Automaton


1. Build trie A from patterns $P[0..p)$.
2. Add *failure links* to A .
3. Add *output links* to A .

Recap: Tries

- efficient dictionary data structure for strings
- name from retrieval, but pronounced “try”
- tree based on symbol comparisons

► Example:

{aa\$, aaab\$, abaab\$, abb\$,
abbab\$, bba\$, bbab\$, bbb\$}



When stored string is a strict prefix of another, internal nodes can correspond to strings.

Aho-Corasick Automaton – Adding Failure Links

Trie for P_r corresponds to match-edges-only NFA.

- ~~ Interpreting the trie as automaton, add ε -edges back to the root.
- ▶ as in KMP, instead of determinizing the automation classically, we again use failure links
- ~~ construction as for KMP using failure state x , repeated for each word.

Aho-Corasick Automaton – Output Links

An automaton state might contain other patterns as suffix \rightsquigarrow must output match(es)! But we are not in an accepting state, so direct use of automaton so far would miss occurrence!

- ▶ **output links:** each state points to longest suffix pattern (if any).
 - ▶ During matching, traverse linked list of matches and output each $\rightsquigarrow O(\text{output})$ cost overall
 - ▶ Computation of output links similar to failure links! (handle patterns by length)

Aho-Corasick Summary

- ▶ Both failure links and output links can be computed in BFS of pattern trie
 - ~~ total construction time $O(m)$
- ▶ for read mapping, split reads into batches, so Aho-Corasick automaton of batch fits in memory
 - ~~ e. g., 300 batches of 1M reads for human genome
 - ~~ substantial speedup over reading the reference genome millions of times!

Part II

Inexact Matches

Robust Mapping

- ▶ so far: only able to find reads that *perfectly* match reference genome
 - ~~ biologically insufficient
- ▶ For the moment, focus again on a single pattern $P[0..m)$

5.6 Alignment-Based Matching

Fitting Alignment

► **Gold Standard:** best *alignment score* over all substrings of T with P .

~~ Can find this as a pairwise alignment! ~~ Recall Unit 3

Semilocal Alignment a.k.a. Fitting Alignment

Slight twist: We know conserved region, but need to find best match in larger sequence.

What substring of $B[0..n]$ is the best match for $A[0..m]$? (typically then $m \ll n$)

~~ only allow $\text{IgnorePrefix}(B[0..j])$ and $\text{IgnoreSuffix}(B[j..n])$

$$\sim D(i, j) = \begin{cases} -i & \text{if } j = 0 \\ 0 & \text{if } i = 0 \\ \max \begin{cases} D(i-1, j) - 1, \\ D(i, j-1) - 1, \\ D(i-1, j-1) + [A[i-1] = B[j-1]] \end{cases} & \text{otherwise} \end{cases}$$

Optimal local alignment score: $\max_{j \in [0..n]} D[m][j]$

13

~~ naive time and space complexity $\Theta(n \cdot m)$ (ignoring exhaustive tabulation)

Slightly more space efficient fitting alignments

Idea: don't rely on minimization over i' at end

- $D(i, j)$ = score of best alignment of $T[i'..i]$ to $P[0..j)$ for any $i' \in [0..i]$

$$D(i, j) = \begin{cases} \mathbf{0} & \text{if } j = 0 \\ j \cdot g & \text{if } i = 0 \\ D(i - 1, j) & \text{if } j = m \\ \min \begin{cases} D(i - 1, j) + g \cdot [\mathbf{j < m}], \\ D(i, j - 1) + g, \\ D(i - 1, j - 1) + p(A[i-1], B[j-1]) \end{cases} & \text{otherwise} \end{cases}$$

↝ same time complexity $\Theta(n \cdot m)$ (ignoring exhaustive tabulation)

- but better space usage
 - suffices to last 2 columns ↝ $\Theta(m)$ space
 - can remember i' in first case ↝ no need to reconstruct alignment for match

👎 still usually prohibitively slow

👎 even limiting to edit distance $\leq k$ (banded alignment) doesn't speed up DP

5.7 Longest Common Extensions

***k*-Mismatch Matching**

Let's first simplify the problem a bit.

- ▶ instead of a match of P with minimal **edit** distance to a substring of T
find a match of P with minimal **Hamming** distance to a substring of T
- ~~> only length- m substrings of $T[i..i + m)$ relevant, $i \in [0..n - m)$

How fast can we solve this problem?

OPEN!

- ▶ survey of techniques
-  **Navarro:** *A Guided Tour to Approximate String Matching*, ACM Comp. Surv. 2001
- ▶ recent work at FOCS 2020, FOCS 2023

Hamming distance – Brute Force

- We can easily adapt the brute-force string matching to
 - find all matches with $\leq k$ -mismatches
 - find a match with fewest mismatches

```
1 procedure kMismatchFind( $T[0..n], P[0..m], k$ ):  
2     for  $i := 0, \dots, n - m - 1$  do  
3          $mismatches := 0$   
4         for  $j := 0, \dots, m - 1$  do  
5             if  $T[i + j] \neq P[j]$  then  $mismatches := mismatches + 1$   
6             if  $mismatches > k$  then break inner loop  
7             if  $j == m$  then return  $i$   
8     return NO_MATCH
```

-  Simple, easy to parallelize
-  No extra space, no preprocessing
-  time $\Theta(nm)$ in the worst case
even when $k \ll m$

Longest Common Extensions

- To do better, we need a way to quickly skip matching parts.

Enter: ***longest common extension (LCE)*** data structure:

- **Given:** String $T[0..n]$
- **Goal:** Answer LCE queries, i. e.,
given positions i, j in T ,
how far can we read the same text from there?
formally: $\text{LCE}(i, j) = \max\{\ell : T[i..i + \ell] = T[j..j + \ell]\}$




Amazing result: Can compute data structure in $\Theta(n)$ time and space
that finds any LCE in **constant(!) time**.

- we will ultimately see the full result in the coming units
- ~~ for now, use $O(1)$ LCA as black box.
- ~~ After linear preprocessing (time & space), we can find LCEs in $O(1)$ time.

Kangaroo Algorithm for approximate matching

- ▶ Note: LCE defined for single string; we will need it for i in T and j in P
- ~~ build LCE data structure on $T' = T\$_1P\$_2$ in time $O(n + m)$



```
1 procedure kMismatch( $T[0..n - 1]$ ,  $P[0..m - 1]$ ):  
2     // build LCE data structure  
3     for  $i := 0, \dots, n - m - 1$  do  
4          $mismatches := 0$ ;  $t := i$ ;  $p := 0$   
5         while  $mismatches \leq k \wedge p < m$  do  
6              $\ell := \text{LCE}_{T'}(t, p)$  // jump over matching part  
7              $t := t + \ell + 1$ ;  $p := p + \ell + 1$   
8              $mismatches := mismatches + 1$   
9         if  $p == m$  then  
10            return  $i$ 
```

- ▶ Analysis: $\Theta(n + m)$ preprocessing + $O(n \cdot k)$ matching

👍 very efficient for small k for a **single pattern**

Kangaroo Algorithm – Discussion

- ▶ as given above, uses $\Omega(n)$ preprocessing **per pattern**,
but can build joint LCE data structure for several patterns
- 👎 (with blackbox reuse), preprocessing on T must be redone for each set of patterns
- 👎 working space $\Theta(n + m)$

5.8 SNPs via filtering

From Exact to Inexact

- ▶ Overhead of LCE data structure can be sizable.
- ▶ For very small k , a much simpler heuristic can be effective
- ▶ **Observation:** If a pattern $P[0..m)$ matches a text T with $\leq k$ mismatches, there is an **exact** match of a substring of length $\geq t := \left\lfloor \frac{m}{k+1} \right\rfloor$.
 - ~~ Break P into $k+1$ substrings P'_1, \dots, P'_{k+1} of length t (except potentially longer last substring)
 - ▶ Find “**seeds**”: exact occurrences of P'_1, \dots, P'_{k+1} (e.g., using Aho-Corasick)
 - ▶ Try to extend each seed to a full $\leq k$ -mismatch match
 - 👎 Worst case: find $\Theta(nk)$ matches of a t -substring of P , but no k -mismatches occurrence
 - ~~ Time to check all spurious matches $\Omega(nm)$
 - 👍 For $m \gg k$ and biological sequences, spurious matches are rare
 - 👍 seed-extension can be generalized to compute best **alignment** around a seed

Practical Heuristics

- ▶ many practical bioinformatics tools make use of seed extension
- ▶ FASTA and BLAST are heuristic methods for optimal local alignments
 - ▶ no approximation guarantee for returned alignments
 - ▶ BLAST widely used in practice, e. g., for GenBank searches
 - ▶ biological goal: find unknown connections, e. g., *homologous sequences*
- ▶ BLAST starts with short seeds with few mismatches
 - ▶ for a k -mer P and alignment score threshold σ , compute *all* strings P' with $dist(P, P') \leq \sigma$
 - ▶ search for exact matches of all these P'

Summary

- ▶ **exact matching:** $O(n + m + \text{output})$ for single or multiple patterns
- ▶ **k -mismatch matching:**
 - ▶ $O(nk)$ after $O(n + m)$ preprocessing
 - ▶ (exact) seed extension often linear time
- ▶ **k -distance matching:**
 - ▶ gold standard semilocal alignment $O(nm)$
 - ▶ improvement in tutorials
 - ▶ (heuristic) seed extension methods often linear time, often “precise enough”