



# ALGORITHMS OF BIOINFORMATICS



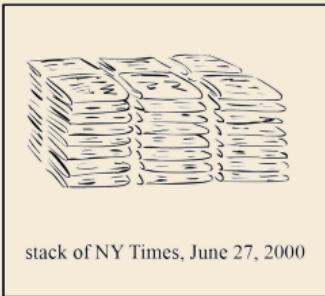
## 4 Assembling Genomes

20 November 2025

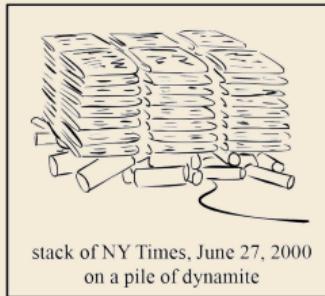
Prof. Dr. Sebastian Wild

## 4.1 Exploding Newspapers

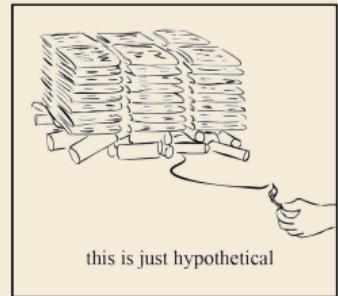
# Exploding Newspapers



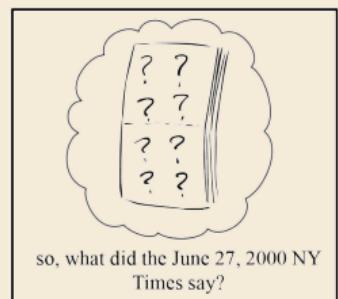
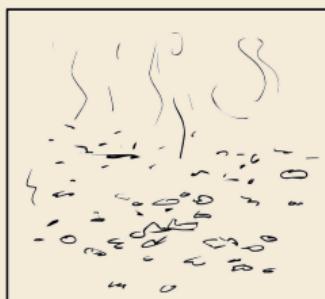
stack of NY Times, June 27, 2000



stack of NY Times, June 27, 2000  
on a pile of dynamite



this is just hypothetical



so, what did the June 27, 2000 NY  
Times say?

Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.1

<https://cogniterra.org/lesson/29884/step/2?unit=21982>

# DNA Sequencing

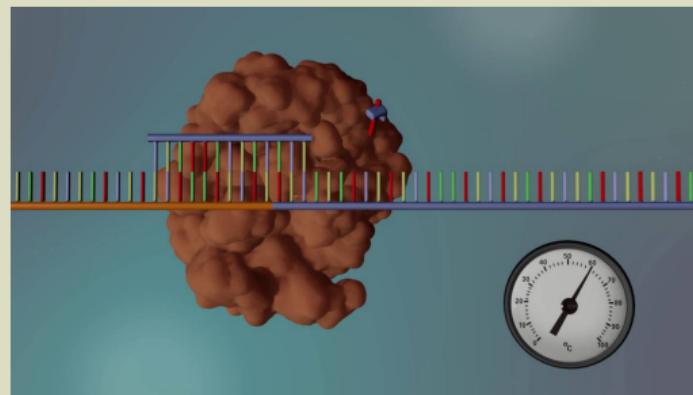
*How can we possibly “read out” DNA sequences?*

# DNA Sequencing

How can we possibly “read out” DNA sequences?

Through clever combinations of chemistry and engineering ...

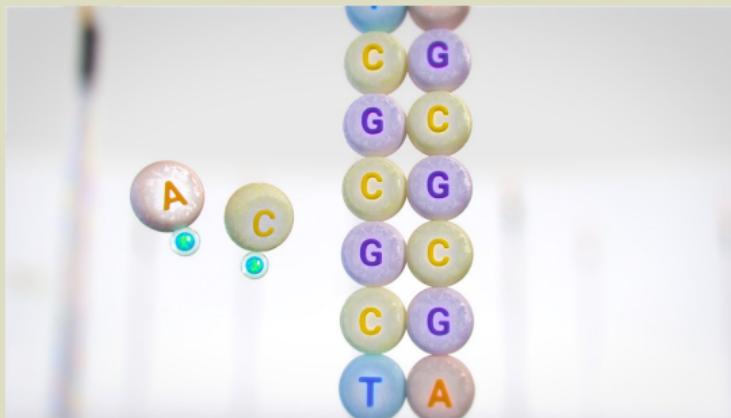
- ▶ Historically first method: *Sanger sequencing*
  - ▶ developed 1977 by Frederick Sanger
  - ▶ earning him his *second* Nobel prize in chemistry  
(the first one was for determining the amino acid sequence of insulin)



► DNA Sequencing - 3D  
<https://youtu.be/ONGdehkB8jU>

# “Next Generation Sequencing”

- ▶ group of methods that are massively parallel
- ~~ easier to automate, much cheaper per read
- ▶ commercially available since 2005



■ Overview of Illumina Sequencing by Synthesis Workflow  
<https://youtu.be/EDVKxSNdSic>

# Limitations fo Sequencing Techniques

- ▶ Sequencing machines can produce somewhat reliable *reads* of a few hundred to a few thousand bases.
- ~~ NGS will produce many reads, covering much of sequence but we need to *assemble* entire DNA sequence

# Overview

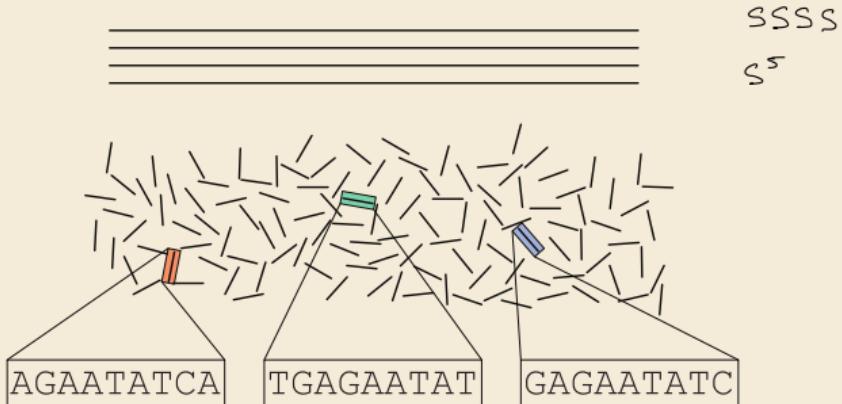
## Overview of Sequencing

Multiple identical copies of a genome

Shatter the genome into reads

Sequence the reads

Assemble the genome using overlapping reads



AGAATATCA  
GAGAATATC  
TGAGAATAT  
... TGAGAATATCA ...

Compeau & Pevzner, *Bioinformatics Algorithms*, Fig. 4.13  
<https://cogniterra.org/lesson/29918/step/2?unit=22015>

In the following: Consider an idealized situation.

## 4.2 The De-Novo Sequencing Problem

# Formalization of Sequence Assembly

- $k$ -mer = length- $k$ -(sub)string

# Formalization of Sequence Assembly

- ▶  $k$ -mer = length- $k$ -(sub)string
- ▶ Given text (genome)  $T[0..N]$  and  $k$ ,  
the  $k$ -mer composition of  $T$  is the multiset of all  $k$ -mers in  $T$ :

$$\text{Composition}_k(T) = \{\{T[0..k), T[1..1+k), T[2..2+k), \dots, T[n-k, n)\}\}$$

# Formalization of Sequence Assembly

- ▶  $k$ -mer = length- $k$ -(sub)string
- ▶ Given text (genome)  $T[0..N]$  and  $k$ ,  
the  $k$ -mer composition of  $T$  is the multiset of all  $k$ -mers in  $T$ :  
$$\text{Composition}_k(T) = \{\{T[0..k], T[1..1+k], T[2..2+k], \dots, T[n-k, n]\}\}$$
- ▶ Example:  $\text{Composition}_3(\text{TAT}\underline{\text{GGGG}}\text{TGC}) = \{\{\text{ATG}, \underline{\text{GGG}}, \underline{\text{GGG}}, \text{GGT}, \text{GTG}, \text{TAT}, \text{TGC}, \text{TGG}\}\}$

# Formalization of Sequence Assembly

- ▶  $k$ -mer = length- $k$ -(sub)string
- ▶ Given text (genome)  $T[0..N]$  and  $k$ ,  
the  *$k$ -mer composition* of  $T$  is the multiset of all  $k$ -mers in  $T$ :  
$$\text{Composition}_k(T) = \{\{T[0..k], T[1..1+k], T[2..2+k], \dots, T[n-k, n]\}\}$$
- ▶ Example:  $\text{Composition}_3(\text{TATGGGGTGC}) = \{\{\text{ATG}, \text{GGG}, \text{GGG}, \text{GGT}, \text{GTG}, \text{TAT}, \text{TGC}, \text{TGG}\}\}$
- ▶ **(De-Novo) Sequence Assembly Problem:**
  - ▶ Given multiset of  $k$ -mers  $\mathcal{R}$
  - ▶ Find a string  $T$  such that  $\text{Composition}_k(T) = \mathcal{R}$

# Formalization of Sequence Assembly

- ▶  $k$ -mer = length- $k$ -(sub)string
- ▶ Given text (genome)  $T[0..N]$  and  $k$ ,  
the  $k$ -mer composition of  $T$  is the multiset of all  $k$ -mers in  $T$ :  
$$\text{Composition}_k(T) = \{\{T[0..k], T[1..1+k], T[2..2+k], \dots, T[n-k, n]\}\}$$
- ▶ Example:  $\text{Composition}_3(\text{TATGGGGTGC}) = \{\{\text{ATG}, \text{GGG}, \text{GGG}, \text{GGT}, \text{GTG}, \text{TAT}, \text{TGC}, \text{TGG}\}\}$
- ▶ **(De-Novo) Sequence Assembly Problem:**
  - ▶ Given multiset of  $k$ -mers  $\mathcal{R}$
  - ▶ Find a string  $T$  such that  $\text{Composition}_k(T) = \mathcal{R}$
- ▶ Example:  $\mathcal{R} = \{\underbrace{\text{AAT}}, \underbrace{\text{ATG}}, \underbrace{\text{GTT}}, \underbrace{\text{TAA}}, \underbrace{\text{TGT}}\} \quad |[\underline{0} \dots \underline{n-k}]| = n-k+1 = 5$   
$$\text{T A A T G T T}$$

# Formalization of Sequence Assembly

- ▶  $k$ -mer = length- $k$ -(sub)string
- ▶ Given text (genome)  $T[0..N]$  and  $k$ ,  
the  *$k$ -mer composition* of  $T$  is the multiset of all  $k$ -mers in  $T$ :  
$$\text{Composition}_k(T) = \{\{T[0..k], T[1..1+k], T[2..2+k], \dots, T[n-k, n]\}\}$$
- ▶ Example:  $\text{Composition}_3(\text{TATGGGGTGC}) = \{\{\text{ATG}, \text{GGG}, \text{GGG}, \text{GGT}, \text{GTG}, \text{TAT}, \text{TGC}, \text{TGG}\}\}$
- ▶ **(De-Novo) Sequence Assembly Problem:**
  - ▶ Given multiset of  $k$ -mers  $\mathcal{R}$
  - ▶ Find a string  $T$  such that  $\text{Composition}_k(T) = \mathcal{R}$
- ▶ Example:  $\mathcal{R} = \{\{\text{AAT}, \text{ATG}, \text{GTT}, \text{TAA}, \text{TGT}\}\}$ 
  - ▶ form overlapping pairs
  - ▶  $T$  must start with **TTA** as no 3-mer ends with **TA** and end with **GTT**

# Formalization of Sequence Assembly

- ▶  $k$ -mer = length- $k$ -(sub)string
  - ▶ Given text (genome)  $T[0..N]$  and  $k$ ,  
the  $k$ -mer composition of  $T$  is the multiset of all  $k$ -mers in  $T$ :  
$$\text{Composition}_k(T) = \{\{T[0..k], T[1..1+k], T[2..2+k], \dots, T[n-k, n]\}\}$$
  - ▶ Example:  $\text{Composition}_3(\text{TATGGGGTGC}) = \{\{\text{ATG}, \text{GGG}, \text{GGG}, \text{GGT}, \text{GTG}, \text{TAT}, \text{TGC}, \text{TGG}\}\}$
  - ▶ (De-Novo) Sequence Assembly Problem:
    - ▶ Given multiset of  $k$ -mers  $\mathcal{R}$
    - ▶ Find a string  $T$  such that  $\text{Composition}_k(T) = \mathcal{R}$
  - ▶ Example:  $\mathcal{R} = \{\{\text{AAT}, \text{ATG}, \text{GTT}, \text{TAA}, \text{TGT}\}\}$ 
    - ▶ form overlapping pairs
    - ▶  $T$  must start with **TTA** as no 3-mer ends with **TA** and end with **GTT**
- TAA  
AAT  
ATG  
TGT  
GTT  
TAATGTT

**... not always as easy**

Now consider  $\mathcal{T} = \{\{\text{AAT}, \text{ATG}, \text{ATG}, \text{ATG}, \text{CAT}, \text{CCA}, \text{GAT}, \text{GCC}, \text{GGA}, \text{GGG}, \text{GTT}, \text{TAA}, \text{TGC}, \text{TGG}, \text{TGT}\}\}$

## ... not always as easy

Now consider  $\mathcal{T} = \{\text{AAT, ATG, ATG, ATG, CAT, CCA, GAT, GCC, GGA, GGG, GTT, TAA, TGC, TGG, TGT}\}$

Let's start again with **TTA**.

The first few steps are forced.

TAA  
AAT  
ATG  
TAATG

Then we have there are 3 options **TG[CGT]**.

Say we pick **TGT**

## *... not always as easy*

Now consider  $\mathcal{T} = \{\{AAT, ATG, ATG, ATG, CAT, CCA, GAT, GCC, GGA, GGG, GTT\}, \{TAA, TGC, TGG, TGT\}\}$

Let's start again with **TTA**.

The first few steps are forced.

TAA  
AAT  
ATG  
TAATG

Then we have there are 3 options **TG[CGT]**.

Say we pick **TGT**

TAA  
AAT  
ATG  
TGT  
GTT  
TAATGTT

## ... not always as easy

Now consider  $\mathcal{T} = \{\text{AAT, ATG, ATG, ATG, CAT, CCA, GAT, GCC, GGA, GGG, GTT, TAA, TGC, TGG, TGT}\}$

Let's start again with **TTA**.

The first few steps are forced.

TAA  
AAT  
ATG  
TAATG

Then we have there are 3 options **TG[CGT]**.

Say we pick **TGT**

TAA  
AAT  
ATG  
TGT  
GTT  
TAATGTT ⚡

We've reached a dead end!

If we instead choose more wisely,  
we can continue:

TAA  
AAT  
ATG  
TGC  
TAATGC

## ... not always as easy

Now consider  $\mathcal{T} = \{\text{AAT, ATG, ATG, ATG, CAT, CCA, GAT, GCC, GGA, GGG, GTT, TAA, TGC, TGG, TGT}\}$

Let's start again with **TTA**.

The first few steps are forced.

TAA  
AAT  
ATG  
TAATG

Then we have there are 3 options **TG[CGT]**.

Say we pick **TGT**

TAA  
AAT  
ATG  
TGT  
GTT  
TAATGTT 

We've reached a dead end!

If we instead choose more wisely,  
we can continue:

TAA  
AAT  
ATG  
TGC  
GCC  
CCA  
CAT  
ATG  
TGG  
GGA  
GAT  
ATG  
TGT  
GTT  
TAATGCCATGGATGTT

## ... not always as easy

Now consider  $\mathcal{T} = \{\text{AAT, ATG, ATG, ATG, CAT, CCA, GAT, GCC, GGA, GGG, GTT, TAA, TGC, TGG, TGT}\}$

Let's start again with **TTA**.

The first few steps are forced.

TAA  
AAT  
ATG  
TAATG

Then we have there are 3 options **TG[CGT]**.

Say we pick **TGT**

TAA  
AAT  
ATG  
TGT  
GTT  
TAATGTT 

We've reached a dead end!

If we instead choose more wisely,  
we can continue:

TAA  
AAT  
ATG  
TGC  
GCC  
CCA  
CAT  
ATG  
TGG  
GGA  
GAT  
ATG  
TGT  
GTT  
TAATGCCATGGATGTT

GGG is missing!

## ... not always as easy

Now consider  $\mathcal{T} = \{\text{AAT, ATG, ATG, ATG, CAT, CCA, GAT, GCC, GGA, GGG, GTT, TAA, TGC, TGG, TGT}\}$

Let's start again with **TTA**.

The first few steps are forced.

TAA  
AAT  
ATG  
TAATG

Then we have there are 3 options TG[CGT].

Say we pick **TGT**

TAA  
AAT  
ATG  
TGT  
GTT  
TAATGTT 

We've reached a dead end!

If we instead choose more wisely,  
we can continue:

TAA  
AAT  
ATG  
TGC  
GCC  
CCA  
CAT  
ATG  
TGG  
GGG  
GGA  
GAT  
ATG  
TGT  
GTT  
TAATGCCATGGGATGTT

# Escaping the Dead Ends

*How can we systematically avoid having to try all options?*

# Escaping the Dead Ends

*How can we systematically avoid having to try all options?*

*Repeats make life challenging.*

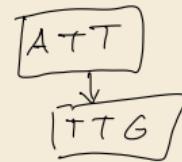


## 4.3 Assembly with Overlap Graphs

# Read Overlap Graphs

Given a  $k$ -mer composition  $\mathcal{R} = \{R_0, \dots, R_{n-k}\}$ ,  
construct the *overlap graph* as directed graph  $G_O(\mathcal{R}) = (V, E)$

- ▶  $V = \{r_0, \dots, r_{n-k}\}$  with  $r_i$  labeled with  $R_i$   
can have duplicate  $k$ -mers
- ▶  $r_i r_j \in E$  if  $R_i[1..k] = R_j[0..k-1]$

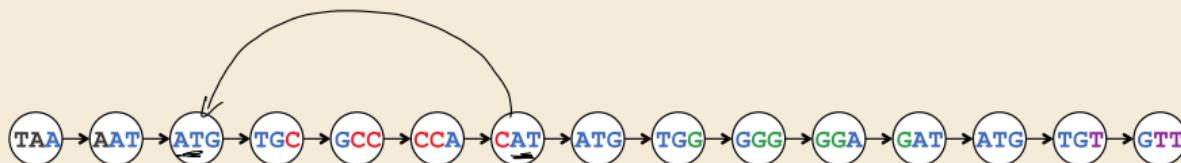


# Read Overlap Graphs

Given a  $k$ -mer composition  $\mathcal{R} = \{R_0, \dots, R_{n-k}\}$ ,  
construct the *overlap graph* as directed graph  $G_O(\mathcal{R}) = (V, E)$

- ▶  $V = \{r_0, \dots, r_{n-k}\}$  with  $r_i$  labeled with  $R_i$   
can have duplicate  $k$ -mers
- ▶  $r_i r_j \in E$  if  $R_i[1..k] = R_j[0..k-1]$

Example Overlap Graph Backbone



TA **A**TGCC**C**ATGG**G**AT**G**T**T**

Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.7  
<https://cogniterra.org/lesson/29886/step/6?unit=21984>

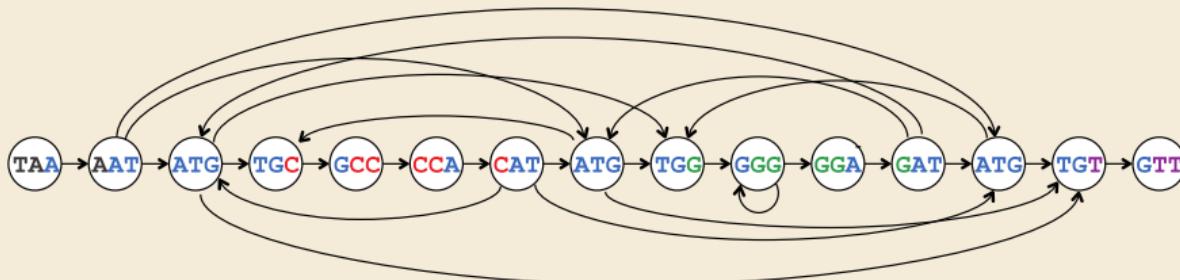
# Read Overlap Graphs

Given a  $k$ -mer composition  $\mathcal{R} = \{R_0, \dots, R_{n-k}\}$ ,  
construct the *overlap graph* as directed graph  $G_O(\mathcal{R}) = (V, E)$

- $V = \{r_0, \dots, r_{n-k}\}$  with  $r_i$  labeled with  $R_i$   
can have duplicate  $k$ -mers
- $r_i r_j \in E$  if  $R_i[1..k] = R_j[0..k-1]$

want: Hamiltonian path  
in  $G_O$

Example Overlap Graph



TAATGCCATGGGATGTT

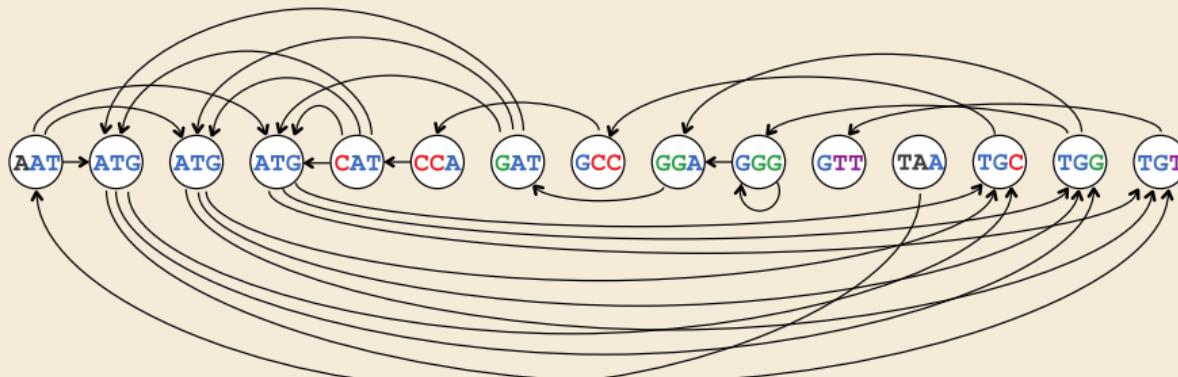
Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.7  
<https://cogniterra.org/lesson/29886/step/6?unit=21984>

# Read Overlap Graphs

Given a  $k$ -mer composition  $\mathcal{R} = \{R_0, \dots, R_{n-k}\}$ ,  
construct the *overlap graph* as directed graph  $G_O(\mathcal{R}) = (V, E)$

- ▶  $V = \{r_0, \dots, r_{n-k}\}$  with  $r_i$  labeled with  $R_i$   
can have duplicate  $k$ -mers
- ▶  $r_i r_j \in E$  if  $R_i[1..k] = R_j[0..k-1]$

Example Overlap Graph (Alphabetical)



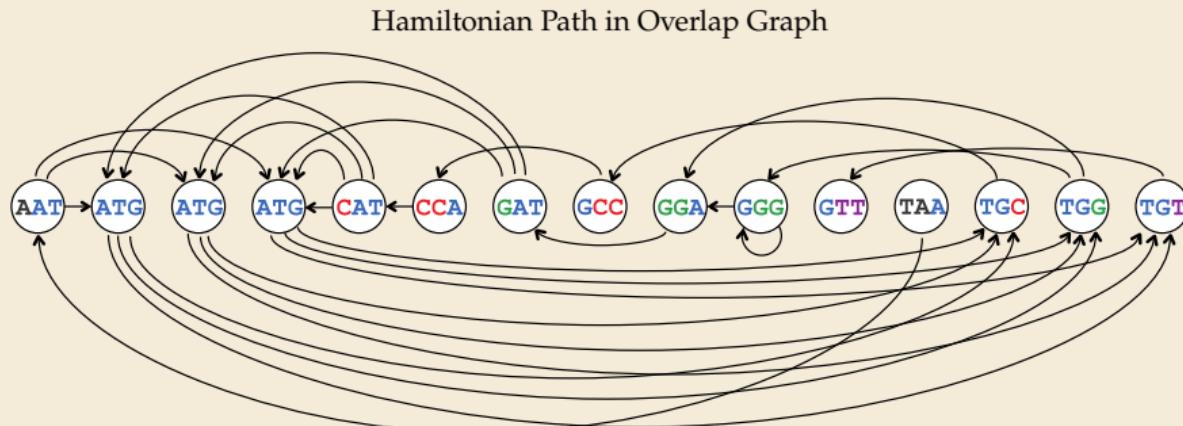
Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.8  
<https://cogniterra.org/lesson/29886/step/7?unit=21984>

## Hamiltonian Paths

Given the overlap graph  $G(\mathcal{R})$  of a  $k$ -mer composition  $\mathcal{R}$   
a sequence assembly solution is a *Hamiltonian path* in  $G(\mathcal{R})$

# Hamiltonian Paths

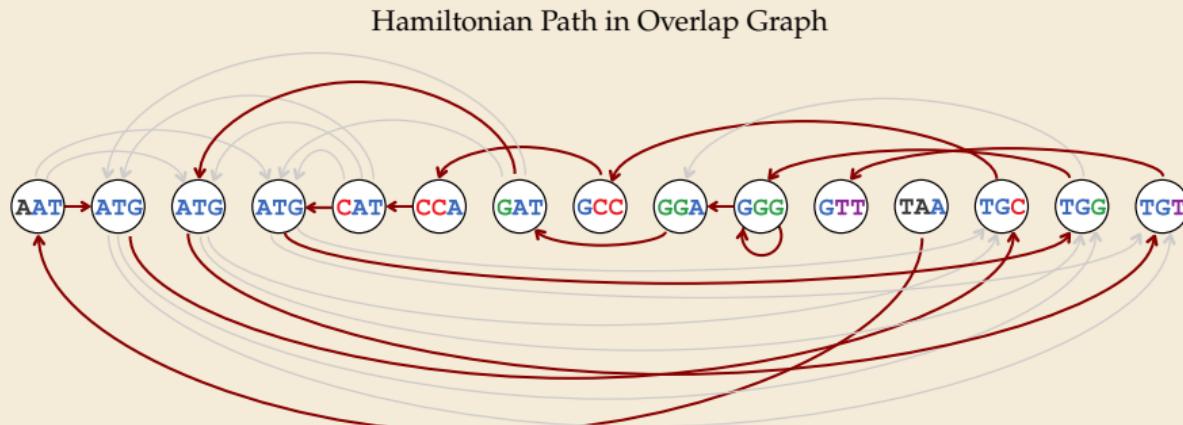
Given the overlap graph  $G(\mathcal{R})$  of a  $k$ -mer composition  $\mathcal{R}$   
a sequence assembly solution is a *Hamiltonian path* in  $G(\mathcal{R})$



Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.8  
<https://cogniterra.org/lesson/29886/step/7?unit=21984>

# Hamiltonian Paths

Given the overlap graph  $G(\mathcal{R})$  of a  $k$ -mer composition  $\mathcal{R}$   
a sequence assembly solution is a *Hamiltonian path* in  $G(\mathcal{R})$

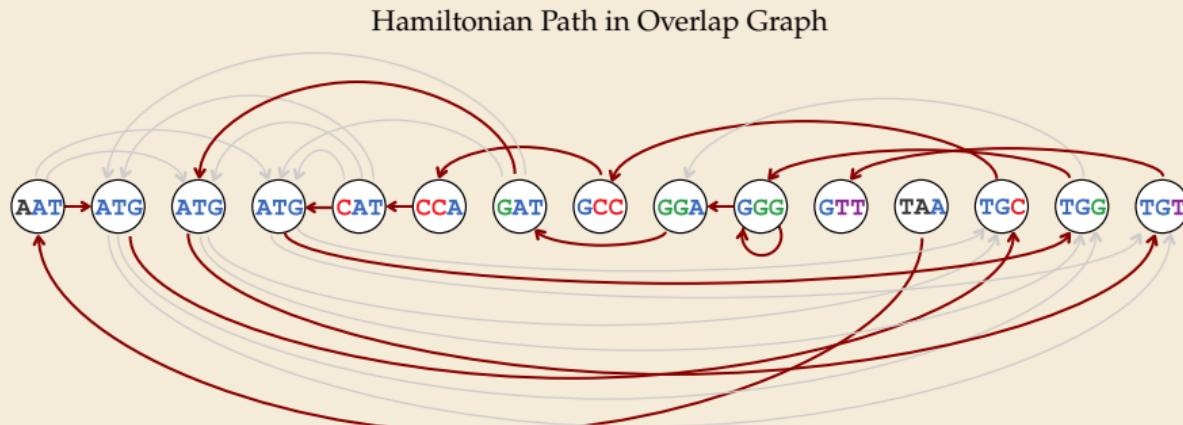


Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.9  
<https://cogniterra.org/lesson/29886/step/8?unit=21984>

- ▶ **Problem 1:** Result may not unique (as in this example)
- ▶ **Problem 2:** DIRECTED HAMILTONIAN PATH is NP-complete.

# Hamiltonian Paths

Given the overlap graph  $G(\mathcal{R})$  of a  $k$ -mer composition  $\mathcal{R}$   
a sequence assembly solution is a *Hamiltonian path* in  $G(\mathcal{R})$



Compeau & Pevzner, *Bioinformatics Algorithms*, Fig 3.9  
<https://cogniterra.org/lesson/29886/step/8?unit=21984>

- ▶ **Problem 1:** Result may not unique (as in this example)
- ▶ **Problem 2:** DIRECTED HAMILTONIAN PATH is NP-complete.

## Bad News Again

*It seems we once again ran into a hard problem.*

(And we haven't even accounted for computing the overlap graph . . . )

## Bad News Again

*It seems we once again ran into a hard problem.*

(And we haven't even accounted for computing the overlap graph . . . )

. . . but sometimes, a different way to look at the problem helps.

## 4.4 De Bruijn Graphs

# Origins

Nicolaas Govert de Bruijn's  $k$ -universal string puzzle:

How long does a binary string need to be to contain all  $2^k$  binary  $k$ -mers as substrings?

How to construct a shortest such string?

$$\leq 2^k \cdot k$$

Examples:

- ▶  $k = 2$ 
  - ▶ must contain 00, 01, 10, 11
  - ~~ a 2-universal string needs at least 5 letters
    - 4 starting positions, each for a 2-letter substring

00011

# Origins

Nicolaas Govert de Bruijn's  $k$ -universal string puzzle:

How long does a binary string need to be to contain all  $2^k$  binary  $k$ -mers as substrings?  
How to construct a shortest such string?

Examples:

- ▶  $k = 2$ 
  - ▶ must contain 00, 01, 10, 11
    - ~~ a 2-universal string needs at least 5 letters
      - 4 starting positions, each for a 2-letter substring
  - ▶ Also achievable: 00110
    - ~~ Answer: 2-universal requires 5 characters.

# Origins

Nicolaas Govert de Bruijn's  $k$ -universal string puzzle:

How long does a binary string need to be to contain all  $2^k$  binary  $k$ -mers as substrings?

How to construct a shortest such string?

Examples:

- ▶  $k = 2$ 
  - ▶ must contain 00, 01, 10, 11
    - ~~ a 2-universal string needs at least 5 letters
      - 4 starting positions, each for a 2-letter substring
  - ▶ Also achievable: 00110
    - ~~ Answer: 2-universal requires 5 characters.
- ▶  $k = 3$     ~~ a shortest 3-universal string is 0001110100

# Origins

Nicolaas Govert de Bruijn's  $k$ -universal string puzzle:

How long does a binary string need to be to contain all  $2^k$  binary  $k$ -mers as substrings?  
How to construct a shortest such string?

Examples:

- ▶  $k = 2$ 
  - ▶ must contain 00, 01, 10, 11
    - ~~ a 2-universal string needs at least 5 letters
      - 4 starting positions, each for a 2-letter substring
  - ▶ Also achievable: 00110
    - ~~ Answer: 2-universal requires 5 characters.
- ▶  $k = 3$  ~~ a shortest 3-universal string is 0001110100

Same question makes sense for strings over alphabet  $\Sigma = [0..\sigma)$ .

~~  $T$  is  $k$ -universal  $\iff$  contains all  $\sigma^k$   $k$ -mers as substring.

# de Bruijn Graphs

Given a  $k$ -mer composition  $\mathcal{R} = \{R_0, \dots, R_{n-k}\}$ ,  
construct the *de Bruijn graph* as directed multigraph  
 $G_B(\mathcal{R}) = (V, E)$

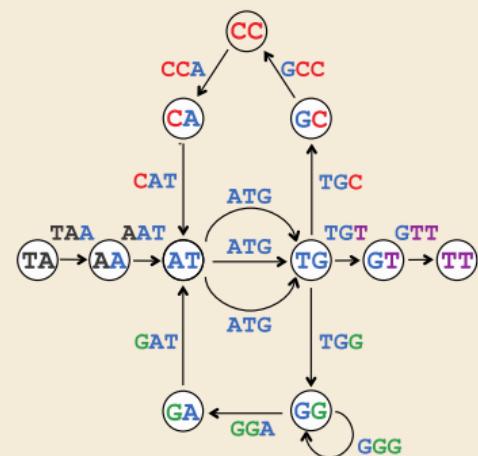
- ▶  $V = \{R[0..k-1] : R \in \mathcal{R}\}$
- ▶ For each read  $R \in \mathcal{R}$ , add the edge  $vw$  (with label  $R[0..k]$ ) to  $E$   
for  $v = R[0..k-1]$  and  $w = R[1..k]$



↔ Sequence using each  $k$ -mer once = Euler path in  $G_B$

- ▶ Euler path efficiently computable!

DEBRUIJN<sub>3</sub>(TAATGCCATGGGATGTT)

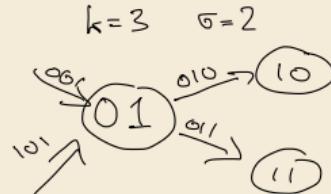


Compeau & Pevzner, *Bioinformatics Algorithms*, Fig. 4.1  
<https://cogniterra.org/lesson/29910/step/2?unit=22007>

# $k$ -universal strings

The de Bruijn graph “solves” the universal string problem for  $\mathcal{R}$  all  $\sigma^k$   $k$ -mers

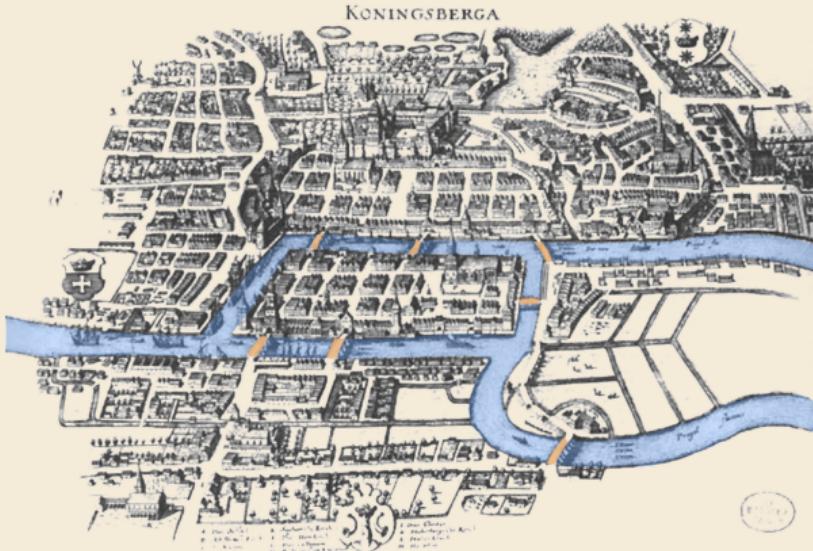
- ▶  $\sigma^{k-1}$  vertices,  $\sigma^k$  edges
- ▶ each vertex has out- and in-degree =  $\sigma$
- ~~  $G_B$  is Eulerian
- ~~ Compute Euler path (see below)



$$\text{total length} = \sigma^k + k - 1$$

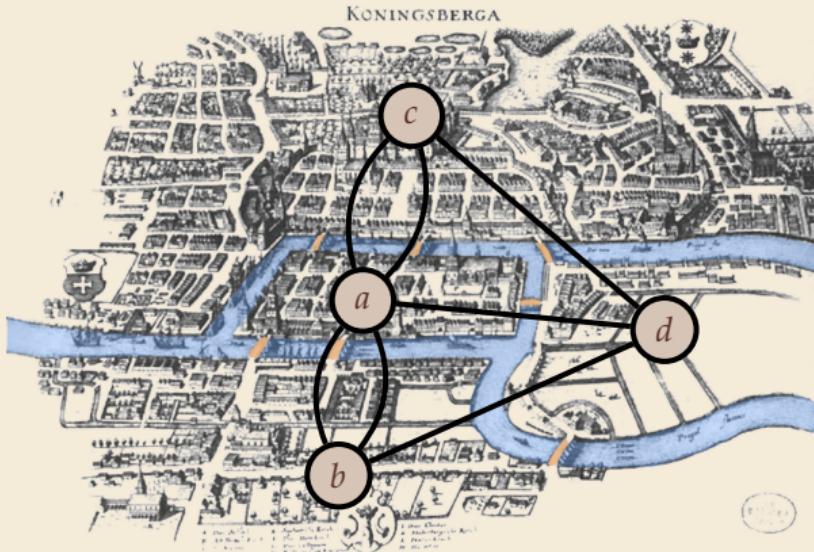
# Euler Cycles

**Euler Walk:** Walk using every edge in  $G = (V, E)$  exactly once.



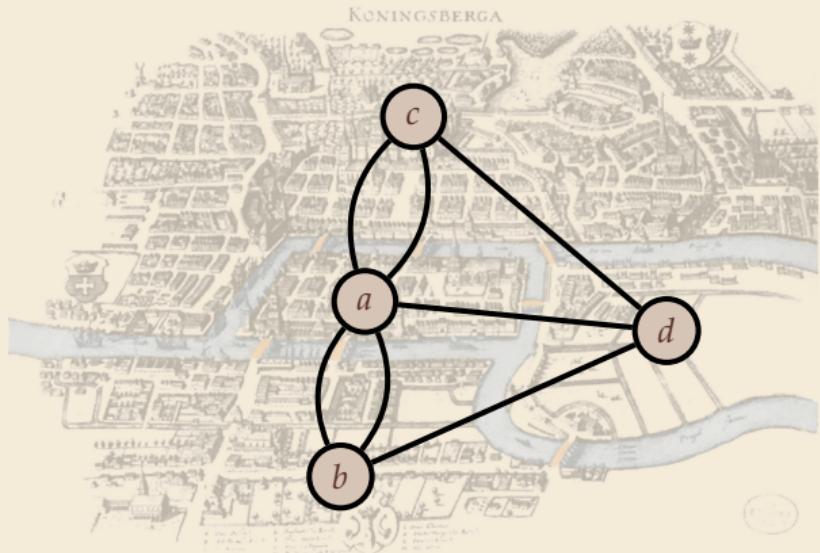
# Euler Cycles

**Euler Walk:** Walk using every edge in  $G = (V, E)$  exactly once.



# Euler Cycles

**Euler Walk:** Walk using every edge in  $G = (V, E)$  exactly once.



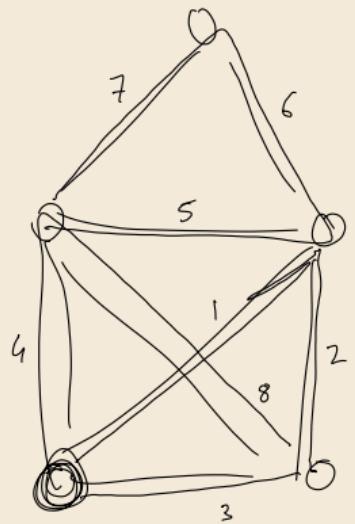
## Euler's Theorem:

Euler walk exists iff  $G$  connected and 0 or 2 vertices have odd degree.

' $\Rightarrow$ ' trivial (need to enter and exit intermediate vertices equally often)

' $\Leftarrow$ ' Following algorithm *constructs* Euler walk under this assumption





# Euler Cycles – Hierholzer’s Algorithm

- ▶ use an *edge-centric DFS*

- ▶ We mark *edges* (not vertices)
  - ~~ stack = **edge-simple walk**
- ▶ We remember iterator *i* globally per *v* to resume traversal

---

```
1 procedure eulerWalk(G):
2     // Assume G = (V, E) is connected (multi)graph
3     Vodd := {v ∈ V : d(v) odd}
4     if |Vodd| ∈ {0, 2} return NOT_EULERIAN
5     if Vodd = {x, y} then s := x else s := 0
6     euler[0..m) := NONE; j := m - 1
7     visited[0..n, 0..n) := false // mark edges as visited
8     for v := 0, ..., n - 1
9         // globally remember next unexplored edge
10        nextEdge[v] := G.adj[w].iterator()
11        edgeDFS(s)
12        return euler
```

---

```
1 procedure edgeDFS(s):
2     frontier := new Stack;
3     frontier.push(s)
4     while ¬frontier.isEmpty()
5         v := frontier.top(); i := nextEdge[v]
6         if ¬i.hasNext() // v has no unused edge
7             frontier.pop()
8             if ¬frontier.isEmpty()
9                 // assign edge leading here largest free index
10                euler[j] := (frontier.top(), v); j := j - 1
11            end if
12        else
13            w := i.next()
14            if ¬visited[v, w]
15                visited[v, w] := true
16                visited[w, v] := true // ✎
17                frontier.push(w)
18            end if
19        end if
20    end while
```

# Directed Euler Cycles

- ▶ de Bruijn graphs are *directed* ... what changes?
- ▶ **Euler's Theorem, directed version:**
  - Euler walk exists iff  $G$  is (a) strongly connected and
    - (b)  $d_{\text{out}}(v) - d_{\text{in}}(v) = 0$  for all vertices  $v$  or
      - 0 for all but 2 vertices where it is +1 and -1, respectively.
- ' $\Rightarrow$ ' need to enter and exit intermediate vertices equally often
- ' $\Leftarrow$ ' Hierholzer's algorithm also works on directed graphs, with changes below.
- ▶ Changes to Hierholzer's algorithm:
  - ▶ Check balance ( $d_{\text{out}}(v) - d_{\text{in}}(v)$ ) of vertices
  - ▶ start at the unique vertex with  $d_{\text{out}}(s) - d_{\text{in}}(s) = +1$  if  $s$  exists;  
otherwise arbitrary
  - ▶ in edgeDFS, do not mark the reverse edge as also visited

## Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>†</sup>   Detlef Plump<sup>‡</sup>   Sebastian Wild<sup>§</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicities.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their computation is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classic algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both continue to be taught in introductory courses on graph theory [1] due to their simplicity and elegance. Being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(n \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg m)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over the standard algorithm. In addition, our algorithm can be easily adapted to directed graphs; we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space required by auxiliary data structures during the computation. Our algorithm runs in linear time on the extra space required by the graph representation, which is natural in many applications. In contrast, for directed graphs, we require that we can iterate over all incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

To our knowledge, we present the first algorithm with working space  $O(n \lg m)$  bits that runs in linear time and that computes an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.

<sup>†</sup>University of Liverpool, United Kingdom (ziad.ismaili-alaeui@liverpool.ac.uk).

<sup>‡</sup>University of York, United Kingdom (detlef.plump@york.ac.uk).

<sup>§</sup>Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>†</sup>   Detlef Plump<sup>‡</sup>   Sebastian Wild<sup>§</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the existence of Eulerian cycles and their enumeration is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both continue to be taught in introductory courses on graph theory [1] and are used in optimization software. Being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(n \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg m)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over the standard algorithm. In addition, our algorithm is more robust. In particular, we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space required by auxiliary data structures during the computation. Our algorithm runs on an assumption that the graph remains static, which is satisfied in most real-life applications. In contrast, for directed graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

To our knowledge, we present the first algorithm with working space  $O(n \lg m)$  that runs in linear time and that produces an Eulerian cycle. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into



\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaoui@liverpool.ac.uk).

‡University of York, United Kingdom (detlef.plump@york.ac.uk).

§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n^2)$ -stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute a spanning tree and then traverse the remaining edges in a backtracking search. Both are slow and space-inefficient, the latter by the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over previous work. Our algorithm is based on the same basic idea as the classical ones; we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time for directed graphs, while the original algorithm can be shown to run in quadratic time for undirected graphs; for directed graphs, we require that we iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

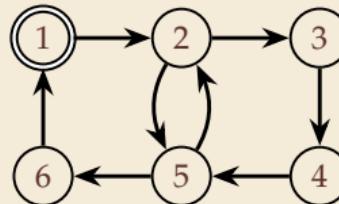
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n^2)$ -stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute a spanning tree and then traverse the remaining edges in a backtracking search. Both are slow and inefficient. In particular, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(n \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over the standard approach. Our algorithm is based on the same basic idea as the standard algorithm: we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time for directed graphs, while the standard algorithm runs in quadratic time. Our algorithm also runs in linear time for directed graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

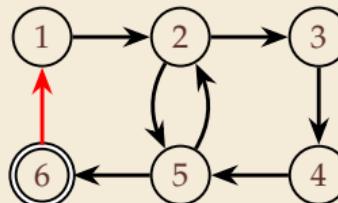
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first one that can work in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

**SOSA 2026**

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n\lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n^2)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1. Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute a spanning tree and then traverse the remaining edges in a backtracking search. They are optimal in terms of being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over the standard algorithm. Our algorithm is also more robust than the classical ones; we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time for directed graphs, we treat the input graph as a directed acyclic graph, which is natural in many applications. For undirected graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

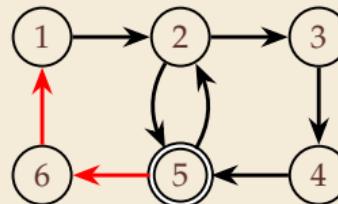
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first one that can work in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

**SOSA 2026**

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n^2)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both concepts and runs in  $O(n + m)$  time. However, they are not always optimal. For example, Fleury's algorithm is both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(n \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over previous work. Our algorithm is based on the same basic idea as the classical ones; we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time for directed graphs, while the classical variants can be quadratic. When restricted to undirected graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

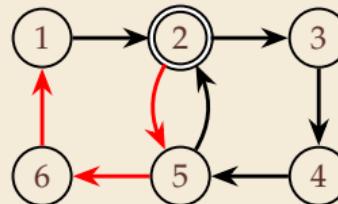
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n\lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both concepts and runs in  $O(m^2)$  time. Hierholzer's algorithm is well-known for its simplicity, but it is not optimal either. Both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(n \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over previous work. Our algorithm is based on the same basic idea as the classical ones; we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time for directed graphs, while the classical one can be quadratic. Our algorithm also runs in linear time for undirected graphs; we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

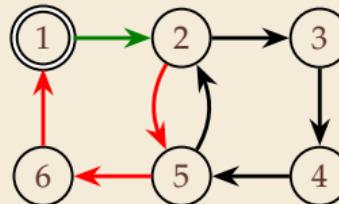
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first that can easily be implemented in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n^2)$ -stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the implementation is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute and store an Eulerian cycle in  $O(n^2)$  time and  $O(n^2)$  space, which is optimal for sparse graphs. Being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(n \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over the standard algorithm. Our algorithm is also more robust than the standard one: we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an unsorted stream of the graph represented as edges, which is natural in many applications. In contrast, for directed graphs, we require that we iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

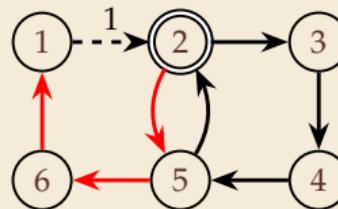
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first that can work in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute and store a stack of vertices, which is quadratic in the number of edges. In contrast, our algorithm is both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. [Here and throughout,  $\lg = \log_2$ .]

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more robust than others. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an adjacency list representation of the graph, where  $n$  is either the number of adjacent vertices or edges for directed graphs; we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

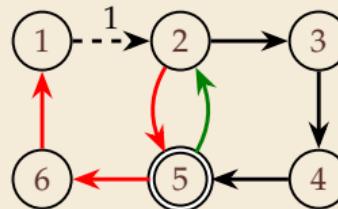
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first one that can easily be implemented in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m \cdot n)$  stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the implementation is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1. Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute and store a stack of vertices, which is quadratic in the number of edges. Our algorithm is both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. [Here and throughout,  $\lg = \log_2$ .]

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more robust than others. In particular, we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time for directed graphs, we treat that case as a special case. Our algorithm runs in linear time for undirected graphs; for directed graphs, we require that we iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

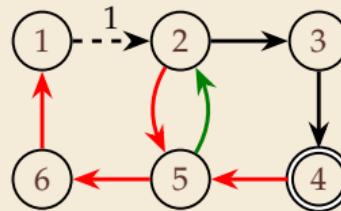
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute and store a stack of vertices, which is quadratic in the number of edges. In contrast, we propose a Being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. [Here and throughout,  $\lg = \log_2$ .]

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more robust than others. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an adjacency list representation of the graph, where  $n$  is either the number of adjacent vertices or edges for directed graphs; we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

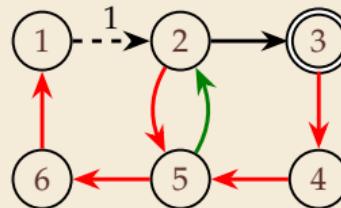
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first that can easily be implemented in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n^2)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute a spanning tree and then traverse the remaining edges in a backtracking search. Both are slow, though. Being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(n \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more robust than others. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an adjacency list representation of the graph, where  $n$  is either the number of adjacent vertices or edges for directed graphs; we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

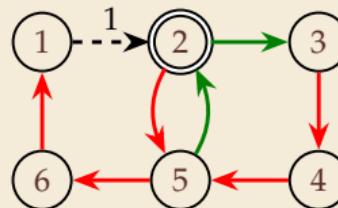
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first one that can easily be implemented in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(n^2)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute and store a stack of vertices, which is quadratic in size. Our algorithm is linear in time and space. Being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(n \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more robust than others. In particular, we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on any graph, even if the graph represents a DAG, which is natural in many applications. In contrast, for directed graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

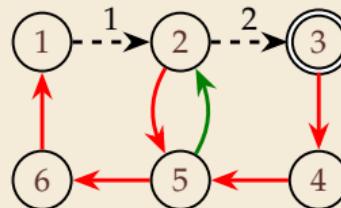
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first one that can easily be implemented in real-world code of computation, such as graph-transformation heuristics like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m \lg n)$  stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1. Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both compute and store a stack of vertices, which is quadratic in the number of edges. Our algorithm is both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more robust than others. In particular, we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an arbitrary graph, provided that it can be stored in memory. As for the adjacency lists, we require that for directed graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

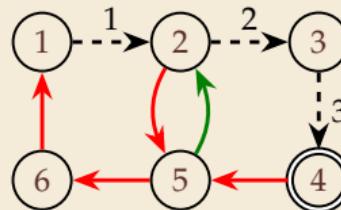
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both concepts and runs in  $O(m^2)$  time. Our algorithm is based on the same basic idea, but is much more efficient, being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. [Here and throughout,  $\lg = \log_2$ .]

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over previous work. Our algorithm is also very easy to implement. In particular, we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an arbitrary graph, provided that it can be stored in adjacent memory locations. For directed graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

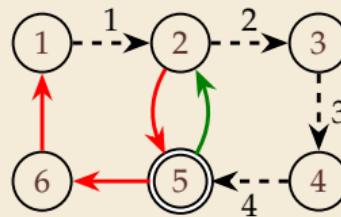
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle. Our algorithm is also the first one that can easily be implemented in real-world code, such as graph-transformation heuristics like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m^2)$ -stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both concepts and running times are similar. However, they are not always optimal in practice. Being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. (Here and throughout,  $\lg = \log_2$ .)

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more robust than the classical ones; we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an arbitrary graph, provided that it can be stored in adjacent memory locations. For directed graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

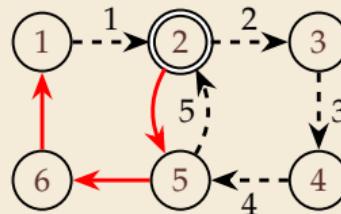
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the method is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both concepts and running times are similar. However, Fleury's algorithm is not optimal, while Hierholzer's is both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. [Here and throughout,  $\lg = \log_2$ .]

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over the standard algorithm. Our algorithm is also more robust than the standard one: we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an arbitrary graph, provided that it can be stored in adjacent memory locations. For directed graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

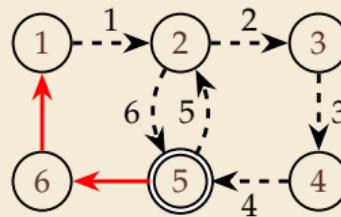
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the implementation is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both concepts are well suited for implementation in imperative languages. However, they are not optimal in terms of both space and efficiency, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. [Here and throughout,  $\lg = \log_2$ .]

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more efficient than Fleury's algorithm: we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an arbitrary graph, provided that the graph remains connected, which is satisfied in the adjacency list representation for directed graphs; we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

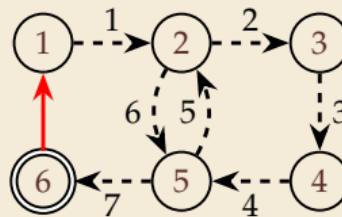
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark DASHED and output edge

# Space-Efficient Hierholzer

*frontier* can use up to  $O(m)$  extra space ( $v$  can appear  $d(v)$  times); can we avoid that?

Space-Efficient Hierholzer:  
Eulerian Cycles in  $O(m)$  Time and  $O(n)$  Space\*

Ziad Ismaili Alaoui<sup>1</sup>   Detlef Plump<sup>2</sup>   Sebastian Wild<sup>3</sup>

**Abstract.** We describe a simple variant of Hierholzer's algorithm that finds an Eulerian cycle in a (multi)graph with  $n$  vertices and  $m$  edges using  $O(n \lg n)$  bits of working memory. This substantially improves the working space compared to standard implementations of Hierholzer's algorithm, which use  $O(m \lg n)$  bits of space. Our algorithm runs in linear time, like the classical version, but avoids an  $O(m)$ -size stack of vertices or storing information for each edge. To our knowledge, this is the first linear-time algorithm to achieve this space bound, and the implementation is very easy to implement. The correctness argument, by contrast, is surprisingly subtle; we give a detailed formal proof. The space savings are particularly relevant for dense graphs or multigraphs with large edge multiplicity.

**1 Introduction.** An *Eulerian cycle* in a graph is a closed walk that traverses every edge exactly once. A graph is called *Eulerian* if it contains an Eulerian cycle. Euler's study of the relation of Eulerian cycles and their applications is often cited as the birth of modern graph theory. Apart from historical significance, efficiently computing them has applications in, e.g., genome assembly [7], routing problems [6], and 3D printing [9].

Classical algorithms such as Fleury's algorithm or Hierholzer's algorithm [5], both dating from the late 19th century, are simple to describe and find an Eulerian cycle in any Eulerian graph in polynomial time. Indeed, both concepts are well suited for a backtracking search. However, they are not optimal in terms of space usage, being both simple and efficient, the latter is the method of choice in practice. Typical implementations of Hierholzer's algorithm (see below) use  $O(m \lg n)$  bits of working memory. [Here and throughout,  $\lg = \log_2$ .]

In this paper, we present a simple variant of Hierholzer's algorithm that uses only  $O(n \lg n)$  bits of working memory and runs in optimal  $O(n + m)$  time. For dense graphs or multigraphs with many parallel edges, this is a substantial improvement over existing implementations. Our algorithm is also more robust than the classical ones; we treat the input graph as given in read-only memory and produce the Eulerian cycle in order to a write-only stream. In particular, our algorithm would be suitable for an application that directly consumes and uses edges of the Eulerian cycle as they are computed, potentially without ever storing the entire cycle. By working memory, we mean the extra space used by auxiliary data structures during the computation. Our algorithm runs in linear time on an arbitrary graph, provided that one can easily iterate through all adjacent edges. For undirected graphs, we require that we can iterate over incoming and outgoing edges; for undirected graphs, we (instead) require a consistent ordering of vertices across all adjacency lists.

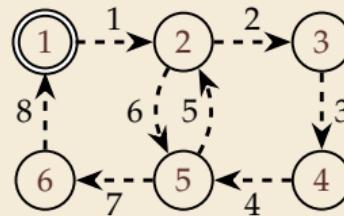
To our knowledge, we present the first algorithm with working space  $O(n \lg n)$  bits that runs in linear time and finds an Eulerian cycle in any Eulerian graph. Our algorithm also works in rule-based models of computation, such as graph-transformation logographs like GP2 [8], which do not natively support stacks or recursion. This is not known to be true for other standard efficient implementations of Eulerian cycle algorithms.

**1.1 Hierholzer's Algorithm.** Hierholzer's algorithm was originally described in 1872 [5], unsurprisingly without detailed information about data structures for efficient execution on a computer. The original algorithm consists of following an arbitrary walk in an Eulerian graph until we get stuck, which can only happen when we closed a cycle. If this cycle has not used all edges yet, starting at such an unused edge, again following an arbitrary walk of unused edges can again only get stuck when closing a cycle; we can then rise this new cycle into

\*Ziad Ismaili Alaoui and Sebastian Wild are supported by EPSRC grant EP/V009427/1.  
†University of Liverpool, United Kingdom (ziad.ismaili-alaeoui@liverpool.ac.uk).  
‡University of York, United Kingdom (detlef.plump@york.ac.uk).  
§Philipps-Universität Marburg, Germany, and University of Liverpool, United Kingdom (wild@informatik.uni-marburg.de).

SOSA 2026

Insight: Euler walk doesn't require full edge DFS  
(backtracking edges in reverse order of forward traversal)  
Suffices to not get stuck!



edges initially **BLACK**  
double line = current vertex

We do a *reverse* edge DFS      ↵ Euler walk in correct order

**RED** = DFS tree edges

vertex already seen      ↵ mark **GREEN**

need to backtrack?      ↵ prefer **GREEN**

whenever we backtrack, mark **DASHED** and output edge

# Space-Efficient Hierholzer – Implementing colors

**Assumption:** Can iterate over outgoing **and** incoming edges

- ▶ Every vertex has at most one outgoing **RED** edge
  - ~~ store explicitly in  $O(n)$  space
- ▶ edges are visited (**RED** or **GREEN**) iff before current iterator in adjacency list
  - ~~ stored implicitly!

total space  $\mathcal{O}(n)$

## Back to Genomes

- ▶ Eulerian walk usually not unique or doesn't exist!
  - ~~ often stop with assembling *contigs*, long contiguous substrings of the genome to be further processed
- ▶ contigs can be found from de Bruijn graph using vertex balances

## 4.5 Practical Assemblers

# Practical Complications

Challenges for DNA sequencing in practice

- ▶ DNA is *not* a single string; one per chromosome!
- ▶ we have (most) chromosomes twice (mom's and dad's), probably similar in large parts!
- ▶ imperfect coverage: some regions might not be covered by any/enough reads
- ▶ sequencing isn't exact; small reading errors need to be tolerated
- ▶ DNA is double stranded ↗ get reads from both strands mixed up

# Some solutions

## Imperfect coverage

- ▶ Rather unlikely to see a read for every possible starting position.
- ▶ But this is harmless; we can break reads of length  $\ell$  into artificial  $k$ -mers for some  $k < \ell$  and obtain perfect coverage!

## Chromosome pairs

- ▶ problematic because of high similarity
- ▶ possible in lab to avoid sequencing both

Inexact Overlaps seem to be the most problematic.

# There's life in the old dog yet!

- ▶ Inexact reads actually much easier to handle in **overlap graph**  
weighted edges for length / similarity of overlap
- ▶ highly effective heuristic: preprocessing overlap graph to transitive reduction
- ▶ Hamiltonian path / TSP tour not the bottleneck      "hard in theory, easy in practice"
- ▶ challenge there: how to compute overlap graph (edge weights) efficiently  
~~ we will come back to this