

# 3

# Efficient Sorting

*18 February 2020*

Sebastian Wild

# Outline

## 3 Efficient Sorting

- 3.1 Mergesort
- 3.2 Quicksort
- 3.3 Comparison-Based Lower Bound
- 3.4 Integer Sorting
- 3.5 Parallel computation
- 3.6 Parallel primitives
- 3.7 Parallel sorting

# Why study sorting?

- ▶ fundamental problem of computer science that is still not solved
- ▶ building brick of many more advanced algorithms
  - ▶ for preprocessing
  - ▶ as subroutine
- ▶ playground of manageable complexity to practice algorithmic techniques

Algorithm with optimal #comparisons in worst case?



Here:

- ▶ “classic” fast sorting method
- ▶ parallel sorting

# Part I

## *The Basics*

# Rules of the game

► **Given:**

- ▶ array  $A[0..n - 1]$  of  $n$  objects
- ▶ a total order relation  $\leq$  among  $A[0], \dots, A[n - 1]$   
(a comparison function)

► **Goal:** rearrange (=permute) elements within  $A$ ,  
so that  $A$  is *sorted*, i. e.,  $A[0] \leq A[1] \leq \dots \leq A[n - 1]$

- for now:  $A$  stored in main memory (*internal sorting*)  
single processor (*sequential sorting*)



## 3.1 Mergesort

# Clicker Question



How does mergesort work?

- A** Split elements around median, then recurse on small / large elements.
- B** Recurse on left / right half, then combine sorted halves.
- C** Grow sorted part on left, repeatedly add next element to sorted range.
- D** Repeatedly choose 2 elements and swap them if they are out of order.
- E** Don't know.

*pingo.upb.de/622222*

# Clicker Question



How does mergesort work?

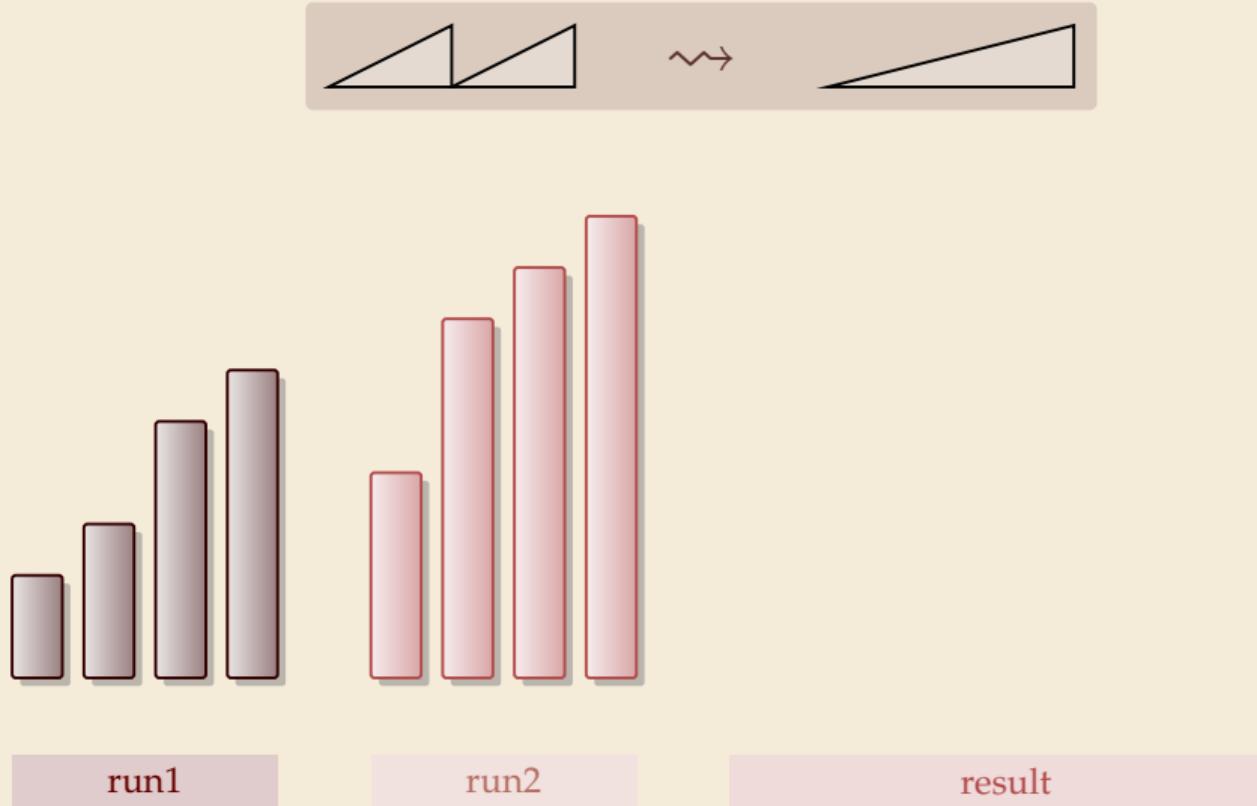
- A** ~~Split elements around median, then recurse on small / large elements.~~
- B** Recurse on left / right half, then combine sorted halves. ✓
- C** ~~Grow sorted part on left, repeatedly add next element to sorted range.~~
- D** ~~Repeatedly choose 2 elements and swap them if they are out of order.~~
- E** ~~Don't know.~~

*pingo.upb.de/622222*

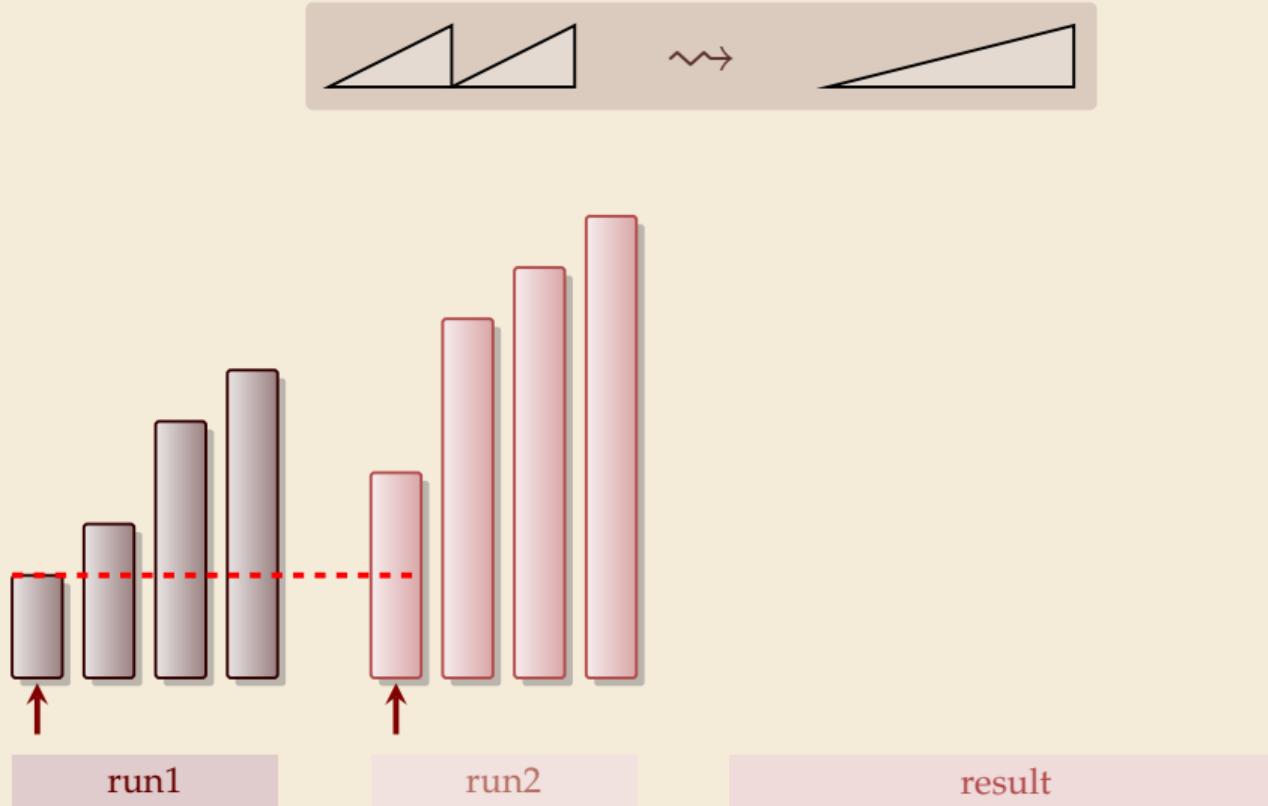
## Merging sorted lists



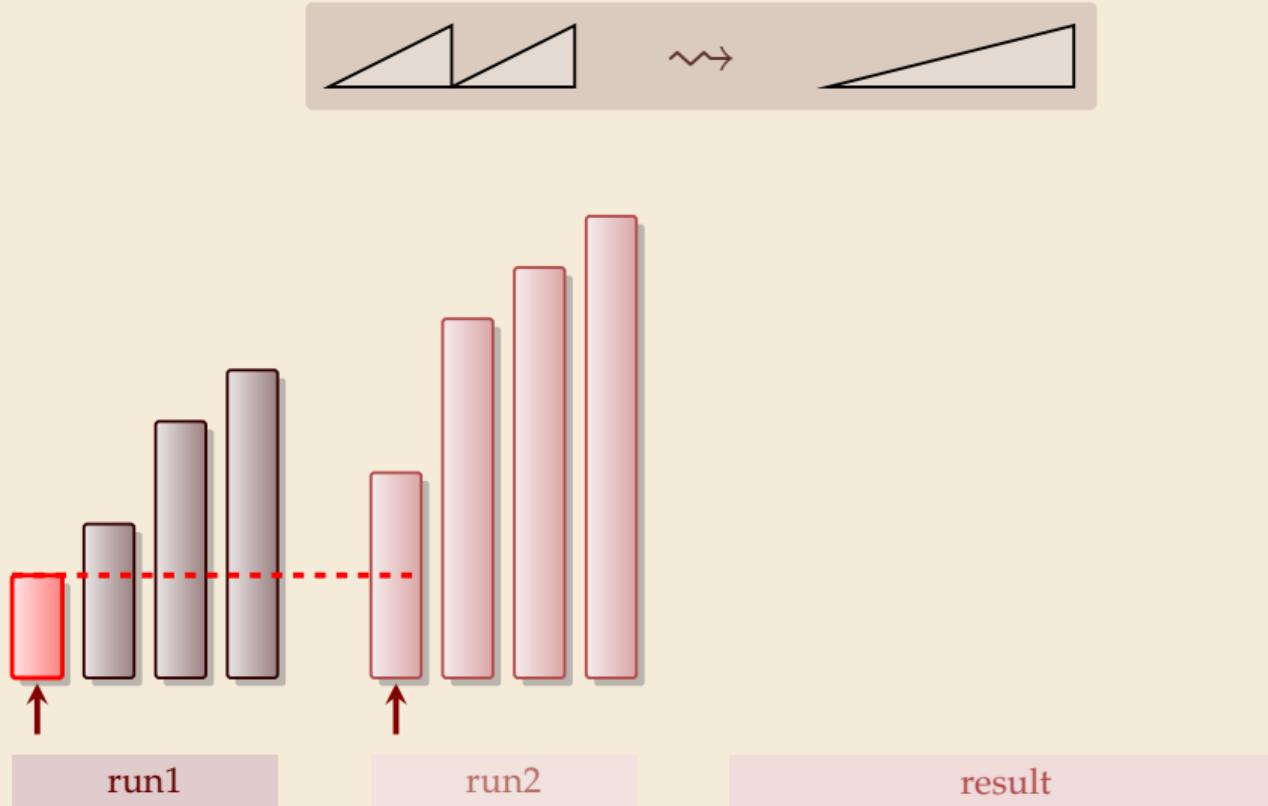
# Merging sorted lists



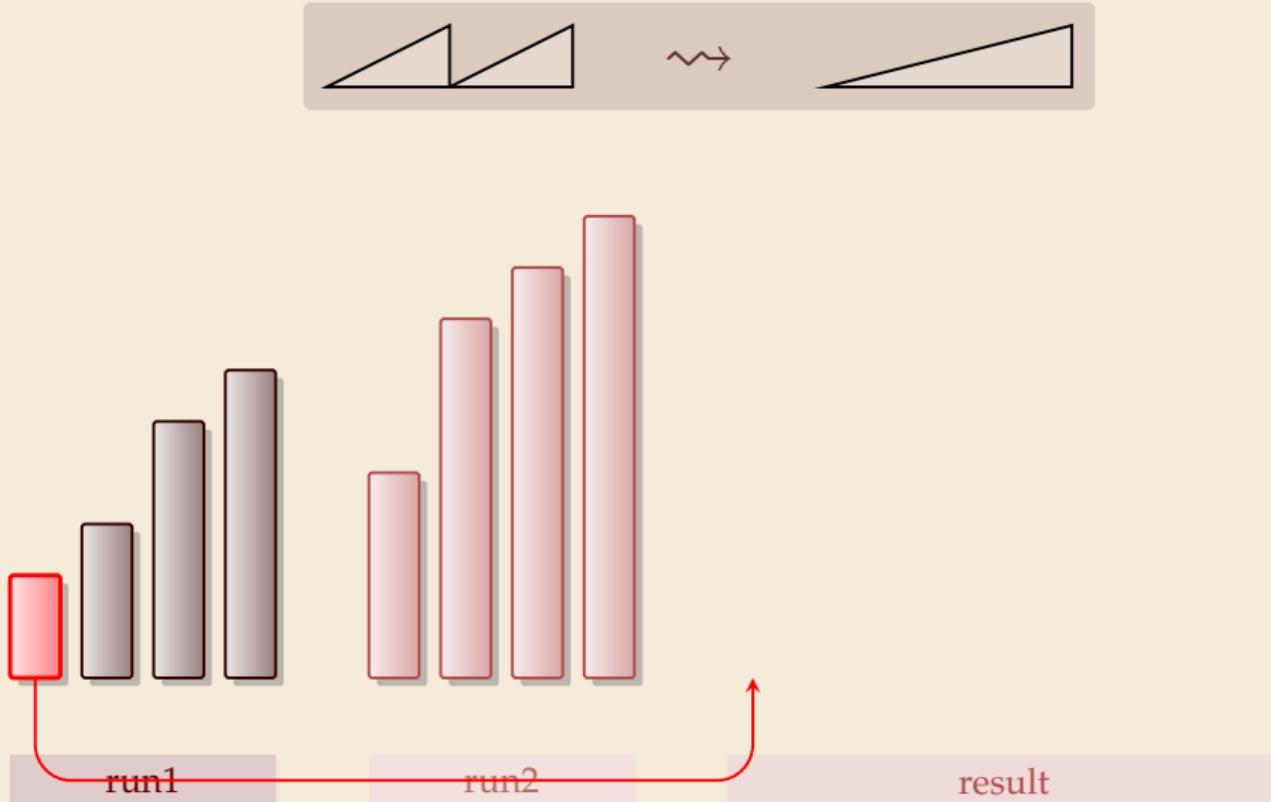
# Merging sorted lists



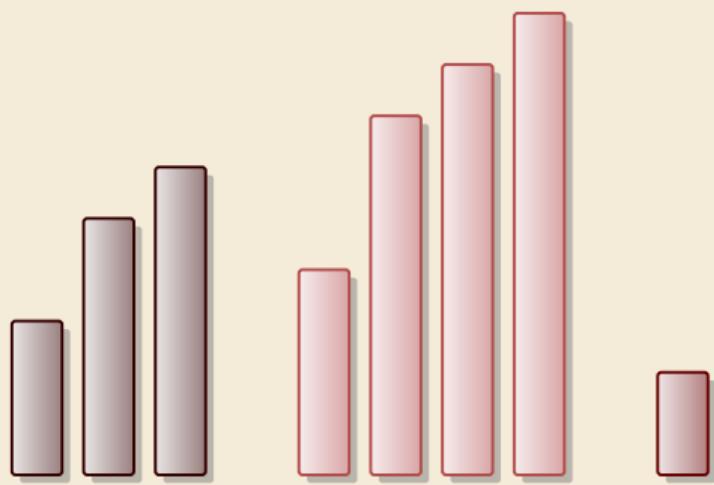
## Merging sorted lists



## Merging sorted lists



# Merging sorted lists

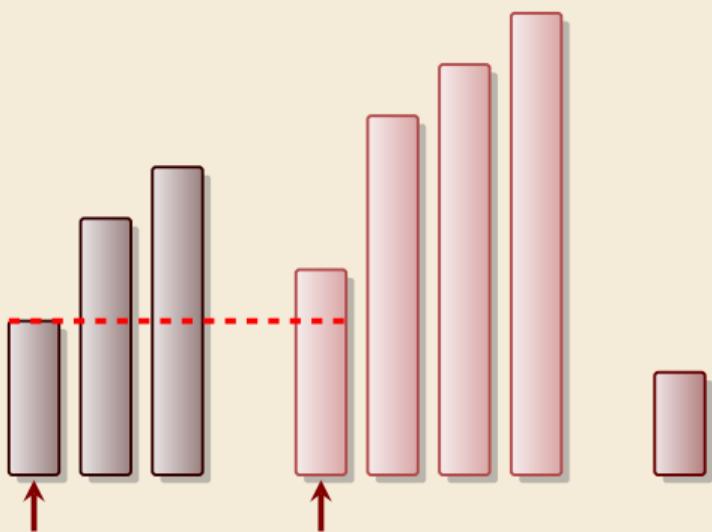


run1

run2

result

# Merging sorted lists

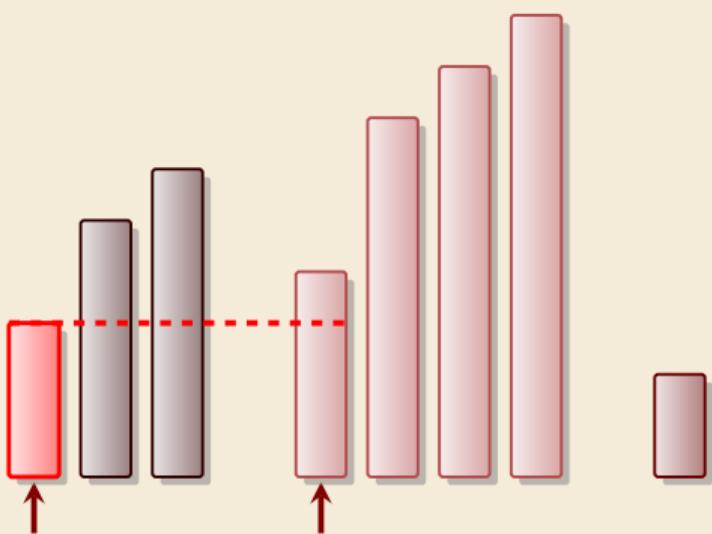


run1

run2

result

# Merging sorted lists

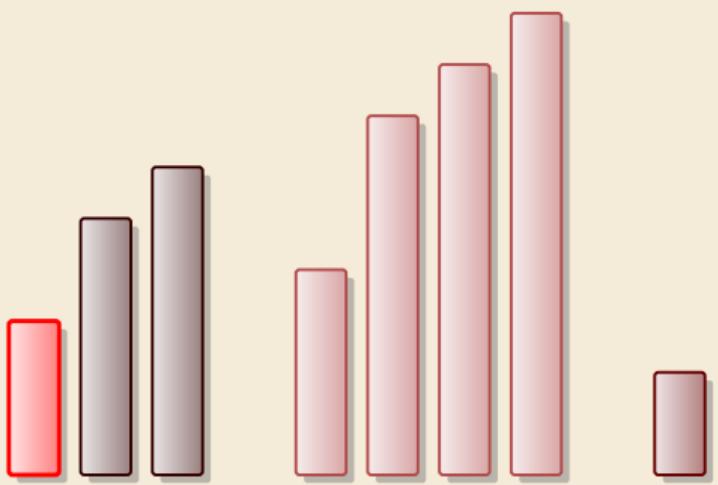


run1

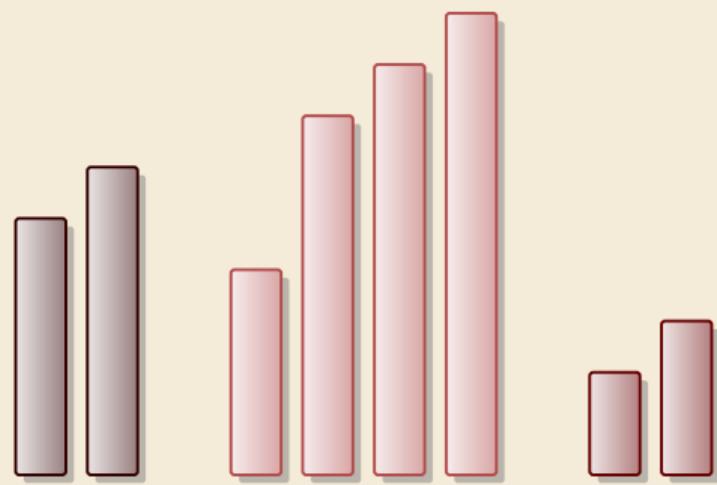
run2

result

# Merging sorted lists



# Merging sorted lists

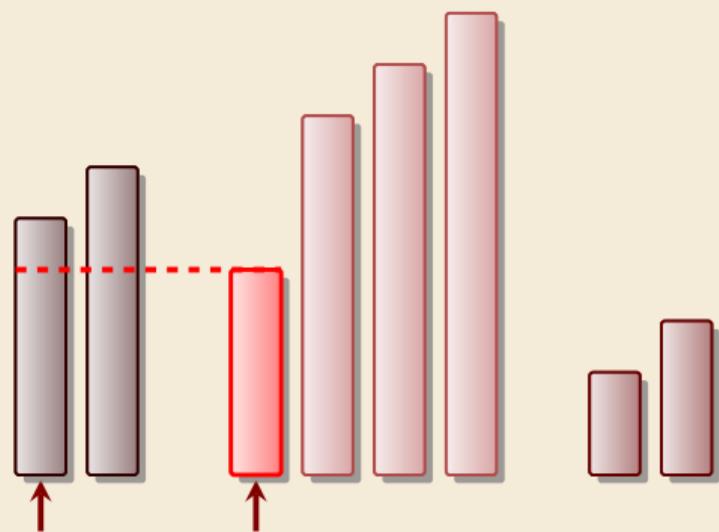


run1

run2

result

# Merging sorted lists

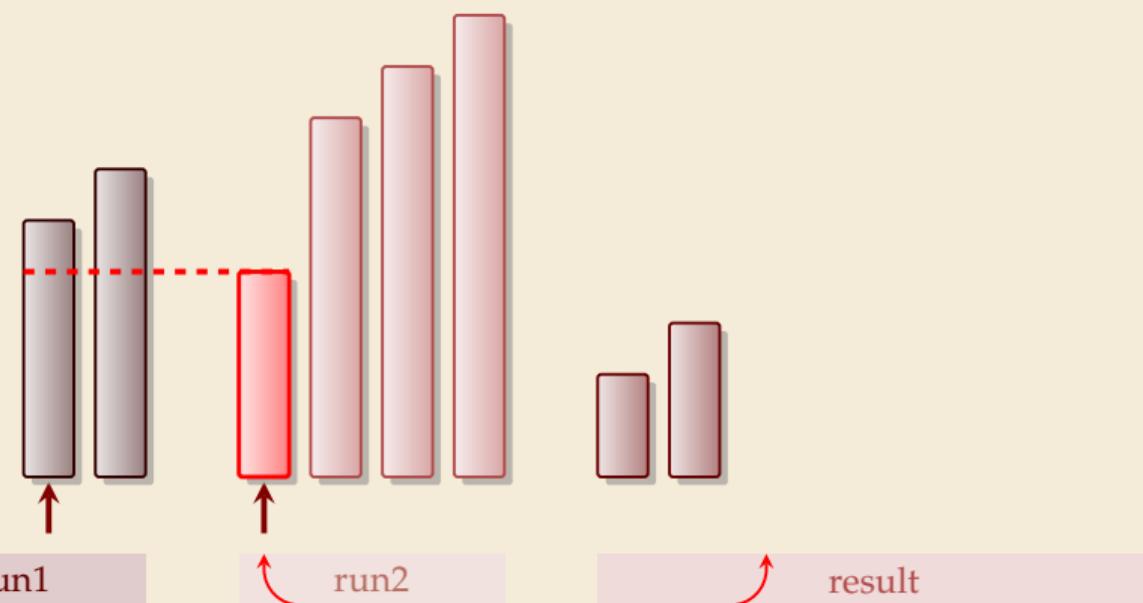


run1

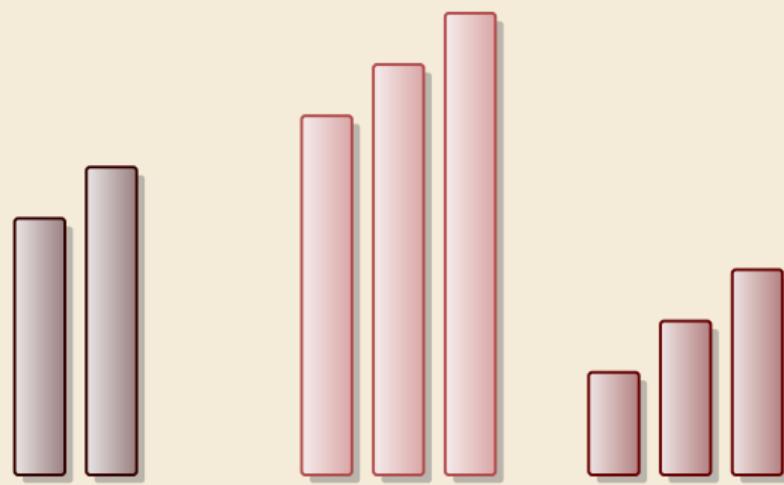
run2

result

# Merging sorted lists



# Merging sorted lists

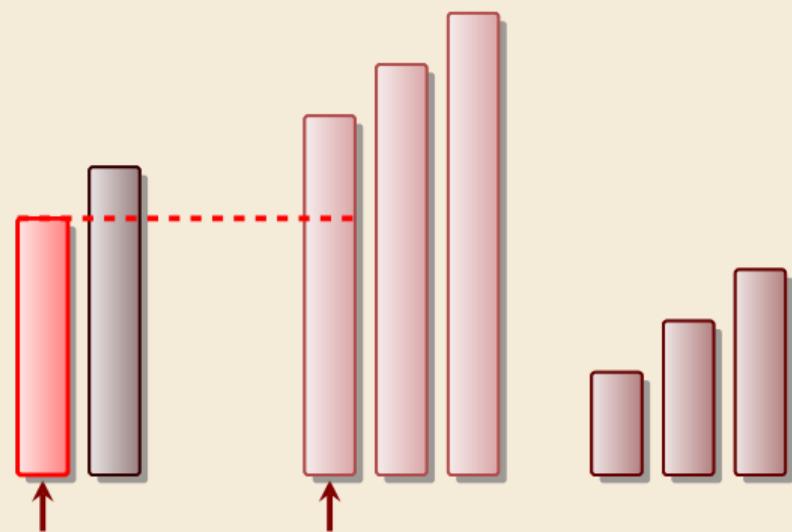


run1

run2

result

## Merging sorted lists

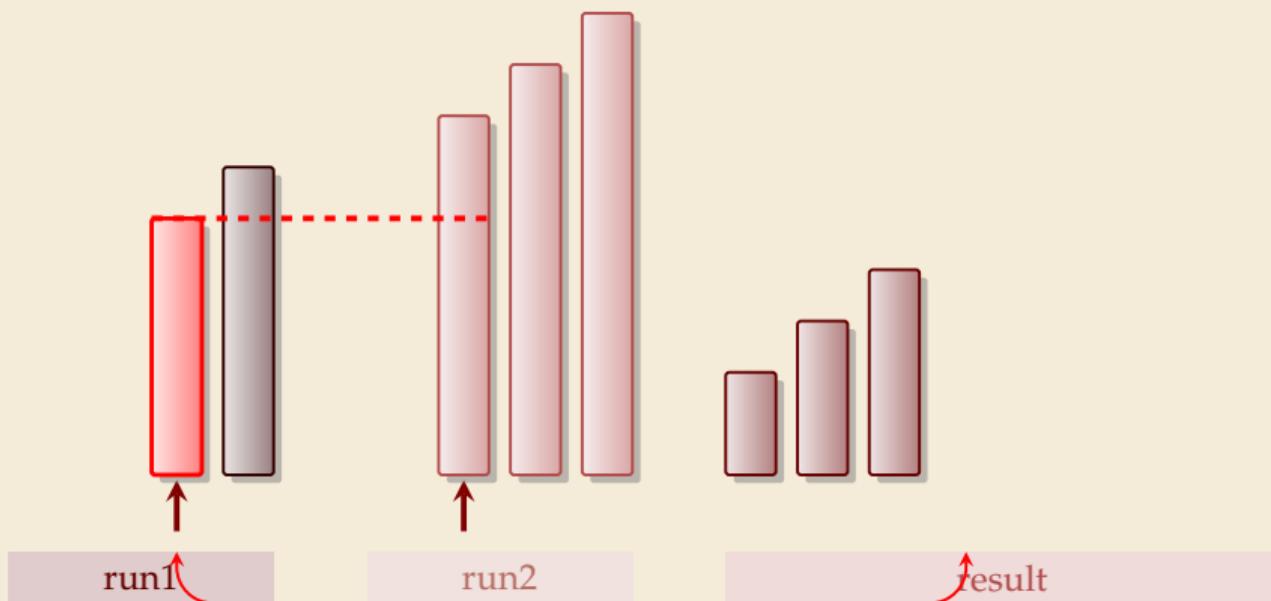


run1

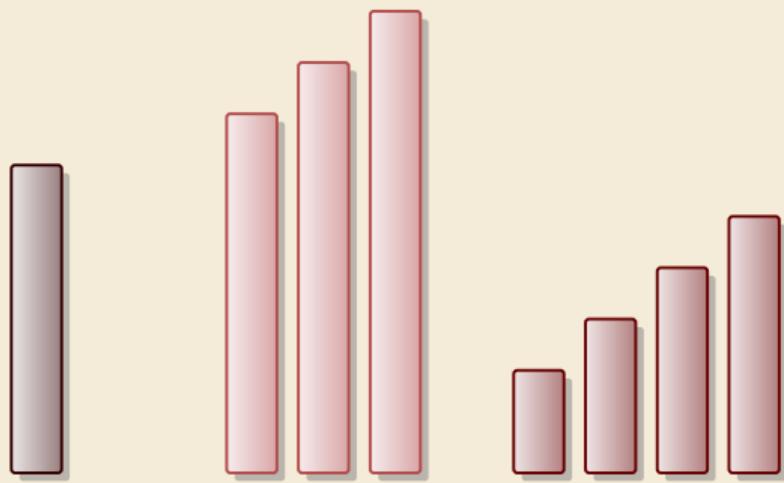
run2

result

## Merging sorted lists



# Merging sorted lists

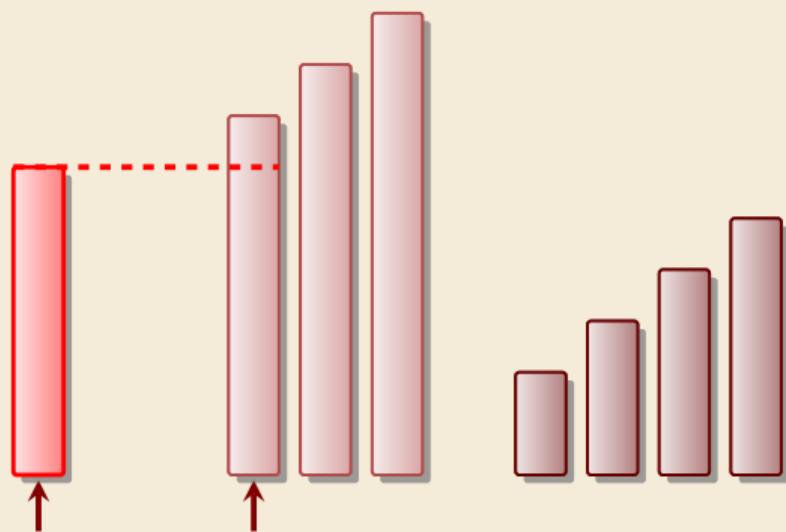


run1

run2

result

## Merging sorted lists

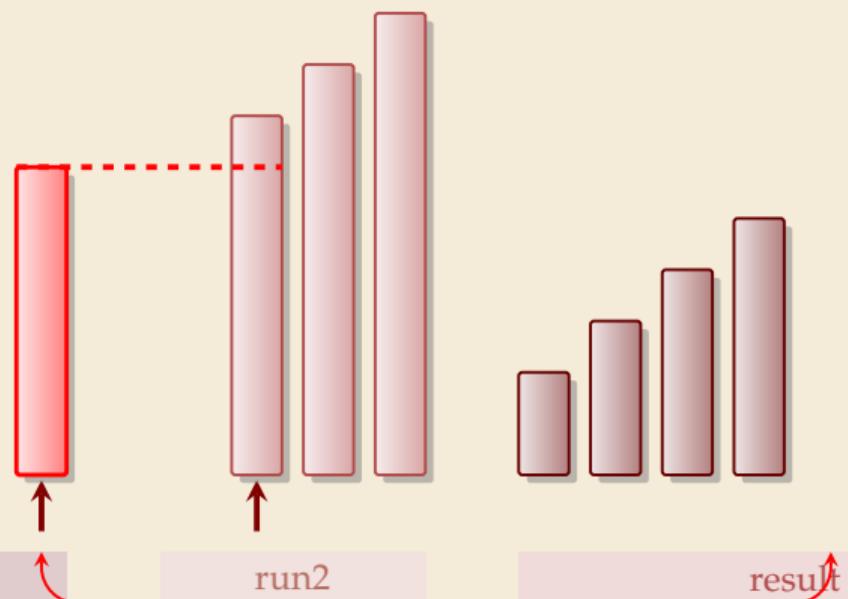


run1

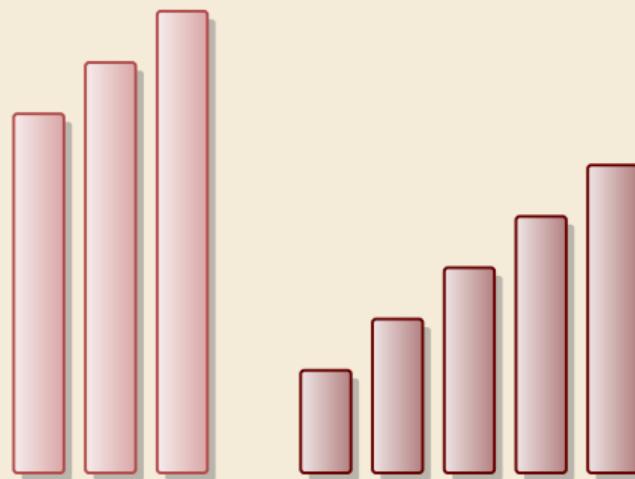
run2

result

# Merging sorted lists



# Merging sorted lists

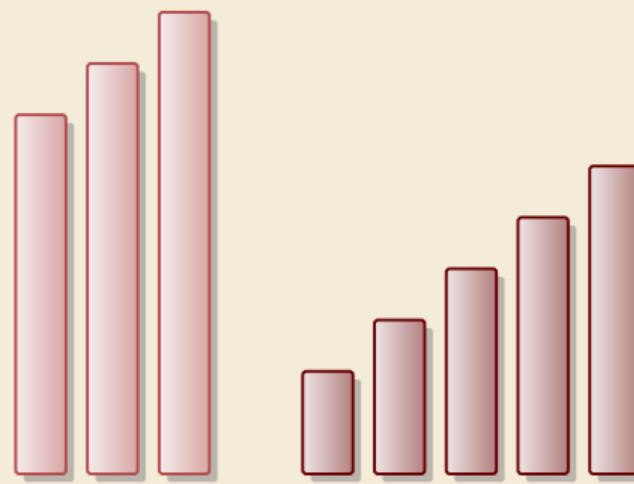


run1

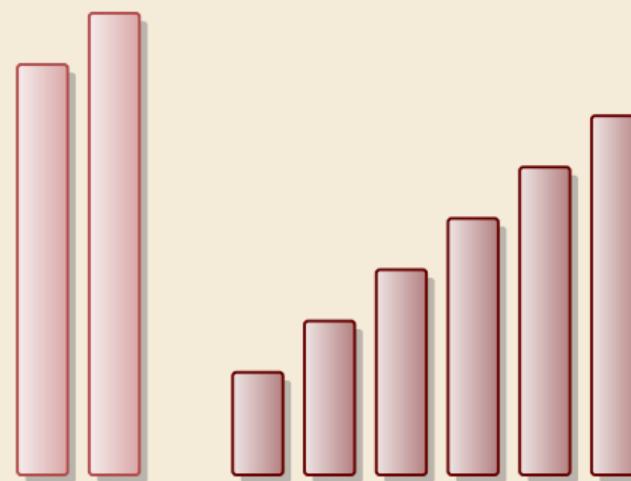
run2

result

## Merging sorted lists



## Merging sorted lists

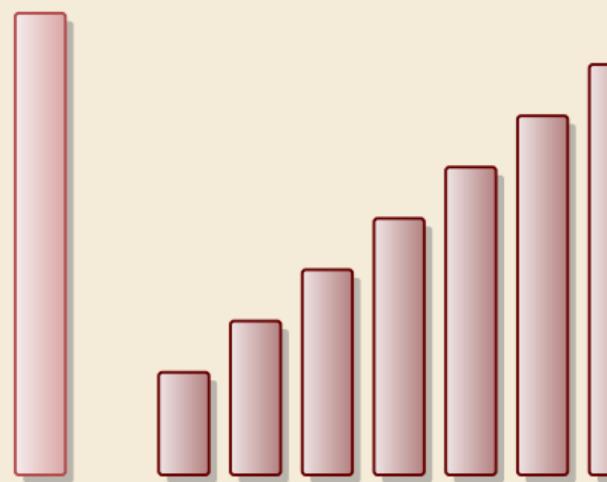


run1

run2

result

## Merging sorted lists

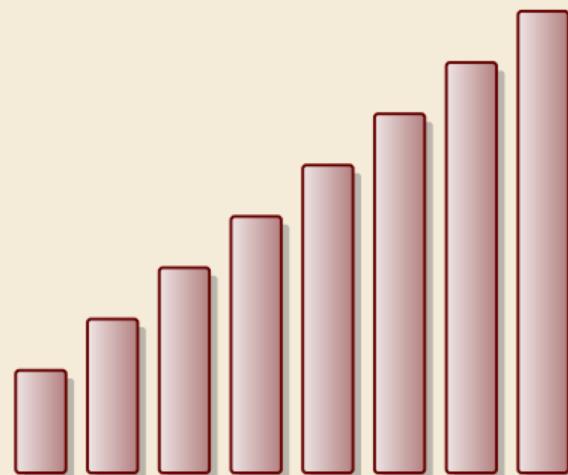


run1

run2

result

# Merging sorted lists



# Mergesort

---

```
1 procedure mergesort(A[l..r])
2     n := r - l + 1
3     if n ≥ 1 return
4         m := l + ⌊ $\frac{n}{2}$ ⌋
5         mergesort(A[l..m - 1])
6         mergesort(A[m..r])
7         merge(A[l..m - 1], A[m..r], buf)
8         copy buf to A[l..r]
```

---

- ▶ recursive procedure; *divide & conquer*
- ▶ merging needs
  - ▶ temporary storage for result      *buf*  
of same size as merged runs
  - ▶ to read and write each element twice  
(once for merging, once for copying back)

# Mergesort

---

```
1 procedure mergesort(A[l..r])
2     n := r - l + 1
3     if n ≥ 1 return
4         m := l + ⌊ n/2 ⌋
5         mergesort(A[l..m - 1])
6         mergesort(A[m..r])
7         merge(A[l..m - 1], A[m..r], buf)
8         copy buf to A[l..r]
```

---

- ▶ recursive procedure; *divide & conquer*
  - ▶ merging needs
    - ▶ temporary storage for result of same size as merged runs
    - ▶ to read and write each element twice (once for merging, once for copying back)
- $2n$

**Analysis:** count “element visits” (read and/or write)

$$C(n) = \begin{cases} 0 & n \leq 1 \\ C(\lfloor n/2 \rfloor) + C(\lceil n/2 \rceil) + 2\underline{n} & n \geq 2 \end{cases}$$

same for best and worst case!

Simplification  $n = 2^k$

$$C(2^k) = \begin{cases} 0 & k \leq 0 \\ 2 \cdot C(2^{k-1}) + 2 \cdot 2^k & k \geq 1 \end{cases} = 2 \cdot 2^k + 2^2 \cdot \underline{2^{k-1}} + 2^3 \cdot 2^{k-2} + \cdots + 2^k \cdot 2^1 = \underbrace{2k \cdot 2^k}_{\lg n}$$

$$C(n) = \underline{2n \lg(n)} = \Theta(n \log n)$$

# Mergesort – Discussion

thumb up optimal time complexity of  $\Theta(n \log n)$  in the worst case

thumb up *stable* sorting method i. e., retains relative order of equal-key items

2 3 1 2 2 5

thumb up memory access is sequential (scans over arrays)

→ 1 2 2 2 3 5 sorted

thumb down requires  $\Theta(n)$  extra space

there are in-place merging methods,  
but they are substantially more complicated  
and not (widely) used

but not stable sorted

## 3.2 Quicksort

# Clicker Question



How does quicksort work?

- A** split elements around median, then recurse on small / large elements.
- B** recurse on left / right half, then combine sorted halves.
- C** grow sorted part on left, repeatedly add next element to sorted range.
- D** repeatedly choose 2 elements and swap them if they are out of order.
- E** Don't know.

*pingo.upb.de/622222*

# Clicker Question



How does quicksort work?

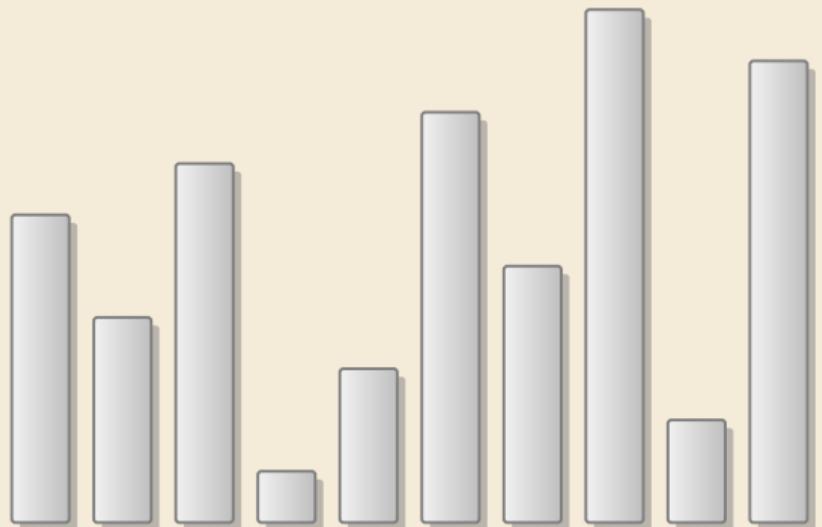
- A** split elements around median, then recurse on small / large elements. ✓
- B** ~~recurse on left / right half, then combine sorted halves.~~
- C** ~~grow sorted part on left, repeatedly add next element to sorted range.~~
- D** ~~repeatedly choose 2 elements and swap them if they are out of order.~~
- E** ~~Don't know.~~

*pingo.upb.de/622222*

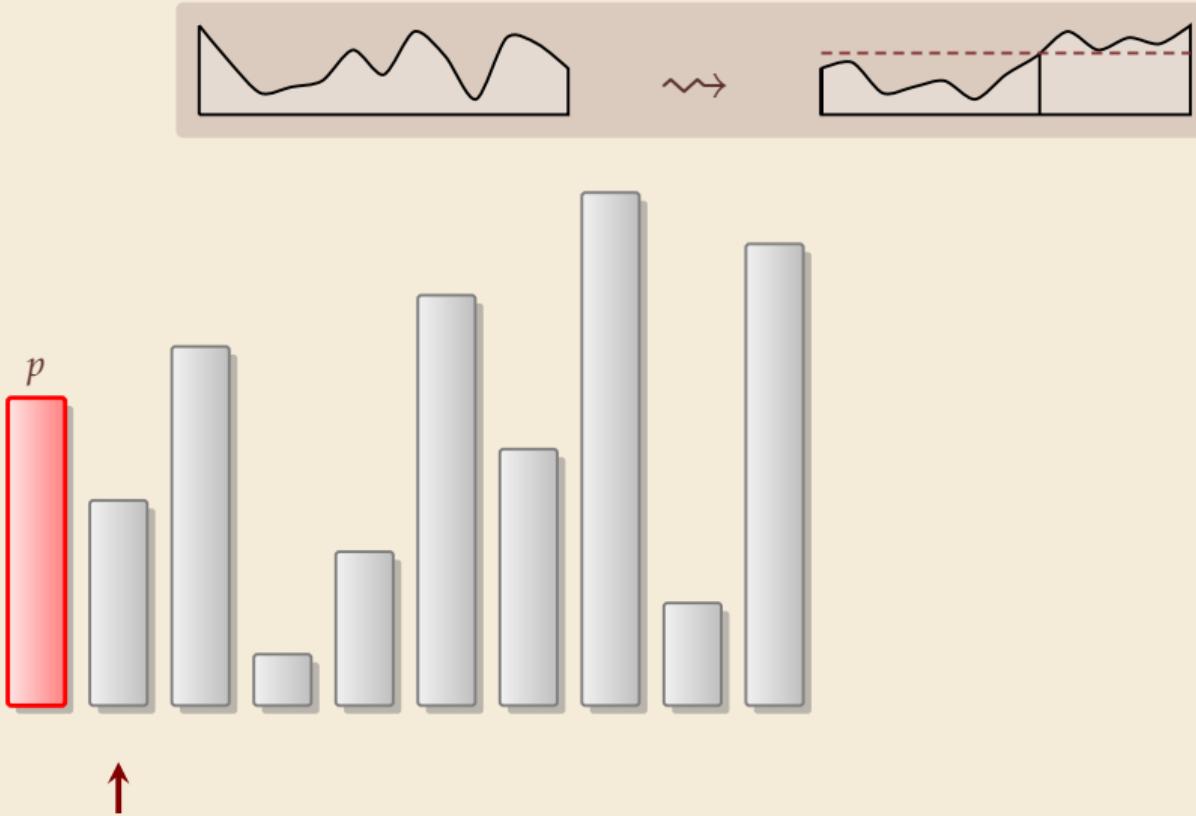
## Partitioning around a pivot



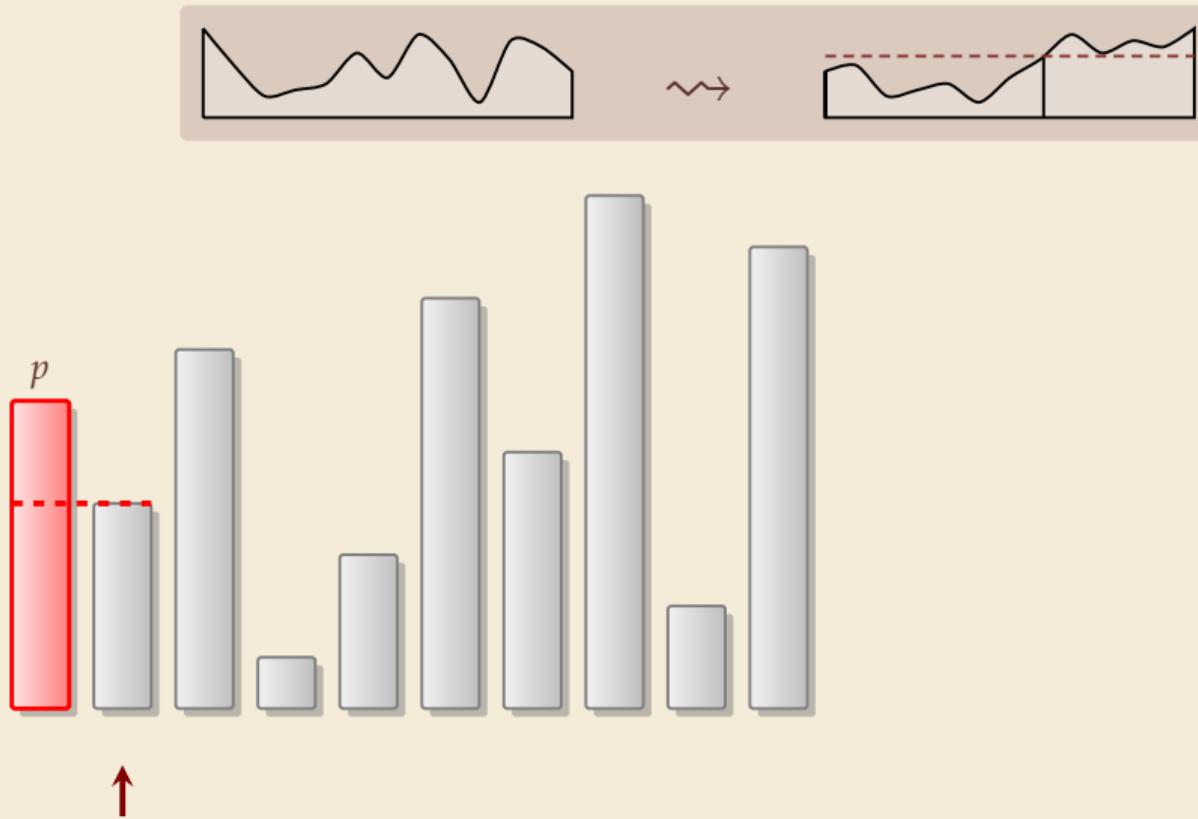
## Partitioning around a pivot



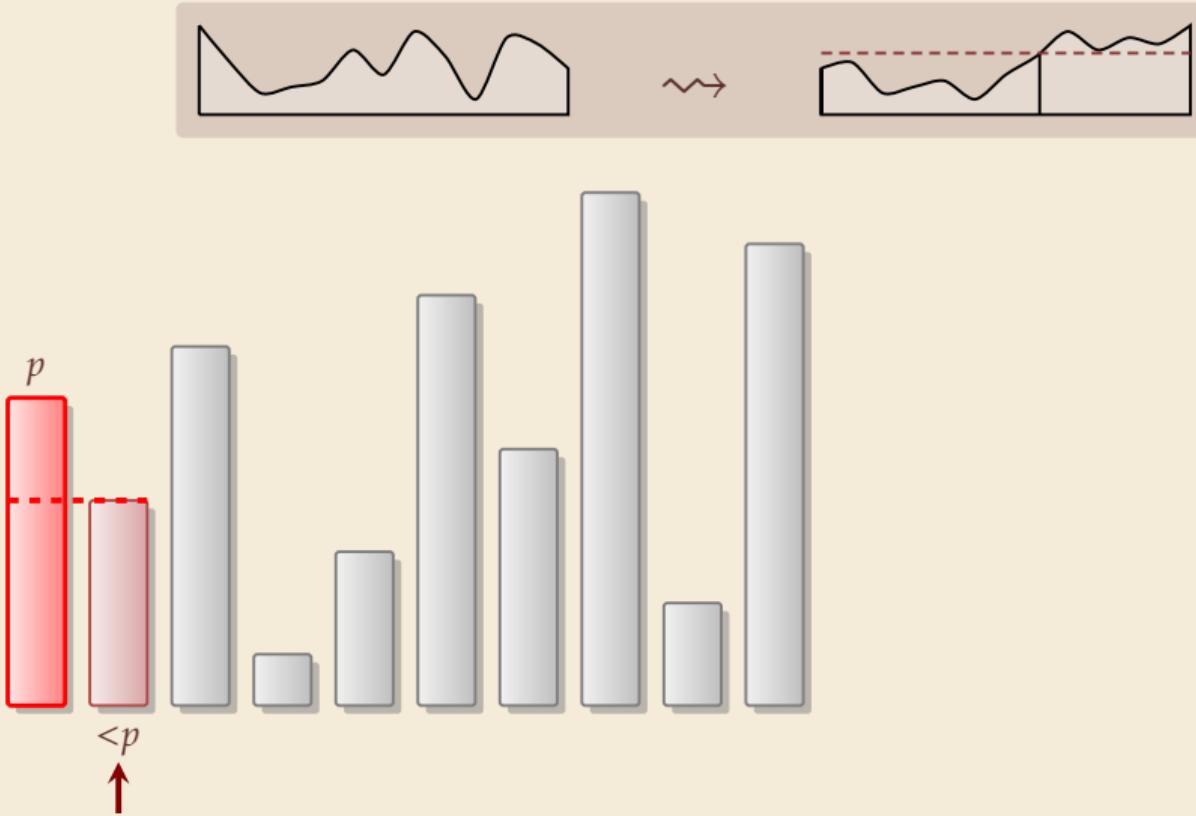
## Partitioning around a pivot



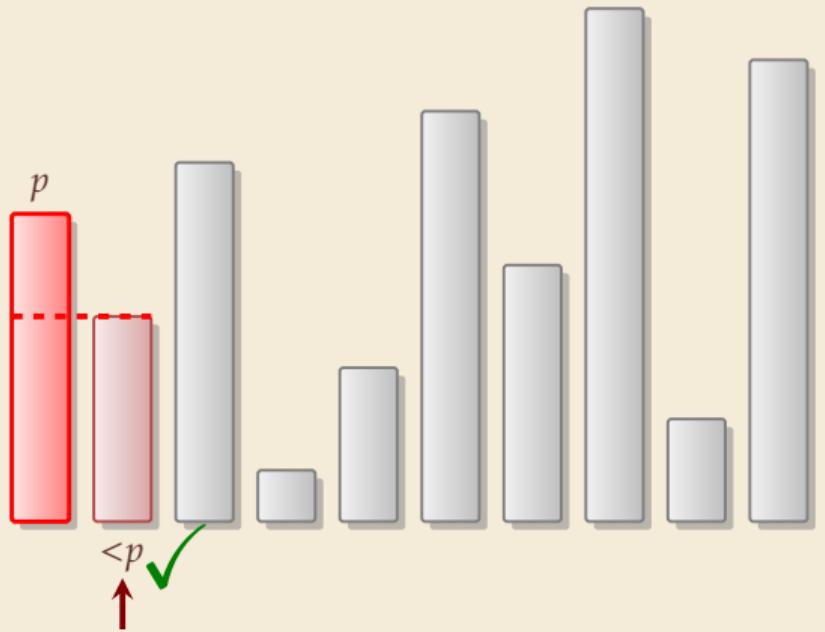
## Partitioning around a pivot



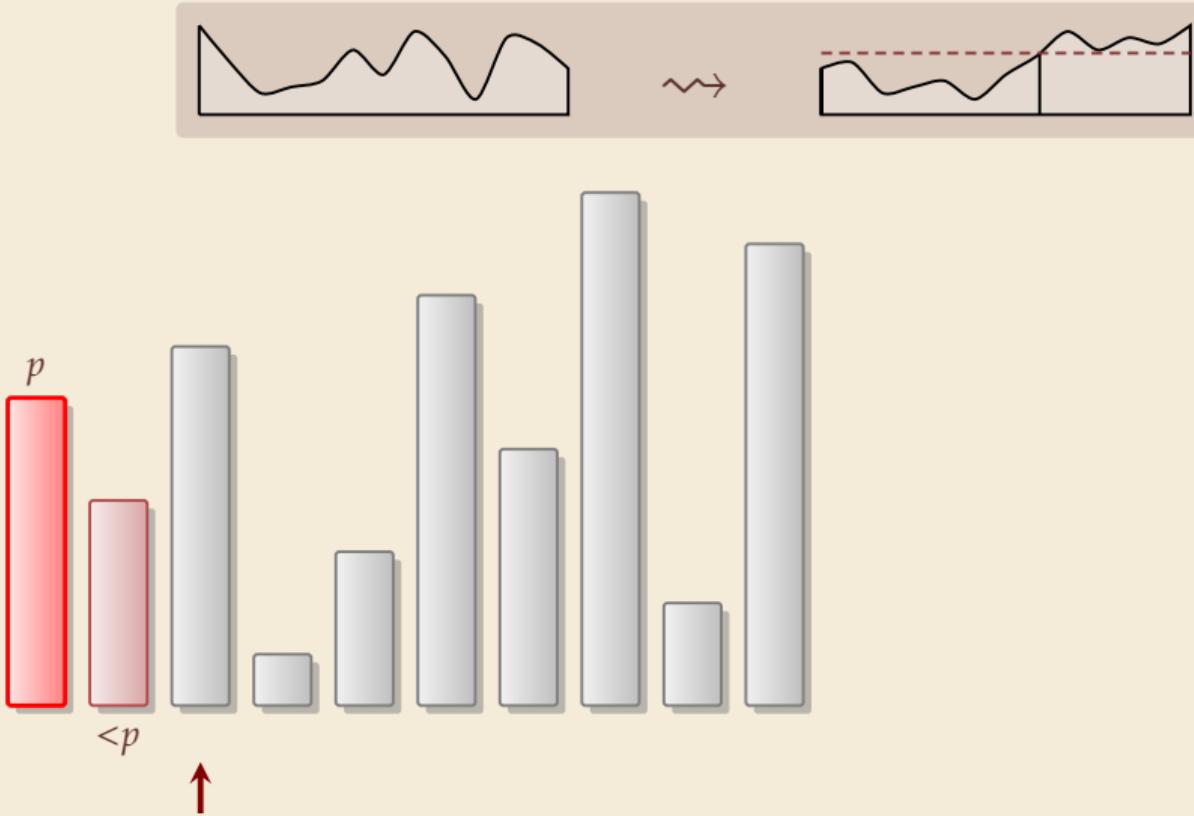
## Partitioning around a pivot



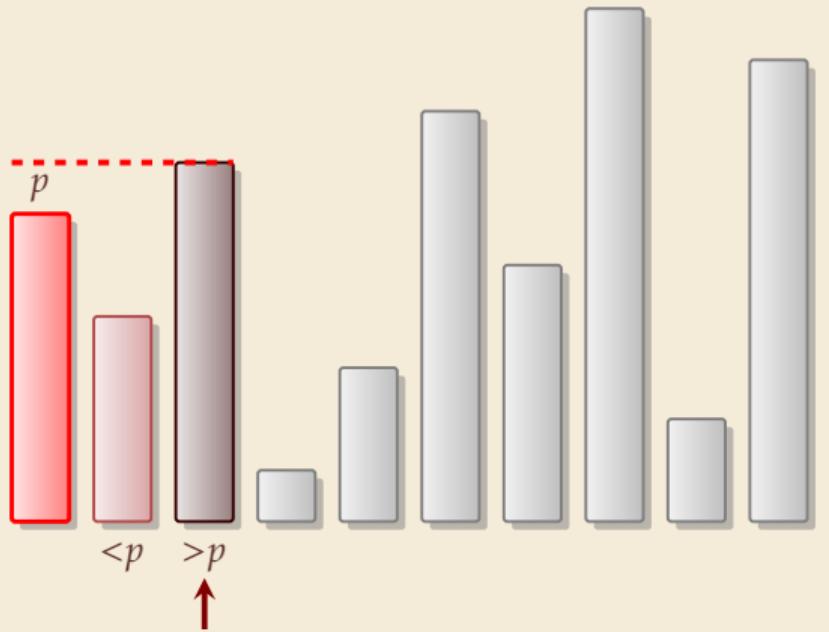
## Partitioning around a pivot



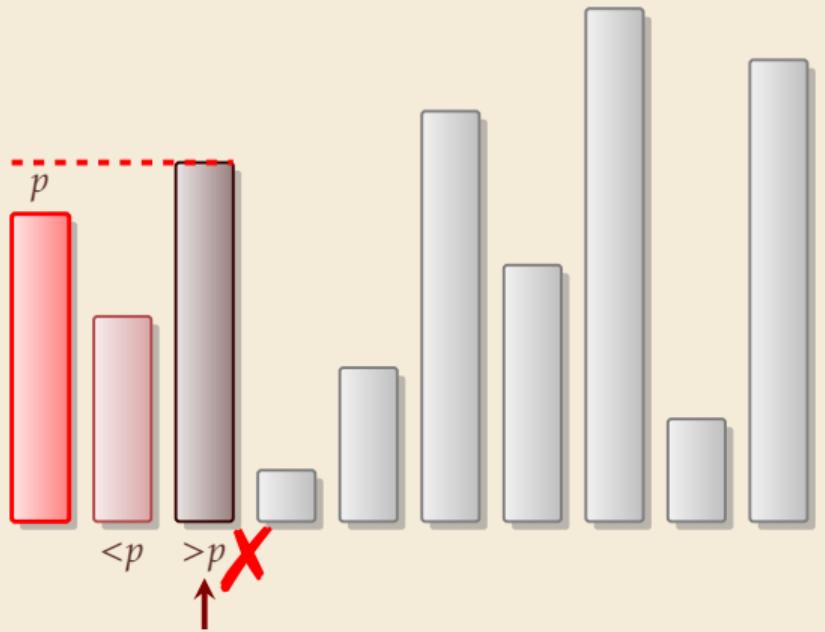
## Partitioning around a pivot



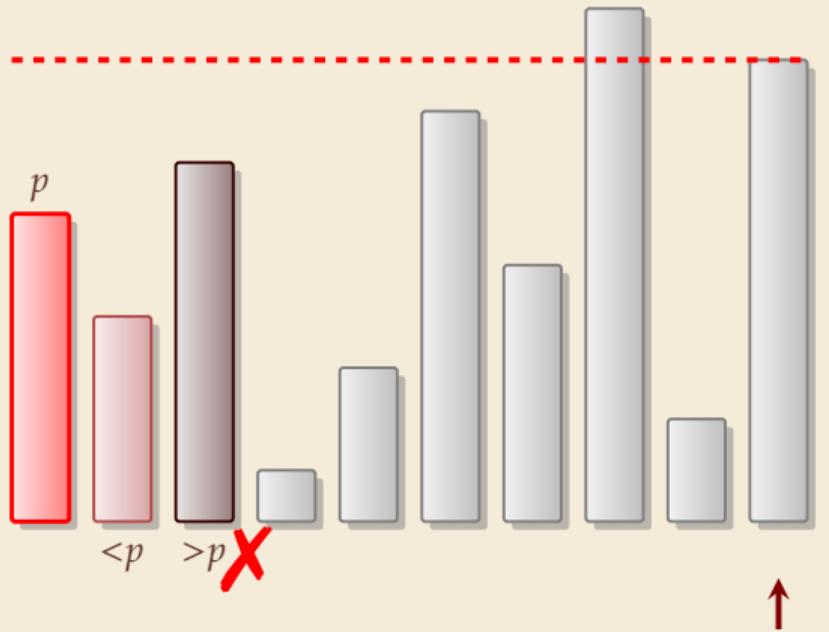
## Partitioning around a pivot



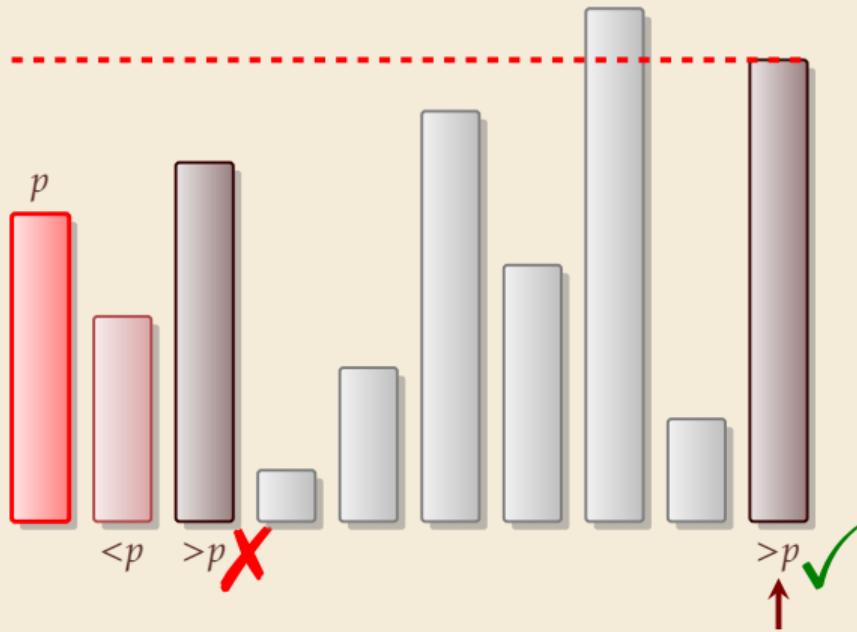
## Partitioning around a pivot



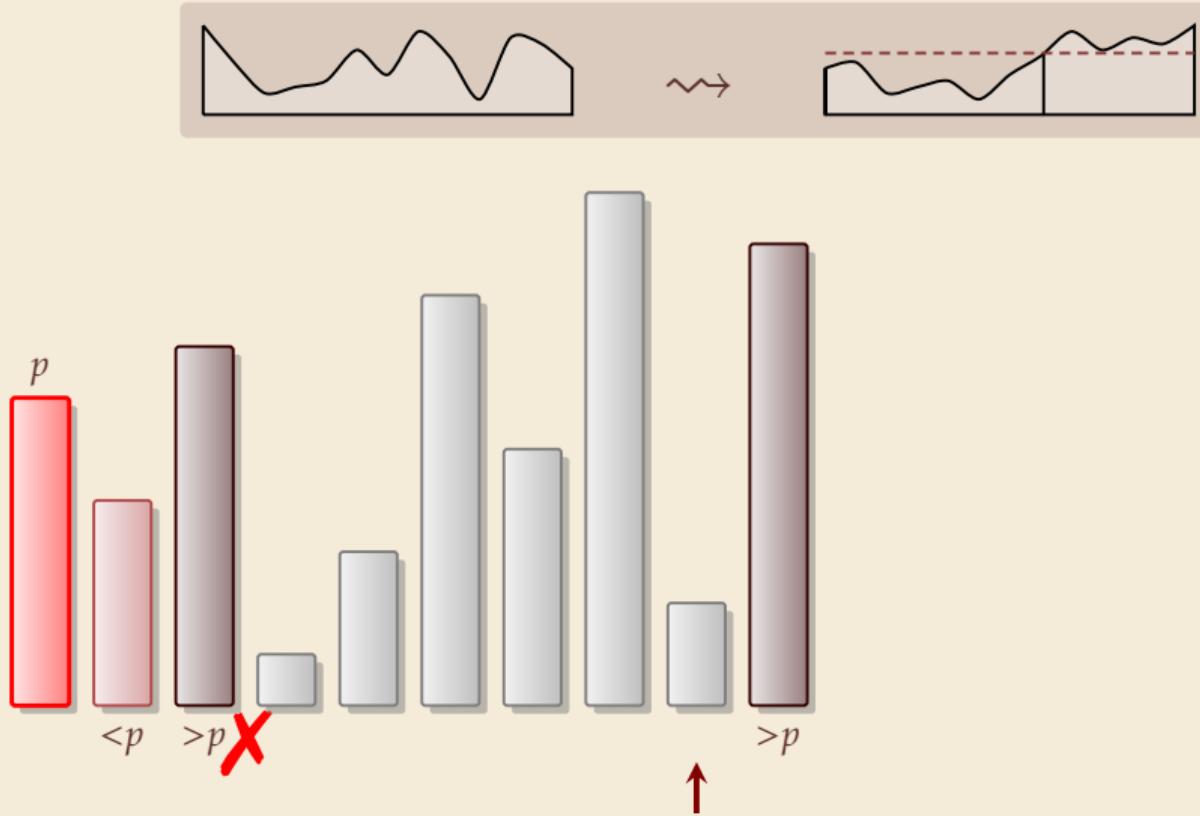
## Partitioning around a pivot



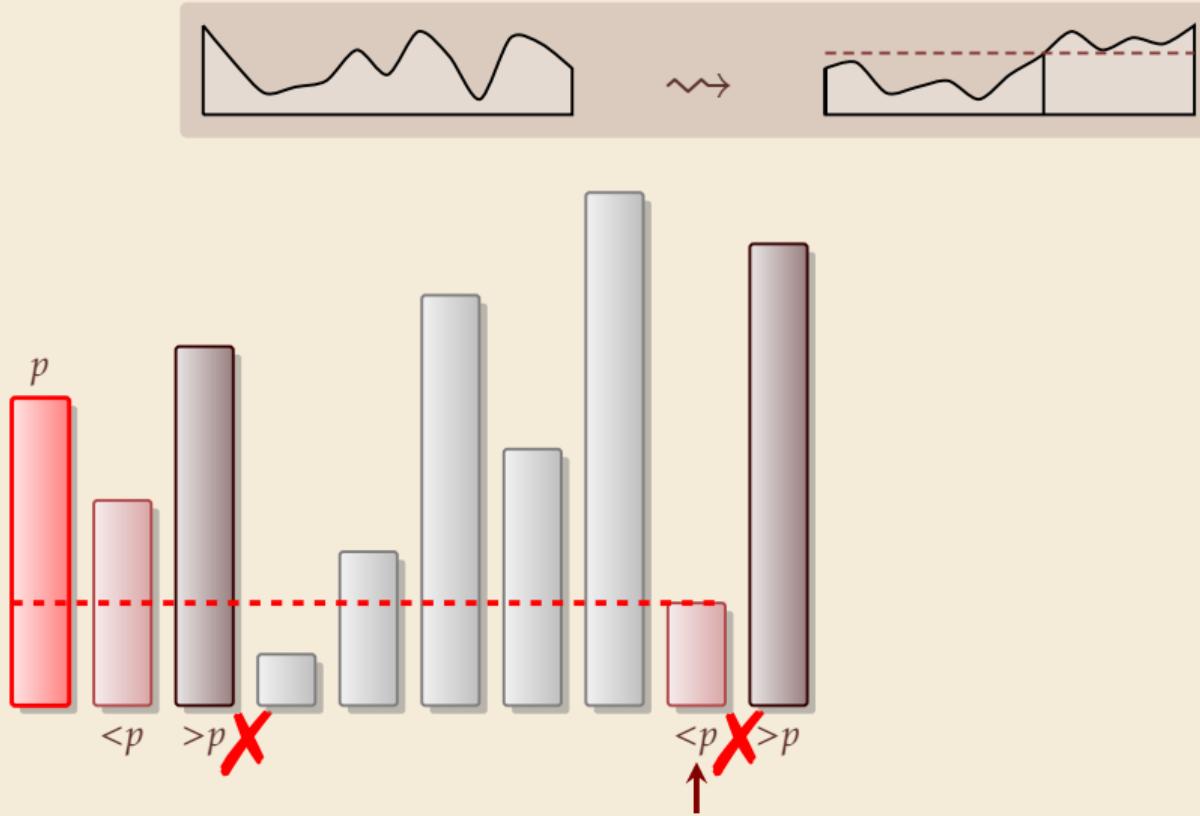
## Partitioning around a pivot



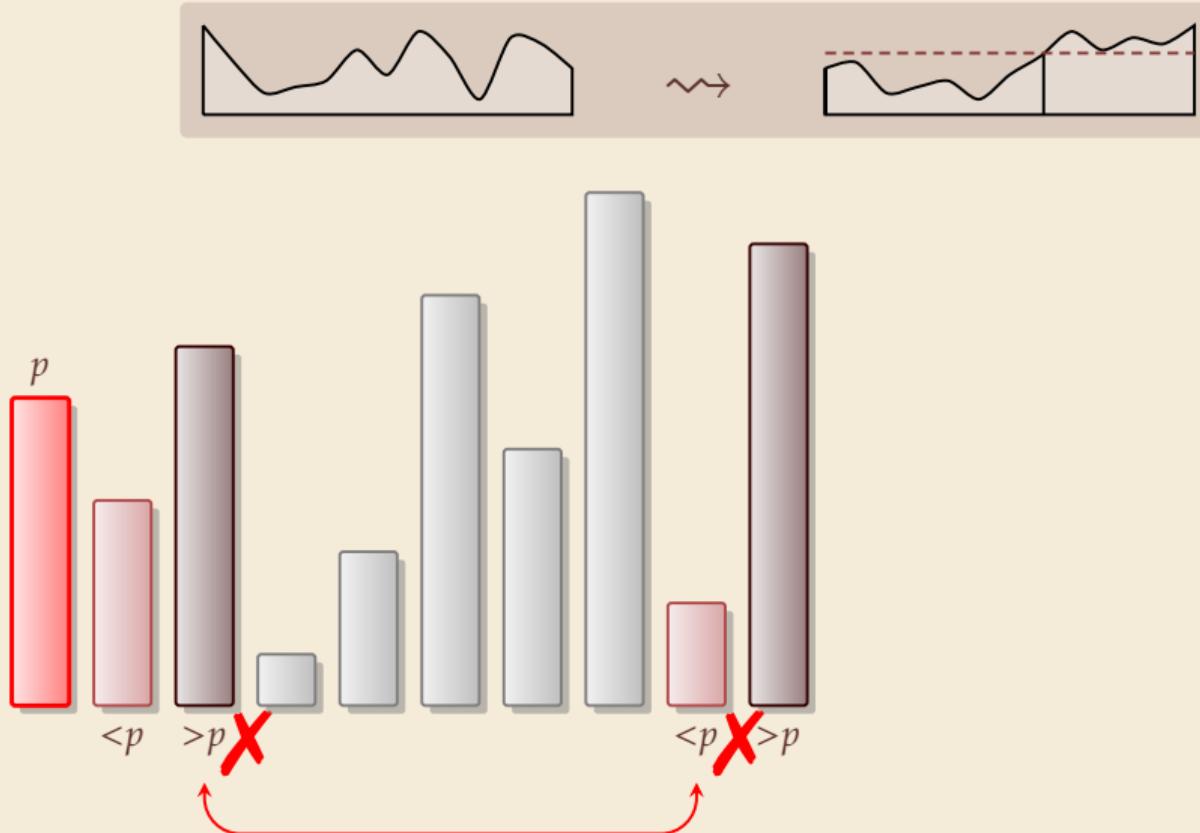
## Partitioning around a pivot



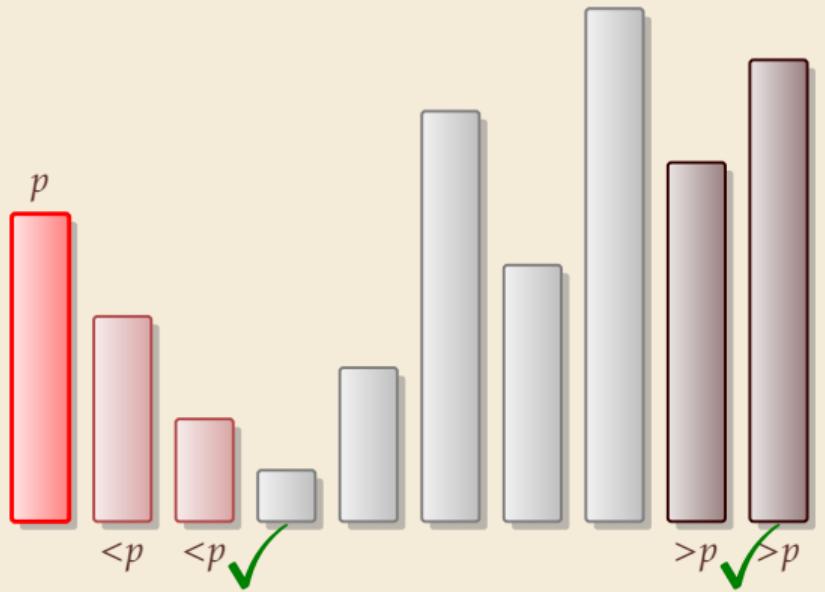
## Partitioning around a pivot



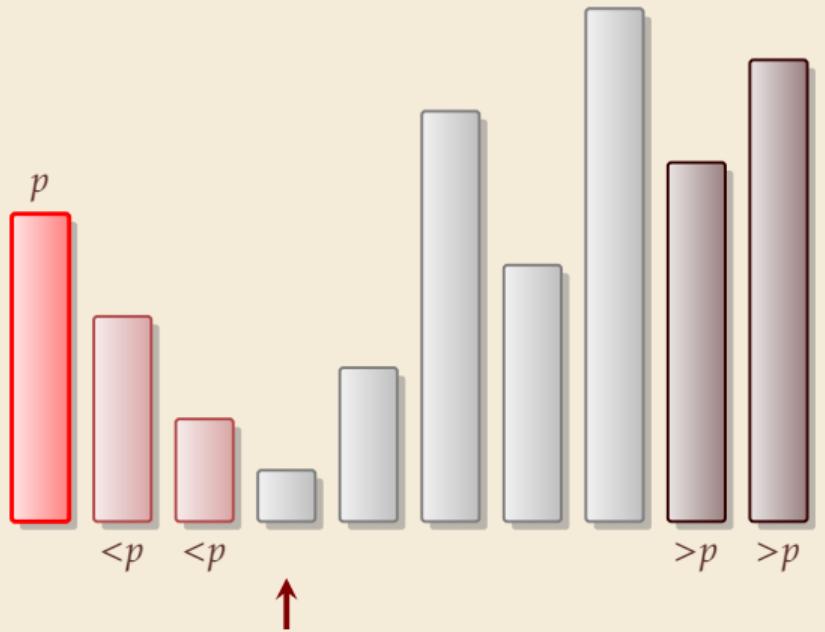
## Partitioning around a pivot



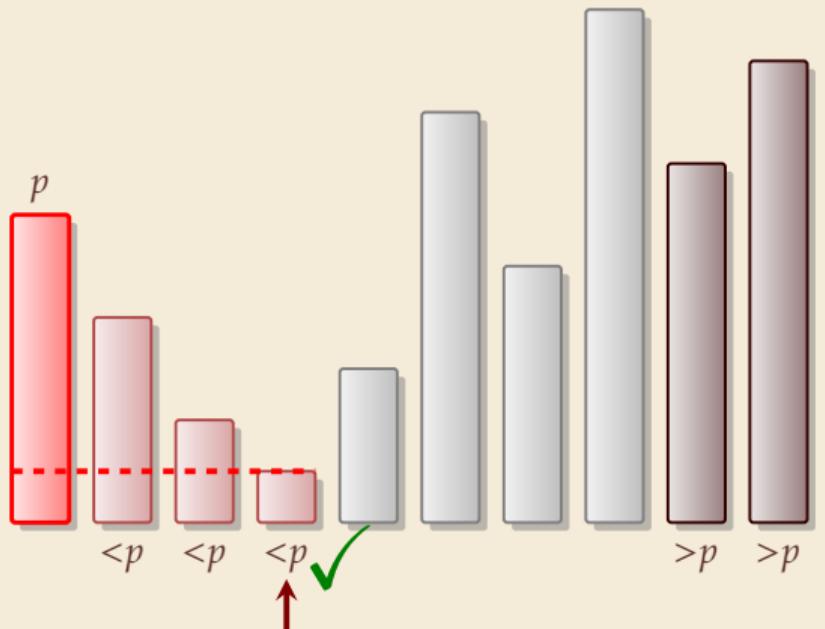
## Partitioning around a pivot



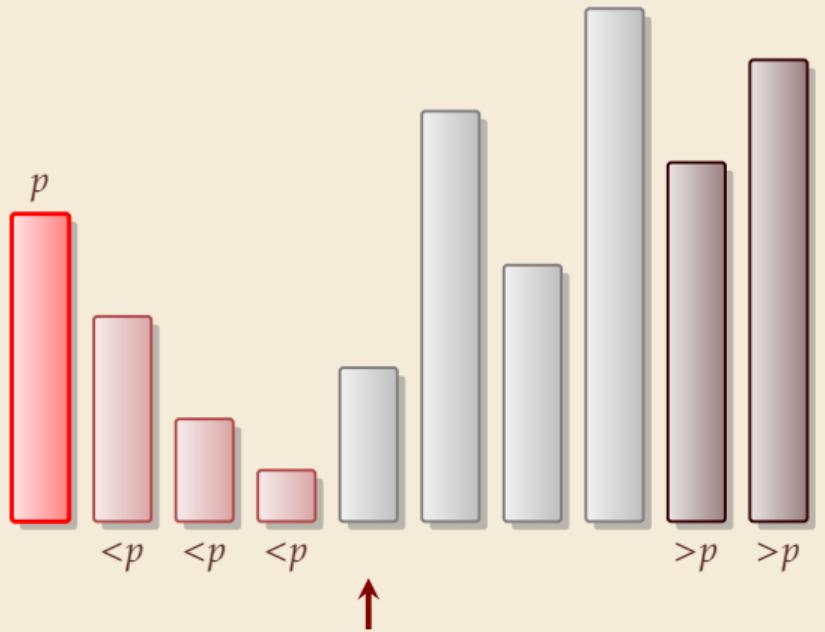
## Partitioning around a pivot



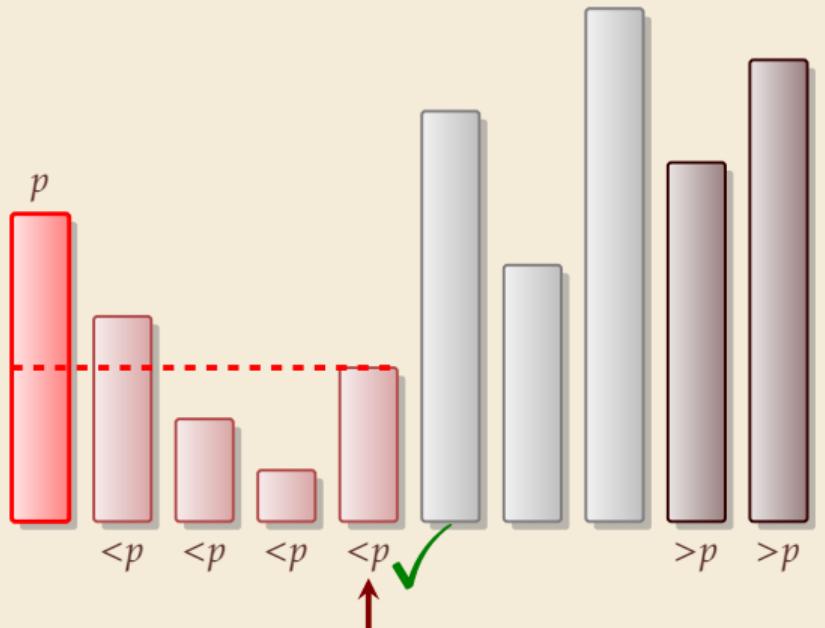
## Partitioning around a pivot



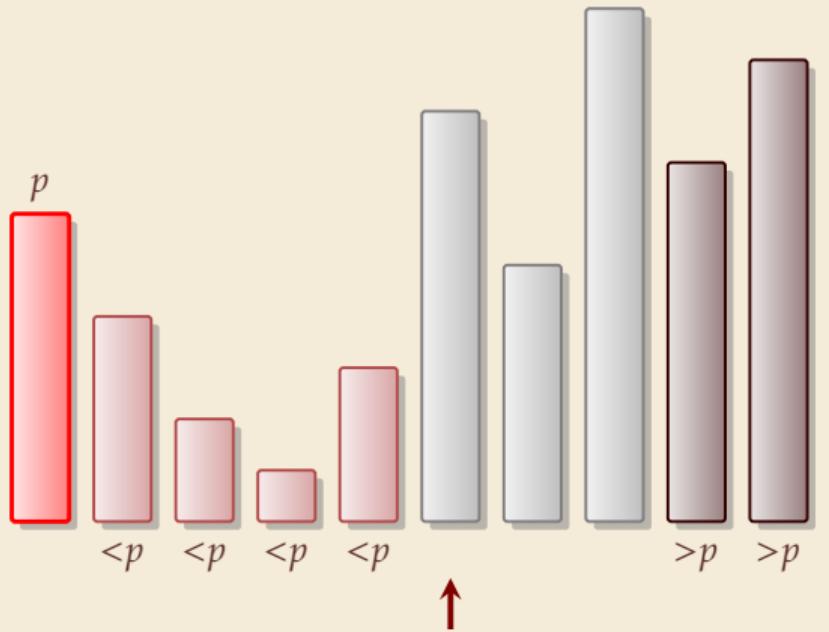
## Partitioning around a pivot



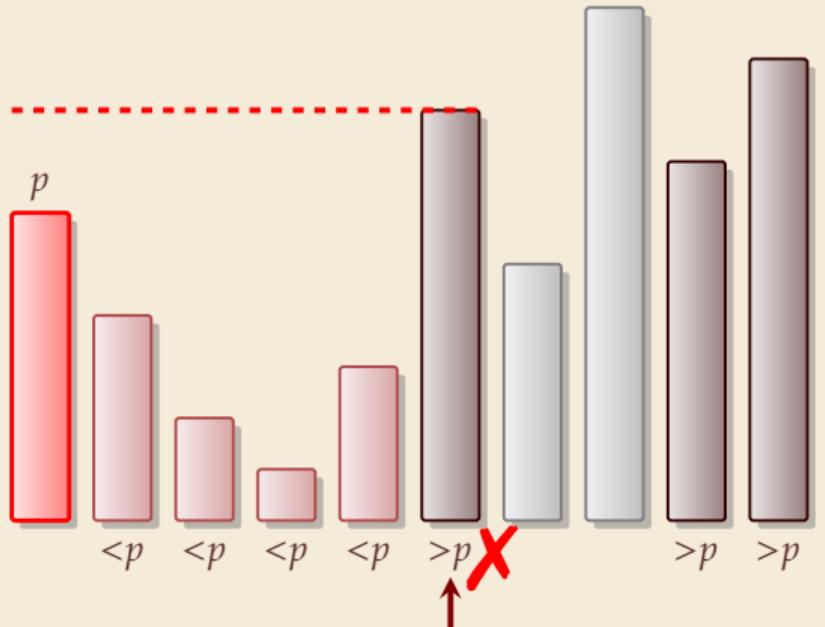
## Partitioning around a pivot



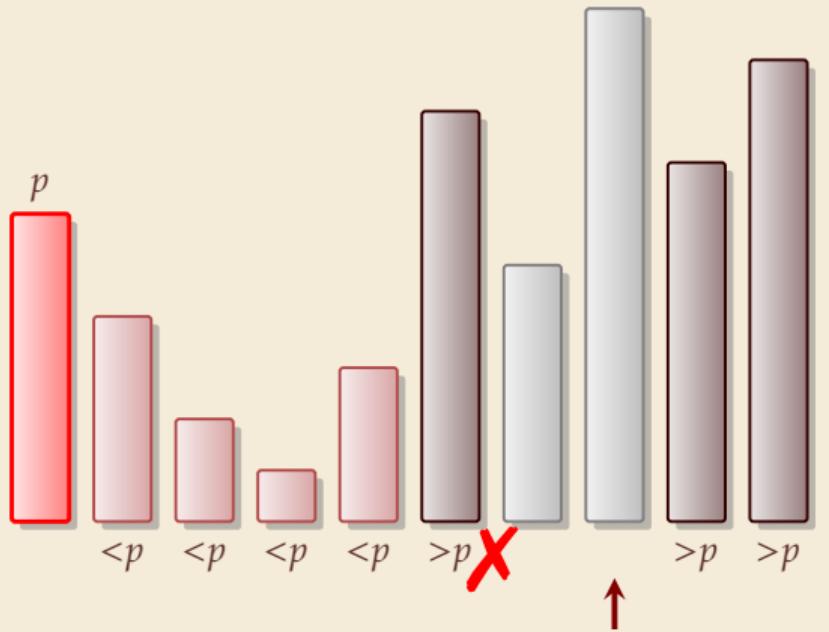
## Partitioning around a pivot



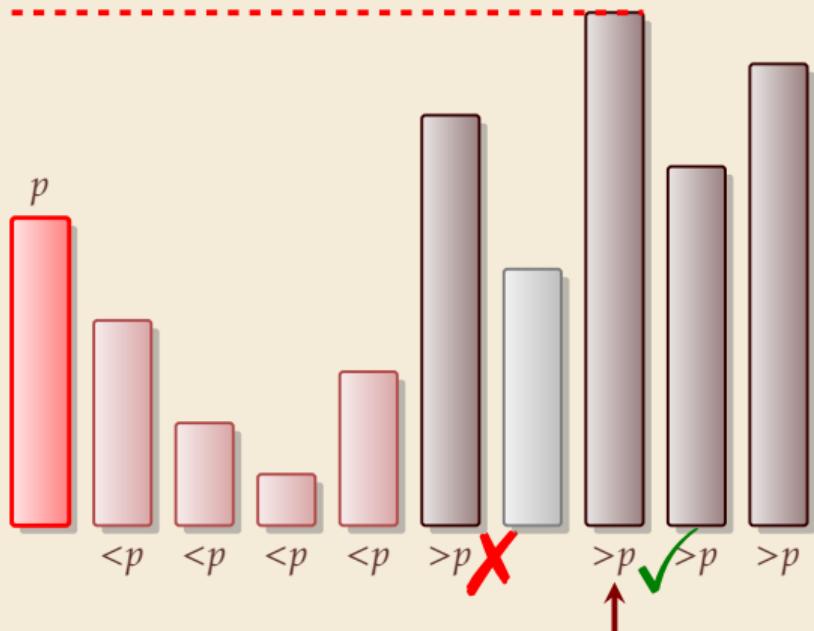
## Partitioning around a pivot



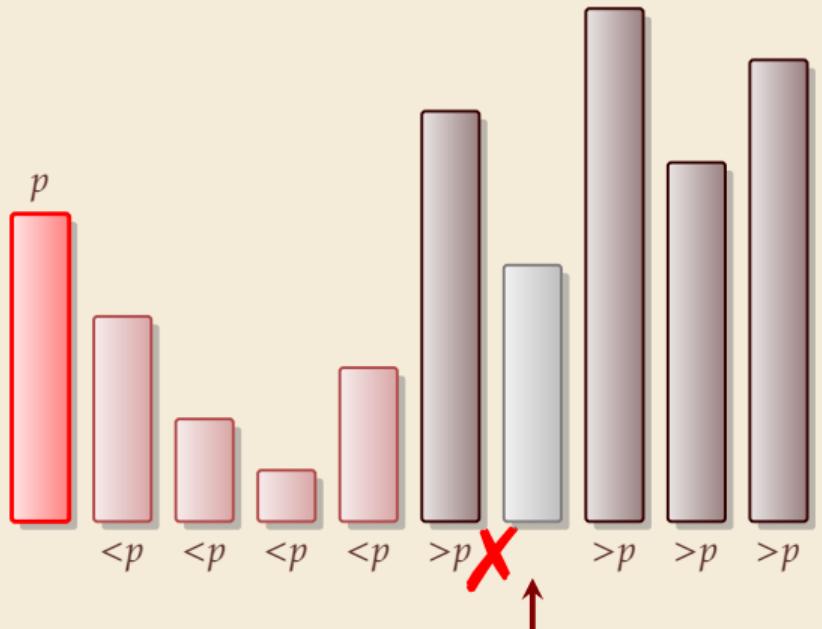
## Partitioning around a pivot



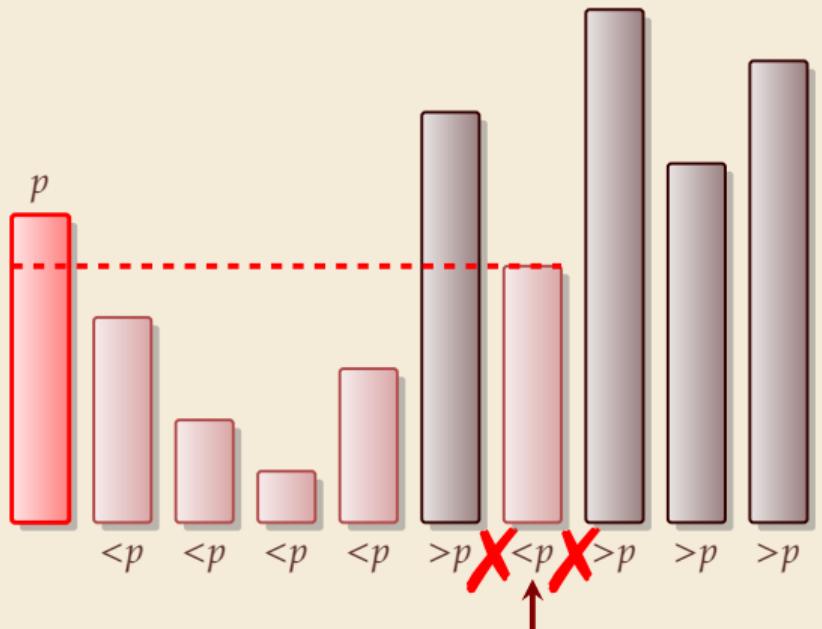
## Partitioning around a pivot



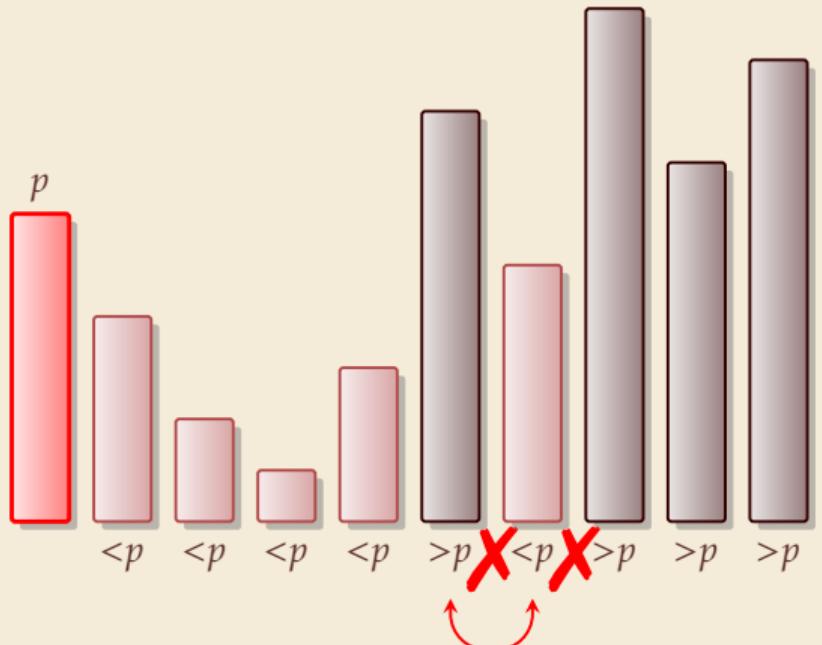
## Partitioning around a pivot



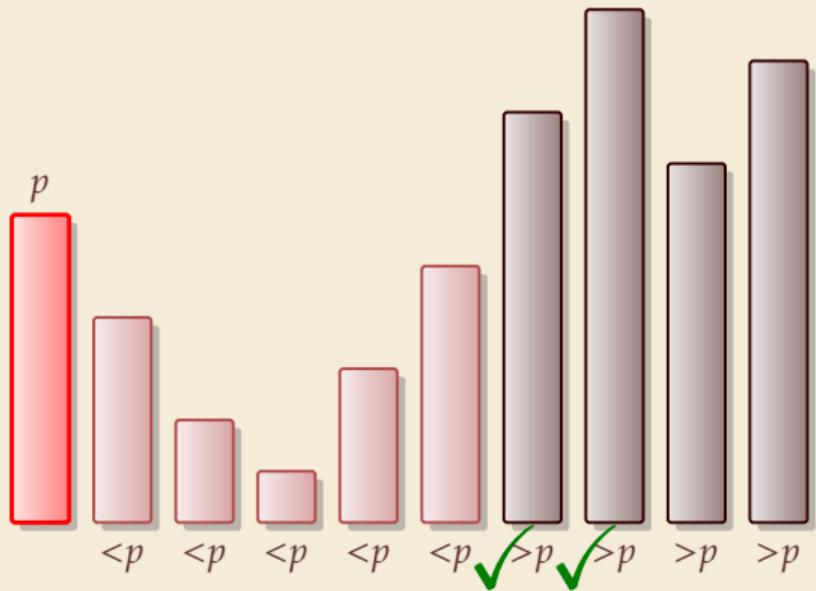
## Partitioning around a pivot



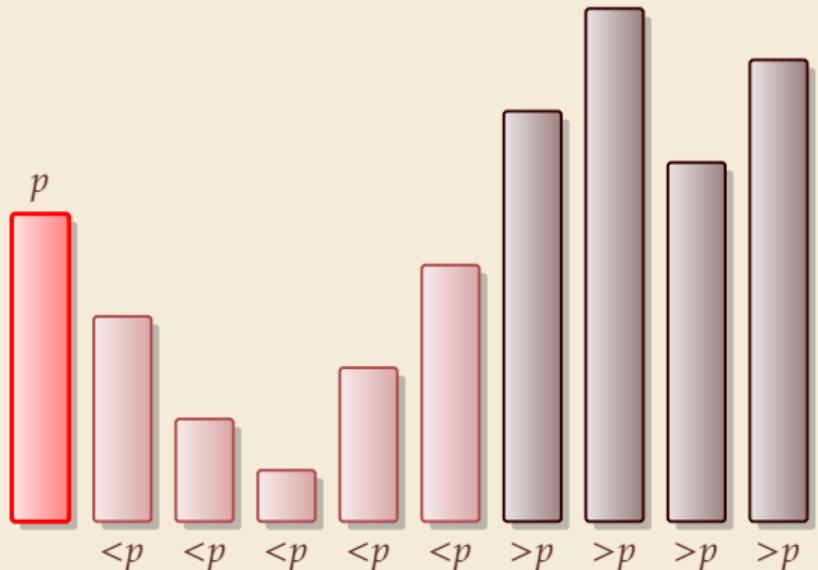
## Partitioning around a pivot



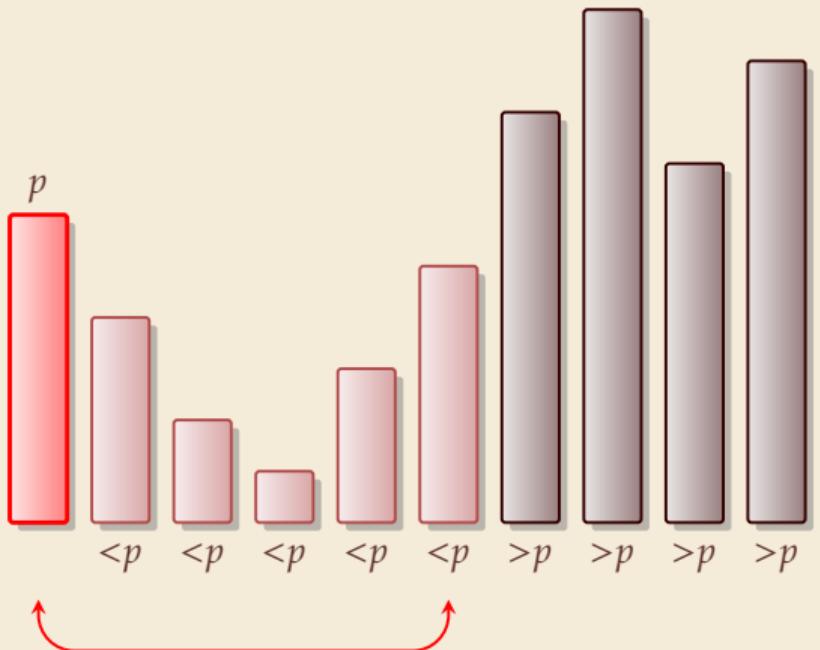
## Partitioning around a pivot



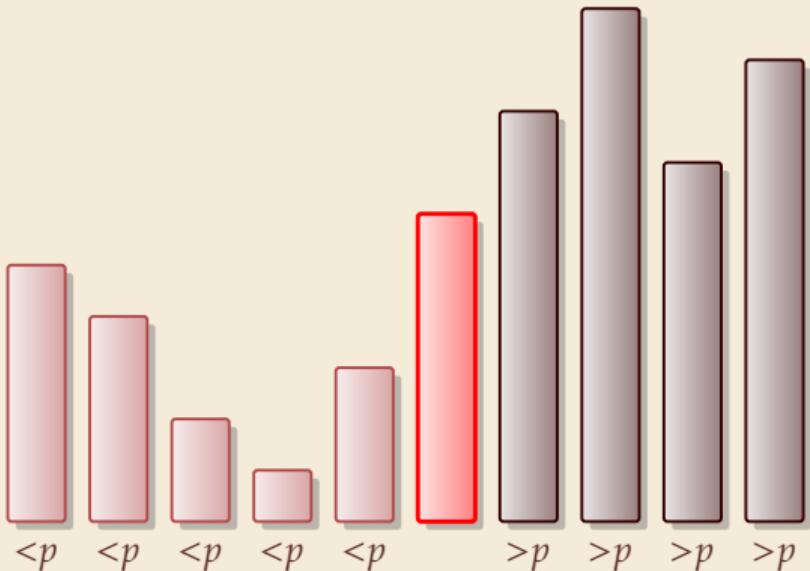
## Partitioning around a pivot



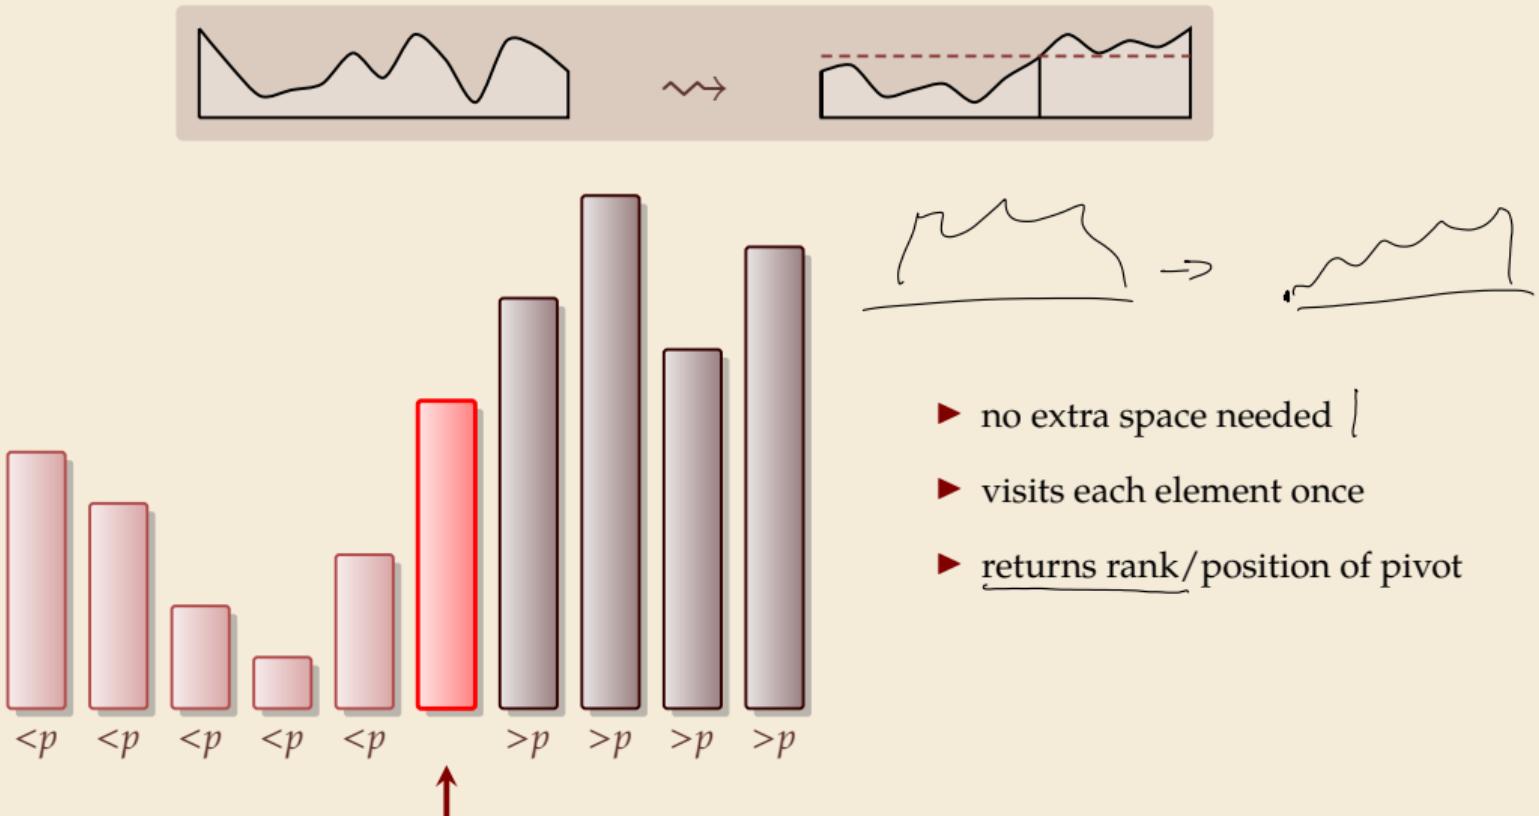
## Partitioning around a pivot



## Partitioning around a pivot



# Partitioning around a pivot



# Partitioning – Detailed code

Beware: details easy to get wrong; use this code!

---

```
1 procedure partition( $A, b$ )
2     // input: array  $A[0..n - 1]$ , position of pivot  $b \in [0..n - 1]$ 
3     swap( $A[0], A[b]$ )
4      $i := 0, j := n$ 
5     while true do
6         do  $i := i + 1$  while  $i < n$  and  $A[i] < A[0]$ 
7         do  $j := j - 1$  while  $j \geq 1$  and  $A[j] > A[0]$ 
8         if  $i \geq j$  then break (goto 8)
9         else swap( $A[i], A[j]$ )
10    end while
11    swap( $A[0], A[j]$ )
12    return  $j$ 
```

---

Loop invariant (5–10):

$A$	$p$	$\leq p$	$?$	$\geq p$
		$i$		$j$

# Quicksort

---

```
1 procedure quicksort( $A[l..r]$ )
2   if  $l \geq r$  then return
3    $b := \text{choosePivot}(A[l..r])$ 
4    $j := \text{partition}(A[l..r], b)$ 
5   quicksort( $A[l..j - 1]$ )
6   quicksort( $A[j + 1..r]$ )
```

---

- ▶ recursive procedure; *divide & conquer*
- ▶ choice of pivot can be
  - ▶ fixed position ↗ dangerous!
  - ▶ random
  - ▶ more sophisticated, e. g., median of 3

# Quicksort & Binary Search Trees

## Quicksort

7	4	2	9	1	3	8	5	6
---	---	---	---	---	---	---	---	---

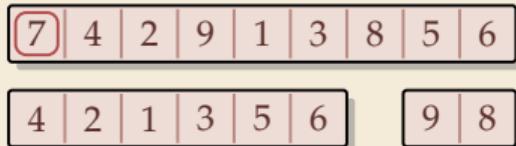
# Quicksort & Binary Search Trees

## Quicksort

7	4	2	9	1	3	8	5	6
---	---	---	---	---	---	---	---	---

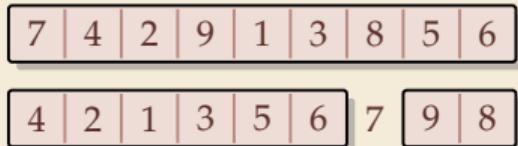
# Quicksort & Binary Search Trees

## Quicksort



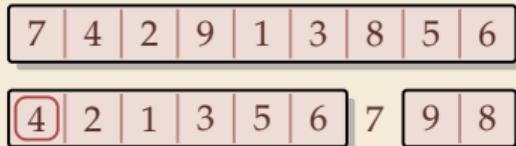
# Quicksort & Binary Search Trees

## Quicksort



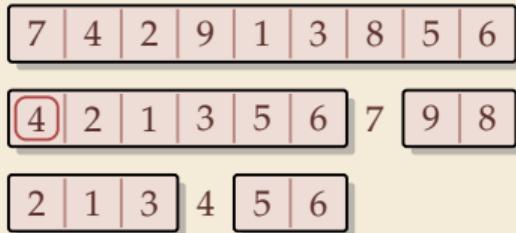
# Quicksort & Binary Search Trees

## Quicksort



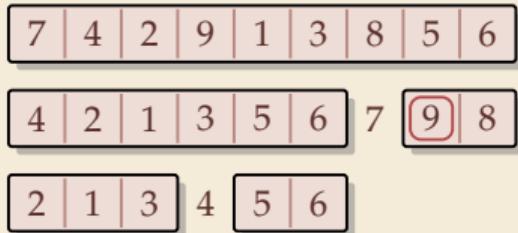
# Quicksort & Binary Search Trees

## Quicksort



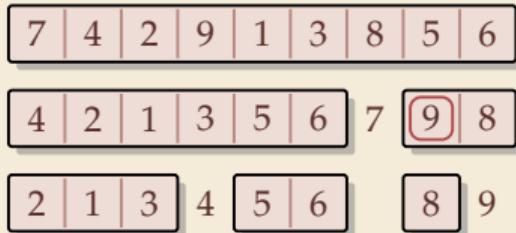
# Quicksort & Binary Search Trees

## Quicksort



# Quicksort & Binary Search Trees

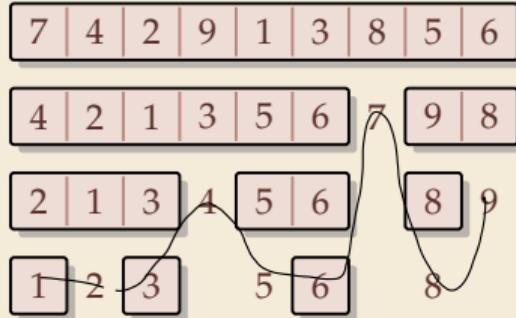
## Quicksort



# Quicksort & Binary Search Trees

## Quicksort

time ↓



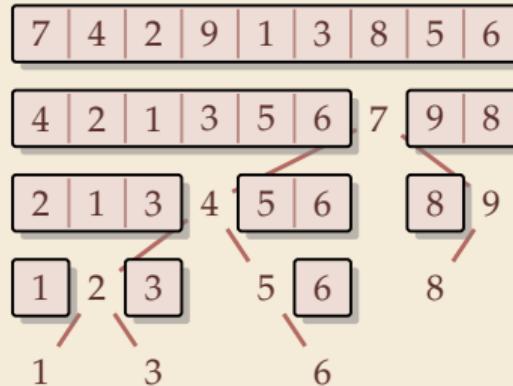
# Quicksort & Binary Search Trees

## Quicksort



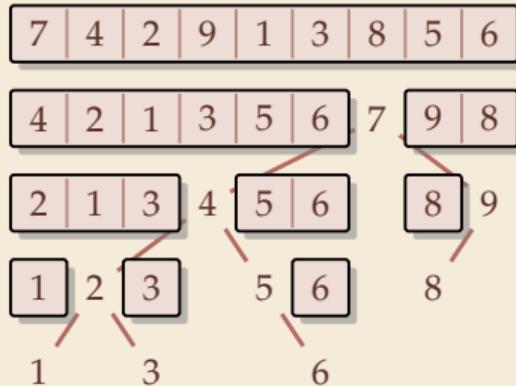
# Quicksort & Binary Search Trees

## Quicksort



# Quicksort & Binary Search Trees

Quicksort

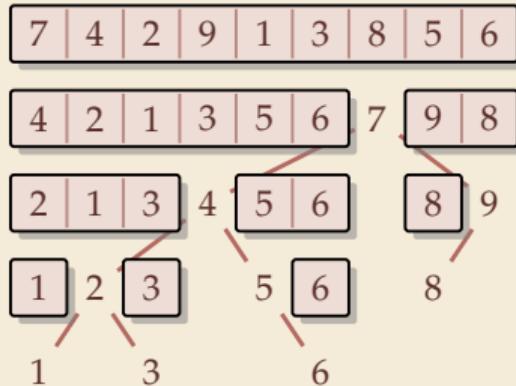


Binary Search Tree (BST)

7 4 2 9 1 3 8 5 6

# Quicksort & Binary Search Trees

Quicksort

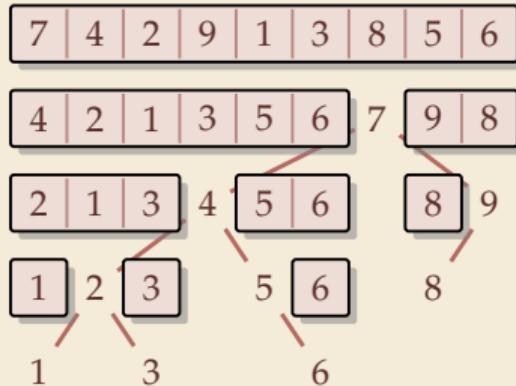


Binary Search Tree (BST)

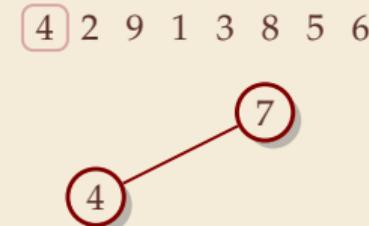


# Quicksort & Binary Search Trees

Quicksort

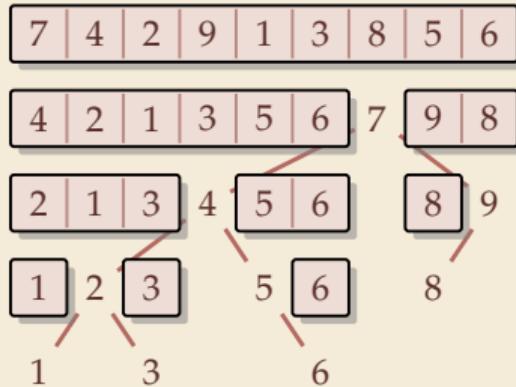


Binary Search Tree (BST)

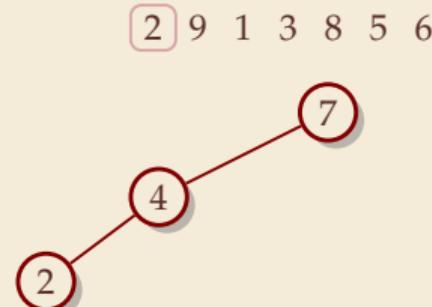


# Quicksort & Binary Search Trees

Quicksort

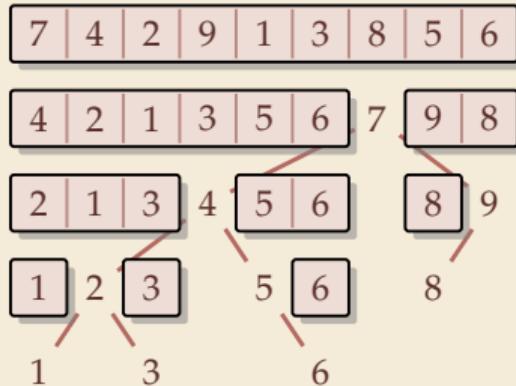


Binary Search Tree (BST)

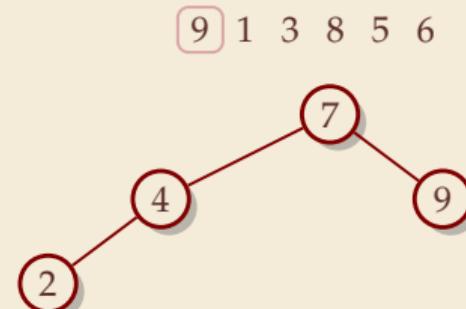


# Quicksort & Binary Search Trees

Quicksort

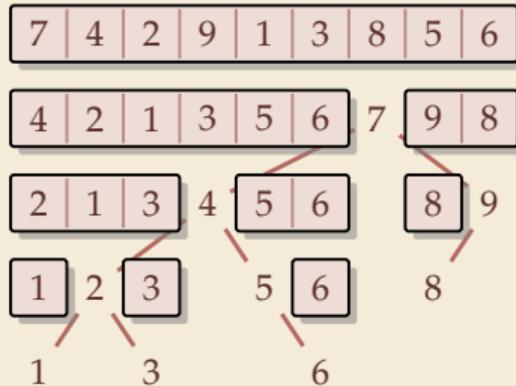


Binary Search Tree (BST)

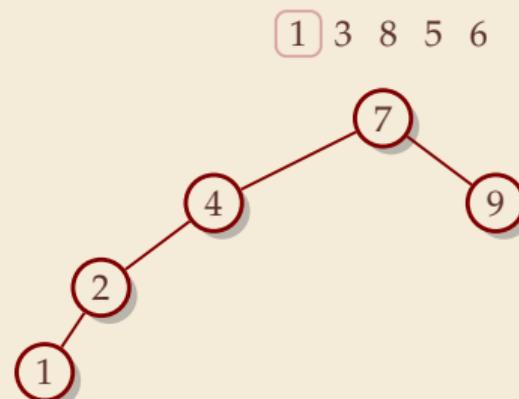


# Quicksort & Binary Search Trees

Quicksort

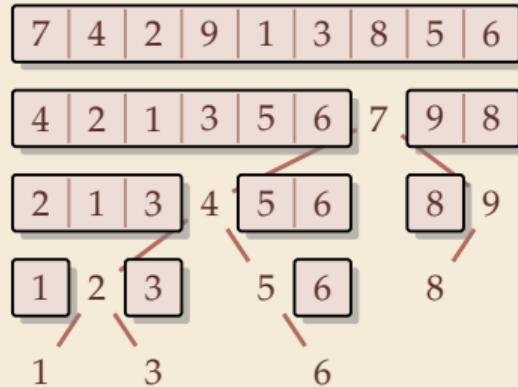


Binary Search Tree (BST)

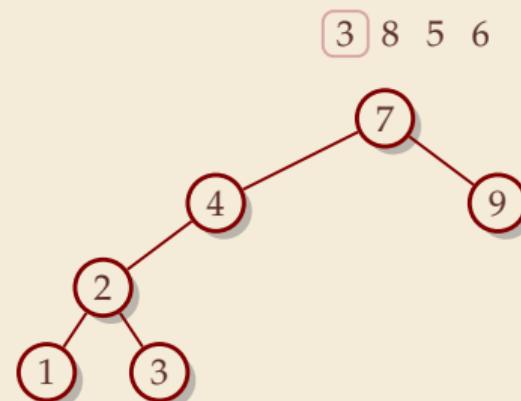


# Quicksort & Binary Search Trees

Quicksort

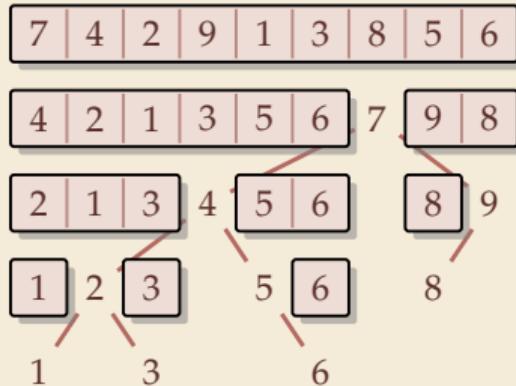


Binary Search Tree (BST)

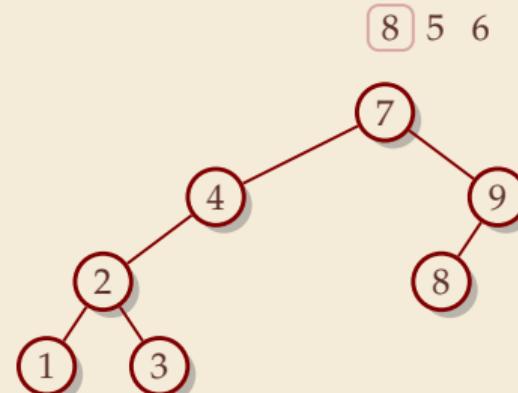


# Quicksort & Binary Search Trees

Quicksort

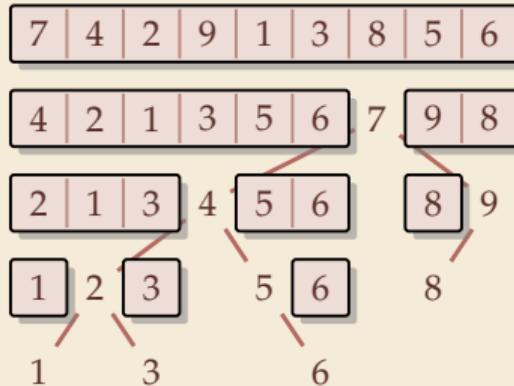


Binary Search Tree (BST)

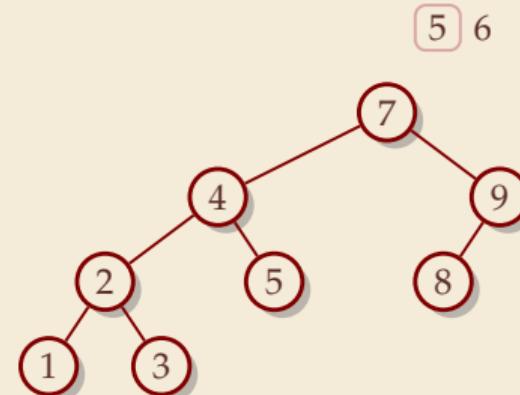


# Quicksort & Binary Search Trees

Quicksort

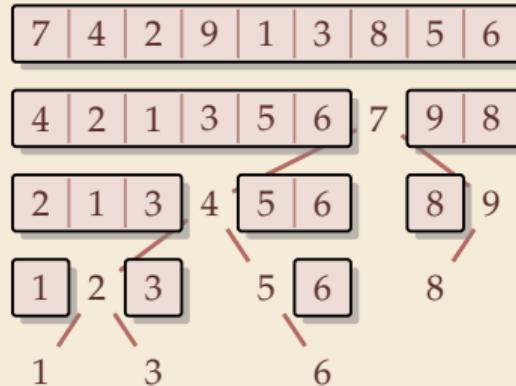


Binary Search Tree (BST)

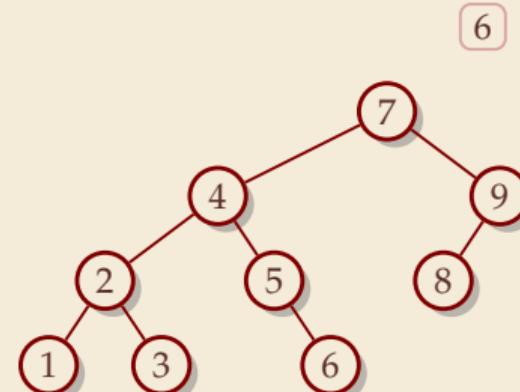


# Quicksort & Binary Search Trees

Quicksort

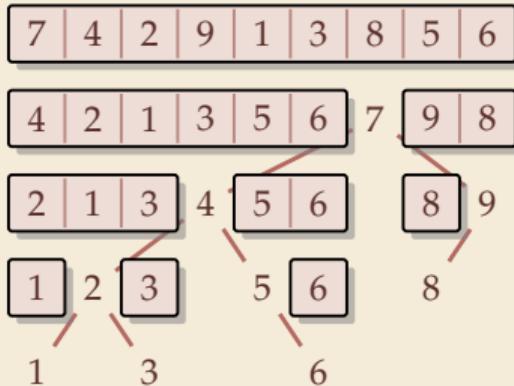


Binary Search Tree (BST)

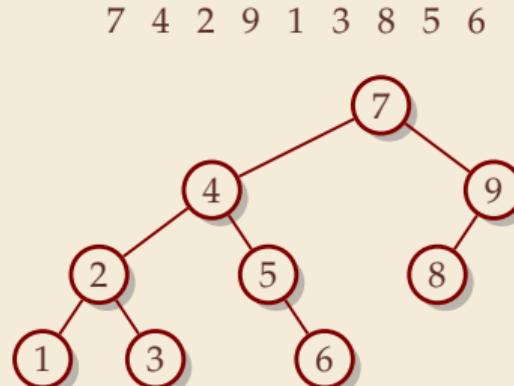


# Quicksort & Binary Search Trees

Quicksort



Binary Search Tree (BST)



- recursion tree of quicksort = binary search tree from successive insertion
- comparisons in quicksort = comparisons to built BST
- comparisons in quicksort  $\approx$  comparisons to search each element in BST

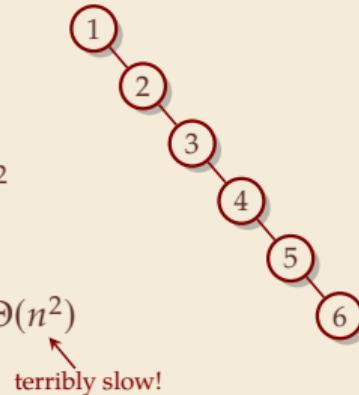
# Quicksort – Worst Case

- ▶ Problem: BSTs can degenerate

- ▶ Cost to search for  $k$  is  $k - 1$

$$\rightsquigarrow \text{Total cost } \sum_{k=1}^n (k - 1) = \frac{n(n - 1)}{2} \sim \frac{1}{2}n^2$$

rightsquigarrow quicksort worst-case running time is in  $\Theta(n^2)$



terribly slow!

But, we can fix this:

## Randomized quicksort:

- ▶ choose a *random pivot* in each step

rightsquigarrow same as randomly shuffling input before sorting    ↴

## Randomized Quicksort – Analysis

- ▶  $C(n)$  = element visits (as for mergesort)
  - ~~ quicksort needs  $\sim 2 \ln(2) \cdot n \lg n \approx \underline{1.39n \lg n}$  *in expectation*
- ▶ also: very unlikely to be much worse:
  - e. g., one can prove:  $\Pr[\text{cost} > 10n \lg n] = O(n^{-2.5})$
  - distribution of costs is “concentrated around mean”
- ▶ intuition: have to be constantly unlucky with pivot choice ]

## Quicksort – Discussion

- thumb up fastest general-purpose method
- thumb up  $\Theta(n \log n)$  average case
- thumb up works *in-place* (no extra space required)
- thumb up memory access is sequential (scans over arrays)
- thumb down  $\Theta(n^2)$  worst case (although extremely unlikely) —————
- thumb down not a *stable* sorting method

Open problem: Simple algorithm that is fast, stable and in-place.

### 3.3 Comparison-Based Lower Bound

# Lower Bounds

- ▶ **Lower bound:** mathematical proof that no algorithm can do better.
  - ▶ very powerful concept: bulletproof *impossibility* result  
≈ *conservation of energy* in physics
  - ▶ **(unique?) feature of computer science:**  
for many problems, solutions are known that (asymptotically) *achieve the lower bound*  
~~ can speak of “*optimal* algorithms”

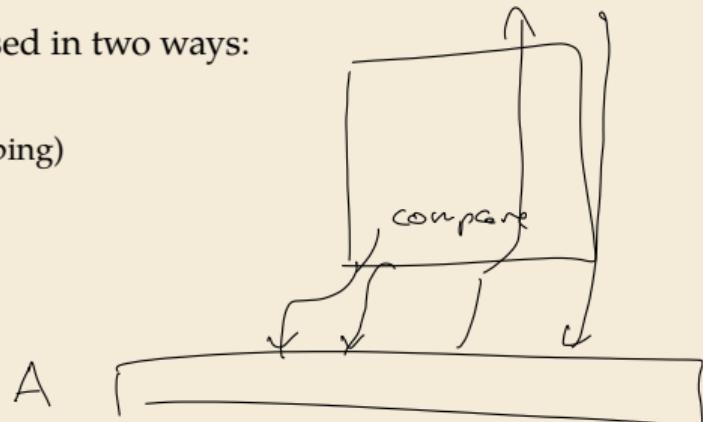
# Lower Bounds

- ▶ **Lower bound:** mathematical proof that no algorithm can do better.
  - ▶ very powerful concept: bulletproof *impossibility* result  
≈ *conservation of energy* in physics
  - ▶ **(unique?) feature of computer science:**  
for many problems, solutions are known that (asymptotically) *achieve the lower bound*  
~~ can speak of “*optimal* algorithms”
- ▶ To prove a statement about *all algorithms*, we must precisely define what that is!
- ▶ already know one option: the word-RAM model
- ▶ Here: use a simpler, more restricted model.

# The Comparison Model

buffer

- ▶ In the *comparison model* data can only be accessed in two ways:
  - ▶ comparing two elements
  - ▶ moving elements around (e.g. copying, swapping)
  - ▶ Cost: number of these operations.



# The Comparison Model

- ▶ In the *comparison model* data can only be accessed in two ways:
  - ▶ comparing two elements
  - ▶ moving elements around (e.g. copying, swapping)
  - ▶ Cost: number of these operations.
- ▶ This makes very few assumptions on the kind of objects we are sorting.
- ▶ Mergesort and Quicksort work in the comparison model.

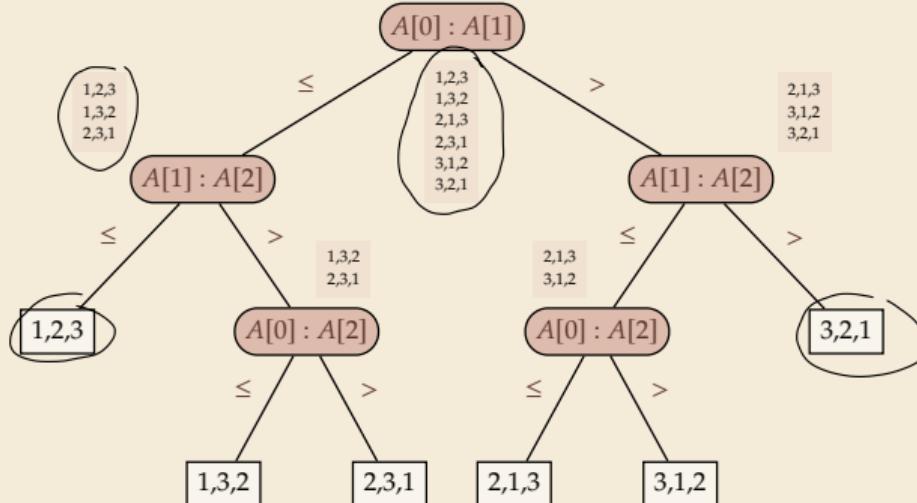
That's good!  
Keeps algorithms general!

# The Comparison Model

- ▶ In the *comparison model* data can only be accessed in two ways:
    - ▶ comparing two elements
    - ▶ moving elements around (e.g. copying, swapping)
    - ▶ Cost: number of these operations.
  - ▶ This makes very few assumptions on the kind of objects we are sorting.
    - That's good!  
Keeps algorithms general!
  - ▶ Mergesort and Quicksort work in the comparison model.
- ~~ Every comparison-based sorting algorithm corresponds to a *decision tree*.
- ▶ only model comparisons    ~~ ignore data movement
  - ▶ nodes = comparisons the algorithm does —
  - ▶ next comparisons can depend on outcomes    ~~ different subtrees
  - ▶ child links = outcomes of comparison
  - ▶ leaf = unique initial input permutation compatible with comparison outcomes

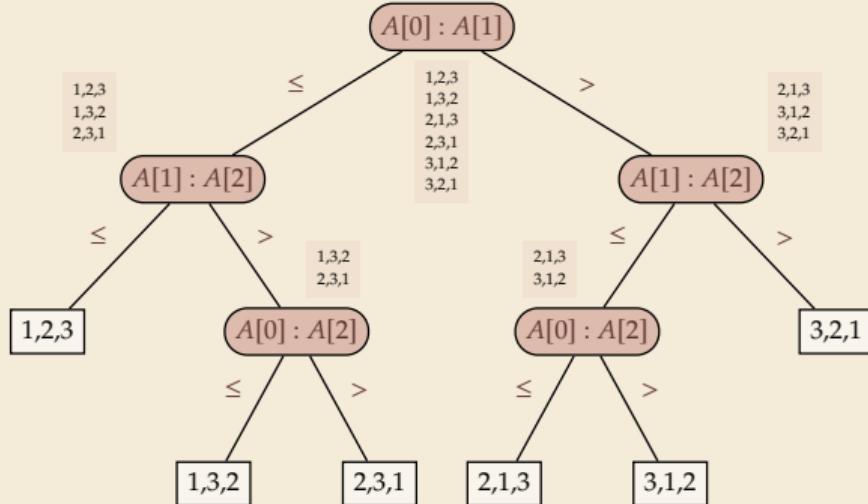
# Comparison Lower Bound

Example: Comparison tree for a sorting method for  $A[0..2]$ :



# Comparison Lower Bound

Example: Comparison tree for a sorting method for  $A[0..2]$ :

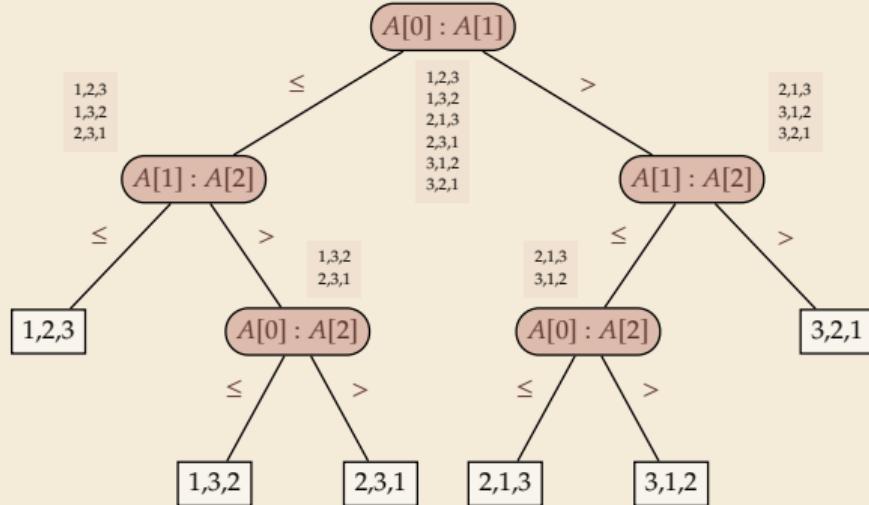


- ▶ Execution = follow a path in comparison tree.
  - ~~ height of comparison tree = worst-case # comparisons
- ▶ comparison trees are *binary* trees
  - ~~  $\ell$  leaves ~~ height  $\geq \lceil \lg(\ell) \rceil$
- ▶ comparison trees for sorting method must have  $\geq \underline{n!}$  leaves
  - ~~ height  $\geq \lg(n!) \sim \underline{n \lg n}$

more precisely:  $\lg(n!) = n \lg n - \lg(e)n + O(\log n)$

# Comparison Lower Bound

Example: Comparison tree for a sorting method for  $A[0..2]$ :



- ▶ Execution = follow a path in comparison tree.  
~~ height of comparison tree = worst-case # comparisons
- ▶ comparison trees are *binary* trees  
~~  $\ell$  leaves ~~ height  $\geq \lceil \lg(\ell) \rceil$
- ▶ comparison trees for sorting method must have  $\geq n!$  leaves  
~~ height  $\geq \lg(n!) \sim n \lg n$   
more precisely:  $\lg(n!) = n \lg n - \lg(e)n + O(\log n)$

- ▶ Mergesort achieves  $\sim n \lg n$  comparisons ~~ asymptotically comparison-optimal!
- ▶ Open (theory) problem: Can we sort with  $n \lg n - \lg(e)n + o(n)$  comparisons?

$$\approx 1.4427$$

## Clicker Question



Does the comparison-tree from the previous slide correspond to a worst-case optimal sorting method?

**A** Yes

**B** No

*pingo.upb.de/622222*

## Clicker Question



Does the comparison-tree from the previous slide correspond to a worst-case optimal sorting method?

A

Yes ✓

B

No

*pingo.upb.de/622222*

## 3.4 Integer Sorting

## How to beat a lower bound

- ▶ Does the above lower bound mean, sorting always takes time  $\Omega(n \log n)$ ?

## How to beat a lower bound

- ▶ Does the above lower bound mean, sorting always takes time  $\Omega(n \log n)$ ?
- ▶ **Not necessarily;** only in the *comparison model*!
  - ~~ Lower bounds show where to change the model!

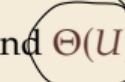
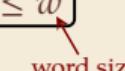
# How to beat a lower bound

- ▶ Does the above lower bound mean, sorting always takes time  $\Omega(n \log n)$ ?
- ▶ **Not necessarily;** only in the *comparison model*!
  - ~~ Lower bounds show where to change the model!
- ▶ Here: sort *n integers*
  - ▶ can do *a lot* with integers: add them up, compute averages, ... (full power of word-RAM)
    - ~~ we are **not** working in the comparison model
    - ~~ *above lower bound does not apply!*

# How to beat a lower bound

- ▶ Does the above lower bound mean, sorting always takes time  $\Omega(n \log n)$ ?
- ▶ **Not necessarily;** only in the *comparison model*!
  - ~~ Lower bounds show where to change the model!
- ▶ Here: sort *n integers*
  - ▶ can do *a lot* with integers: add them up, compute averages, ... (full power of word-RAM)
    - ~~ we are **not** working in the comparison model
    - ~~ *above lower bound does not apply!*
  - ▶ but: a priori unclear how much arithmetic helps for sorting ...

# Counting sort

- ▶ Important parameter: size/range of numbers
  - ▶ numbers in range  $[0..U] = \{0, \dots, U-1\}$  typically  $U = 2^b \rightsquigarrow b$ -bit binary numbers
- ▶ We can sort  $n$  integers in  $\Theta(n + U)$  time and  $\Theta(U)$  space when  $b \leq w$   
  


## Counting sort

```
1 procedure countingSort( $A[0..n - 1]$ )
2   //  $A$  contains integers in range  $[0..U]$ .
3    $C[0..U - 1] :=$  new integer array, initialized to 0
4   // Count occurrences
5   for  $i := 0, \dots, n - 1$ 
6      $C[A[i]] := C[A[i]] + 1$ 
7     i := 0 // Produce sorted list
8   for  $k := 0, \dots, U - 1$ 
9     for  $j := 1, \dots, C[k]$ 
10     $A[i] := k; i := i + 1$ 
```

count  
#occurrences  
of all  $i \in [0..U]$

- Java uses this `sort(byte[])`
- ▶ count how often each possible value occurs
  - ▶ produce sorted result directly from counts
  - ▶ circumvents lower bound by using integers as array index / pointer offset

~ Can sort  $n$  integers in range  $[0..U]$  with  $U = O(n)$  in time and space  $\Theta(n)$ . 

# Integer Sorting – State of the art

- ▶  $O(n)$  time sorting also possible for numbers in range  $U = O(n^c)$  for constant  $c$ .
  - ▶ radix sort with radix  $2^w$

- ▶ **algorithm theory**
  - ▶ suppose  $U = 2^w$ , but  $w$  can be arbitrary function of  $n$
  - ▶ how fast can we sort  $n$  such  $w$ -bit integers on a  $w$ -bit word-RAM?
    - ▶ for  $w = O(\log n)$ : linear time (*radix/counting sort*)
    - ▶ for  $w = \Omega(\log^{2+\epsilon} n)$ : linear time (*signature sort*)
    - ▶ for  $w$  in between: can do  $O(n\sqrt{\lg \lg n})$  (very complicated algorithm)  
don't know if that is best possible!

*Unit 1 :  $w = \Theta(\log n)$*

outside  
of exam

# Integer Sorting – State of the art

- ▶  $O(n)$  time sorting also possible for numbers in range  $U = O(n^c)$  for constant  $c$ .
  - ▶ *radix sort* with radix  $2^w$
- ▶ **algorithm theory**
  - ▶ suppose  $U = 2^w$ , but  $w$  can be arbitrary function of  $n$
  - ▶ how fast can we sort  $n$  such  $w$ -bit integers on a  $w$ -bit word-RAM?
    - ▶ for  $w = O(\log n)$ : linear time (*radix/counting sort*)
    - ▶ for  $w = \Omega(\log^{2+\varepsilon} n)$ : linear time (*signature sort*)
    - ▶ for  $w$  in between: can do  $O(n\sqrt{\lg \lg n})$  (very complicated algorithm)  
don't know if that is best possible!

\* \* \*

- ▶ for the rest of this unit: back to the comparisons model!

# Part II

## *Sorting with many processors*

## 3.5 Parallel computation

# Types of parallel computation

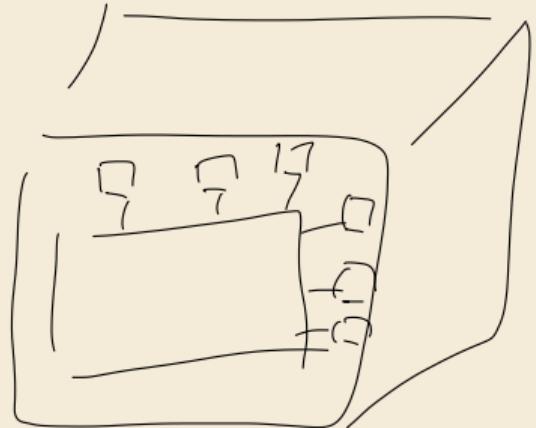
£££ can't buy you more time, but more computers!

~~ Challenge: Algorithms for parallel computation.

There are two main forms of parallelism

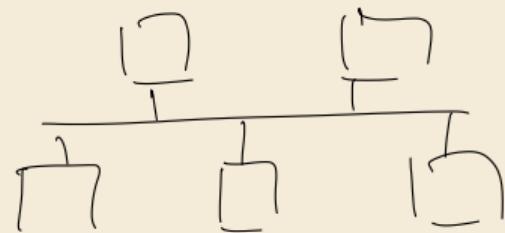
## 1. shared-memory parallel computer $\leftarrow$ focus of today

- ▶  $p$  processing elements (PEs, processors) working in parallel
- ▶ single big memory, accessible from every PE
- ▶ communication via shared memory
- ▶ think: a big server, 128 CPU cores, terabyte of main memory



## 2. distributed computing

- ▶  $p$  PEs working in parallel
- ▶ each PE has private memory
- ▶ communication by sending messages via a network
- ▶ think: a cluster of individual machines



# PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 1)
- ▶ the  $p$  PEs are identified by ids  $0, \dots, p - 1$ 
  - ▶ like  $w$  (the word size),  $p$  is a parameter of the model that can grow with  $n$
  - ▶  $\underline{p = \Theta(n)}$  is not unusual      many processors!
- ▶ the PEs all **independently** run a RAM-style program  
(they can use their id there) |
- ▶ each PE has its own registers, but MEM is shared among all PEs
- ▶ computation runs in synchronous steps:  
in each time step, every PE executes one instruction

↙ questionable  
assumption  
in practice ...

# PRAM – Conflict management



**Problem:** What if several PEs simultaneously overwrite a memory cell?

- ▶ **EREW-PRAM** (exclusive read, exclusive write)  
any **parallel access** to same memory cell is **forbidden** (crash if happens)

- ▶ **CREW-PRAM** (concurrent read, exclusive write)  
parallel **write** access to same memory cell is *forbidden*, but reading is fine

- ▶ **CRCW-PRAM** (concurrent read, concurrent write)  
concurrent access is allowed,  
need a rule for write conflicts:

- ▶ common CRCW-PRAM:  
all concurrent writes to same cell must write *same* value
- ▶ arbitrary CRCW-PRAM:  
some unspecified concurrent write wins
- ▶ (more exist ...)

- ▶ no single model is always adequate, but our default is CREW

# PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory      same as for RAM
- ▶ **time:** number of steps till all PEs finish      assuming sufficiently many PEs!  
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

Holy grail of PRAM algorithms:

- ▶ minimal time
- ▶ work (asymptotically) no worse than running time of best sequential algorithm
  - ▶ *work-efficient* algorithm: work in same  $\Theta$ -class as best sequential

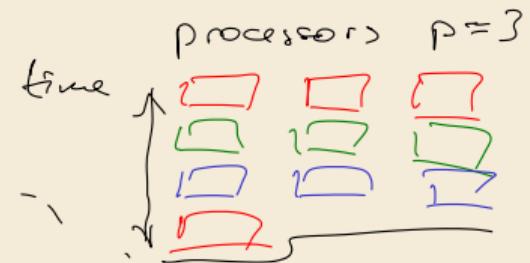
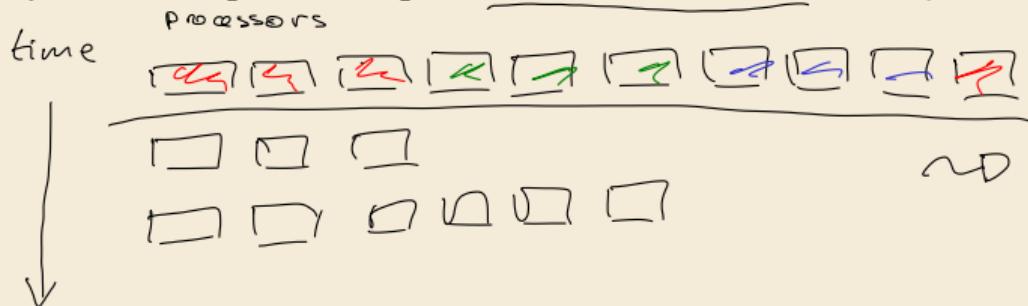
# The number of processors

Hold on, my computer does not have  $\Theta(n)$  processors! Why should I care for span and work!?

## Theorem 3.1 (Brent's Theorem):

If an algorithm has span  $T$  and work  $W$  (for an arbitrarily large number of processors), it can be run on a PRAM with  $p$  PEs in time  $O(T + \frac{W}{p})$  (and using  $O(W)$  work). ◀

Proof: schedule parallel steps in round-robin fashion on the  $p$  PEs.



~~ span and work give guideline for *any* number of processors

## 3.6 Parallel primitives

# Prefix sums

Before we come to parallel sorting, we study some useful building blocks.

**Prefix-sum problem** (also: cumulative sums, running totals)

- ▶ Given: array  $A[0..n - 1]$  of numbers
- ▶ Goal: compute all prefix sums  $A[0] + \dots + A[i]$  for  $i = 0, \dots, n - 1$   
may be done “in-place”, i. e., by overwriting  $A$

**Example:**

input:

3		0		0		5		7		0		0		2		0		0		0		4		0		8		0		1
---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--	---

$\Sigma$

output:

3		3		3		8		15		15		15		17		17		17		17		17		21		21		29		29		30
---	--	---	--	---	--	---	--	----	--	----	--	----	--	----	--	----	--	----	--	----	--	----	--	----	--	----	--	----	--	----	--	----

# Clicker Question



What is the *sequential* running time achievable for prefix sums?

A  $O(n^3)$

B  $O(n^2)$

C  $O(n \log n)$

D  $O(n)$

E  $O(\sqrt{n})$

F  $O(\log n)$

*pingo.upb.de/622222*

# Clicker Question

What is the *sequential* running time achievable for prefix sums?



A  $\cancel{O(n^3)}$

D  $O(n)$  ✓

B  $\cancel{O(n^2)}$

E  $\cancel{O(\sqrt{n})}$

C  $\cancel{O(n \log n)}$

F  $\cancel{O(\log n)}$

*pingo.upb.de/622222*

## Prefix sums – Sequential

- ▶ sequential solution does  $n - 1$  additions
- ▶ but: cannot parallelize them
  - data dependencies!
- ~~ need a different approach

---

```
1 procedure prefixSum( $A[0..n - 1]$ )
2     for  $i := 1, \dots, n - 1$  do
3          $A[i] := A[i - 1] + A[i]$ 
```

---