

4

String Matching – What's behind Ctrl+F?

24 February 2020

Sebastian Wild

4 String Matching

- 4.1 Introduction
- 4.2 Brute Force
- 4.3 String Matching with Finite Automata
- 4.4 The Knuth-Morris-Pratt algorithm
- 4.5 Beyond Optimal? The Boyer-Moore Algorithm
- 4.6 The Rabin-Karp Algorithm

4.1 Introduction

Ubiquitous strings

string = sequence of characters

- ▶ universal data type for ... everything!
 - ▶ natural language texts
 - ▶ programs (source code)
 - ▶ websites
 - ▶ XML documents
 - ▶ DNA sequences
 - ▶ bitstrings
 - ▶ ... a computer's memory \rightsquigarrow ultimately any data is a string

\rightsquigarrow many different tasks and algorithms

Ubiquitous strings

string = sequence of characters

- ▶ universal data type for ... everything!
 - ▶ natural language texts
 - ▶ programs (source code)
 - ▶ websites
 - ▶ XML documents
 - ▶ DNA sequences
 - ▶ bitstrings
 - ▶ ... a computer's memory ~→ ultimately any data is a string


~→ many different tasks and algorithms

- ▶ This unit: finding (exact) **occurrences of a pattern** text.
 - ▶ Ctrl+F
 - ▶ grep
 - ▶ computer forensics (e. g. find signature of file on disk)
 - ▶ virus scanner
- ▶ basis for many advanced applications

Notations

- ▶ *alphabet* Σ : finite set of allowed **characters**; $\sigma = |\Sigma|$ “a string over alphabet Σ ”
 - ▶ letters (Latin, Greek, Arabic, Cyrillic, Asian scripts, ...)
 - ▶ “what you can type on a keyboard”, Unicode characters
 - ▶ $\{0, 1\}$; nucleotides $\{A, C, G, T\}$; ...
- ↖ comprehensive standard character set
including emoji and all known symbols

Notations

- ▶ *alphabet* Σ : finite set of allowed **characters**; $|\Sigma|$ “a string over alphabet Σ ”
 - ▶ letters (Latin, Greek, Arabic, Cyrillic, Asian scripts, ...)
 - ▶ “what you can type on a keyboard”, Unicode characters
 - ▶ $\{0, 1\}$; nucleotides $\{A, C, G, T\}$; ...
comprehensive standard character set including emoji and all known symbols
- ▶ $\Sigma^n = \Sigma \times \dots \times \Sigma$: strings of **length** $n \in \mathbb{N}_0$ (n -tuples)
- ▶ $\Sigma^* = \bigcup_{n \geq 0} \Sigma^n$: set of **all** (finite) strings over Σ
- ▶ $\Sigma^+ = \bigcup_{n \geq 1} \Sigma^n$: set of **all** (finite) **nonempty** strings over Σ
- ▶ ε $\in \Sigma^0$: the *empty* string (same for all alphabets)

Notations

- ▶ *alphabet* Σ : finite set of allowed **characters**; $\sigma = |\Sigma|$ “a string over alphabet Σ ”
 - ▶ letters (Latin, Greek, Arabic, Cyrillic, Asian scripts, ...)
 - ▶ “what you can type on a keyboard”, Unicode characters
 - ▶ $\{0, 1\}$; nucleotides $\{A, C, G, T\}$; ...
 - comprehensive standard character set including emoji and all known symbols
- ▶ $\Sigma^n = \Sigma \times \cdots \times \Sigma$: strings of **length** $n \in \mathbb{N}_0$ (n -tuples)
- ▶ $\Sigma^* = \bigcup_{n \geq 0} \Sigma^n$: set of **all** (finite) strings over Σ
- ▶ $\Sigma^+ = \bigcup_{n \geq 1} \Sigma^n$: set of **all** (finite) **nonempty** strings over Σ
- ▶ $\varepsilon \in \Sigma^0$: the *empty* string (same for all alphabets)
 - \swarrow $\Sigma^0 \dots \Sigma^n - \{\}$
- ▶ for $\underline{S \in \Sigma^n}$, write $S[i]$ (other sources: S_i) for ***i*th** character ($0 \leq i < n$)
 - zero-based (like arrays!)
- ▶ for $S, T \in \Sigma^*$, write $ST = S \cdot T$ for **concatenation** of S and T
- ▶ for $S \in \Sigma^n$, write $S[i..j]$ or $S_{i,j}$ for the **substring** $S[i] \cdot S[i+1] \cdots S[j]$ ($0 \leq i \leq j < n$)
 - ▶ $S[0..j]$ is a **prefix** of S ; $S[i..n-1]$ is a **suffix** of S
 - ▶ $S[i..j) = S[i..j \text{ } \overline{\neq} 1]$ (endpoint exclusive) $\rightsquigarrow S = \underline{S[0..n)}$

Clicker Question



True or false: $\Sigma^* = \Sigma^+ \cup \{\varepsilon\}$

A True

B False

pingo.upb.de/622222

Clicker Question



True or false: $\Sigma^* = \Sigma^+ \cup \{\varepsilon\}$

A True ✓

B ~~False~~

pingo.upb.de/622222

String matching – Definition

Search for a string (pattern) in a large body of text

► Input:

- $T \in \Sigma^n$: The text (haystack) being searched within
- $P \in \Sigma^m$: The pattern (needle) being searched for; typically $n \gg m$

► Output:

- the *first occurrence (match)* of P in T : $\min\{i \in [0..n - m) : T[i..i + m) = P\}$
- or NO_MATCH if there is no such i (“ P does not occur in T ”)

► Variant: Find **all** occurrences of P in T .

↪ Can do that iteratively (update T to $T[i + 1..n)$ after match at i)

► Example:

- $T = \text{“Where is he?”}$
- $P_1 = \text{“he”} \rightsquigarrow i = 1$
- $P_2 = \text{“who”} \rightsquigarrow \text{NO_MATCH}$

► string matching is implemented in Java in String.indexOf

4.2 Brute Force

Abstract idea of algorithms

Pattern matching algorithms consist of *guesses* and *checks*:

- ▶ A **guess** is a position i such that P might start at $T[i]$.
Possible guesses (initially) are $0 \leq i \leq n - m$.
- ▶ A **check** of a guess is a pair $(\underline{i}, \underline{j})$ where we compare $T[i + j]$ to $P[j]$.
- ▶ Note: need all m checks to verify a single **correct** guess \underline{i} ,
but it may take (many) fewer checks to recognize an **incorrect** guess.
- ▶ Cost measure: #character comparisons = #checks

$\leadsto \text{cost} \leq n \cdot m$ (number of possible checks)

Brute-force method

```

1 procedure bruteForceSM( $T[0..n], P[0..m]$ )
2   for  $i := 0, \dots, n - m - 1$  do
3     for  $j := 0, \dots, m - 1$  do
4       if  $T[i + j] \neq P[j]$  then break inner loop
5     if  $j == m$  then return  $i$ 
6   return NO MATCH

```

- ▶ try all guesses i
- ▶ check each guess (left to right); stop early on mismatch
- ▶ essentially the implementation in Java!

► Example:

$T = \text{abbbababbab}$

$P = \text{abba}$

~ 15 char cmps
(vs $n \cdot m = 44$)
not too bad!

[illegible]

Brute-force method

```
1 procedure bruteForceSM( $T[0..n]$ ,  $P[0..m]$ )  
2   for  $i := 0, \dots, n - m - 1$  do  
3     for  $j := 0, \dots, m - 1$  do  
4       if  $T[i + j] \neq P[j]$  then break inner loop  
5       if  $j == m$  then return  $i$   
6   return NO_MATCH
```

- try all guesses i
- check each guess (left to right); stop early on mismatch
- essentially the implementation in Java!

► **Example:**

$T = \text{abbbababbab}$

$P = \text{abba}$

↪ 15 char cmps
(vs $n \cdot m = 44$)
not too bad!

	a	b	b	b	a	b	a	b	b	a	b
a	a	b	b	a							
		a									
			a								
				a							
					a	b	b				
						a					
							a	b	b	a	

Brute-force method – Discussion



Brute-force method can be good enough

- ▶ typically works well for natural language text
- ▶ also for random strings



but: can be as bad as it gets!

	a	a	a	a	a	a	a	a	a	a	a
a	a	a	b								
	a	a	a	b							
		a	a	a	b						
			a	a	a	b					
				a	a	a	b				
					a	a	a	b			
						a	a	a	b		
							a	a	a	b	

▶ Worst possible input: $P = a^{m-1}b$,
 $T = a^n$

▶ Worst-case performance: $(n - m + 1) \cdot m$

\rightsquigarrow for $m \leq n/2$ that is $\Theta(mn)$

Brute-force method – Discussion



Brute-force method can be good enough

- ▶ typically works well for natural language text
- ▶ also for random strings



but: can be as bad as it gets!

	a	a	a	a	a	a	a	a	a	a	a
a	a	a	b								
	a	a	a	b							
		a	a	a	b						
			a	a	a	b					
				a	a	a	b				
					a	a	a	b			
						a	a	a	b		
							a	a	a	b	

▶ Worst possible input: $P = a^{m-1}b$,
 $T = a^n$

▶ Worst-case performance: $(n - m + 1) \cdot m$

\rightsquigarrow for $m \leq n/2$ that is $\Theta(mn)$

- ▶ Bad input: lots of self-similarity in T ! \rightsquigarrow can we exploit that?
- ▶ brute force does ‘obviously’ stupid repetitive comparisons \rightsquigarrow can we avoid that?

Roadmap

- ▶ **Approach 1** (this week): Use *preprocessing* on the pattern P to eliminate guesses (avoid 'obvious' redundant work)
 - ▶ Deterministic finite automata (DFA)
 - ▶ Knuth-Morris-Pratt algorithm
 - ▶ Boyer-Moore algorithm
 - ▶ Rabin-Karp algorithm
- ▶ **Approach 2** (↪ Unit 6): Do preprocessing on the text T
Can find matches in time *independent of text size(!)*
 - ▶ inverted indices
 - ▶ Suffix trees
 - ▶ Suffix arrays

4.3 String Matching with Finite Automata

Clicker Question



Do you know what regular expressions, NFAs and DFAs are, and how to convert between them?

- ☐ A Never heard of this; are these new emoji?
- ☐ B Heard the terms, but don't remember conversion methods.
- ☐ C Had that in my undergrad course (memories fading a bit).
- ☐ D Sure, I could do that blindfolded!

pingo.upb.de/622222

Theoretical Computer Science to the rescue!

- ▶ string matching = deciding whether $T \in \Sigma^* \cdot P \cdot \Sigma^*$
- ▶ $\Sigma^* \cdot P \cdot \Sigma^*$ is *regular* formal language
- $\rightsquigarrow \exists$ *deterministic finite automaton* (DFA) to recognize $\Sigma^* \cdot P \cdot \Sigma^*$
- \rightsquigarrow can check for occurrence of P in $|T| = n$ steps!

Theoretical Computer Science to the rescue!

► string matching = deciding whether $T \in \Sigma^* \cdot P \cdot \Sigma^*$

► $\Sigma^* \cdot P \cdot \Sigma^*$ is *regular* formal language

↪ \exists *deterministic finite automaton* (DFA) to recognize $\Sigma^* \cdot P \cdot \Sigma^*$

↪ can check for occurrence of P in $|T| = n$ steps!



Job done!

Theoretical Computer Science to the rescue!

► string matching = deciding whether $T \in \Sigma^* \cdot P \cdot \Sigma^*$

► $\Sigma^* \cdot P \cdot \Sigma^*$ is *regular* formal language

↪ \exists *deterministic finite automaton* (DFA) to recognize $\Sigma^* \cdot P \cdot \Sigma^*$

↪ can check for occurrence of P in $|T| = n$ steps!



Job done!



WTF!?

Theoretical Computer Science to the rescue!

- ▶ string matching = deciding whether $T \in \Sigma^* \cdot P \cdot \Sigma^*$
- ▶ $\Sigma^* \cdot P \cdot \Sigma^*$ is *regular* formal language
- $\rightsquigarrow (\exists)$ *deterministic finite automaton* (DFA) to recognize $\Sigma^* \cdot P \cdot \Sigma^*$
- \rightsquigarrow can check for occurrence of P in $|T| = n$ steps!



Job done!



WTF!?

We are not quite done yet.

- ▶ (Problem 0: programmer might not know automata and formal languages ...)
- ▶ Problem 1: existence alone does not give an algorithm!
- ▶ Problem 2: automaton could be very big!

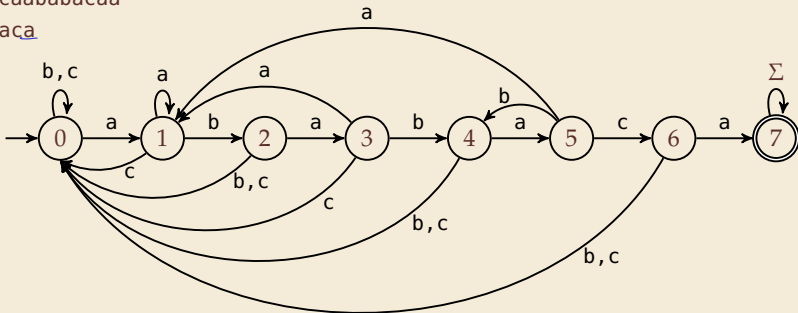
String matching with DFA

- Assume first, we already have a deterministic automaton
- How does string matching work?

Example:

$T = \text{aabacaababacaa}$

$P = \text{ababaca}$



text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3										

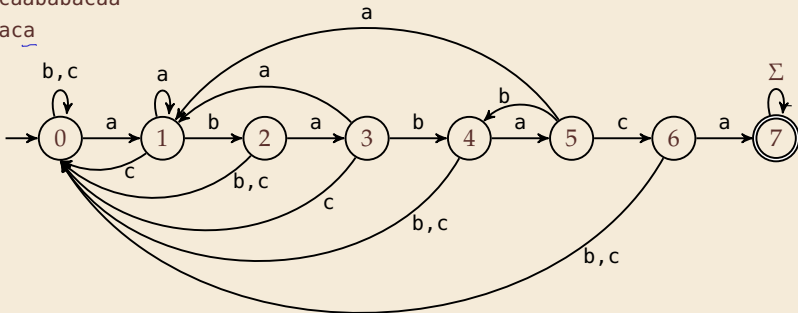
String matching with DFA

- Assume first, we already have a deterministic automaton
- How does string matching work?

Example:

$T = \text{aabacaababacaa}$

$P = \text{ababaca}$



text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3	0	1	<u>1</u>	2	3	4	5	6	7	7

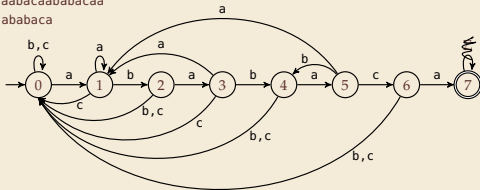
String matching DFA – Intuition

Why does this work?

- Main insight: *Invariant*

State q means:
*“we have seen $P[0..q]$ until here
 (but not any longer prefix of P)”*

$T = \text{aabacaababacaa}$
 $P = \text{ababaca}$



text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3	0	1	<u>1</u>	2	3	4	5	6	7	7

$P[q]$
 \neq

- If the next text character c does not match, we know:

- (i) text seen so far ends with $\underline{P[0..q]} \cdot c$
- (ii) $\underline{P[0..q]} \cdot c$ is not a prefix of \underline{P}
- (iii) without reading c , $\underline{P[0..q]}$ was the *longest* prefix of P that ends here.

$T = \dots \boxed{P[0..q]} \cdot c$
 $\boxed{P[0..q']}$


$$q' \leq q$$

⇒ New longest matched prefix will be (weakly) shorter than q

⇒ All information about the text needed to determine it is contained in $P[0..q] \cdot c$!

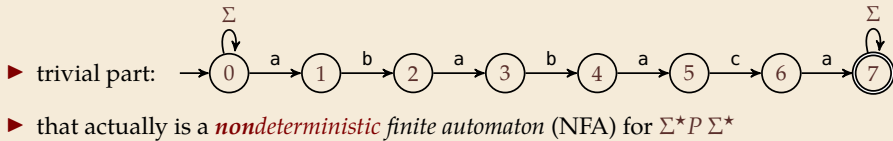
NFA instead of DFA?

It remains to *construct* the DFA.



NFA instead of DFA?

It remains to *construct* the DFA.



~> We *could* use the NFA directly for string matching:

- at any point in time, we are in a **set of states**
- accept when one of them is final state

Example: *Previous versions of this example were missing states; this is the correct version:*

text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	0,1	0,1	0,2	0,1,3	0	0,1	0,1	0,2	0,1,3	0,2,4	0,1,3,5	0,6	0,1,7	0,1,7

But maintaining a whole set makes this slow ...

Computing DFA directly



You have an NFA and want a DFA?
Simply apply the power-set construction
(and maybe DFA minimization)!

The powerset method has exponential state blow up!
I guess I might as well use brute force ...



Computing DFA directly



You have an NFA and want a DFA?
Simply apply the power-set construction
(and maybe DFA minimization)!

The powerset method has exponential state blow up!
I guess I might as well use brute force ...



Ingenious algorithm by Knuth, Morris, and Pratt: construct DFA *inductively*:

Suppose we add character $P[j]$ to automaton A_{j-1} for $P[0..j-1]$

- ▶ add new state and matching transition \rightsquigarrow easy
- ▶ for each $c \neq P[j]$, we need $\delta(j, c)$ (transition from j when reading c)



Computing DFA directly



You have an NFA and want a DFA?
Simply apply the power-set construction
(and maybe DFA minimization)!

The powerset method has exponential state blow up!
I guess I might as well use brute force ...





Ingenious algorithm by Knuth, Morris, and Pratt: construct DFA *inductively*:

Suppose we add character $P[j]$ to automaton A_{j-1} for $P[0..j-1]$

- ▶ add new state and matching transition \rightsquigarrow easy
- ▶ for each $c \neq P[j]$, we need $\delta(j, c)$ (transition from j when reading c)
- ▶ $\delta(j, c) =$ length of the longest prefix of $P[0..j]c$ that is a suffix of $P[1..j]c$
= state of automaton after reading $P[1..j]c$
 $\leq j \rightsquigarrow$ can use known automaton A_{j-1} for that!

\rightsquigarrow can directly compute A_j from A_{j-1} !

 seems to require simulating automata $m \cdot \sigma$ times



State q means:
"we have seen $P[0..q]$ until here
(but not any longer prefix of P)"



Computing DFA efficiently

- KMP's second insight: simulations in one step differ only in last symbol

↪ simply maintain state x , the state after reading $P[1..j-1]$.

- copy its transitions
- update x by following transitions for $P[j]$

Demo: Algorithms videos of Sedgewick and Wayne

Knuth-Morris-Pratt construction demo (in linear time)

Mismatch transition. For each state j and char $c \neq \text{pat.charAt}(j)$, set $\text{dfa}[c][j] = \text{dfa}[c][X]$; then update $X = \text{dfa}[\text{pat.charAt}(j)][X]$.

$X = \text{simulation of B A B A}$

	0	1	2	3	4	5
pat.charAt(j)	A	B	A	B	A	C
A	1	1	3	1	5	1
B	0	2	0	4	0	4
C	0	0	0	0	0	6

Constructing the DFA for KMP substring search for A B A B A C

<https://cuvds.io/app/video/194/watch>


String matching with DFA – Discussion


► Time:


- Matching: n table lookups for DFA transitions
 - building DFA: $\Theta(m\sigma)$ time (constant time per transition edge).
- $\rightsquigarrow \Theta(m\sigma + n)$ time for string matching.

► Space:

- $\Theta(m\sigma)$ space for transition matrix.

 **fast matching** time actually: hard to beat!

 total time asymptotically optimal for small alphabet (for $\sigma = O(n/m)$)

 substantial **space overhead**, in particular for large alphabets


4.4 The Knuth-Morris-Pratt algorithm

Failure Links

- ▶ Recall: String matching with is DFA fast,
but needs table of $m \times \sigma$ transitions.
- ▶ in fast DFA construction, we used that all simulations differ only by *last* symbol
- ~> **KMP's third insight:** do this last step of simulation from state x during *matching*!
... but how?

Failure Links


- ▶ Recall: String matching with DFA is fast, but needs table of $m \times \sigma$ transitions.
- ▶ in fast DFA construction, we used that all simulations differ only by *last* symbol
- ~> **KMP's third insight:** do this last step of simulation from state x during *matching*! ... but how?
- ▶ **Answer:** Use a new type of transition, the *failure links*
 - ▶ Use this transition (only) if no other one fits.
 - ▶ \times does not consume a character. ~> might follow several failure links



~> Computations are deterministic (but automaton is not a real DFA.)

Failure link automaton – Example

Example: $T = \text{abababaaaca}$, $P = \text{ababaca}$




$T :$	a	b	a	b	a	b	a	a	b	a	b
	0	1	2	3	4	5,3	4	5,3,1,0	1		

Failure link automaton – Example

Example: $T = \text{abababaaaca}$, $P = \text{ababaca}$

for failure link construction,
simulate on $P[1..j) = \text{babac}..$



T : a b a b a b a a b a b

P :	a	b	a	b	a	×					
			(a)	(b)	(a)	b	a	×			
								a	b	a	b

to state 3
to state 1

q :	1	2	3	4	5	3,4	5	3,1,0,1	2	3	4
-------	---	---	---	---	---	-----	---	---------	---	---	---

(after reading this character)

Clicker Question



What is the worst-case time to process one character in a failure-link automaton for $P[0..m]$?

A $\Theta(1)$


C $\Theta(m)$

B $\Theta(\log m)$

D $\Theta(m^2)$

[*pingo.upb.de/622222*](https://pingo.upb.de/622222)

Clicker Question



What is the worst-case time to process one character in a failure-link automaton for $P[0..m)$?

A ~~$\Theta(1)$~~

C $\Theta(m)$ ✓

B ~~$\Theta(\log m)$~~

D ~~$\Theta(m^2)$~~

pingo.upb.de/622222

The Knuth-Morris-Pratt Algorithm

```
1 procedure KMP( $T[0..n - 1]$ ,  $P[0..m - 1]$ )
2    $fail[0..m] := failureLinks(P)$ 
3    $i := 0$  // current position in  $T$ 
4    $q := 0$  // current state of KMP automaton
5   while  $i < n$  do
6     if  $T[i] == P[q]$  then
7        $i := i + 1$ ;  $q := q + 1$ 
8     if  $q == m$  then
9       return  $i - q$  // occurrence found
10    else // i.e.  $T[i] \neq P[q]$ 
11      if  $q \geq 1$  then
12         $q := fail[q]$  // follow one  $\times$ 
13      else
14         $i := i + 1$ 
15  end while
16  return NO_MATCH
```

► only need single array *fail* for failure links

► (procedure failureLinks later)

The Knuth-Morris-Pratt Algorithm

```
1 procedure KMP( $T[0..n-1]$ ,  $P[0..m-1]$ )
2    $fail[0..m] := failureLinks(P)$ 
3    $i := 0$  // current position in  $T$ 
4    $q := 0$  // current state of KMP automaton
5   while  $i < n$  do
6     if  $T[i] == P[q]$  then
7        $i := i + 1$ ;  $q := q + 1$ 
8       if  $q == m$  then
9         return  $i - q$  // occurrence found
10      else // i.e.  $T[i] \neq P[q]$ 
11        if  $q \geq 1$  then
12           $q := fail[q]$  // follow one  $\times$ 
13        else
14           $i := i + 1$ 
15   end while
16   return NO_MATCH
```

► only need single array *fail* for failure links

► (procedure failureLinks later)

Analysis: (matching part)

► always have $fail[j] < j$ for $j \geq 1$

↪ in each iteration

► either advance position in text ($i := i + 1$)

► or shift pattern forward (guess $i - j$)

► each can happen at most n times

↪ $\leq 2n$ symbol comparisons!

⇒ $O(1)$ time per character on average

Computing failure links

► failure links point to error state x (from DFA construction)

↪ run same algorithm, but store $fail[j] := x$ instead of copying all transitions

```
1 procedure failureLinks( $P[0..m-1]$ )
2    $fail[0] := 0$ 
3    $x := 0$ 
4   for  $j := 1, \dots, m-1$  do
5      $fail[j] := x$ 
6     // update failure state using failure links:
7     while  $P[x] \neq P[j]$ 
8       if  $x == 0$  then
9          $x := -1$ ; break
10      else
11         $x := fail[x]$ 
12      end while
13      $x := x + 1$ 
14 end for
```

simulates FLA on $P[1..j]$

Computing failure links

- ▶ failure links point to error state x (from DFA construction)

↪ run same algorithm, but store $fail[j] := x$ instead of copying all transitions

```
1 procedure failureLinks( $P[0..m-1]$ )
2    $fail[0] := 0$ 
3    $x := 0$ 
4   for  $j := 1, \dots, m-1$  do
5      $fail[j] := x$ 
6     // update failure state using failure links:
7     while  $P[x] \neq P[j]$ 
8       if  $x == 0$  then
9          $x := -1$ ; break
10      else
11         $x := fail[x]$ 
12      end while
13       $x := x + 1$ 
14  end for
```

Analysis:

- ▶ m iterations of for loop
- ▶ while loop always decrements x
- ▶ x is incremented only once per iteration of for loop

↪ $\leq m$ iterations of while loop *in total*

↪ $\leq 2m$ symbol comparisons

Knuth-Morris-Pratt – Discussion

► Time:

► $\leq 2n + 2m = O(n + m)$ character comparisons

► clearly must at least *read* both T and P

~> KMP has optimal worst-case complexity!

► Space:

► $\Theta(m)$ space for failure links

👍 total time asymptotically optimal (for any alphabet size)

👍 reasonable extra space

Clicker Question



What are the main advantages of the KMP string matching (using the failure-link automaton) over string matching with DFAs? Check all that apply.

- ☐ **A** faster preprocessing on pattern
- ☐ **B** faster matching in text
- ☐ **C** fewer character comparisons
- ☐ **D** uses less space
- ☐ **E** makes running time independent of σ
- ☐ **F** I don't have to do automata theory

pingo.upb.de/622222

Clicker Question



What are the main advantages of the KMP string matching (using the failure-link automaton) over string matching with DFAs? Check all that apply.

- ☒ A faster preprocessing on pattern ✓
- ☐ B ~~faster matching in text~~
- ☐ C ~~fewer character comparisons~~
- ☒ D uses less space ✓
- ☒ E makes running time independent of σ ✓
- ☐ F ~~I don't have to do automata theory~~

pingo.upb.de/622222

The KMP prefix function

- ▶ It turns out that the failure links are useful beyond KMP
- ▶ a slight variation is more widely used: (for historic reasons)
the (KMP) *prefix function* $F : [1..m - 1] \rightarrow [0..m - 1]$:
 *$F[j]$ is the length of the longest prefix of $P[0..j]$
that is a suffix of $P[1..j]$.*
- ▶ Can show: $fail[j] = F[j - 1]$ for $j \geq 1$, and hence

$fail[j] = \text{length of the longest prefix of } P[0..j] \text{ that is a suffix of } P[1..j].$

