6 Text Indexing – Searching whole genomes

9 March 2020

Sebastian Wild

version 2020-03-09 15:54

Outline

6 Text Indexing

6.1 Motivation

Text indexing

- *Text indexing* (also: *offline text search*):
  - case of string matching: find $P[0..m-1]$ in $T[0..n-1]$
  - but with *fixed* text $\rightsquigarrow$ preprocess T (instead of P)
  - $\rightsquigarrow$ expect many queries P, answer them without looking at all of T
  - $\rightsquigarrow$ essentially a data structuring problem: "building an *index* of T"

Latin: "one who points out"

- application areas
  - web search engines
  - online dictionaries
  - online encyclopedia
  - DNA/RNA data bases
  - . . . searching in any collection of text documents (that grows only moderately)

Inverted indices

same as "indexes"

- original indices in books: list of (key) words $\mapsto$ page numbers where they occur
- assumption: searches are only for **whole** (key) **words**

$\rightsquigarrow$ often reasonable for natural language text

Inverted indices

same as "indexes"

- original indices in books: list of (key) words $\mapsto$ page numbers where they occur
- assumption: searches are only for **whole** (key) **words**

$\rightsquigarrow$ often reasonable for natural language text

Inverted index:

- collect all words in T
  - can be as simple as splitting T at whitespace
  - actual implementations typically support *stemming* of words
    `goes` $\rightarrow$ `go`, `cats` $\rightarrow$ `cat`
- store mapping from words to a list of occurrences $\rightsquigarrow$ *how?* — not here — BST

go $\mapsto$ {5, 10, 20}

cat $\mapsto$ {4, 21}

Clicker Question

Do you know what a *trie* is?

- **A** A what? No!
- **B** I have heard the term, but don't quite remember.
- **C** I remember hearing about it in a module.
- **D** Sure.

`pingo.upb.de/622222`

Tries

{aa, a}

- efficient dictionary data structure for strings
- name from re**trie**val, but pronounced "try"
- tree based on symbol comparisons
- **Assumption:** stored strings are *prefix-free* (no string is a prefix of another)
  - strings of same length ✓
  - strings have "end-of-string" marker $ ✓ (some character $\notin \Sigma$)
- **Example**:
  {aa$, aaab$, abaab$, abb$, abbab$, bba$, bbab$, bbb$}

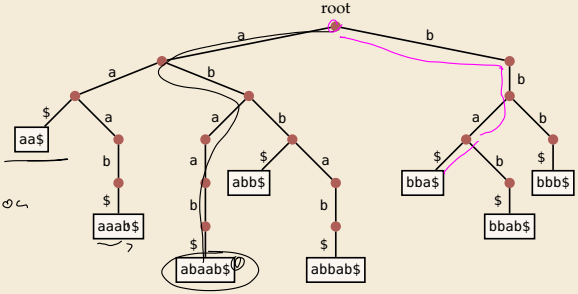

- construction: top-down
  independent of order of insertion
- query: (get) ex: bba$

Construction {aaa$, babb$, aa$, abbba$}

NOT our standard tries, but version with compacted paths to leaves (see next page)

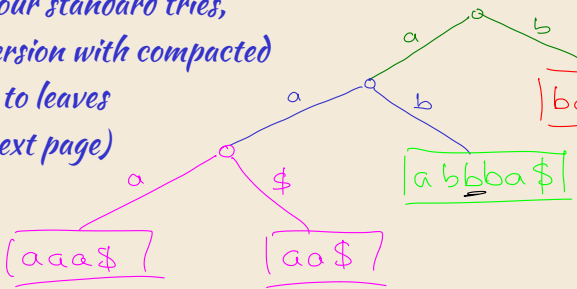

add: aba$

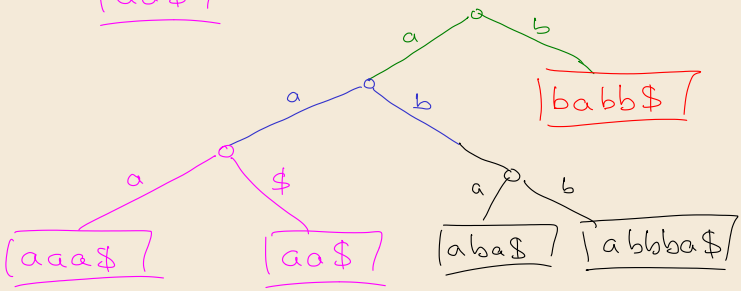

Trie construction (correct version)

{aaa$, bab$, aa$, abbba$}

standard trie

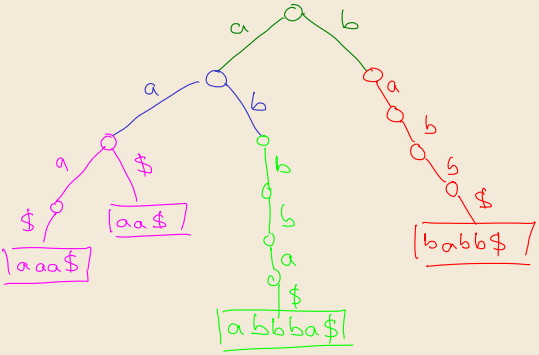

trie with compacted paths to leaves

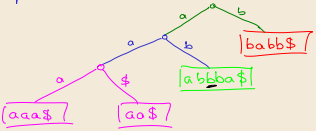

Clicker Question

Suppose we have a trie that stores n strings over $\Sigma = \{\texttt{A}, \ldots, \texttt{Z}\}$.
Each stored string consists of m characters.
We now search for a query string Q with $|Q| = q$.
How many **nodes** in the trie are **visited** during this **query**?

A $\Theta(\log n)$

B $\Theta(\log(nm))$

C $\Theta(m \cdot \log n)$

D $\Theta(m + \log n)$

E $\Theta(m)$

F $\Theta(\log m)$

G $\Theta(q)$

H $\Theta(\log q)$

I $\Theta(q \cdot \log n)$

J $\Theta(q + \log n)$

pingo.upb.de/622222

Clicker Question

Suppose we have a trie that stores n strings over $\Sigma = \{\texttt{A}, \ldots, \texttt{Z}\}$.
Each stored string consists of m characters.
We now search for a query string Q with $|Q| = q$.
How many **nodes** in the trie are **visited** during this **query**?

- A ~~$\Theta(\log n)$~~
- B ~~$\Theta(\log(nm))$~~
- C ~~$\Theta(m \cdot \log n)$~~
- D ~~$\Theta(m + \log n)$~~
- E ~~$\Theta(m)$~~
- F ~~$\Theta(\log m)$~~
- G $\Theta(q)$ ✓
- H ~~$\Theta(\log q)$~~
- I ~~$\Theta(q \cdot \log n)$~~
- J ~~$\Theta(q + \log n)$~~

pingo.upb.de/622222

Clicker Question

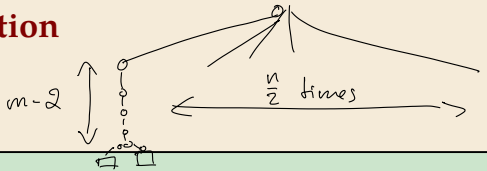

$S_1 = \$_1 aaaaaaa\, b$
$S_2 = \$_1 aaaaaaa$
$S_3 = \$_2$
$S_4 = \$_2$

Suppose we have a trie that stores n strings over $\Sigma = \{\texttt{A}, \ldots, \texttt{Z}\}$.
Each stored string consists of m characters.
How many **nodes** does the trie have **in total** *in the worst case*?

A $\Theta(n)$

B $\Theta(n + m)$

C $\Theta(n \cdot m)$

D $\Theta(n \log m)$

E $\Theta(m)$

F $\Theta(m \log n)$

pingo.upb.de/622222

Clicker Question

Suppose we have a trie that stores n strings over $\Sigma = \{\texttt{A}, \ldots, \texttt{Z}\}$.
Each stored string consists of m characters.
How many **nodes** does the trie have **in total** *in the worst case*?

- **A** ~~$\Theta(n)$~~
- **B** ~~$\Theta(n + m)$~~
- **C** $\Theta(n \cdot m)$ ✓
- **D** ~~$\Theta(n \log m)$~~
- **E** ~~$\Theta(m)$~~
- **F** ~~$\Theta(m \log n)$~~

pingo.upb.de/622222

Compact tries

=1 child

- compress paths of unary nodes into single edge
- nodes store index of next character

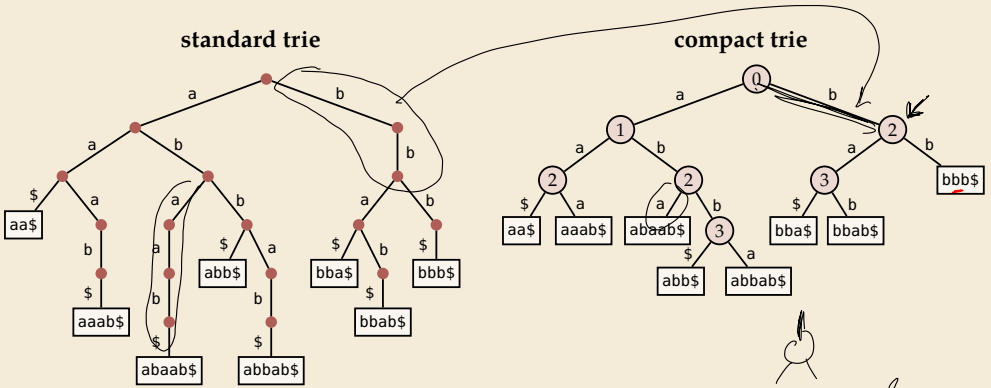

- searching slightly trickier, but same time complexity as in trie
- all nodes ≥ 2 children $\rightsquigarrow$ #nodes $\leq$ #leaves = #strings $\rightsquigarrow$ linear space