

6 Text Indexing – Searching whole genomes

16 March 2021

Sebastian Wild

Outline

6 Text Indexing

- 6.1 Motivation
- 6.2 Suffix Trees
- 6.3 Applications
- 6.4 Longest Common Extensions
- 6.5 Suffix Arrays
- 6.6 Linear-Time Suffix Sorting
- 6.7 The LCP Array

6.1 Motivation

Text indexing

- ▶ *Text indexing* (also: *offline text search*):

- ▶ case of string matching: find $P[0..m)$ in $T[0..n)$

- ▶ but with *fixed* text \rightsquigarrow preprocess T (instead of P)

- \rightsquigarrow expect many queries P , answer them without looking at all of T

- \rightsquigarrow essentially a data structuring problem: “building an *index* of T ”

Latin: “one who points out”

- ▶ application areas

- ▶ web search engines

- ▶ online dictionaries

- ▶ online encyclopedia

- ▶ DNA/RNA data bases

- ▶ ... searching in any collection of text documents (that grows only moderately)

Inverted indices

same as "indexes"

- ▶ original indices in books: list of (key) words \mapsto page numbers where they occur
- ▶ assumption: searches are only for **whole** (key) **words**
- ~> often reasonable for natural language text

Inverted indices

same as "indexes"

- ▶ original indices in books: list of (key) words \mapsto page numbers where they occur
 - ▶ assumption: searches are only for **whole** (key) **words**
- \rightsquigarrow often reasonable for natural language text

Inverted index:

- ▶ collect all words in T
 - ▶ can be as simple as splitting T at whitespace
- (▶ actual implementations typically support *stemming* of words)
goes \rightarrow go, cats \rightarrow cat
- ▶ store mapping from words to a list of occurrences \rightsquigarrow how?

like a dictionary!

keys = words

values = list of occurrences,

BST
but $O(\log n)$
time

Clicker Question



Do you know what a *trie* is?

- ☐ **A** A what? No!
- ☐ **B** I have heard the term, but don't quite remember.
- ☐ **C** I remember hearing about it in a module.
- ☐ **D** Sure.

sli.do/comp526


Click on "Polls" tab

Tries

- ▶ efficient dictionary data structure for strings
- ▶ name from retrieval, but pronounced “try” \approx free
- ▶ tree based on symbol comparisons
- ▶ **Assumption:** stored strings are prefix-free (no string is a prefix of another)
 - ▶ strings of same length ✓ some character $\notin \Sigma$
 - ▶ strings have “end-of-string” marker \$ ✓

▶ Example:

$\Sigma = \{a, b\}$
{aa\$, aaab\$, abaab\$, abb\$,
abbab\$, bba\$, bbab\$, bbb\$, \$}



Clicker Question

Suppose we have a trie that stores n strings over $\Sigma = \{A, \dots, Z\}$. Each stored string consists of m characters.

We now search for a query string Q with $|Q| = q$. ($q \leq m$)

How many **nodes** in the trie are **visited** during this **query**?



A $\Theta(\log n)$

F $\Theta(\log m)$

B $\Theta(\log(nm))$

G $\Theta(q)$

C $\Theta(m \cdot \log n)$

H $\Theta(\log q)$

D $\Theta(m + \log n)$

I $\Theta(q \cdot \log n)$

E $\Theta(m)$

J $\Theta(q + \log n)$

sli.do/comp526

Click on "Polls" tab

Clicker Question

Suppose we have a trie that stores n strings over $\Sigma = \{A, \dots, Z\}$. Each stored string consists of m characters.

We now search for a query string Q with $|Q| = q$.

How many **nodes** in the trie are **visited** during this **query**?



A ~~$\Theta(\log n)$~~

F ~~$\Theta(\log m)$~~

B ~~$\Theta(\log(nm))$~~

G $\Theta(q)$ ✓

C ~~$\Theta(m \log n)$~~

H ~~$\Theta(\log q)$~~

D ~~$\Theta(m + \log n)$~~

I ~~$\Theta(q \log n)$~~

E ~~$\Theta(m)$~~

J ~~$\Theta(q + \log n)$~~

sli.do/comp526

Click on "Polls" tab

Clicker Question



Suppose we have a trie that stores n strings over $\Sigma = \{A, \dots, Z\}$. Each stored string consists of m characters.

How many **nodes** does the trie have **in total** in the worst case?

A $\Theta(n)$

D $\Theta(n \log m)$

B $\Theta(n + m)$

E $\Theta(m)$

C $\Theta(n \cdot m)$

F $\Theta(m \log n)$

sli.do/comp526

Click on "Polls" tab

Clicker Question



Suppose we have a trie that stores n strings over $\Sigma = \{A, \dots, Z\}$. Each stored string consists of m characters.

How many **nodes** does the trie have **in total** *in the worst case*?

A ~~$\Theta(n)$~~

D ~~$\Theta(n \log m)$~~

B ~~$\Theta(n + m)$~~

E ~~$\Theta(m)$~~

C $\Theta(n \cdot m)$ ✓

F ~~$\Theta(m \log n)$~~

sli.do/comp526


Click on "Polls" tab

Compact tries


- ▶ compress paths of unary nodes into single edge
- ▶ nodes store index of next character

=1 child

standard trie



compact trie



↪ searching slightly trickier, but same time complexity as in trie

not $O(n \cdot m)$

- ▶ all nodes ≥ 2 children ↪ $\#nodes \leq \#leaves = \#strings$ ↪ linear space $O(n)$

Tries as inverted index



simple



fast lookup



cannot handle more general queries:

- ▶ search part of a word
- ▶ search phrase (sequence of words)

Tries as inverted index



simple



fast lookup



cannot handle more general queries:

- ▶ search part of a word
- ▶ search phrase (sequence of words)



what if the 'text' does not even have words to begin with?!

- ▶ biological sequences

```
ACAAGATGCCATTGTCCCCGGCCTCTGCTGCTGCTGCTCTCCGGGGCCACGGCCACCGCTGCCCTGCCCTGGAGGGTGGCCCCACCGGC  
CGAGACAGCGAGCATATGCAGGAAGCGGCAGGAATAAGGAAAAGCAGCCTCCTGACTTTCTCGCTTGGTGGTTTGAGTGGACCTCCAGGC  
CAGTGCCGGGCCCCCTCATAGGAGAGGAAGCTCGGGAGGTTGGCCAGGCGGCAGGAAGGCGCACCCCCCAGCAATCCGCGCGCCGGGACAGAA  
TGCCCTGCAGGAACCTTCTTCTGGAAGACCTTCTCCTCTGCAAATAAAACCTCACCCATGAATGCTCACGCAAGTTTAATTACAGACCTGAA
```

- ▶ binary streams

```
00000010101001111010111000001111100011111011111001101101000011100010011011110000010001101010  
011011000011010110100000001000000011101011000001000011110101110110010001100101101110111111  
110001010001011001010000001110101010011000000001101100001100111110000101 0101011101111000011  
10101110010010101010100000111110100110000001111001101010000000100100100000101100011000110111
```



need new ideas

6.2 Suffix Trees

Suffix trees – A ‘magic’ data structure

Appetizer: Longest common substring problem

- ▶ Given: strings S_1, \dots, S_k **Example:** $S_1 = \text{superiorcalifornialives}$, $S_2 = \text{sealiver}$
- ▶ Goal: find the longest substring that occurs in all k strings

Suffix trees – A ‘magic’ data structure

Appetizer: Longest common substring problem

► Given: strings S_1, \dots, S_k **Example:** $S_1 = \text{superiorcalifornialives}$, $S_2 = \text{sealiver}$

► Goal: find the longest substring that occurs in all k strings \rightsquigarrow alive



Can we do this in time $O(|S_1| + \dots + |S_k|)$? How??

Suffix trees – A ‘magic’ data structure

Appetizer: Longest common substring problem

- ▶ Given: strings S_1, \dots, S_k **Example:** $S_1 = \text{superiorcalifornialives}$, $S_2 = \text{sealiver}$
- ▶ Goal: find the longest substring that occurs in all k strings \rightsquigarrow alive



Can we do this in time $O(|S_1| + \dots + |S_k|)$? How??

Enter: *suffix trees*

- ▶ versatile data structure for index with full-text search
- ▶ linear time (for construction) and linear space
- ▶ allows efficient solutions for many advanced string problems




“Although the longest common substring problem looks trivial now, given our knowledge of suffix trees, it is very interesting to note that in 1970 Don Knuth conjectured that a linear-time algorithm for this problem would be impossible.”

[Gusfield: Algorithms on Strings, Trees, and Sequences (1997)]

Suffix trees – Definition

- suffix tree \mathcal{T} for text $T = T[0..n)$ = compact trie of all suffixes of T (set $T[n] := \$$)




1
anana \$

2
nana\$

3
ana \$

4 na \$

5-
a. 9

6
\$

Suffix trees – Definition


- suffix tree \mathcal{T} for text $T = T[0..n)$ = compact trie of all suffixes of $T\$$ (set $T[n] := \$$)

Example:

$T = \text{bananaban\$}$

suffixes: { $\text{bananaban\$}$, $\text{ananaban\$}$, $\text{nanaban\$}$,
 $\text{anaban\$}$, $\text{naban\$}$, $\text{aban\$}$, $\text{ban\$}$, $\text{an\$}$, $\text{n\$}$, $\text{\$}$ }

	0	1	2	3	4	5	6	7	8	9
$T =$	b	a	n	a	n	a	b	a	n	\$



Suffix trees – Definition

- ▶ suffix tree \mathcal{T} for text $T = T[0..n)$ = compact trie of all suffixes of $T\$$ (set $T[n] := \$$)
- ▶ except: in leaves, store *start index* (instead of actual string)


Example:

$T = \text{bananaban\$}$

suffixes: { $\text{bananaban\$}$, $\text{ananaban\$}$, $\text{nanaban\$}$,
 $\text{anaban\$}$, $\text{naban\$}$, $\text{aban\$}$, $\text{ban\$}$, $\text{an\$}$, $\text{n\$}$, $\text{\$}$ }

0	1	2	3	4	5	6	7	8	9
b	a	n	a	n	a	b	a	n	\$

$T =$



Suffix trees – Definition

- ▶ suffix tree \mathcal{T} for text $T = T[0..n)$ = compact trie of all suffixes of $T\$$ (set $T[n] := \$$)
- ▶ except: in leaves, store *start index* (instead of actual string)


Example:

$T = \text{bananaban\$}$

suffixes: { $\text{bananaban\$}$, $\text{ananaban\$}$, $\text{nanaban\$}$,
 $\text{anaban\$}$, $\text{naban\$}$, $\text{aban\$}$, $\text{ban\$}$, $\text{an\$}$, $\text{n\$}$, $\text{\$}$ }

0	1	2	3	4	5	6	7	8	9
b	a	n	a	n	a	b	a	n	\$

- ▶ also: edge labels like in compact trie
- ▶ (more readable form on slides to explain algorithms)



Suffix trees – Construction

- ▶ $T[0..n)$ has $n + 1$ suffixes (starting at character $i \in [0..n)$)
- ▶ We can build the suffix tree by inserting each suffix of T into a compressed trie.
But that takes time $\Theta(n^2)$. \rightsquigarrow not interesting!

Suffix trees – Construction

- ▶ $T[0..n)$ has $n + 1$ suffixes (starting at character $i \in [0..n)$)
- ▶ We can build the suffix tree by inserting each suffix of T into a compressed trie. But that takes time $\Theta(n^2)$. \rightsquigarrow not interesting!



same order of growth as reading the text!

Amazing result: Can construct the suffix tree of T in $\Theta(n)$ time!

- ▶ algorithms are a bit tricky to understand
- ▶ but were a theoretical breakthrough
- ▶ and they are efficient in practice (and heavily used)!


\rightsquigarrow for now, take linear-time construction for granted. What can we do with them?

6.3 Applications


Applications of suffix trees

- In this section, always assume suffix tree \mathcal{T} for T given.

Recall: \mathcal{T} stored like this:




but think about this:



- Moreover: assume internal nodes store pointer to leftmost leaf in subtree.
- Notation: $T_i = T[i..n]$ (including \$)

Clicker Question



What does T 's suffix tree (on the left) tell you about the question whether T contains the pattern $P = \text{ana}$?


Check all that apply to this example.

- ☐ **A** Nothing.
- ☐ **B** P occurs in T .
- ☐ **C** P does not occur in T .
- ☐ **D** P occurs once in T .
- ☐ **E** P occurs twice in T .
- ☐ **F** P starts at index 0.
- ☐ **G** P starts at index 1.
- ☐ **H** P starts at index 2.
- ☐ **I** P starts at index 3.
- ☐ **J** P starts at index 4.
- ☐ **K** P starts at index 7.

sli.do/comp526

Click on "Polls" tab

Clicker Question



What does T 's suffix tree (on the left) tell you about the question whether T contains the pattern $P = \text{ana}$?

Check all that apply to this example.


- ☒ **A** ~~Nothing.~~
- ☒ **B** P occurs in T . ✓
- ☒ **C** ~~P does not occur in T .~~
- ☒ **D** ~~P occurs once in T .~~
- ☒ **E** P occurs twice in T . ✓
- ☒ **F** ~~P starts at index 0.~~
- ☒ **G** P starts at index 1. ✓
- ☒ **H** ~~P starts at index 2.~~
- ☒ **I** P starts at index 3. ✓
- ☒ **J** ~~P starts at index 4.~~
- ☒ **K** ~~P starts at index 7.~~

sli.do/comp526

Click on "Polls" tab

Application 1: Text Indexing / String Matching

- ▶ P occurs in $T \iff \underline{P}$ is a prefix of a suffix of T
- ▶ we have all suffixes in \mathcal{T} !



Application 1: Text Indexing / String Matching

- ▶ P occurs in $T \iff P$ is a prefix of a suffix of T
- ▶ we have all suffixes in \mathcal{T} !

↪ (try to) follow path with label P , until

1. we get stuck

at internal node (no node with next character of P) n^b
or inside edge (mismatch of next characters) $b a_\alpha$

$\rightsquigarrow P$ does not occur in T

2. we run out of pattern

we run out of pattern
reach end of P at internal node v or inside edge towards v

→ P occurs at all leaves in subtree of v


- ### 3. we run out of tree

reach a leaf ℓ with part of P left \rightsquigarrow compare P to ℓ .



This cannot happen when testing edge labels since $\$ \notin \Sigma$, but needs check(s) in compact trie implementation!

- Finding first match (or NO MATCH) takes $O(|P|)$ time!



not possible/relevant
text indexing

Application 1: Text Indexing / String Matching

► P occurs in $T \iff P$ is a prefix of a suffix of T

► we have all suffixes in \mathcal{T} !

↪ (try to) follow path with label P , until

1. **we get stuck**

at internal node (no node with next character of P)
or inside edge (mismatch of next characters)

↪ P does not occur in T

2. **we run out of pattern**

reach end of P at internal node v or inside edge towards v

↪ P occurs at all leaves in subtree of v


3. **we run out of tree**

reach a leaf ℓ with part of P left ↪ compare P to ℓ .



This cannot happen when testing edge labels since $\$ \notin \Sigma$,
but needs check(s) in compact trie implementation!

► Finding first match (or NO_MATCH) takes $O(|P|)$ time!



Examples:

► $P = \text{ann}$

► $P = \text{ana}$

► $P = \text{briar}$


Application 2: Longest repeated substring

► **Goal:** Find longest substring $T[i..i + \ell)$ that occurs also at $j \neq i$: $T[j..j + \ell) = T[i..i + \ell)$.



How can we efficiently check *all possible substrings*?

e.g. for compression \rightsquigarrow Unit 7



Application 2: Longest repeated substring

- **Goal:** Find longest substring $T[i..i + \ell)$ that occurs also at $j \neq i$: $T[j..j + \ell) = T[i..i + \ell)$.




How can we efficiently check *all possible substrings*?

e.g. for compression \rightsquigarrow Unit 7




Repeated substrings = shared paths in *suffix tree*



- $T_5 = \underline{a}ban\$$ and $T_7 = \underline{a}n\$$ have *longest common prefix* 'a'

$\rightsquigarrow \exists$ internal node with path label 'a'

here single edge, can be longer path



Application 2: Longest repeated substring

- **Goal:** Find longest substring $T[i..i + \ell)$ that occurs also at $j \neq i$: $T[j..j + \ell) = T[i..i + \ell)$.



How can we efficiently check *all possible substrings*?



Repeated substrings = shared paths in *suffix tree*



e.g. for compression \rightsquigarrow Unit 7

- $T_5 = \text{aban\$}$ and $T_7 = \text{an\$}$ have *longest common prefix* 'a'

$\rightsquigarrow \exists$ internal node with path label 'a'

here single edge, can be longer path


\rightsquigarrow longest repeated substring = longest common prefix (LCP) of two suffixes

actually: adjacent leaves

- Algorithm:

1. Compute string depth (=length of path label) of nodes
2. Find internal nodes with maximal string depth

- Both can be done in depth-first traversal $\rightsquigarrow \underline{\Theta(n)}$ time



$T = \text{bananaban\$}$

Generalized suffix trees

- ▶ longest *repeated* substring (of one string) feels very similar to longest *common* substring of several strings $T^{(1)}, \dots, T^{(k)}$ with $T^{(j)} \in \Sigma^{n_j}$
- ▶ can we solve that in the same way?
- ▶ could build the suffix tree for each $T^{(j)}$... but doesn't seem to help

Generalized suffix trees

- ▶ longest *repeated* substring (of one string) feels very similar to longest *common* substring of several strings $T^{(1)}, \dots, T^{(k)}$ with $T^{(j)} \in \Sigma^{n_j}$
 - ▶ can we solve that in the same way?
 - ▶ could build the suffix tree for each $T^{(j)}$... but doesn't seem to help
- \rightsquigarrow need a *single/joint* suffix tree for *several* texts

Generalized suffix trees

- ▶ longest *repeated* substring (of one string) feels very similar to longest *common* substring of several strings $T^{(1)}, \dots, T^{(k)}$ with $T^{(j)} \in \Sigma^{n_j}$
 - ▶ can we solve that in the same way?
 - ▶ could build the suffix tree for each $T^{(j)} \dots$ but doesn't seem to help
- \rightsquigarrow need a *single/joint* suffix tree for *several* texts

Enter: *generalized suffix tree*

- ▶ Define $T := T^{(1)}\$1 T^{(2)}\$2 \dots T^{(k)}\$k$ for k new end-of-word symbols
- ▶ Construct suffix tree \mathcal{T} for T

\rightsquigarrow $\$j$ -edges always leads to leaves $\rightsquigarrow \exists$ leaf (j, i) for each suffix $T_i^{(j)} = T^{(j)}[i..n_j]$



Clicker Question



What is the longest common substring of the strings
bcabcac, aabca and bcaa?

sli.do/comp526

Click on "Polls" tab

Application 3: Longest common substring

► With that new idea, we can find longest common superstrings:

1. Compute generalized suffix tree \mathcal{T} .
2. Store with each node the *subset of strings* that contain its path label:
 - 2.1. Traverse \mathcal{T} bottom-up.
 - 2.2. For a leaf (j, i) , the subset is $\{j\}$.
 - 2.3. For an internal node, the subset is the union of its children.
3. In top-down traversal, compute string depths of nodes. (as above)
4. Report deepest node (by string depth) whose subset is $\{1, \dots, k\}$.

② stores set of j so that there is a leaf (j, i) in the subtree


► Each step takes time $\Theta(n)$ for $n = n_1 + \dots + n_k$ the total length of all texts.

“Although the longest common substring problem looks trivial now, given our knowledge of suffix trees, it is very interesting to note that in 1970 Don Knuth conjectured that a linear-time algorithm for this problem would be impossible.”

[Gusfield: Algorithms on Strings, Trees, and Sequences (1997)]

Longest common substring – Example

$T^{(1)} = \text{bcabcac}$, $T^{(2)} = \text{aabca}$, $T^{(3)} = \text{bcaa}$



6.4 Longest Common Extensions

Application 4: Longest Common Extensions

- We implicitly used a special case of a more general, versatile idea:

The *longest common extension (LCE)* data structure:

- **Given:** String $T[0..n)$
- **Goal:** Answer LCE queries, i. e.,
given positions i, j in T ,
how far can we read the same text from there?
formally: $\text{LCE}(i, j) = \max\{\ell : T[i..i + \ell) = T[j..j + \ell)\}$




Application 4: Longest Common Extensions


- We implicitly used a special case of a more general, versatile idea:

The *longest common extension (LCE)* data structure:

- **Given:** String $T[0..n)$
- **Goal:** Answer LCE queries, i. e.,
 given positions i, j in T ,
 how far can we read the same text from there?
 formally: $\text{LCE}(i, j) = \max\{\ell : T[i..i + \ell) = T[j..j + \ell)\}$



gban\$
an\$



$T = \text{bananabans}$



↪ use suffix tree of T !

- $\text{LCE}(i, j)$ = longest common prefix of i th and j th suffix
 In \mathcal{T} : $\text{LCE}(i, j) = \text{LCP}(T_i, T_j) \rightsquigarrow$ same thing, different name!
 = string depth of *lowest common ancestor (LCA)* of leaves i and j
- in short: $\text{LCE}(i, j) = \text{LCP}(T_i, T_j) = \text{stringDepth}(\text{LCA}(i, j))$



Efficient LCA

How to find lowest common ancestors?

- ▶ Could walk up the tree to find LCA $\rightsquigarrow \Theta(n)$ worst case 
- ▶ Could store all LCAs in big table $\rightsquigarrow \Theta(n^2)$ space and preprocessing 

Efficient LCA

How to find lowest common ancestors?

- ▶ Could walk up the tree to find LCA $\rightsquigarrow \Theta(n)$ worst case 🙄
- ▶ Could store all LCAs in big table $\rightsquigarrow \Theta(n^2)$ space and preprocessing 🙄



Amazing result: Can compute data structure in $\Theta(n)$ time and space that finds any LCA is **constant(!) time**.

- ▶ a bit tricky to understand
- ▶ but a theoretical breakthrough
- ▶ and useful in practice

and suffix tree construction inside ...



\rightsquigarrow for now, use $O(1)$ LCA as black box.

\rightsquigarrow After linear preprocessing (time & space), we can find LCEs in $O(1)$ time.

Application 5: Approximate matching

k -mismatch matching:

- ▶ **Input:** text $T[0..n)$, pattern $P[0..m)$, $k \in [0..m)$ # mismatched characters
- ▶ **Output:** "Hamming distance $\leq k$ "
 - ▶ smallest i so that $T[i..i+m)$ and P differ in at most k characters
 - ▶ or NO_MATCH if there is no such i

↪ searching with typos

- ▶ Assume longest common extensions in $T_1 P_2$ can be found in $O(1)$
 - ↪ generalized suffix tree \mathcal{T} has been built
 - ↪ string depths of all internal nodes have been computed
 - ↪ constant-time LCA data structure for \mathcal{T} has been built

Clicker Question



What is the Hamming distance between heart and beard?

2


sli.do/comp526

Click on "Polls" tab

Kangaroo Algorithm for approximate matching



```
1 procedure kMismatch( $T[0..n - 1]$ ,  $P[0..m - 1]$ )  
2   // build LCE data structure  
3   for  $i := 0, \dots, n - m - 1$  do  
4     mismatches := 0;  $t := i$ ;  $p := 0$   
5     while mismatches  $\leq k \wedge p < m$  do  
6        $\ell := \text{LCE}(t, p)$  // jump over matching part  
7        $t := t + \ell + 1$ ;  $p := p + \ell + 1$   
8       mismatches := mismatches + 1  
9     if  $p == m$  then  
10      return  $i$ 
```



► **Analysis:** $\Theta(n + m)$ preprocessing + $O(n \cdot k)$ matching

↪ very efficient for small k

► State of the art

- $O\left(n \frac{k^2 \log k}{m}\right)$ possible with complicated algorithms
- extensions for edit distance $\leq k$ possible

Application 6: Matching with wildcards

- ▶ Allow a wildcard character in pattern

stands for arbitrary (single) character

unit*	P
in_unit5_we_will	T

- ▶ similar algorithm as for k -mismatch $\rightsquigarrow O(n \cdot k + m)$ when P has k wildcards

* * *

Many more applications, in particular for problems on biological sequences

20+ described in Gusfield, *Algorithms on strings, trees, and sequences* (1999)