

ALGORITHMS \$ EFFICIENT
CIENTALGORITHMS \$ EFFI
EFFICIENTALGORITHMS \$
FFICIENTALGORITHMS \$
FICIENTALGORITHMS \$
GORITHMS \$ EFFICIENTAL
HMS \$ EFFICIENTALGORIT
IENTALGORITHMS \$ EFF
ITHMS \$ EFFICIENTALGOR
LGORITHMS \$ EFFICIENTALG
MS \$ EFFICIENTALGORITH
NTALGORITHMS \$ EFFICIENTALG
ORITHMS \$ EFFICIENTALGO
RITHMS \$ EFFICIENTALGO
S \$ EFFICIENTALGORITHM
TALGORITHMS \$ EFFICIENT
THMS \$ EFFICIENTALGORI

7

Parallel Algorithms

17 November 2023

Sebastian Wild

Learning Outcomes

1. Know and apply *parallelization strategies* for embarrassingly parallel problems.
2. Identify *limits of parallel speedups*.
3. Understand and use the *parallel random-access-machine* model in its different variants.
4. Be able to *analyze* and compare simple shared-memory parallel algorithms by determining *parallel time and work*.
5. Understand efficient parallel *prefix sum* algorithms.
6. Be able to devise high-level description of *parallel quicksort and mergesort* methods.

Unit 7: *Parallel Algorithms*



7 Parallel Algorithms

- 7.1 Parallel Computation
- 7.2 Parallel String Matching
- 7.3 Parallel Primitives
- 7.4 Parallel Sorting

7.1 Parallel Computation

Clicker Question



Have you ever written a concurrent program (explicit threads, job pools library, or using a framework for distributed computing)?

- ☐ **A** Yes
- ☐ **B** No
- ☐ **C** Concur... what?



→ *sli.do/comp526*

Types of parallel computation

£££ can't buy you more time . . . but more computers!

↪ Challenge: Algorithms for *parallel* computation.

Types of parallel computation


£££ can't buy you more time ... but more computers!

↪ Challenge: Algorithms for *parallel* computation.

There are two main forms of parallelism:


1. shared-memory parallel computer ← focus of today

- ▶ *p processing elements* (PEs, processors) working in parallel
- ▶ **single** big memory, **accessible from every PE**
- ▶ communication via shared memory
- ▶ think: a big server, 128 CPU cores, terabyte of main memory




2. distributed computing

- ▶ *p* PEs working in parallel
- ▶ each PE has **private** memory
- ▶ communication by sending **messages** via a network
- ▶ think: a cluster of individual machines



PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 1)
- ▶ the p PEs are identified by ids $0, \dots, p-1$
 - ▶ like w (the word size), p is a parameter of the model that can grow with n
 - ▶ $p = \Theta(n)$ is not unusual maaany processors!
- ▶ the PEs all **independently** run the same RAM-style program (they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:
in each time step, every PE executes one instruction



PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 1)
- ▶ the p PEs are identified by ids $0, \dots, p - 1$
 - ▶ like w (the word size), p is a parameter of the model that can grow with n
 - ▶ $p = \Theta(n)$ is not unusual maaany processors!
- ▶ the PEs all **independently** run the same RAM-style program (they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:
in each time step, every PE executes one instruction
- ▶ As for RAM:
 - ▶ assume a basic “operating system”
 - ↪ write algorithms in pseudocode instead of RAM assembly
 - ▶ **NEW:** loops and commands can be run “in parallel” (examples coming up)

PRAM – Conflict management



Problem: What if several PEs simultaneously overwrite a memory cell?

- ▶ **EREW-PRAM** (exclusive read, exclusive write)
any **parallel access** to same memory cell is **forbidden** (crash if happens)
- ▶ **CREW-PRAM** (concurrent read, exclusive write)
parallel **write** access to same memory cell is **forbidden**, *but reading is fine*
- ▶ **CRCW-PRAM** (concurrent read, concurrent write)
concurrent access is allowed,
need a rule for write conflicts:
 - ▶ common CRCW-PRAM:
all concurrent writes to same cell must write **same** value
 - ▶ arbitrary CRCW-PRAM:
some unspecified concurrent write wins ↗
 - ▶ (more exist . . .)
- ▶ no single model is always adequate, but our default is **CREW**

PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish assuming sufficiently many PEs!
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish assuming sufficiently many PEs!
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

Holy grail of PRAM algorithms:

- ▶ minimal time (=span)
- ▶ work (asymptotically) no worse than running time of best sequential algorithm
 \rightsquigarrow “*work-efficient*” algorithm: work in same Θ -class as best sequential