

# 8

# Error-Correcting Codes

*28 April 2021*

Sebastian Wild

## Outline

# 8 Error-Correcting Codes

8.1 Introduction

8.2 Lower Bounds

8.3 Hamming Codes

## 8.1 Introduction

# Noisy Communication

- ▶ most forms of communication are “noisy”
  - ▶ humans: acoustic noise, unclear pronunciation, misunderstanding, foreign languages

# Noisy Communication

- ▶ most forms of communication are “noisy”
  - ▶ humans: acoustic noise, unclear pronunciation, misunderstanding, foreign languages
- ▶ How do humans cope with that?
  - ▶ slow down and/or speak up
  - ▶ ask to repeat if necessary



# Noisy Communication

- ▶ most forms of communication are “noisy”
  - ▶ humans: acoustic noise, unclear pronunciation, misunderstanding, foreign languages
- ▶ How do humans cope with that?
  - ▶ slow down and/or speak up
  - ▶ ask to repeat if necessary
- ▶ But how is it possible (for us) to decode a message in the presence of noise & errors?



*Because it seems that natural language has a lot of **redundancy** built into it!*

# Noisy Communication

- ▶ most forms of communication are “noisy”
  - ▶ humans: acoustic noise, unclear pronunciation, misunderstanding, foreign languages

- ▶ How do humans cope with that?

- ▶ slow down and/or speak up
  - ▶ ask to repeat if necessary



- ▶ But how is it possible (for us) to decode a message in the presence of noise & errors?

*Bcaesue it semes taht ntaurul lanaguge has a lots fo **redundancy** bilt itno it!*

→ We can

1. **detect errors**      “This sentence has aao pi dgsdho gioasghds.”
2. **correct (some) errors**      “Tiny errs ar corrected automaticly.”  
(sometimes too eagerly as in the Chinese Whispers / Telephone)



# Noisy Channels

- ▶ computers: copper cables & electromagnetic interference
  - ▶ transmit a binary string
  - ▶ but occasionally bits can “flip”
- ⇒ want a robust code





# Noisy Channels

- ▶ computers: copper cables & electromagnetic interference
  - ▶ transmit a binary string
  - ▶ but occasionally bits can “flip”
- ~> want a robust code



- ▶ We can aim at
  1. **error detection** ~> can request a re-transmit
  2. **error correction** ~> avoid re-transmit for common types of errors

# Noisy Channels

- ▶ computers: copper cables & electromagnetic interference
  - ▶ transmit a binary string  $= \text{message}$
  - ▶ but occasionally bits can “flip”
- ~> want a robust code



- ▶ We can aim at
    1. **error detection** ~> can request a re-transmit
    2. **error correction** ~> avoid re-transmit for common types of errors
  - ▶ This will require *redundancy*: sending *more* bits than plain message
    - ~> **goal**: robust code with lowest redundancy
- that's the opposite of compression!

## Clicker Question



What do you think, how many extra bits do we need to **detect** a **single bit error** in a message of 100 bits?

*[sli.do/comp526](https://sli.do/comp526)*

Click on “Polls” tab

## Clicker Question



What do you think, how many extra bits do we need to **correct** a **single bit error** in a message of 100 bits?

*[sli.do/comp526](https://sli.do/comp526)*

Click on "Polls" tab

## 8.2 Lower Bounds

# Block codes

## ► model:

- want to send message  $S \in \{0, 1\}^*$  (bitstream) across a (*communication*) channel
  - any bit transmitted through the channel might *flip* ( $0 \rightarrow 1$  resp.  $1 \rightarrow 0$ )  
**no other errors** occur (no bits lost, duplicated, inserted, etc.)
  - instead of  $S$ , we send *encoded bitstream*  $C \in \{0, 1\}^*$   
sender *encodes*  $S$  to  $C$ , receiver *decodes*  $C$  to  $S$  (hopefully)
- ~> what errors can be detected and/or corrected?

# Block codes

## ► model:

- want to send message  $S \in \{0, 1\}^*$  (bitstream) across a (*communication*) channel
- any bit transmitted through the channel might *flip* ( $0 \rightarrow 1$  resp.  $1 \rightarrow 0$ )  
**no other errors** occur (no bits lost, duplicated, inserted, etc.)
- instead of  $S$ , we send *encoded bitstream*  $C \in \{0, 1\}^*$   
sender *encodes*  $S$  to  $C$ , receiver *decodes*  $C$  to  $S$  (hopefully)

↪ what errors can be detected and/or corrected?

## ► all codes discussed here are block codes

- divide  $S$  into *messages*  $m \in \{0, 1\}^k$  of  $k$  bits each ( $k = \text{message length}$ )
- encode each message (separately) as  $C(m) \in \{0, 1\}^n$  ( $n = \text{block length}$ ,  $n \geq k$ )

↪ can analyze everything block-wise

# Block codes

## ► model:

- want to send message  $S \in \{0, 1\}^*$  (bitstream) across a (*communication*) *channel*
- any bit transmitted through the channel might *flip* ( $0 \rightarrow 1$  resp.  $1 \rightarrow 0$ )  
**no other errors** occur (no bits lost, duplicated, inserted, etc.)
- instead of  $S$ , we send *encoded bitstream*  $C \in \{0, 1\}^*$   
sender *encodes*  $S$  to  $C$ , receiver *decodes*  $C$  to  $S$  (hopefully)

~> what errors can be detected and/or corrected?

## ► all codes discussed here are *block codes*

- divide  $S$  into *messages*  $m \in \{0, 1\}^k$  of  $k$  bits each ( $k = \text{message length}$ )
- encode each message (separately) as  $C(m) \in \{0, 1\}^n$  ( $n = \text{block length}$ ,  $n \geq k$ )

~> can analyze everything block-wise

## ► between 0 and $n$ bits might be flipped

invalid code

- how many flipped bits can we definitely **detect**?
- how many flipped bits can we **correct** without retransmit?

i. e. decoding  $m$  still possible



## Code distance

$$m \neq m' \implies C(m) \neq C(m')$$

- ▶ each block code is an *injective* function  $C : \{0, 1\}^k \rightarrow \{0, 1\}^n$

# Code distance

$$m \neq m' \implies C(m) \neq C(m')$$


► each block code is an *injective* function  $C : \{0, 1\}^k \rightarrow \{0, 1\}^n$

► define  $\mathcal{C}$  = set of all codewords =  $C(\{0, 1\}^k) = \{b \in \{0, 1\}^n : \exists m \in \{0, 1\}^k : b = C(m)\}$

$$\rightsquigarrow \mathcal{C} \subseteq \{0, 1\}^n$$

$|\mathcal{C}| = 2^k$  out of  $2^n$   $n$ -bit strings are valid codewords

► decoding = finding closest valid codeword



# Code distance

$$m \neq m' \implies C(m) \neq C(m')$$

- ▶ each block code is an *injective* function  $C : \{0, 1\}^k \rightarrow \{0, 1\}^n$
- ▶ define  $\mathcal{C}$  = set of all codewords =  $C(\{0, 1\}^k)$

$$\leadsto \mathcal{C} \subseteq \{0, 1\}^n$$

$|\mathcal{C}| = 2^k$  out of  $2^n$   $n$ -bit strings are valid codewords

- ▶ decoding = finding closest valid codeword

- ▶ *distance of code:*

$$d = \text{minimal Hamming distance of any two codewords} = \min_{x, y \in \mathcal{C}} d_H(x, y)$$

# Code distance

$$m \neq m' \implies C(m) \neq C(m')$$


► each block code is an *injective* function  $C : \{0, 1\}^k \rightarrow \{0, 1\}^n$

► define  $\mathcal{C}$  = set of all codewords =  $C(\{0, 1\}^k)$

$$\leadsto \mathcal{C} \subseteq \{0, 1\}^n$$

$|\mathcal{C}| = 2^k$  out of  $2^n$   $n$ -bit strings are valid codewords

► decoding = finding closest valid codeword



► *distance of code:*

$d$  = minimal Hamming distance of any two codewords =  $\min_{x, y \in \mathcal{C}} d_H(x, y)$

## Implications for codes

1. Need distance  $d$  to detect all errors flipping up to  $d - 1$  bits.
2. Need distance  $d$  to **correct** all errors flipping up to  $\lfloor \frac{d-1}{2} \rfloor$  bits.

$\Rightarrow$  for detecting 1 bit errors  $\leadsto$  need distance 2  
 for correcting 1 bit errors  $\leadsto$  need distance 3

## Lower Bounds

- ▶ Main advantage of concept of code distance:  
can *prove* lower bounds on block length

# Lower Bounds

- ▶ Main advantage of concept of code distance:  
can *prove* lower bounds on block length


$$k = 100$$

$$d = 2$$

$$n \geq 101$$

- ▶ **Singleton bound:**  $2^k \leq 2^{n-(d-1)} \rightsquigarrow \boxed{n \geq k + d - 1}$

- ▶ *proof sketch:* We have  $2^k$  codeswords with distance  $d$   
after deleting the first  $d - 1$  bits, all are still distinct  
but there are only  $\underline{2^{n-(d-1)}}$  such shorter bitstrings.



ℓ

# Lower Bounds

- ▶ Main advantage of concept of code distance:  
can *prove* lower bounds on block length

- ▶ **Singleton bound:**  $2^k \leq 2^{n-(d-1)} \rightsquigarrow n \geq k + d - 1$

- ▶ *proof sketch:* We have  $2^k$  codeswords with distance  $d$   
after deleting the first  $d - 1$  bits, all are still distinct  
but there are only  $2^{n-(d-1)}$  such shorter bitstrings.



- ▶ **Hamming bound:**  $2^k \leq \frac{2^n}{\sum_{f=0}^{\lfloor (d-1)/2 \rfloor} \binom{n}{f}}$

- ▶ *proof idea:* consider “balls” of bitstrings around codewords  
count bitstrings with Hamming-distance  $\leq t = \lfloor (d - 1)/2 \rfloor$   
correcting  $t$  errors means all these balls are disjoint  
so  $2^k \cdot \text{ball size} \leq 2^n$




$\rightsquigarrow$  We will come back to these.

## 8.3 Hamming Codes




# Parity Bit

- ▶ simplest possible error-detecting code: add a **parity bit**



# Parity Bit

- ▶ simplest possible error-detecting code: add a **parity bit**




~> code distance 2

- ▶ can detect any single-bit error (actually, any odd number of flipped bits)
- ▶ used in many hardware (communication) protocols
  - ▶ PCI buses, serial buses
  - ▶ caches
  - ▶ early forms of main memory

# Parity Bit

- ▶ simplest possible error-detecting code: add a **parity bit**



~> code distance 2

- ▶ can detect any single-bit error (actually, any odd number of flipped bits)
- ▶ used in many hardware (communication) protocols
  - ▶ PCI buses, serial buses
  - ▶ caches
  - ▶ early forms of main memory



very simple and cheap



cannot correct any errors

## Clicker Question



What do you think, how many extra bits do we need to **detect** a **single bit error** in a message of 100 bits?

*[sli.do/comp526](https://sli.do/comp526)*


Click on “Polls” tab

# Error-correcting codes

any downtime is expensive!

- ▶ typical application: heavy-duty server RAM
  - ▶ bits can randomly flip (e.g., by cosmic rays)
  - ▶ individually very unlikely, but in always-on server with lots of RAM, it happens!

<https://blogs.oracle.com/linux/attack-of-the-cosmic-rays-v2>




# Error-correcting codes

any downtime is expensive!

- ▶ typical application: heavy-duty server RAM
  - ▶ bits can randomly flip (e.g., by cosmic rays)
  - ▶ individually very unlikely, but in always-on server with lots of RAM, it happens!

<https://blogs.oracle.com/linux/attack-of-the-cosmic-rays-v2>




Can we **correct** a bit error without knowing where it occurred? How?

# Error-correcting codes

any downtime is expensive!

- ▶ typical application: heavy-duty server RAM
  - ▶ bits can randomly flip (e.g., by cosmic rays)
  - ▶ individually very unlikely, but in always-on server with lots of RAM, it happens!

<https://blogs.oracle.com/linux/attack-of-the-cosmic-rays-v2>



Can we **correct** a bit error without knowing where it occurred? How?


- ▶ Yes! store every bit *three times*!
  - ▶ upon read, do majority vote
  - ▶ if only one bit flipped, the other two (correct) will still win

# Error-correcting codes

any downtime is expensive!

- ▶ typical application: heavy-duty server RAM
  - ▶ bits can randomly flip (e.g., by cosmic rays)
  - ▶ individually very unlikely, but in always-on server with lots of RAM, it happens!

<https://blogs.oracle.com/linux/attack-of-the-cosmic-rays-v2>



Can we **correct** a bit error without knowing where it occurred? How?

- ▶ Yes! store every bit *three times*!
  - ▶ upon read, do majority vote
  - ▶ if only one bit flipped, the other two (correct) will still win

👎 *triples* the cost!



*You want WHAT!?!*




# Error-correcting codes

- ▶ typical application: heavy-duty server RAM
  - ▶ bits can randomly flip (e.g., by cosmic rays)
  - ▶ individually very unlikely, but in always-on server with lots of RAM, it happens!

<https://blogs.oracle.com/linux/attack-of-the-cosmic-rays-v2>

any downtime is expensive!



Can we **correct** a bit error without knowing where it occurred? How?

- ▶ Yes! store every bit *three times*!
  - ▶ upon read, do majority vote
  - ▶ if only one bit flipped, the other two (correct) will still win



*triples* the cost!



*You want WHAT!?!*



instead of 200% (!)

Can do it with 11% extra memory!