



Clever Codes

1 December 2025

Prof. Dr. Sebastian Wild

Learning Outcomes

Unit 8: *Clever Codes*

1. Know the principles and performance characteristics of *arithmetic coding*.
2. Judge the use of arithmetic coding in applications.
3. Understand the context of *error-prone communication*.
4. Understand concepts of *error-detecting codes* and *error-correcting codes*.
5. Know and understand *Hamming codes*, in particular (7,4) Hamming code.
6. Reason about the *suitability of a code* for an application.

Outline

8 Clever Codes

- 8.1 Arithmetic Coding
- 8.2 Practical Arithmetic Coding
- 8.3 Error Correcting Codes
- 8.4 Coding Theory
- 8.5 Hamming Codes

8.1 Arithmetic Coding

Stream Codes

- ▶ **Recall:** (binary) character encoding $E : \Sigma \rightarrow \{0, 1\}^*$
 - ▶ Huffman codes *optimal* for any given character frequencies
 - ↪ encoding all characters with that code *minimizes* compressed size
- ... *if we assume that all characters must be encoded individually by a codeword!*
- ▶ Stream codes instead compress entire **sequence** of characters
 - ▶ RLE and LZW are examples of stream codes ↪ can sometimes do better
- ▶ Two indicative examples
 1. **“Low entropy bits:”** $\Sigma = \{0, 1\}$, highly skewed: $p_0 = 0.99$
 - ↪ entropy $\mathcal{H}(\frac{1}{100}, \frac{99}{100}) \approx 0.08$ bits per character,
Huffman code must use 1 bit per character!
 - ↪ “optimal” Huffman code gives 12-fold space increase over entropy!
 - ▶ Can certainly do better here (RLE!)
 2. **“Trits”:** $\Sigma = \{0, 1, 2\}$, equally likely
 - ↪ entropy $\mathcal{H}(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}) = \lg(3) \approx 1.58$ bits per character,
Huffman code uses average of $\frac{1}{3} \cdot 1 + \frac{2}{3} \cdot 2 = \frac{5}{3} \approx 1.67$
- ▶ Can we do better?

A Decent Hack: Block Codes

- ▶ Huffman on trits wastes ≈ 0.0817 bits per character and over 5 % of space
- ▶ A simple trick can reduce this substantially!
 - ▶ treat 5 trits as one “supercharacter”, e. g., 21101
 - ↪ $3^5 = 243$ possible combinations
 - ↪ encode these using 8 bits (with $2^8 = 256$ possible combinations)
 - ▶ entropy $\lg(3^5) \approx 7.92$ bits, so less than 0.1 % wasted space!
- ▶ We can even use a Huffman code for the supercharacters to handle nonuniformity!
- ▶ For the low-entropy bits, could use 3 bits
 - ↪ probabilities:
 - 000 : 0.97
 - 001, 010, 100 : 0.0098
 - 011, 101, 110 : 0.000099
 - 111: 0.000001
 - ↪ with Huffman code, 1.06 bits per superchar of 3 input bits
 - ↪ almost factor 3 better; can improve with larger blocks!

Block Codes – A Panacea?

- ▶ Using supercharacters works well in our examples.



*Hmmm . . . so why don't we treat the entire source text as one large block?
Wouldn't that be even better!?*

↪ *We can optimally compress any text, without doing anything intelligent!?*



⚡ For general case, need to *communicate* the supercharacter encoding

- ▶ Blocks of k characters need $\Omega(\sigma^k)$ space for code
- ▶ Huffman code has to be part of coded message

↪ Can only sensibly use block codes for small σ and k



There is no such thing as a free lunch . . .

Arithmetic Coding

except in isolated lucky cases



- ▶ Also: Block codes still had $\Theta(n)$ wasted space for sequences of n symbols

▶ *Arithmetic Coding:*

0. Maintain $[\ell, \ell + p) \subseteq [0, 1)$; initially $\ell = 0, p = 1$

1. Zoom into subinterval for each character


2. Output dyadic encoding of final interval

▶ *Step 1: “Zoom” for each character (trit) in $S[0..n)$:*

- ▶ Of the current subinterval $[\ell, \ell + p)$,
take first, second or last third
depending whether $S[i] = 0, 1$, resp. 2:

$$\ell := \ell + S[i] \cdot \frac{1}{3} \cdot p$$

$$p := p \cdot \frac{1}{3}$$



▶ *Step 2: Dyadic encoding*

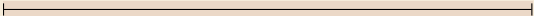





- ▶ Find smallest m so that $\exists x \in \mathbb{N}_0$ with $\left[\frac{x}{2^m}, \frac{x+1}{2^m} \right) \subseteq [\ell, \ell + p)$
- ▶ Output x in binary using m bits.

\rightsquigarrow Encode n trits in $n \lg(3) + 2$ bits(!) without cheating

Arithmetic Coding – Encode Trits Example

► $S[0..n) = 21101$ ($n = 5$)

► **Step 1:** Zoom into subintervals

Iteration	ℓ	p	Interval (rounded)	
0	0	1	[0.00000, 1.00000)	
1	$\frac{2}{3}$	$\frac{1}{3}$	[0.66667, 1.00000)	
2	$\frac{7}{9}$	$\frac{1}{9}$	[0.77778, 0.88889)	
3	$\frac{22}{27}$	$\frac{1}{27}$	[0.81482, 0.85185)	
4	$\frac{66}{81}$	$\frac{1}{81}$	[0.81482, 0.82716)	
5	$\frac{199}{243}$	$\frac{1}{243}$	[0.81893, 0.82305)	

► **Step 2:** Dyadic encoding for interval $[\ell, \ell + p) = \left[\frac{199}{243}, \frac{200}{243}\right)$

► Must have $m \geq \lg(1/p) > 7$

► $m = 8$: smallest $x/2^m \geq \frac{199}{243}$ is $x = 210$, but $[210/256, 211/256) \approx [0.82031, 0.82422) \not\subset [\ell, \ell + p)$

► $m = 9$: smallest $x/2^m \geq \frac{199}{243}$ is $x = 420$ and $[420/512, 421/512) \approx [0.82031, 0.82227) \subset [\ell, \ell + p)$ ✓

↪ Output $x = 420$ in binary with $m = 9$ digits: 110100100

The diagram illustrates the construction of a binary tree for the Huffman coding algorithm. The tree is built from the bottom up, starting with individual characters 'a' and 'b' and their frequencies. The root node is '00000' with a frequency of 1000. The tree structure is shown with nodes and edges, and the final Huffman code is listed on the right.

Tree Structure and Frequencies:

- Root: 00000 (1000)
- Level 1: 00001 (000), 00010 (000)
- Level 2: 00011 (001), 00100 (00)
- Level 3: 00101 (0010), 00110 (001), 00111 (0)
- Level 4: 01000 (0100), 01001 (010), 01010 (0101), 01011 (01), 01100 (0110), 01101 (011), 01110 (011), 01111 (011)
- Level 5: 10000 (1000), 10001 (100), 10010 (1001), 10011 (10), 10100 (1010), 10101 (101), 10110 (1011), 10111 (1)
- Level 6: 11000 (1100), 11001 (110), 11010 (1101), 11011 (11), 11100 (1110), 11101 (111), 11110 (1111), 11111 (1111)

Huffman Code:

- a: 01101
- b: 10011

[illegible]

7

Arithmetic Coding – General framework

- ▶ Note: Arithmetic coder *doesn't care* if probabilities or even σ change all the time!
 - ▶ As long as encoder and decoder know from context what they are!

General stochastic sequence:

Sequence of random variables X_0, X_1, X_2, \dots such that

1. $X_i \in [0..U_i) \cup \{\$ \}$ (We use \$ to signal “end of text”)
2. $\mathbb{P}[X_i = j] = P_{ij}$
3. both U_i and P_{ij} are random variables as they *depend* on X_0, \dots, X_{i-1} , but conditioned on X_0, \dots, X_{i-1} , they are fixed and known:
$$P_{ij} = P_{ij}(X_0, \dots, X_{i-1}) = \mathbb{P}[X_i = j \mid X_0, \dots, X_{i-1}]$$
$$U_i = U_i(X_0, \dots, X_{i-1}) = \max\{j : P_{ij}(X_0, \dots, X_{i-1}) > 0\}$$

- ▶ Can model arbitrary dependencies on previous outcomes
- ▶ Assume here that random process is known by both encoder and decoder (fixed coding) otherwise extra space needed to encode model!

Arithmetic Coding – Encoding

```
1 procedure arithmeticEncode( $X_0, \dots, X_n$ ):  
2   // Assume model  $U_i$  and  $P_{ij}$  are fixed.  
3   // Assume  $X_i \in [0..U_i)$  for  $i < n$  and  $X_n = \$$   
4   // Step 1: Interval zooming  
5    $\ell := 0$ ;  $p := 1$   
6   for  $i := 0, \dots, n - 1$  do  
7        $q := \sum_{j=0}^{X_i-1} P_{ij}$ ;  
8        $\ell := \ell + q \cdot p$ ;  $p := p \cdot P_{i, X_i}$   
9   end for  
10   $q := 1 - P_{n, \$}$  // encode $ as last character  
11   $\ell := \ell + q \cdot p$ ;  $p := p \cdot P_{n, \$}$   
12  // Step 2: Dyadic encoding  
13   $m := \lceil \lg(1/p) \rceil - 1$   
14  do  
15       $m := m + 1$ ;  $x := \lceil \ell \cdot 2^m \rceil$   
16  while  $(x + 1)/2^m > \ell + p$   
17  return  $x$  in binary using  $m$  bits
```

Arithmetic Coding – Decoding

```
1 procedure arithmeticDecode( $C[0..m]$ ):
2   // Assume model  $U_i$  and  $P_{ij}$  are fixed.
3   //  $C[0..m]$  bit string produced by arithmeticEncode
4    $x = \sum_{i=0}^{m-1} C[i] \cdot 2^{m-1-i}$  // final interval  $[x/2^m, (x+1)/2^m]$ 
5    $\ell := 0; p := 1; i := 0$ 
6   while true
7      $c := 0; q := 0$  // Decode next character  $c$ 
8     while  $\ell + q \cdot p < x/2^m$  // Iterate through characters until final interval
9       if  $c == U_i + 1$  // reached $
10          $X[i] := \$$ 
11         return  $X[0..i]$ 
12       else
13          $q := q + P_{i,c}; c := c + 1$ 
14     end while
15      $c := c - 1; q := q - P_{i,c}$  // we overshoot by 1
16      $X[i] := c$ 
17      $\ell := \ell + q \cdot p; p := p \cdot P_{i,c}$ 
18      $i := i + 1$ 
19 end for
```

8.2 Practical Arithmetic Coding

Arithmetic Coding – Numerics

- ▶ As implemented above, p usually gets smaller by a constant factor with *each character*

↪ p gets exponentially small in n !

- ▶ ℓ does not get smaller in absolute terms, but we need it to ever higher accuracy

↪ requires $\Omega(n)$ bit precision and exact arithmetic!

- ▶ *With a clever trick, this can be avoided!*

- ▶ If $[\ell, \ell + p) \subseteq [0, \frac{1}{2})$, we know:

- ▶ Our final x with $[\frac{x}{2^m}, \frac{x+1}{2^m}) \subseteq [\ell, \ell + p)$ *must start with a 0-bit!*

↪ Output a 0 and renormalize interval:

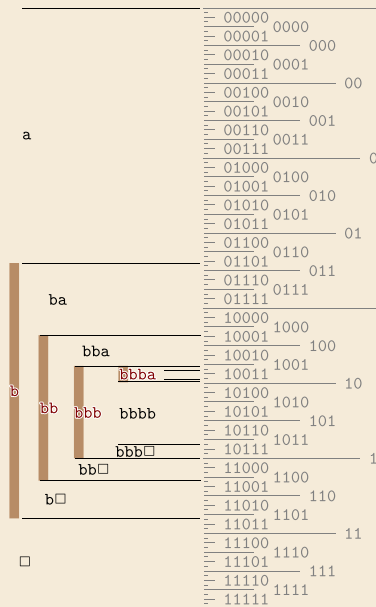
$$\ell := 2\ell; \quad p := 2p$$

- ▶ If $[\ell, \ell + p) \subseteq [\frac{1}{2}, 1)$, similarly:

- ▶ Output 1 and renormalize:

$$\ell := \ell - \frac{1}{2}$$

$$\ell := 2\ell; \quad p := 2p$$



Arithmetic Coding – Renormalization

Does this guarantee ℓ and p stay in a reasonable range?

- ▶ No! Consider (uniform) trits in $\{0, 1, 2\}$ again and encode $1111111111111111 \dots$

$$\rightsquigarrow p = \left(\frac{1}{3}\right)^n, \quad \ell = \frac{1}{3} + \frac{1}{9} + \frac{1}{27} + \dots = \sum_{i=1}^n 3^{-i} = \frac{1}{2} - \frac{3^{-n}}{2}$$

$$\rightsquigarrow \ell < \frac{1}{2} \text{ and } \ell + p > \frac{1}{2} \rightsquigarrow \text{next bit unknown as of yet}$$

But: If $[\ell, \ell + p) \subseteq [\frac{1}{4}, \frac{3}{4})$, next **two** bits are either 01 or 10

- ▶ Remember an “*outstanding opposite bit*” (increment counter)

- ▶ Renormalize:

$$\ell := \ell - \frac{1}{4}$$

$$\ell := 2\ell; \quad p := 2p$$

$$\rightsquigarrow \ell \text{ and } p \text{ remain in range of } P_{ij}$$

$$\rightsquigarrow \text{round } P_{ij} \text{ to integer multiple of } 2^{-F} \rightsquigarrow \text{fixed-precision arithmetic}$$

00000	0000
00001	000
00010	0001
00011	00
00100	0010
00101	001
00110	0011
00111	0
01000	0100
01001	010
01010	0101
01011	01
01100	0110
01101	011
01110	0111
01111	
10000	1000
10001	100
10010	1001
10011	10
10100	1010
10101	101
10110	1011
10111	1
11000	1100
11001	110
11010	1101
11011	11
11100	1110
11101	111
11110	1111
11111	

Fixed Precision Arithmetic Encode

Detailed code from Moffat, Neal, Witten, *Arithmetic Coding Revisited*, ACM Trans. Inf. Sys. 1998

Note: L is our ℓ , R is our p , $b \leq w$ is #bits for variables

```
arithmetic_encode( $l, h, t$ )
```

```
/* Arithmetically encode the range  $[l/t, h/t)$  using low-precision arithmetic.  
The state variables  $R$  and  $L$  are modified to reflect the new range, and then  
renormalized to restore the initial and final invariants  $2^{b-2} < R \leq 2^{b-1}$ ,  
 $0 \leq L < 2^b - 2^{b-2}$ , and  $L + R \leq 2^b$  */
```

```
(1) Set  $r \leftarrow R \text{ div } t$ 
```

```
(2) Set  $L \leftarrow L + r$  times  $l$ 
```

```
(3) If  $h < t$  then
```

```
    set  $R \leftarrow r$  times  $(h - l)$ 
```

```
else
```

```
    set  $R \leftarrow R - r$  times  $l$ 
```

```
(4) While  $R \leq 2^{b-2}$  do
```

```
    Use Algorithm ENCODER RENORMALIZATION (Figure 7) to renormalize  $R$ ,  
    adjust  $L$ , and output one bit
```

Fixed Precision Renormalize

In *arithmetic_encode()*

/* Reestablish the invariant on R , namely that $2^{b-2} < R \leq 2^{b-1}$. Each doubling of R corresponds to the output of one bit, either of known value, or of value opposite to the value of the next bit actually output */

- (4) While $R \leq 2^{b-2}$ do
 If $L + R \leq 2^{b-1}$ then
 bit_plus_follow(0)
 else if $2^{b-1} \leq L$ then
 bit_plus_follow(1)
 Set $L \leftarrow L - 2^{b-1}$
 else
 Set *bits_outstanding* \leftarrow *bits_outstanding* + 1 and $L \leftarrow L - 2^{b-2}$
 Set $L \leftarrow 2L$ and $R \leftarrow 2R$

bit_plus_follow(x)

/* Write the bit x (value 0 or 1) to the output bit stream, plus any outstanding following bits, which are known to be of opposite polarity */

- (1) *write_one_bit*(x).
(2) While *bits_outstanding* > 0 do
 write_one_bit($1 - x$)
 Set *bits_outstanding* \leftarrow *bits_outstanding* - 1

Fixed Precision Arithmetic Decode

Functions *decode_target* and *arithmetic_decode* to be called alternatingly.

decode_target(*t*)

/* Returns an integer *target*, $0 \leq \textit{target} < t$ that is guaranteed to lie in the range $[l, h)$ that was used at the corresponding call to *arithmetic_encode*() */







- (1) Set $r \leftarrow R \text{ div } t$
- (2) Return $(\min\{t - 1, D \text{ div } r\})$

arithmetic_decode(*l, h, t*)

/* Adjusts the decoder's state variables *R* and *D* to reflect the changes made in the encoder during the corresponding call to *arithmetic_encode*(). Note that, compared with Algorithm CACM CODER (Figure 6), the transformation $D = V - L$ is used. It is also assumed that *r* has been set by a prior call to *decode_target*() */

- (1) Set $D \leftarrow D - r \text{ times } l$
- (2) If $h < t$ then
 set $R \leftarrow r \text{ times } (h - l)$
else
 set $R \leftarrow R - r \text{ times } l$
- (3) While $R \leq 2^{b-2}$ do
 Set $R \leftarrow 2R$ and $D \leftarrow 2D + \textit{read_one_bit}()$

Arithmetic Coding Discussion

-  Subtle code (↪ libraries!)
-  Typically slower to encode/decode than Huffman codes
-  Encoded bits can be produced/consumed in bursts
-  Extremely versatile w. r. t. random process
-  Almost optimal space usage / compression
-  Widely used (instead of Huffman) in JPEG, zip variants, . . .

8.3 Error Correcting Codes

Noisy Communication

- ▶ most forms of communication are “noisy”
 - ▶ humans: acoustic noise, unclear pronunciation, misunderstanding, foreign languages

- ▶ How do humans cope with that?

- ▶ slow down and/or speak up
 - ▶ ask to repeat if necessary

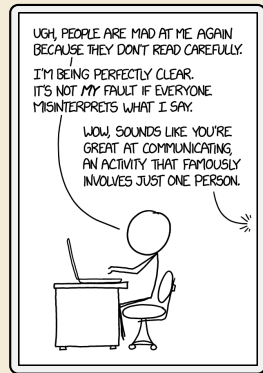


- ▶ But how is it possible (for us) to decode a message in the presence of noise & errors?

*Bcaesue it semes taht ntaurul lanaguge has a lots fo **redundancy** bilt itno it!*

⇒ We can

- 1. detect errors** “This sentence has aao pi dgsdho gioasghds.”
- 2. correct (some) errors** “Tiny errs ar corrected automaticly.”
(sometimes too eagerly as in the Chinese Whispers / Telephone)



Noisy Channels

- ▶ computers: copper cables & electromagnetic interference
 - ▶ transmit a binary string
 - ▶ but occasionally bits can “flip”
- ~> want a robust code



- ▶ We can aim at
 1. **error detection** ~> can request a re-transmit
 2. **error correction** ~> avoid re-transmit for common types of errors
 - ▶ This will require *redundancy*: sending *more* bits than plain message
 - ~> **goal**: robust code with lowest redundancy
- that's the opposite of compression!

8.4 Coding Theory

Block codes

► model:

- want to send message $S \in \{0, 1\}^*$ (bitstream) across a (*communication*) *channel*
- any bit transmitted through the channel might *flip* ($0 \rightarrow 1$ resp. $1 \rightarrow 0$)
no other errors occur (no bits lost, duplicated, inserted, etc.)
- instead of S , we send *encoded bitstream* $C \in \{0, 1\}^*$
sender *encodes* S to C , receiver *decodes* C to S (hopefully)

~> what errors can be detected and/or corrected?

► all codes discussed here are *block codes*

- divide S into *messages* $m \in \{0, 1\}^k$ of k bits each ($k = \text{message length}$)
- encode each message (separately) as $C(m) \in \{0, 1\}^n$ ($n = \text{block length}$, $n \geq k$)

~> can analyze everything block-wise

► between 0 and n bits might be flipped invalid code

- how many flipped bits can we definitely **detect**?
- how many flipped bits can we **correct** without retransmit?

i. e. decoding m still possible

Code distance

$$m \neq m' \implies C(m) \neq C(m')$$

► each block code is an *injective* function $C : \{0, 1\}^k \rightarrow \{0, 1\}^n$

► define \mathcal{C} = set of all codewords = $C(\{0, 1\}^k)$

$$\rightsquigarrow \mathcal{C} \subseteq \{0, 1\}^n$$

$|\mathcal{C}| = 2^k$ out of 2^n n -bit strings are valid codewords

► decoding = finding closest valid codeword

► *distance of code:*

d = minimal Hamming distance of any two codewords = $\min_{x, y \in \mathcal{C}} d_H(x, y)$

Implications for codes

1. Need distance d to **detect** all errors flipping up to $d - 1$ bits.
2. Need distance d to **correct** all errors flipping up to $\lfloor \frac{d-1}{2} \rfloor$ bits.

Lower Bounds

- ▶ Main advantage of concept of code distance:
can *prove* lower bounds on block length

otherwise no such code exists

Given block length n , message length k , code distance d , we must have:

- ▶ **Singleton bound:** $2^k \leq 2^{n-(d-1)} \rightsquigarrow n \geq k + d - 1$

- ▶ *proof sketch:* We have 2^k codewords with distance d
after deleting the first $d - 1$ bits, all are still distinct
but there are only $2^{n-(d-1)}$ such shorter bitstrings.

- ▶ **Hamming bound:** $2^k \leq \frac{2^n}{\sum_{f=0}^{\lfloor (d-1)/2 \rfloor} \binom{n}{f}}$


- ▶ *proof idea:* consider “balls” of bitstrings around codewords
count bitstrings with Hamming-distance $\leq t = \lfloor (d - 1)/2 \rfloor$
correcting t errors means all these balls are disjoint
so $2^k \cdot \text{ball size} \leq 2^n$

\rightsquigarrow We will come back to these.

8.5 Hamming Codes

Parity Bit

- ▶ simplest possible error-detecting code: add a **parity bit**



⇒ code distance 2

- ▶ can detect any single-bit error (actually, any odd number of flipped bits)
- ▶ used in many hardware (communication) protocols
 - ▶ PCI buses, serial buses
 - ▶ caches
 - ▶ early forms of main memory

👍 very simple and cheap


👎 cannot correct any errors

Error-correcting codes

- ▶ typical application: heavy-duty server RAM
 - ▶ bits can randomly flip (e. g., by cosmic rays)
 - ▶ individually very unlikely, but in always-on server with lots of RAM, it happens!

<https://blogs.oracle.com/linux/attack-of-the-cosmic-rays-v2>

any downtime is expensive!



Can we **correct** a bit error without knowing where it occurred? How?

- ▶ Yes! store every bit *three times*!
 - ▶ upon read, do majority vote
 - ▶ if only one bit flipped, the other two (correct) will still win

👎 triples the cost!



You want WHAT!?!




instead of 200% (!)

Can do it with 11% extra memory!

How to locate errors?

- ▶ **Idea:** Use several parity bits
 - ▶ each covers a **subset** of bits
 - ▶ clever subsets \rightsquigarrow violated/valid parity bit pattern narrows down error
- ⚠ flipped bit can be one of the parity bits!
- ▶ Consider $n = 7$ bits B_1, \dots, B_7 with the following constraints:




Observe:

- ▶ No error (all 7 bits correct) $\rightsquigarrow C = C_2C_1C_0 = 000_2 = 0$ ✓
- ▶ What happens if (exactly) 1 bit, say B_i flips?

$C_j = 1$ iff j th bit in binary representation of i is 1 $\rightsquigarrow C$ encodes **position of error!**

(7, 4) Hamming Code

► How can we turn this into a code?



► B_4, B_2 and B_1 occur only in one constraint each \rightsquigarrow **define** them based on rest!

► (7, 4) Hamming Code – Encoding

1. **Given:** message $D_3D_2D_1D_0$ of length $k = 4$
2. copy $D_3D_2D_1D_0$ to $B_7B_6B_5B_3$
3. compute $P_2P_1P_0 = B_4B_2B_1$ so that $C = 0$
4. send $D_3D_2D_1P_2D_0P_1P_0$

(7, 4) Hamming Code – Decoding

► (7, 4) Hamming Code – Decoding

1. **Given:** block $B_7B_6B_5B_4B_3B_2B_1$ of length $n = 7$
2. compute C (as above)
3. if $C = 0$ no (detectable) error occurred
otherwise, flip B_C (the C th bit was twisted)
4. return 4-bit message $B_7B_6B_5B_3$

(7, 4) Hamming Code – Properties

- ▶ **Hamming bound:**

- ▶ 2^4 valid 7-bit codewords (one per message)
- ▶ any of the 7 single-bit errors corrected towards valid codeword

↪ each codeword covers 8 of all possible 7-bit strings

- ▶ $2^4 \cdot 2^3 = 2^7$ ↪ exactly cover space of 7-bit strings

- ▶ distance $d = 3$

- ▶ can *correct* any 1-bit error

- ▶ How about 2-bit errors?

- ▶ We can *detect* that *something* went wrong.


- ▶ **But:** above decoder mistakes it for a (different!) 1-bit error and “corrects” that

- ▶ Variant: store one additional parity bit for entire block

↪ Can *detect* any 2-bit error, but *not correct* it.

Hamming Codes – General recipe

- ▶ construction can be generalized:
 - ▶ Start with $n = 2^\ell - 1$ bits for $\ell \in \mathbb{N}$ (we had $\ell = 3$)
 - ▶ use the ℓ bits whose index is a power of 2 as parity bits
 - ▶ the other $n - \ell$ are data bits
- ▶ Choosing $\ell = 7$ we can encode entire word of memory (64 bit) with 11% overhead (using only 64 out of the 120 possible data bits)

 simple and efficient coding / decoding

 fairly space-efficient

Outlook

- ▶ Indeed: $(2^\ell - 1, 2^\ell - \ell - 1)$ Hamming Code is “*perfect*” code

↪ cannot use fewer bits ...

= matches Hamming lower bound

- ▶ if message length is $2^\ell - \ell - 1$ for $\ell \in \mathbb{N}_{\geq 2}$
i. e., one of 1, 4, 11, 26, 57, 120, 247, 502, 1013, ...
- ▶ **and** we want to correct 1-bit errors
- ▶ For other scenarios, finding good codes is an active research area
 - ▶ information theory predicts that *almost all* randomly chosen codes are good(!)
 - ▶ but these are inefficient to decode
- ↪ clever tricks and constructions needed
e. g. *low density parity check codes*