

8

Clever Codes

1 December 2025

Prof. Dr. Sebastian Wild

Learning Outcomes

Unit 8: *Clever Codes*

1. Know the principles and performance characteristics of *arithmetic coding*.
2. Judge the use of arithmetic coding in applications.
3. Understand the context of *error-prone communication*.
4. Understand concepts of *error-detecting codes* and *error-correcting codes*.
5. Know and understand *Hamming codes*, in particular (7,4) Hamming code.
6. Reason about the *suitability of a code* for an application.

Outline

8 Clever Codes

- 8.1 Arithmetic Coding
- 8.2 Arithmetic Coding Beyond Trits
- 8.3 Practical Arithmetic Coding
- 8.4 Error Correcting Codes
- 8.5 Coding Theory
- 8.6 Hamming Codes

8.1 Arithmetic Coding

Stream Codes

- ▶ Recall: (binary) character encoding $E : \Sigma \rightarrow \{0, 1\}^*$
 - ▶ Huffman codes *optimal* for any given character frequencies
 - ~~ encoding all characters with that code *minimizes* compressed size
 - ... *if we assume that all characters must be encoded individually by a codeword!*
- ▶ Stream codes instead compress entire **sequence** of characters
 - ▶ RLE and LZW are examples of stream codes ~~ can sometimes do better
- ▶ Two indicative examples
 1. **"Low entropy bits:"** $\Sigma = \{0, 1\}$, highly skewed: $p_0 = 0.99$
 - ~~ entropy $\mathcal{H}(\frac{1}{100}, \frac{99}{100}) \approx 0.08$ bits per character,
Huffman code must use 1 bit per character!
 - ~~ "optimal" Huffman code gives 12-fold space increase over entropy!
 - ▶ Can certainly do better here (RLE!)
 2. **"Trits":** $\Sigma = \{0, 1, 2\}$, equally likely
 - ~~ entropy $\mathcal{H}(\frac{1}{3}, \frac{1}{3}, \frac{1}{3}) = \lg(3) \approx 1.58$ bits per character,
Huffman code uses average of $\frac{1}{3} \cdot 1 + \frac{2}{3} \cdot 2 = \frac{5}{3} \approx 1.67$
- ▶ Can we do better?

A Decent Hack: Block Codes

- ▶ Huffman on trits wastes ≈ 0.0817 bits per character and over 5 % of space
- ▶ A simple trick can reduce this substantially!
 - ▶ treat 5 trits as one “supercharacter”, e.g., **21101**
 - ~~ $3^5 = 243$ possible combinations
 - ~~ encode these using 8 bits (with $2^8 = 256$ possible combinations)
 - ▶ entropy $\lg(3^5) \approx 7.92$ bits, so less than 0.1 % wasted space!
- ▶ We can even use a Huffman code for the supercharacters to handle nonuniformity!
- ▶ For the low-entropy bits, could use 3 bits
 - ~~ probabilities:
 - 000** : 0.97
 - 001**, **010**, **100** : 0.0098
 - 011**, **101**, **110** : 0.000099
 - 111**: 0.000001
 - ~~ with Huffman code, 1.06 bits per superchar of 3 input bits
 - ~~ almost factor 3 better; can improve with larger blocks!

Block Codes – A Panacea?

- ▶ Using supercharacters works well in our examples.



*Hmmm . . . so why don't we treat the entire source text as one large block?
Wouldn't that be even better!?*

~~ We can optimally compress any text, without doing anything intelligent!?



- ⚡ For general case, need to *communicate* the supercharacter encoding

- ▶ Blocks of k characters need $\Omega(\sigma^k)$ space for code
- ▶ Huffman code has to be part of coded message
- ~~ Can only sensibly use block codes for small σ and k



There is no such thing as a free lunch . . .

Arithmetic Coding


except in isolated lucky cases

- Also: Block codes still had $\Theta(n)$ wasted space for sequences of n symbols

Arithmetic Coding:

- Maintain $[\ell, \ell + p) \subseteq [0, 1]$; initially $\ell = 0, p = 1$
 - Zoom into subinterval for each character
 - Output dyadic encoding of final interval
- Step 1:** “Zoom” for each character (trit) in $S[0..n]$:

- Of the current subinterval $[\ell, \ell + p)$,
take first, second or last third
depending whether $S[i] = 0, 1$, resp. 2:
 $\ell := \ell + S[i] \cdot \frac{1}{3} \cdot p$
 $p := p \cdot \frac{1}{3}$



Step 2: Dyadic encoding

- Find smallest m so that $\exists x \in \mathbb{N}_0$ with $\left[\frac{x}{2^m}, \frac{x+1}{2^m} \right) \subseteq [\ell, \ell + p)$
- Output x in binary using m bits.

\rightsquigarrow Encode n trits in $n \lg(3) + 2$ bits(!) without cheating

Arithmetic Coding – Encode Trits Example

- ▶ $S[0..n) = 21101 \quad (n = 5)$
- ▶ Step 1: Zoom into subintervals

Iteration	ℓ	p	Interval (rounded)	
0	0	1	[0.00000, 1.00000)	
1	$\frac{2}{3}$	$\frac{1}{3}$	[0.66667, 1.00000)	
2	$\frac{7}{9}$	$\frac{1}{9}$	[0.77778, 0.88889)	
3	$\frac{22}{27}$	$\frac{1}{27}$	[0.81482, 0.85185)	
4	$\frac{66}{81}$	$\frac{1}{81}$	[0.81482, 0.82716)	
5	$\frac{199}{243}$	$\frac{1}{243}$	[0.81893, 0.82305)	

- ▶ Step 2: Dyadic encoding for interval $[\ell, \ell + p) = \left[\frac{199}{243}, \frac{200}{243} \right)$
 - ▶ Must have $m \geq \lg(1/p) > 7$
 - ▶ $m = 8$: smallest $x/2^m \geq \frac{199}{243}$ is $x = 210$, but $[210/256, 211/256) \approx [0.82031, 0.82422) \not\subset [\ell, \ell + p)$
 - ▶ $m = 9$: smallest $x/2^m \geq \frac{199}{243}$ is $x = 420$ and $[420/512, 421/512) \approx [0.82031, 0.82227) \subset [\ell, \ell + p)$ ✓
 - ~~ Output $x = 420$ in binary with $m = 9$ digits: 110100100

8.2 Arithmetic Coding Beyond Trits


Beyond Trits

In the example above, we always split the interval into thirds.


But there's nothing special about thirds.

- ~~> Any subdivision of $[0, 1)$ works!

Versatility of Arithmetic Coding – Adaptive Model



Context (sequence thus far)	Probability of next symbol		
	$P(a) = 0.425$	$P(b) = 0.425$	$P(\square) = 0.15$
b	$P(a b) = 0.28$	$P(b b) = 0.57$	$P(\square b) = 0.15$
bb	$P(a bb) = 0.21$	$P(b bb) = 0.64$	$P(\square bb) = 0.15$
bbb	$P(a bbb) = 0.17$	$P(b bbb) = 0.68$	$P(\square bbb) = 0.15$
bbba	$P(a bbba) = 0.28$	$P(b bbba) = 0.57$	$P(\square bbba) = 0.15$



adapted from Figure 6.4 of MacKay: *Information Theory, Inference, and Learning Algorithms* 2003

Arithmetic Coding – General framework

- ▶ Note: Arithmetic coder *doesn't care* if probabilities or even σ change all the time!
 - ▶ As long as encoder and decoder know from context what they are!

General stochastic sequence:

Sequence of random variables X_0, X_1, X_2, \dots such that

1. $X_i \in [0..U_i) \cup \{\$\}$ (We use \$ to signal “end of text”)
2. $\mathbb{P}[X_i = j] = P_{ij}$
3. both U_i and P_{ij} are random variables as they *depend* on X_0, \dots, X_{i-1} ,
but conditioned on X_0, \dots, X_{i-1} , they are fixed and known:

$$P_{ij} = P_{ij}(X_0, \dots, X_{i-1}) = \mathbb{P}[X_i = j | X_0, \dots, X_{i-1}]$$

$$U_i = U_i(X_0, \dots, X_{i-1}) = \max\{j : P_{ij}(X_0, \dots, X_{i-1}) > 0\}$$

- ▶ Can model arbitrary dependencies on previous outcomes
- ▶ Assume here that random process is known by both encoder and decoder (fixed coding)
otherwise extra space needed to encode model!

Arithmetic Coding – Encoding

```
1 procedure arithmeticEncode( $X_0, \dots, X_n$ ):  
2     // Assume model  $U_i$  and  $P_{ij}$  are fixed.  
3     // Assume  $X_i \in [0..U_i)$  for  $i < n$  and  $X_n = \$$   
4     // Step 1: Interval zooming  
5      $\ell := 0$ ;  $p := 1$   
6     for  $i := 0, \dots, n - 1$  do  
7          $q := \sum_{j=0}^{X_i-1} P_{ij};$   
8          $\ell := \ell + q \cdot p$ ;  $p := p \cdot P_{i,X_i}$   
9     end for  
10     $q := 1 - P_{n,\$}$  // encode  $\$$  as last character  
11     $\ell := \ell + q \cdot p$ ;  $p := p \cdot P_{n,\$}$   
12    // Step 2: Dyadic encoding  
13     $m := \lceil \lg(1/p) \rceil - 1$   
14    do  
15         $m := m + 1$ ;  $x := \lceil \ell \cdot 2^m \rceil$   
16        while  $(x + 1)/2^m > \ell + p$   
17        return  $x$  in binary using  $m$  bits
```

Arithmetic Coding – Decoding

```
1 procedure arithmeticDecode(C[0..m]):  
2     // Assume model  $U_i$  and  $P_{ij}$  are fixed.  
3     // C[0..m) bit string produced by arithmeticEncode  
4      $x = \sum_{i=0}^{m-1} C[i] \cdot 2^{m-1-i}$  // final interval  $[x/2^m, (x+1)/2^m)$   
5      $\ell := 0; p := 1; i := 0$   
6     while true  
7          $c := 0; q := 0$  // Decode next character c  
8         while  $\ell + q \cdot p < x/2^m$  // Iterate through characters until final interval  
9             if  $c == U_i + 1$  // reached $  
10                 $X[i] := \$$   
11                return X[0..i]  
12            else  
13                 $q := q + P_{i,c}; c := c + 1$   
14        end while  
15         $c := c - 1; q := q - P_{i,c}$  // we overshot by 1  
16         $X[i] := c$   
17         $\ell := \ell + q \cdot p; p := p \cdot P_{i,c}$   
18         $i := i + 1$   
19    end for
```

8.3 Practical Arithmetic Coding

Arithmetic Coding – Numerics

- As implemented above, p usually gets smaller by a constant factor with *each character*
 - ~~ p gets exponentially small in $n!$

- ℓ does not get smaller in absolute terms, but we need it to ever higher accuracy

~~ requires $\Omega(n)$ bit precision and exact arithmetic!

- With a clever trick, this can be avoided!

- If $[\ell, \ell + p) \subseteq [0, \frac{1}{2}]$, we know:

- Our final x with $[\frac{x}{2^m}, \frac{x+1}{2^m}) \subseteq [\ell, \ell + p)$
must start with a 0-bit!


- ~~ Output a 0 and renormalize interval:
 $\ell := 2\ell; p := 2p$

- If $[\ell, \ell + p) \subseteq [\frac{1}{2}, 1]$, similarly:

- Output 1 and renormalize:

$$\ell := \ell - \frac{1}{2}$$

$$\ell := 2\ell; p := 2p$$



Arithmetic Coding – Renormalization

Does this guarantee ℓ and p stay in a reasonable range?

- ▶ No! Consider (uniform) trits in $\{0, 1, 2\}$ again and encode
 $1111111111111111\dots$

$$\rightsquigarrow p = \left(\frac{1}{3}\right)^n, \quad \ell = \frac{1}{3} + \frac{1}{9} + \frac{1}{27} + \dots = \sum_{i=1}^n 3^{-i} = \frac{1}{2} - \frac{3^{-n}}{2}$$

$\rightsquigarrow \ell < \frac{1}{2}$ and $\ell + p > \frac{1}{2}$ \rightsquigarrow next bit unknown as of yet

But: If $[\ell, \ell + p) \subseteq [\frac{1}{4}, \frac{3}{4})$, next **two** bits are either **01** or **10**

- ▶ Remember an “*outstanding opposite bit*” (increment counter)

- ▶ Renormalize:

$$\ell := \ell - \frac{1}{4}$$

$$\ell := 2\ell; \quad p := 2p$$

$\rightsquigarrow \ell$ and p remain in range of P_{ij}

\rightsquigarrow round P_{ij} to integer multiple of 2^{-F} \rightsquigarrow fixed-precision arithmetic

-	00000	0000
-	00001	0000
-	00010	000
-	00011	0001
-	00100	00
-	00101	0010
-	00110	001
-	00111	0011
-	01000	0
-	01001	0100
-	01010	010
-	01011	0101
-	01100	01
-	01101	0110
-	01110	011
-	01111	0111
-	10000	1000
-	10001	1000
-	10010	100
-	10011	1001
-	10100	10
-	10101	1010
-	10110	101
-	10111	1011
-	11000	1
-	11001	1100
-	11010	110
-	11011	1101
-	11100	11
-	11101	1110
-	11110	111
-	11111	1111

Fixed Precision Arithmetic Encode

Detailed code from



Moffat, Neal, Witten: *Arithmetic Coding Revisited*, ACM Trans. Inf. Sys. 1998

Note: Their L is our ℓ , R is our p , $b \leq w$ is #bits for variables

arithmetic_encode(l, h, t)

```
/* Arithmetically encode the range [l/t, h/t) using low-precision arithmetic.  
The state variables R and L are modified to reflect the new range, and then  
renormalized to restore the initial and final invariants  $2^{b-2} < R \leq 2^{b-1}$ ,  
 $0 \leq L < 2^b - 2^{b-2}$ , and  $L + R \leq 2^b$  */
```

- (1) Set $r \leftarrow R \text{ div } t$
- (2) Set $L \leftarrow L + r \text{ times } l$
- (3) If $h < t$ then
 - set $R \leftarrow r \text{ times } (h - l)$
 - else
 - set $R \leftarrow R - r \text{ times } l$
- (4) While $R \leq 2^{b-2}$ do
 - Use Algorithm ENCODER RENORMALIZATION (Figure 7) to renormalize R ,
 - adjust L , and output one bit

Fixed Precision Renormalize

In *arithmetic_encode()*

/* Reestablish the invariant on R , namely that $2^{b-2} < R \leq 2^{b-1}$. Each doubling of R corresponds to the output of one bit, either of known value, or of value opposite to the value of the next bit actually output */

- (4) While $R \leq 2^{b-2}$ do
 - If $L + R \leq 2^{b-1}$ then
 - bit_plus_follow(0)*
 - else if $2^{b-1} \leq L$ then
 - bit_plus_follow(1)*
 - Set $L \leftarrow L - 2^{b-1}$
 - else
 - Set $bits_outstanding \leftarrow bits_outstanding + 1$ and $L \leftarrow L - 2^{b-2}$
 - Set $L \leftarrow 2L$ and $R \leftarrow 2R$

bit_plus_follow(x)

/* Write the bit x (value 0 or 1) to the output bit stream, plus any outstanding following bits, which are known to be of opposite polarity */

- (1) *write_one_bit(x).*
- (2) While $bits_outstanding > 0$ do
 - write_one_bit(1 - x)*
 - Set $bits_outstanding \leftarrow bits_outstanding - 1$

Fixed Precision Arithmetic Decode

Functions *decode_target* and *arithmetic_decode* to be called alternately.

decode_target(t)

/* Returns an integer *target*, $0 \leq \text{target} < t$ that is guaranteed to lie in the range $[l, h]$ that was used at the corresponding call to *arithmetic_encode()* */

- (1) Set $r \leftarrow R \text{ div } t$
- (2) Return ($\min\{t - 1, D \text{ div } r\}$)

arithmetic_decode(l, h, t)

/* Adjusts the decoder's state variables R and D to reflect the changes made in the encoder during the corresponding call to *arithmetic_encode()*. Note that, compared with Algorithm CACM CODER (Figure 6), the transformation $D = V - L$ is used. It is also assumed that r has been set by a prior call to *decode_target()* */

- (1) Set $D \leftarrow D - r \text{ times } l$
- (2) If $h < t$ then
 - set $R \leftarrow r \text{ times } (h - l)$
 - else
 - set $R \leftarrow R - r \text{ times } l$
- (3) While $R \leq 2^{b-2}$ do
 - Set $R \leftarrow 2R$ and $D \leftarrow 2D + \text{read_one_bit}()$

Arithmetic Coding Discussion

-  Subtle code (\rightsquigarrow libraries!)
-  Typically slower to encode/decode than Huffman codes
-  Encoded bits can be produced/consumed in bursts
-  Extremely versatile w. r. t. random process
-  Almost optimal space usage / compression
-  Widely used (instead of Huffman) in JPEG, zip variants, . . .

8.4 Error Correcting Codes

Noisy Communication

- ▶ most forms of communication are “noisy”
 - ▶ humans: acoustic noise, unclear pronunciation, misunderstanding, foreign languages

- ▶ How do humans cope with that?

- ▶ slow down and/or speak up
 - ▶ ask to repeat if necessary



- ▶ But how is it possible (for us) to decode a message in the presence of noise & errors?

Bcaesue it semes taht ntaurul lanaguge has a lots fo redundancy bilt itno it!

~~ We can

- 1. detect errors** “This sentence has aao pi dgsdho gioasghds.”
- 2. correct (some) errors** “Tiny errs ar corrrected automatically.”
(sometimes too eagerly as in the Chinese Whispers / Telephone)



UGH, PEOPLE ARE MAD AT ME AGAIN
BECAUSE THEY DONT READ CAREFULLY.

I'M BEING PERFECTLY CLEAR.
IT'S NOT MY FAULT IF EVERYONE
MISINTERPRETS WHAT I SAY.

WOW, SOUNDS LIKE YOU'RE
GREAT AT COMMUNICATING,
AN ACTIVITY THAT FAMOUSLY
INVOLVES JUST ONE PERSON.

Noisy Channels

- ▶ computers: copper cables & electromagnetic interference
- ▶ transmit a binary string
- ▶ but occasionally bits can “flip”
- ~~ want a robust code



- ▶ We can aim at
 1. **error detection** ~~ can request a re-transmit
 2. **error correction** ~~ avoid re-transmit for common types of errors
- ▶ This will require *redundancy*: sending *more* bits than plain message
~~ **goal**: robust code with lowest redundancy that's the opposite of compression!

8.5 Coding Theory

Block codes

► model:

- ▶ want to send message $S \in \{0, 1\}^*$ (bitstream) across a (*communication*) *channel*
- ▶ any bit transmitted through the channel might *flip* ($0 \rightarrow 1$ resp. $1 \rightarrow 0$)
no other errors occur (no bits lost, duplicated, inserted, etc.)
- ▶ instead of S , we send *encoded bitstream* $C \in \{0, 1\}^*$
sender *encodes* S to C , receiver *decodes* C to S (hopefully)
 - ~~ what errors can be detected and/or corrected?
- ▶ all codes discussed here are **block codes**
 - ▶ divide S into messages $m \in \{0, 1\}^k$ of k bits each ($k = \text{message length}$)
 - ▶ encode each message (separately) as $C(m) \in \{0, 1\}^n$ ($n = \text{block length}$, $n \geq k$)
 - ~~ can analyze everything block-wise
- ▶ between 0 and n bits might be flipped
 - ▶ how many flipped bits can we definitely **detect**?
 - invalid code
 - ▶ how many flipped bits can we **correct** without retransmit?
 - i. e. decoding m still possible

Code distance

$$m \neq m' \implies C(m) \neq C(m')$$

► each block code is an *injective* function $C : \{0, 1\}^k \rightarrow \{0, 1\}^n$

► define \mathcal{C} = set of all codewords = $C(\{0, 1\}^k)$

~ \sim $\mathcal{C} \subseteq \{0, 1\}^n$ $|\mathcal{C}| = 2^k$ out of 2^n n -bit strings are valid codewords

► decoding = finding closest valid codeword

► *distance of code:*

d = minimal Hamming distance of any two codewords = $\min_{x, y \in \mathcal{C}} d_H(x, y)$

Implications for codes

1. Need distance d to detect all errors flipping up to $d - 1$ bits.
2. Need distance d to correct all errors flipping up to $\lfloor \frac{d-1}{2} \rfloor$ bits.

Lower Bounds


- ▶ Main advantage of concept of code distance:
can *prove* lower bounds on block length
otherwise no such code exists
- Given block length n , message length k , code distance d , we must have:
- ▶ **Singleton bound:** $2^k \leq 2^{n-(d-1)} \rightsquigarrow n \geq k + d - 1$
 - ▶ *proof sketch:* We have 2^k codeswords with distance d
after deleting the first $d - 1$ bits, all are still distinct
but there are only $2^{n-(d-1)}$ such shorter bitstrings.
- ▶ **Hamming bound:** $2^k \leq \frac{2^n}{\sum_{f=0}^{\lfloor(d-1)/2\rfloor} \binom{n}{f}}$
 - ▶ *proof idea:* consider “balls” of bitstrings around codewords
count bitstrings with Hamming-distance $\leq t = \lfloor(d-1)/2\rfloor$
correcting t errors means all these balls are disjoint
so $2^k \cdot \text{ball size} \leq 2^n$

rightsquigarrow We will come back to these.

8.6 Hamming Codes

Parity Bit

- ▶ simplest possible error-detecting code: add a **parity bit**



~~ code distance 2

- ▶ can detect any single-bit error (actually, any odd number of flipped bits)
- ▶ used in many hardware (communication) protocols
 - ▶ PCI buses, serial buses
 - ▶ caches
 - ▶ early forms of main memory

👍 very simple and cheap


👎 cannot correct any errors

Error-correcting codes

- ▶ typical application: heavy-duty server RAM
 - ▶ bits can randomly flip (e.g., by cosmic rays)
 - ▶ individually very unlikely,
but in always-on server with lots of RAM, it happens!

<https://blogs.oracle.com/linux/attack-of-the-cosmic-rays-v2>

any downtime is expensive!



Can we **correct** a bit error without knowing where it occurred? How?

- ▶ Yes! store every bit *three times!*
 - ▶ upon read, do majority vote
 - ▶ if only one bit flipped, the other two (correct) will still win
- ◀ *triples the cost!*



instead of 200% (!)


Can do it with 11% extra memory!



You want WHAT?!?

How to locate errors?

- Idea: Use several parity bits
 - each covers a **subset** of bits
 - clever subsets \rightsquigarrow violated/valid parity bit pattern narrows down error
- ⚠️ flipped bit can be one of the parity bits!
- Consider $n = 7$ bits B_1, \dots, B_7 with the following constraints:




Observe:

- No error (all 7 bits correct) $\rightsquigarrow C = C_2C_1C_0 = 000_2 = 0$ ✓
- What happens if (exactly) 1 bit, say B_i flips?

$C_j = 1$ iff j th bit in binary representation of i is 1 $\rightsquigarrow C$ encodes **position of error!**

(7, 4) Hamming Code

- ▶ How can we turn this into a code?



- ▶ B_4, B_2 and B_1 occur only in one constraint each \rightsquigarrow **define** them based on rest!
- ▶ (7, 4) Hamming Code – Encoding
 - 1. Given: message $D_3D_2D_1D_0$ of length $k = 4$
 - 2. copy $D_3D_2D_1D_0$ to $B_7B_6B_5B_3$
 - 3. compute $P_2P_1P_0 = B_4B_2B_1$ so that $C = 0$
 - 4. send $D_3D_2D_1P_2D_0P_1P_0$

(7, 4) Hamming Code – Decoding

► (7, 4) Hamming Code – Decoding

1. Given: block $B_7B_6B_5B_4B_3B_2B_1$ of length $n = 7$
2. compute C (as above)
3. if $C = 0$ no (detectable) error occurred
otherwise, flip B_C (the C th bit was twisted)
4. return 4-bit message $B_7B_6B_5B_3$

(7, 4) Hamming Code – Properties

- ▶ **Hamming bound:**
 - ▶ 2^4 valid 7-bit codewords (on per message)
 - ▶ any of the 7 single-bit errors corrected towards valid codeword
 - ~~ each codeword covers 8 of all possible 7-bit strings
 - ▶ $2^4 \cdot 2^3 = 2^7$ ~~ exactly cover space of 7-bit strings
- ▶ distance $d = 3$
- ▶ can *correct* any 1-bit error
- ▶ How about 2-bit errors?
 - ▶ We can *detect* that *something* went wrong.
 - ▶ **But:** above decoder mistakes it for a (different!) 1-bit error and “corrects” that
 - ▶ Variant: store one additional parity bit for entire block
 - ~~ Can *detect* any 2-bit error, but *not correct* it.

Hamming Codes – General recipe

- ▶ construction can be generalized:
 - ▶ Start with $n = 2^\ell - 1$ bits for $\ell \in \mathbb{N}$ (we had $\ell = 3$)
 - ▶ use the ℓ bits whose index is a power of 2 as parity bits
 - ▶ the other $n - \ell$ are data bits
- ▶ Choosing $\ell = 7$ we can encode entire word of memory (64 bit) with 11% overhead (using only 64 out of the 120 possible data bits)

-  simple and efficient coding / decoding
-  fairly space-efficient

Outlook

- ▶ Indeed: $(2^\ell - 1, 2^\ell - \ell - 1)$ Hamming Code is “*perfect*” code
 - ~~ cannot use fewer bits . . .
 - = matches Hamming lower bound
- ▶ if message length is $2^\ell - \ell - 1$ for $\ell \in \mathbb{N}_{\geq 2}$
 - i. e., one of $1, 4, 11, 26, 57, 120, 247, 502, 1013, \dots$
- ▶ and we want to correct 1-bit errors
- ▶ For other scenarios, finding good codes is an active research area
 - ▶ information theory predicts that *almost all* randomly chosen codes are good(!)
 - ▶ but these are inefficient to decode
 - ~~ clever tricks and constructions needed
 - e. g. *low density parity check codes*