

# 10

## Parallel Algorithms

12 January 2026

Prof. Dr. Sebastian Wild

# Learning Outcomes

## Unit 10: *Parallel Algorithms*

1. Know and apply *parallelization strategies* for embarrassingly parallel problems.
2. Identify *limits of parallel speedups*.
3. Understand and use the *parallel random-access-machine* model in its different variants.
4. Be able to *analyze* and compare simple shared-memory parallel algorithms by determining *parallel time and work*.
5. Understand efficient parallel *prefix sum* algorithms.
6. Be able to devise high-level description of *parallel quicksort and mergesort* methods.

## Outline

# 10 Parallel Algorithms

- 10.1 Parallel Computation
- 10.2 Parallel String Matching
- 10.3 Parallel Primitives
- 10.4 Parallel Sorting

## 10.1 Parallel Computation

# Types of parallel computation

€€€ can't buy you more time . . . but more computers!

~ Challenge: Algorithms for *parallel* computation.

There are two main forms of parallelism:

## 1. shared-memory parallel computer $\leftarrow$ focus of today

- ▶  $p$  *processing elements* (PEs, processors) working in parallel
- ▶ **single big memory, accessible from every PE**
- ▶ communication via shared memory
- ▶ think: a big server, 128 CPU cores, terabyte of main memory

## 2. distributed computing

- ▶  $p$  PEs working in parallel
- ▶ each PE has **private** memory
- ▶ communication by sending **messages** via a network
- ▶ think: a cluster of individual machines

# PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 2)
- ▶ the  $p$  PEs are identified by ids  $0, \dots, p - 1$ 
  - ▶ like  $w$  (the word size),  $p$  is a parameter of the model that can grow with  $n$
  - ▶  $p = \Theta(n)$  is not unusual      maaany processors!
- ▶ the PEs all **independently** run the same RAM-style program  
(they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:  
in each time step, every PE executes one instruction
- ▶ As for RAM:
  - ▶ assume a basic “operating system”
    - ~~ write algorithms in pseudocode instead of RAM assembly
  - ▶ **NEW:** loops and commands can be run “**in parallel**” (examples coming up)

# PRAM – Conflict management



**Problem:** What if several PEs simultaneously overwrite a memory cell?

- ▶ **EREW-PRAM** (exclusive read, exclusive write)  
any **parallel access** to same memory cell is **forbidden** (crash if happens)
- ▶ **CREW-PRAM** (concurrent read, exclusive write)  
parallel **write** access to same memory cell is **forbidden**, *but reading is fine*
- ▶ **CRCW-PRAM** (concurrent read, concurrent write)  
concurrent access is allowed,  
need a rule for write conflicts:
  - ▶ common CRCW-PRAM:  
all concurrent writes to same cell must write **same** value
  - ▶ arbitrary CRCW-PRAM:  
some unspecified concurrent write wins
  - ▶ (more exist . . . )
- ▶ no single model is always adequate, but our default is CREW

# PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish      assuming sufficiently many PEs!  
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

Holy grail of PRAM algorithms:

- ▶ minimal time (=span)
- ▶ work (asymptotically) no worse than running time of best sequential algorithm  
~~ “*work-efficient*” algorithm: work in same  $\Theta$ -class as best sequential

# The number of processors

*Hold on, my computer does not have  $\Theta(n)$  processors! Why should I care for span and work!?*

## Theorem 10.1 (Brent's Theorem)

If an algorithm has span  $T$  and work  $W$  (for an arbitrarily large number of processors), it can be run on a PRAM with  $p$  PEs in time  $O(T + \frac{W}{p})$  (and using  $O(W)$  work). ◀

*Proof:* schedule parallel steps in round-robin fashion on the  $p$  PEs.

~~ span and work give guideline for *any* number of processors

## 10.2 Parallel String Matching

# Embarrassingly Parallel

- ▶ A problem is called "*embarrassingly parallel*"  
if it can immediately be split into *many, small subtasks*  
that can be solved completely *independently* of each other
- ▶ Typical example: sum of two large matrices      (all entries independent)  
~~ best case for parallel computation      (simply assign each processor one subtask)
- ▶ Sorting is not embarrassingly parallel
  - ▶ no obvious way to define many *small* (= efficiently solvable) subproblems
  - ▶ but: some subtasks of our algorithms are (stay tuned . . . )

# Parallel string matching – Easy?

- ▶ We have seen a plethora of string matching methods in Unit 6
- ▶ But all efficient methods seem inherently sequential  
*Indeed, they became efficient only after building on knowledge from previous steps!*  
Sounds like the *opposite* of parallel!

~~ How well can we parallelize string matching?

Here: string matching = find *all* occurrences of  $P$  in  $T$       (more natural problem for parallel)  
always assume  $m \leq n$

## Subproblems in string matching:

- ▶ string matching = check all guesses  $i = 0, \dots, n - m - 1$
- ▶ checking one guess is a subtask!

# Parallel string matching – Brute force

- ▶ Check all guesses in parallel

---

```
1 procedure parallelBruteForce( $T[0..n], P[0..m]$ ):  
2   for  $i := 0, \dots, n - m - 1$  do in parallel ← only difference to normal brute force!  
3     for  $j := 0, \dots, m - 1$  do  
4       if  $T[i + j] \neq P[j]$  then break inner loop  
5       if  $j == m$  then report match at  $i$   
6   end parallel for
```

---

- ▶ PE  $k$  is executing the loop iteration where  $i = k$ .
  - ~~ requires that all iterations can be done **independently**!
    - ▶ Different PEs work **in lockstep** (synchronized after each instruction)
    - ▶ similar to OpenMP `#pragma omp parallel for`
- ▶ checking whether *no* match was found by *any* PE more effort ~~ ... stay tuned

~~ **Time:**  $\Theta(m)$  using sequential checks      **Work:**  $\Theta((n - m)m)$  ~~ not great  
                 $\Theta(\log m)$  on CREW-PRAM (~~ tutorials)      ... much more than best sequential  
                 $\Theta(1)$       on CRCW-PRAM (~~ tutorials)

# Parallel string matching – Blocking



Divide  $T$  into **overlapping** blocks of  $2m - 1$  characters:  
 $T[0..2m - 1], T[m..3m - 1], T[2m..4m - 1], T[3m..5m - 1] \dots$

- ▶ Search all blocks in parallel, each using efficient *sequential* method

---

```
1 procedure blockingStringMatching( $T[0..n], P[0..m]$ ):  
2     for  $b := 0, \dots, \lceil n/m \rceil$  do in parallel  
3         result := KMP( $T[bm .. (b+1)m - 1], P$ )  
4         if result  $\neq$  NO_MATCH then report match at result  
5     end parallel for
```

---

~~~ **Time:**

- ▶ loop body has text of length  $n' = 2m - 1$  and pattern of length  $m$

~~~ KPM runtime  $\Theta(n' + m) = \Theta(m)$

~~~ **Work:**  $\Theta(\frac{n}{m} \cdot m) = \Theta(n)$  ~~ work efficient!

# Parallel string matching – Discussion

-  very simple methods
-  could even run distributed with access to part of  $T$
-  parallel speedup only for  $m \ll n$

- work-efficient methods with better parallel time possible?
  - ~~ must genuinely parallelize the matching process! (and the preprocessing of the pattern)
  - ~~ needs new ideas (much more complicated, but possible!)
- **Parallel string matching – State of the art:**
  - $O(\log m)$  time & work-efficient parallel string matching (very complicated)
    - this is optimal for CREW-PRAM
  - on CRCW-PRAM: matching part even in  $O(1)$  time (easy)  
but preprocessing requires  $\Theta(\log \log m)$  time (very complicated)

## 10.3 Parallel Primitives

# Building blocks



- ▶ Most nontrivial problems need tricks to be parallelized
- ▶ Some versatile building blocks are known that help in many problems
  - ~~ We study some of them now, before we apply them to *parallel sorting*

*The following problems might not look natural at first sight . . . but turn out to be good abstractions.*  
~~ bear with me

# Prefix sums

**Prefix-sum problem** (also: cumulative sums, running totals)

- ▶ Given: array  $A[0..n]$  of numbers
- ▶ Goal: compute all prefix sums  $A[0] + \dots + A[i]$  for  $i = 0, \dots, n - 1$   
may be done “in-place”, i. e., by overwriting  $A$

**Example:**

input:

|   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 3 | 0 | 0 | 5 | 7 | 0 | 0 | 2 | 0 | 0 | 0 | 4 | 0 | 8 | 0 | 1 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

$\Sigma$

output:

|   |   |   |   |    |    |    |    |    |    |    |    |    |    |    |    |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|
| 3 | 3 | 3 | 8 | 15 | 15 | 15 | 17 | 17 | 17 | 17 | 21 | 21 | 29 | 29 | 30 |
|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|

# Prefix sums – Sequential

- ▶ sequential solution does  $n - 1$  additions
  - ▶ but: cannot parallelize them!
    - ⚡ data dependencies!
- ~~ need a different approach

Let's try a simpler problem first.

## Excuse: Sum

- ▶ Given: array  $A[0..n]$  of numbers
- ▶ Goal: compute  $A[0] + A[1] + \dots + A[n - 1]$   
(solved by prefix sums)

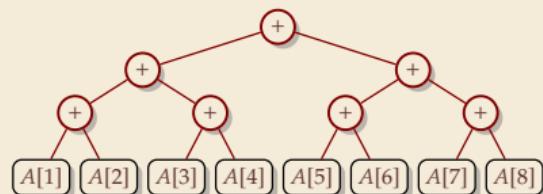
Any algorithm *must* do  $n - 1$  binary additions

- ~~ Height of tree = parallel time!

---

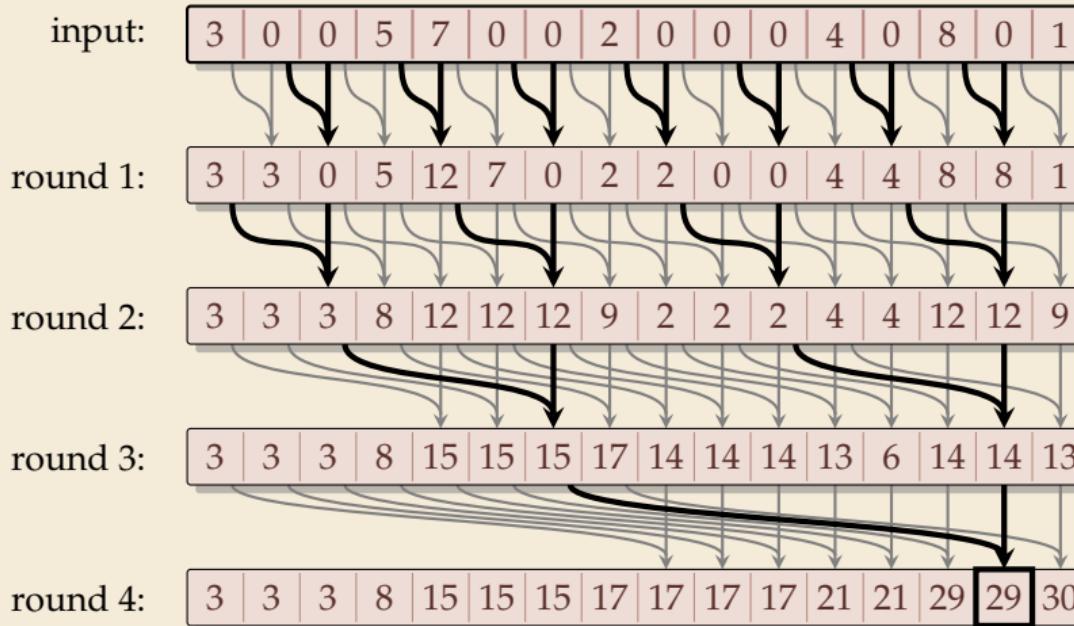
```
1 procedure prefixSum(A[0..n)):  
2   for i := 1, ..., n - 1 do  
3     A[i] := A[i - 1] + A[i]
```

---



# Parallel prefix sums

- Idea: Compute all prefix sums with balanced trees in parallel  
Remember partial results for reuse



## Parallel prefix sums – Code

- ▶ can be realized in-place (overwriting  $A$ )
- ▶ assumption: in each parallel step, all reads precede all writes

---

```
1 procedure parallelPrefixSums( $A[0..n]$ ):  
2     for  $r := 1, \dots, \lceil \lg n \rceil$  do  
3          $step := 2^{r-1}$   
4         for  $i := step, \dots, n-1$  do in parallel  
5              $x := A[i] + A[i-step]$   
6              $A[i] := x$   
7         end parallel for  
8     end for
```

---

# Parallel prefix sums – Analysis

## ► Time:

- ▶ all additions of one round run in parallel
- ▶  $\lceil \lg n \rceil$  rounds
- $\rightsquigarrow \Theta(\log n)$  time      best possible!

## ► Work:

- ▶  $\geq \frac{n}{2}$  additions in all rounds (except maybe last round)
  - $\rightsquigarrow \Theta(n \log n)$  work
  - ▶ more than the  $\Theta(n)$  sequential algorithm!
- 
- Typical trade-off: greater parallelism at the expense of more overall work
  - For prefix sums:
    - ▶ can actually get  $\Theta(n)$  work in *twice* that time!  
 $\rightsquigarrow$  algorithm is slightly more complicated
    - ▶ instead here: linear work in *thrice* the time using “blocking trick”

# Work-efficient parallel prefix sums

standard trick to improve work: compute small blocks sequentially

recall string matching!

1. Set  $b := \lceil \lg n \rceil$
2. For blocks of  $b$  consecutive indices, i. e.,  $A[0..b), A[b..2b), \dots$  do in parallel:
  - ▶ compute local prefix sums with fast **sequential** algorithm
3. Use previous work-inefficient parallel algorithm only on **rightmost elements** of blocks, i. e., to compute prefix sums of  $A[b - 1], A[2b - 1], A[3b - 1], \dots$
4. For blocks  $A[0..b), A[b..2b), \dots$  do in parallel:  
Add block-prefix sums to local prefix sums

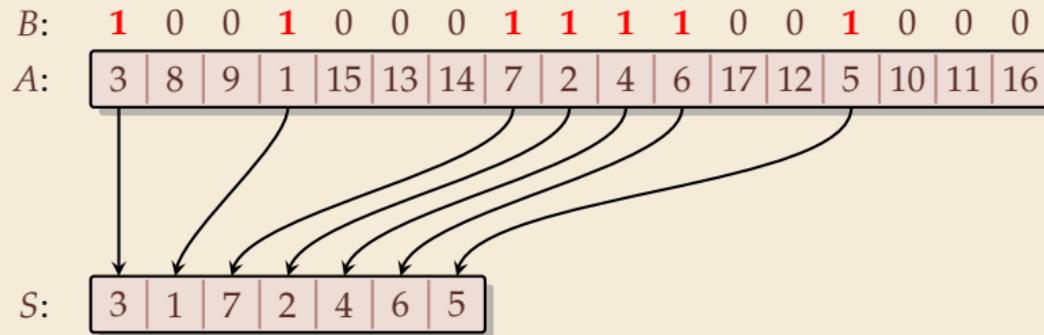
Analysis:

- ▶ Time:
  - ▶ 2. & 4.:  $\Theta(b) = \Theta(\log n)$  time
  - ▶ 3.  $\Theta(\log(n/b)) = \Theta(\log n)$  time
- ▶ Work:
  - ▶ 2. & 4.:  $\Theta(b)$  per block  $\times \lceil \frac{n}{b} \rceil$  blocks  $\rightsquigarrow \Theta(n)$
  - ▶ 3.  $\Theta\left(\frac{n}{b} \log\left(\frac{n}{b}\right)\right) = \Theta(n)$

# Compacting subsequences

How do prefix sums help with sorting? one more step to go ...

**Goal:** *Compact* a subsequence of an array



Use prefix sums on bitvector  $B$

~ offset of selected cells in  $S$

---

```
1 procedure compactArray(A[0..n], B[0..n])
2     C[0..n] := B[0..n] // deep copy of B
3     parallelPrefixSums(C)
4     for j := 0, ..., n - 1 do in parallel
5         if B[j] == 1 then S[C[j] - 1] := A[j]
6     end parallel for
```

---

## 10.4 Parallel Sorting

# Parallel Mergesort

- ▶ Recursive calls can run in parallel (data independent)!
- ▶ how about merging sorted halves  $A[l..m]$  and  $A[m..r]$ ?
- ▶ Our pointer-based sequential method seems hard to parallelize
  - ~~ Must treat all elements independently.
    - ▶ correct position of  $x$  in sorted output = *rank* of  $x$  breaking ties by position in  $A$
    - ▶  $\# \text{elements} \leq x = \# \text{elements from } A[l..m] \text{ that are } \leq x + \# \text{elements from } A[m..r] \text{ that are } \leq x$
  - ▶ rank in **own run** is simply the **index** of  $x$  in that run!
  - ▶ find rank in **other** run by *binary search*
  - ~~ can move  $x$  directly to correct position

# Parallel Mergesort – Code

---

```
1 procedure parMergesort( $A[l..r]$ ,  $buf$ ):  
2      $m := l + \lfloor (r - l)/2 \rfloor$   
3     in parallel { parMergesort( $A[l..m]$ ,  $buf$ ), parMergesort( $A[m..r]$ ,  $buf$ ) }  
4     parallelMerge( $A[l..m]$ ,  $A[m..r]$ ,  $buf$ )  
5     for  $i = l, \dots, r - 1$  do in parallel // copy back in parallel  
6          $A[i] := buf[i]$   
7     end parallel for  
8  
9 procedure parallelMerge( $A[l..m]$ ,  $A[m..r]$ ,  $buf$ ):  
10    for  $i = l, \dots, m - 1$  do in parallel  
11         $r := (i - l) + \text{binarySearch}(A[m..r], A[i])$  //  $\text{binarySearch}(A, x)$  returns #elements <  $x$  in  $A$   
12         $buf[r] = A[i]$   
13    end parallel for  
14    for  $j = m, \dots, r - 1$  do in parallel  
15         $r := \text{binarySearch}(A[l..m], A[j]) + (j - m)$   
16         $buf[r] = A[j]$   
17    end parallel for
```

---

# Parallel mergesort – Analysis

## ► Time:

- ▶ merge:  $\Theta(\log n)$  from binary search, rest  $O(1)$
- ▶ mergesort: depth of recursion tree is  $\Theta(\log n)$ 
  - ~~ total time  $O(\log^2(n))$

## ► Work:

- ▶ merge:  $n$  binary searches    ~~  $\Theta(n \log n)$ 
  - ~~ mergesort:  $O(n \log^2(n))$  work
- ▶ work can be reduced to  $\Theta(n)$  for merge (complicated!)
  - ▶ do full binary searches only for regularly sampled elements
  - ▶ ranks of remaining elements are sandwiched between sampled ranks
  - ▶ use a sequential method for small blocks, treat blocks in parallel
  - ▶ (details omitted)

# Parallel Quicksort

Let's try to parallelize Quicksort

- ▶ As for Mergesort, recursive calls can run in parallel ✓
- ▶ our sequential partitioning algorithm seems hard to parallelize
- ▶ but can split partitioning into *phases*:
  1. **comparisons:** compare all elements to pivot (in parallel), store result in bitvectors
  2. compute prefix sums of bit vectors (in parallel as above)
  3. **compact** subsequences of small and large elements (in parallel as above)

# Parallel Quicksort – Code

---

```
1 procedure parQuicksort( $A[l..r]$ ):
2      $b := \text{choosePivot}(A[l..r])$ 
3      $j := \text{parallelPartition}(A[l..r], b)$ 
4     in parallel { parQuicksort( $A[l..j]$ ), parQuicksort( $A[j + 1..r]$ ) }
5
6 procedure parallelPartition( $A[0..n]$ ,  $b$ ):
7     swap( $A[n - 1], A[b]$ );  $p := A[n - 1]$ 
8     for  $i = 0, \dots, n - 2$  do in parallel
9          $S[i] := [A[i] \leq p]$  //  $S[i]$  is 1 or 0
10         $L[i] := 1 - S[i]$ 
11    end parallel for
12    in parallel { parallelPrefixSum( $S[0..n - 2]$ ); parallelPrefixSum( $L[0..n - 2]$ ) }
13     $j := S[n - 2] + 1$ 
14    for  $i = 0, \dots, n - 2$  do in parallel
15         $x := A[i]$ 
16        if  $x \leq p$  then  $A[S[i] - 1] := x$ 
17        else  $A[j + L[i]] := x$ 
18    end parallel for
19     $A[j] := p$ 
20    return  $j$ 
```

---

# Parallel Quicksort – Analysis

## ► Time:

- ▶ partition: all  $O(1)$  time except prefix sums  $\rightsquigarrow \Theta(\log n)$  time
- ▶ Quicksort: expected depth of recursion tree is  $\Theta(\log n)$ 
  - $\rightsquigarrow$  total time  $O(\log^2(n))$  in expectation

## ► Work:

- ▶ partition:  $O(n)$  time except prefix sums  $\rightsquigarrow \Theta(n)$  work (with work-efficient prefix-sums algorithm)
  - $\rightsquigarrow$  Quicksort  $O(n \log(n))$  work in expectation
- ▶ (expected) work-efficient parallel sorting!

## Parallel sorting – State of the art

- ▶ more sophisticated methods can sort in  $O(\log n)$  parallel time on CREW-PRAM  
(very complicated algorithm based on parallel mergesort with interleaved merges)
  - ▶ practical challenge: small units of work add overhead
  - ▶ need a lot of PEs to see improvement from  $O(\log n)$  parallel time
- ~~ implementations tend to use simpler methods above
- ▶ check the Java library sources for interesting examples!  
`java.util.Arrays.parallelSort(int[])`