

# 6 String Matching – What's behind Ctrl+F?

18 November 2024

Prof. Dr. Sebastian Wild

# Learning Outcomes

## Unit 6: *String Matching*

1. Know and use typical notions for *strings* (substring, prefix, suffix, etc.).
2. Understand principles and implementation of the *KMP*, *BM*, and *RK* algorithms.
3. Know the *performance characteristics* of the KMP, BM, and RK algorithms.
4. Be able to solve simple *stringology problems* using the *KMP failure function*.

## 6 String Matching

- 6.1 String Notation
- 6.2 Brute Force
- 6.3 String Matching with Finite Automata
- 6.4 Constructing String Matching Automata
- 6.5 The Knuth-Morris-Pratt algorithm
- 6.6 Beyond Optimal? The Boyer-Moore Algorithm
- 6.7 The Rabin-Karp Algorithm

## 6.1 String Notation

# Ubiquitous strings

*string* = sequence of characters

- ▶ universal data type for ... everything!
  - ▶ natural language texts
  - ▶ programs (source code)
  - ▶ websites
  - ▶ XML documents
  - ▶ DNA sequences
  - ▶ bitstrings
  - ▶ ... a computer's memory  $\rightsquigarrow$  ultimately any data is a string

$\rightsquigarrow$  many different tasks and algorithms

# Ubiquitous strings

*string* = sequence of characters

- ▶ universal data type for ... everything!
  - ▶ natural language texts
  - ▶ programs (source code)
  - ▶ websites
  - ▶ XML documents
  - ▶ DNA sequences
  - ▶ bitstrings
  - ▶ ... a computer's memory ~→ ultimately any data is a string



~→ many different tasks and algorithms

- ▶ This unit: finding (exact) **occurrences of a pattern** text.
  - ▶ Ctrl+F
  - ▶ grep
  - ▶ computer forensics (e. g. find signature of file on disk)
  - ▶ virus scanner
- ▶ basis for many advanced applications

# Notations

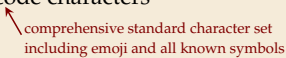
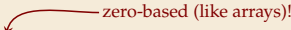
- ▶ *alphabet*  $\Sigma$ : finite set of allowed **characters**;  $\sigma = |\Sigma|$       “a string over alphabet  $\Sigma$ ”
    - ▶ letters (Latin, Greek, Arabic, Cyrillic, Asian scripts, ...)
    - ▶ “what you can type on a keyboard”,    Unicode characters
    - ▶  $\{0, 1\}$ ;    nucleotides  $\{A, C, G, T\}$ ; ...
- ↖ comprehensive standard character set  
including emoji and all known symbols

# Notations

- ▶ *alphabet*  $\Sigma$ : finite set of allowed **characters**;  $|\Sigma|$  “a string over alphabet  $\Sigma$ ”
  - ▶ letters (Latin, Greek, Arabic, Cyrillic, Asian scripts, ...)
  - ▶ “what you can type on a keyboard”, Unicode characters
  - ▶  $\{0, 1\}$ ; nucleotides  $\{A, C, G, T\}$ ; ...  
comprehensive standard character set including emoji and all known symbols
- ▶  $\Sigma^n = \Sigma \times \dots \times \Sigma$ : strings of **length**  $n \in \mathbb{N}_0$  ( $n$ -tuples)
- ▶  $\Sigma^* = \bigcup_{n \geq 0} \Sigma^n$ : set of **all** (finite) strings over  $\Sigma$
- ▶  $\Sigma^+ = \bigcup_{n \geq 1} \Sigma^n$ : set of **all** (finite) **nonempty** strings over  $\Sigma$
- ▶  $\varepsilon \in \Sigma^0$ : the *empty* string (same for all alphabets)  
  
 $\varepsilon$



# Notations

- ▶ *alphabet*  $\Sigma$ : finite set of allowed **characters**;  $\sigma = |\Sigma|$       “a string over alphabet  $\Sigma$ ”
  - ▶ letters (Latin, Greek, Arabic, Cyrillic, Asian scripts, ...)
  - ▶ “what you can type on a keyboard”, Unicode characters
  - ▶  $\{0, 1\}$ ; nucleotides  $\{A, C, G, T\}$ ; ...  
comprehensive standard character set including emoji and all known symbols
- ▶  $\Sigma^n = \Sigma \times \cdots \times \Sigma$ : strings of **length**  $n \in \mathbb{N}_0$  ( $n$ -tuples)
- ▶  $\Sigma^* = \bigcup_{n \geq 0} \Sigma^n$ : set of **all** (finite) strings over  $\Sigma$
- ▶  $\Sigma^+ = \bigcup_{n \geq 1} \Sigma^n$ : set of **all** (finite) **nonempty** strings over  $\Sigma$
- ▶  $\varepsilon \in \Sigma^0$ : the *empty* string (same for all alphabets)
- ▶ for  $S \in \Sigma^n$ , write  $S[i]$  (other sources:  $S_i$ ) for ***i*th** character       $(0 \leq i < n)$   
zero-based (like arrays!)
- ▶ for  $S, T \in \Sigma^*$ , write  $ST = S \cdot T$  for **concatenation** of  $S$  and  $T$
- ▶ for  $S \in \Sigma^n$ , write  $S[i..j]$  or  $S_{i,j}$  for the **substring**  $S[i] \cdot S[i+1] \cdots S[j]$        $(0 \leq i \leq j < n)$ 
  - ▶  $S[0..j]$  is a **prefix** of  $S$ ;  $S[i..n-1]$  is a **suffix** of  $S$
  - ▶  $S[i..j) = S[i..j-1]$  (endpoint exclusive)  $\rightsquigarrow S = S[0..n)$

## Clicker Question



True or false:

$$\bigcup_{n \geq 0} \Sigma^n = \Sigma^+ \cup \{\varepsilon\}$$

A

True

B

False



→ [sli.do/cs566](https://sli.do/cs566)

## Clicker Question



True or false:  $\Sigma^* = \Sigma^+ \cup \{\varepsilon\}$

A

True



B

~~False~~



→ [sli.do/cs566](https://sli.do/cs566)

# String matching – Definition

Search for a string (pattern) in a large body of text

## ► Input:

- $T \in \Sigma^n$ : The text (haystack) being searched within
- $P \in \Sigma^m$ : The pattern (needle) being searched for; typically  $n \gg m$

## ► Output:

- the first occurrence (match) of  $P$  in  $T$ :  $\min\{i \in [0..n - m) : T[i..i + m) = P\}$
  - or NO\_MATCH if there is no such  $i$  (“ $P$  does not occur in  $T$ ”)
- Variant: Find **all** occurrences of  $P$  in  $T$ .  
↪ Can do that iteratively (update  $T$  to  $T[i + 1..n)$  after match at  $i$ )

## ► Example:

- $T = \text{“Where is he?”}$
  - $P_1 = \text{“he”} \rightsquigarrow i = 1$
  - $P_2 = \text{“who”} \rightsquigarrow \text{NO\_MATCH}$
- string matching is implemented in Java in String.indexOf, in Python as `str.find`

## 6.2 Brute Force

# Abstract idea of algorithms

String matching algorithms typically use *guesses* and *checks*:

- ▶ A **guess** is a position  $i$  such that  $P$  might start at  $T[i]$ .  
Possible guesses (initially) are  $0 \leq i \leq n - m$ .
- ▶ A **check** of a guess is a comparison of  $T[i + j]$  to  $P[j]$ .

# Abstract idea of algorithms

String matching algorithms typically use *guesses* and *checks*:

- ▶ A **guess** is a position  $i$  such that  $P$  might start at  $T[i]$ .  
Possible guesses (initially) are  $0 \leq i \leq n - m$ .
- ▶ A **check** of a guess is a comparison of  $T[i + j]$  to  $P[j]$ .
- ▶ Note: need all  $m$  checks to verify a single *correct* guess  $i$ ,  
but it may take (many) fewer checks to recognize an *incorrect* guess.
- ▶ Cost measure: #character comparisons

$\rightsquigarrow$  #checks  $\leq n \cdot m$  (number of possible checks)

# Brute-force method

```
1 procedure bruteForceSM( $T[0..n]$ ,  $P[0..m]$ )  
2   for  $i := 0, \dots, n - m - 1$  do  
3     for  $j := 0, \dots, m - 1$  do  
4       if  $T[i + j] \neq P[j]$  then break inner loop  
5       if  $j == m$  then return  $i$   
6   return NO_MATCH
```

- try all guesses  $i$
- check each guess (left to right); stop early on mismatch
- essentially the implementation in Java!

► **Example:**

$T = \text{abbbababbab}$

$P = \text{abba}$

	a	b	b	b	a	b	a	b	b	a	b
a	a	b	b	<del>b</del>							



# Brute-force method

```
1 procedure bruteForceSM( $T[0..n]$ ,  $P[0..m]$ )  
2   for  $i := 0, \dots, n - m - 1$  do  
3     for  $j := 0, \dots, m - 1$  do  
4       if  $T[i + j] \neq P[j]$  then break inner loop  
5       if  $j == m$  then return  $i$   
6   return NO_MATCH
```

► try all guesses  $i$

► check each guess (left to right);  
stop early on mismatch

► essentially the implementation  
in Java!

► **Example:**

$T = \text{abbbababbab}$

$P = \text{abba}$

↪ 15 char cmps  
(vs  $n \cdot m = 44$ )  
not too bad!

	a	b	b	b	a	b	a	b	b	a	b
a	a	b	b	a							
		a									
			a								
				a							
					a	b	b				
						a					
							a	b	b	a	

# Brute-force method – Discussion



Brute-force method can be good enough

- ▶ typically works well for natural language text
- ▶ also for random strings



but: can be as bad as it gets!

a	a	a	a	a	a	a	a	a	a	a
a	a	a	b							
	a	a	a	b						
		a	a	a	b					
			a	a	a	b				
				a	a	a	b			
					a	a	a	b		
						a	a	a	b	
							a	a	a	b

▶ Worst possible input:  $P = a^{m-1}b$ ,  
 $T = a^n$

▶ Worst-case performance:  $(n - m + 1) \cdot m$

$\rightsquigarrow$  for  $m \leq n/2$  that is  $\Theta(mn)$

# Brute-force method – Discussion



Brute-force method can be good enough

- ▶ typically works well for natural language text
- ▶ also for random strings



but: can be as bad as it gets!

	a	a	a	a	a	a	a	a	a	a	
a	a	a	b								
	a	a	a	b							
		a	a	a	b						
			a	a	a	b					
				a	a	a	b				
					a	a	a	b			
						a	a	a	b		
							a	a	a	b	

▶ Worst possible input:  $P = a^{m-1}b$ ,  
 $T = a^n$

▶ Worst-case performance:  $(n - m + 1) \cdot m$

$\rightsquigarrow$  for  $m \leq n/2$  that is  $\Theta(mn)$

- ▶ Bad input: lots of self-similarity in  $T$ !  $\rightsquigarrow$  can we exploit that?
- ▶ brute force does 'obviously' stupid repetitive comparisons  $\rightsquigarrow$  can we avoid that?

# Roadmap

- ▶ **Approach 1** (this week): Use *preprocessing* on the **pattern**  $P$  to eliminate guesses (avoid 'obvious' redundant work)
  - ▶ Deterministic finite automata (DFA)
  - ▶ Knuth-Morris-Pratt algorithm
  - ▶ Boyer-Moore algorithm
  - ▶ Rabin-Karp algorithm
- ▶ **Approach 2** ( $\rightsquigarrow$  Unit 13): Do *preprocessing* on the **text**  $T$   
Can find matches in time *independent of text size(!)*
  - ▶ inverted indices
  - ▶ Suffix trees
  - ▶ Suffix arrays

## 6.3 String Matching with Finite Automata

## Clicker Question



Do you know what regular expressions, NFAs and DFAs are, and how to convert between them?

- ☐ A Never heard of this; are these new emoji?
- ☐ B Heard the terms, but don't remember conversion methods.
- ☐ C Had that in my undergrad course (memories fading a bit).
- ☐ D Sure, I could do that blindfolded!



→ *[sli.do/cs566](https://sli.do/cs566)*

# Theoretical Computer Science to the rescue!

► string matching = deciding whether  $T \in \Sigma^* \cdot P \cdot \Sigma^*$

►  $\Sigma^* \cdot P \cdot \Sigma^*$  is *regular* formal language

$\rightsquigarrow \exists$  *deterministic finite automaton* (DFA) to recognize  $\Sigma^* \cdot P \cdot \Sigma^*$

$\rightsquigarrow$  can check for occurrence of  $P$  in  $|T| = n$  steps!

# Theoretical Computer Science to the rescue!

► string matching = deciding whether  $T \in \Sigma^* \cdot P \cdot \Sigma^*$

►  $\Sigma^* \cdot P \cdot \Sigma^*$  is *regular* formal language

$\rightsquigarrow \exists$  *deterministic finite automaton* (DFA) to recognize  $\Sigma^* \cdot P \cdot \Sigma^*$

$\rightsquigarrow$  can check for occurrence of  $P$  in  $|T| = n$  steps!



Job done!



# Theoretical Computer Science to the rescue!

► string matching = deciding whether  $T \in \Sigma^* \cdot P \cdot \Sigma^*$

►  $\Sigma^* \cdot P \cdot \Sigma^*$  is *regular* formal language

$\rightsquigarrow \exists$  *deterministic finite automaton* (DFA) to recognize  $\Sigma^* \cdot P \cdot \Sigma^*$

$\rightsquigarrow$  can check for occurrence of  $P$  in  $|T| = n$  steps!



Job done!



WTF!?

# Theoretical Computer Science to the rescue!

► string matching = deciding whether  $T \in \Sigma^* \cdot P \cdot \Sigma^*$

►  $\Sigma^* \cdot P \cdot \Sigma^*$  is *regular* formal language

$\rightsquigarrow \exists$  *deterministic finite automaton* (DFA) to recognize  $\Sigma^* \cdot P \cdot \Sigma^*$

$\rightsquigarrow$  can check for occurrence of  $P$  in  $|T| = n$  steps!



Job done!



WTF!?

We are not quite done yet.

- (Problem 0: programmer might not know automata and formal languages ...)
- Problem 1: existence alone does not give an algorithm!
- Problem 2: automaton could be very big!


# String matching with DFA

- Assume first, we already have a deterministic automaton
- How does string matching work?

**Example:**

$T = \text{aabacaababacaa}$

$P = \underline{\text{ababaca}}$



text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3	0	1	1	2	3	4	5	6	7	7


# String matching with DFA

- Assume first, we already have a deterministic automaton
- How does string matching work?

**Example:**

$T = \text{aabacaababacaa}$

$P = \text{ababaca}$



	c	a	b	c
q				
0		1	0	0
1				
2				

text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3	0	1	1	2	3	4	5	6	7	7


# String matching DFA – Intuition

Why does this work?

► Main insight:

State  $q$  means:  
*“we have seen  $P[0..q)$  until here  
 (but not any longer prefix of  $P$ )”*


$T = \text{aabacaababacaa}$   
 $P = \text{ababaca}$



text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3	0	1	1	2	3	4	5	6	7	7

► If the next text character  $c$  does not match, we know:

- (i) text seen so far ends with  $P[0...q) \cdot c$
- (ii)  $P[0...q) \cdot c$  is not a prefix of  $P$
- (iii) without reading  $c$ ,  $P[0..q)$  was the *longest* prefix of  $P$  that ends here.




# String matching DFA – Intuition

Why does this work?

► Main insight:

State  $q$  means:  
*“we have seen  $P[0..q)$  until here  
 (but not any longer prefix of  $P$ )”*


$T = \text{aabacaababacaa}$   
 $P = \text{ababaca}$



text:		a	a	b	a	c	a	a	b	a	b	a	c	a	a
state:	0	1	1	2	3	0	1	1	2	3	4	5	6	7	7

► If the next text character  $c$  does not match, we know:

- (i) text seen so far ends with  $P[0...q) \cdot c$
- (ii)  $P[0...q) \cdot c$  is not a prefix of  $P$
- (iii) without reading  $c$ ,  $P[0..q)$  was the *longest* prefix of  $P$  that ends here.



↪ New longest matched prefix will be (weakly) shorter than  $q$

↪ All information about the text needed to determine it is contained in  $P[0...q) \cdot c$ !