

10 Parallel Algorithms

12 January 2026

Prof. Dr. Sebastian Wild

Learning Outcomes

Unit 10: *Parallel Algorithms*

1. Know and apply *parallelization strategies* for embarrassingly parallel problems.
2. Identify *limits of parallel speedups*.
3. Understand and use the *parallel random-access-machine* model in its different variants.
4. Be able to *analyze* and compare simple shared-memory parallel algorithms by determining *parallel time and work*.
5. Understand efficient parallel *prefix sum* algorithms.
6. Be able to devise high-level description of *parallel quicksort and mergesort* methods.

Outline

10 Parallel Algorithms

10.1 Parallel Computation

10.2 Parallel String Matching

10.3 Parallel Primitives

10.4 Parallel Sorting

10.1 Parallel Computation

Clicker Question



Have you ever written a concurrent program (explicit threads, job pools library, or using a framework for distributed computing)?

- A Yes
- B No
- C Concur... what?



→ *sli.do/cs566*

Types of parallel computation

€€€ can't buy you more time . . . but more computers!

↝ Challenge: Algorithms for *parallel* computation.

Types of parallel computation

€€€ can't buy you more time . . . but more computers!

↝ Challenge: Algorithms for *parallel* computation.

There are two main forms of parallelism:

1. shared-memory parallel computer ← *focus of today*

- ▶ p *processing elements* (PEs, processors) working in parallel
- ▶ single big memory, **accessible from every PE**
- ▶ communication via shared memory
- ▶ think: a big server, 128 CPU cores, terabyte of main memory

2. distributed computing

- ▶ p PEs working in parallel
- ▶ each PE has **private** memory
- ▶ communication by sending **messages** via a network
- ▶ think: a cluster of individual machines

PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 2)
- ▶ the p PEs are identified by ids $0, \dots, p - 1$
 - ▶ like w (the word size), p is a parameter of the model that can grow with n
 - ▶ $p = \Theta(n)$ is not unusual many processors!
- ▶ the PEs all **independently** run the same RAM-style program
(they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:
in each time step, every PE executes one instruction

PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 2)
- ▶ the p PEs are identified by ids $0, \dots, p - 1$
 - ▶ like w (the word size), p is a parameter of the model that can grow with n
 - ▶ $p = \Theta(n)$ is not unusual maaany processors!
- ▶ the PEs all **independently** run the same RAM-style program
(they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:
in each time step, every PE executes one instruction
- ▶ As for RAM:
 - ▶ assume a basic “operating system”
 - ~~ write algorithms in pseudocode instead of RAM assembly
 - ▶ **NEW:** loops and commands can be run “**in parallel**” (examples coming up)

PRAM – Conflict management



Problem: What if several PEs simultaneously overwrite a memory cell?

- ▶ **EREW-PRAM** (exclusive read, exclusive write)
any **parallel access** to same memory cell is **forbidden** (crash if happens)
- ▶ **CREW-PRAM** (concurrent read, exclusive write)
parallel **write** access to same memory cell is **forbidden**, *but reading is fine*
- ▶ **CRCW-PRAM** (concurrent read, concurrent write)
concurrent access is allowed,
need a rule for write conflicts:
 - ▶ common CRCW-PRAM:
all concurrent writes to same cell must write **same** value
 - ▶ arbitrary CRCW-PRAM:
some unspecified concurrent write wins
 - ▶ (more exist ...)
- ▶ no single model is always adequate, but our default is CREW

PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish assuming sufficiently many PEs!
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish assuming sufficiently many PEs!
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

Holy grail of PRAM algorithms:

- ▶ minimal time (=span)
- ▶ work (asymptotically) no worse than running time of best sequential algorithm
~~ “*work-efficient*” algorithm: work in same Θ -class as best sequential

Clicker Question



Does every computational problem allow a work-efficient algorithm?

A

Yes

B

No



→ *sli.do/cs566*

Clicker Question



Does every computational problem allow a work-efficient algorithm?

A

Yes ✓

B

No



→ *sli.do/cs566*

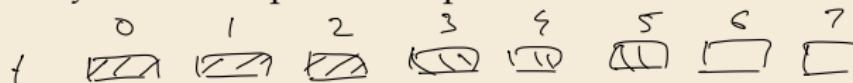
The number of processors

Hold on, my computer does not have $\Theta(n)$ processors! Why should I care for span and work!?

Theorem 10.1 (Brent's Theorem)

If an algorithm has span T and work W (for an arbitrarily large number of processors), it can be run on a PRAM with p PEs in time $O(T + \frac{W}{p})$ (and using $O(W)$ work). 

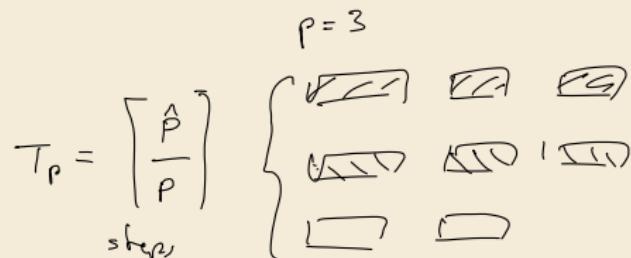
Proof: schedule parallel steps in round-robin fashion on the p PEs.



$$\hat{p} = \# \text{ PEs in PRAM}$$

$$\#\text{rounds} = \Theta(T + \frac{W}{\hat{p}})$$

$$T \cdot \left\lceil \frac{W}{T_p} \right\rceil \leq T \left(\frac{W}{T_p} + 1 \right) = \frac{W}{p} + T$$



~ span and work give guideline for *any* number of processors

10.2 Parallel String Matching

Embarrassingly Parallel

- ▶ A problem is called "*embarrassingly parallel*"
if it can immediately be split into *many, small subtasks*
that can be solved completely *independently* of each other
- ▶ Typical example: sum of two large matrices (all entries independent)
 - ~~ best case for parallel computation (simply assign each processor one subtask)
- ▶ Sorting is not embarrassingly parallel
 - ▶ no obvious way to define many *small* (= efficiently solvable) subproblems
 - ▶ but: some subtasks of our algorithms are (stay tuned ...)

Clicker Question

Is the string-matching problem “embarrassingly parallel”?



- A Yes
- B No
- C Only for $n \gg m$
- D Only for $n \approx m$



→ *sli.do/cs566*

Parallel string matching – Easy?

- ▶ We have seen a plethora of string matching methods in Unit 6
- ▶ But all efficient methods seem inherently sequential
Indeed, they became efficient only after building on knowledge from previous steps!
Sounds like the *opposite* of parallel!

~~ How well can we parallelize string matching?

Here: string matching = find *all* occurrences of P in T (more natural problem for parallel)
always assume $m \leq n$

Parallel string matching – Easy?

- ▶ We have seen a plethora of string matching methods in Unit 6
- ▶ But all efficient methods seem inherently sequential

Indeed, they became efficient only after building on knowledge from previous steps!

Sounds like the *opposite* of parallel!


- ~~ How well can we parallelize string matching?

Here: string matching = find *all* occurrences of P in T (more natural problem for parallel)
always assume $m \leq n$

Subproblems in string matching:

- ▶ string matching = check all guesses $i = 0, \dots, n - m - 1$
- ▶ checking one guess is a subtask!

Parallel string matching – Brute force

- ▶ Check all guesses in parallel

```
1 procedure parallelBruteForce( $T[0..n], P[0..m]$ ):  
2   for  $i := 0, \dots, n - m - 1$  do in parallel ← only difference to normal brute force!  
3     for  $j := 0, \dots, m - 1$  do  
4       if  $T[i + j] \neq P[j]$  then break inner loop  
5       if  $j == m$  then report match at  $i$   
6   end parallel for
```

- ▶ PE k is executing the loop iteration where $i = k$.
 - ~~ requires that all iterations can be done **independently**!
 - ▶ Different PEs work **in lockstep** (synchronized after each instruction)
 - ▶ similar to OpenMP `#pragma omp parallel for`
- ▶ checking whether *no* match was found by *any* PE more effort ~~ ... stay tuned

Parallel string matching – Brute force

- ▶ Check all guesses in parallel

```
1 procedure parallelBruteForce( $T[0..n], P[0..m]$ ):  
2   for  $i := 0, \dots, n - m - 1$  do in parallel ← only difference to normal brute force!  
3     for  $j := 0, \dots, m - 1$  do  
4       if  $T[i + j] \neq P[j]$  then break inner loop  
5       if  $j == m$  then report match at  $i$   
6   end parallel for
```

- ▶ PE k is executing the loop iteration where $i = k$.
 - ~~ requires that all iterations can be done **independently**!
 - ▶ Different PEs work **in lockstep** (synchronized after each instruction)
 - ▶ similar to OpenMP `#pragma omp parallel for`
- ▶ checking whether *no* match was found by *any* PE more effort ~~ ... stay tuned

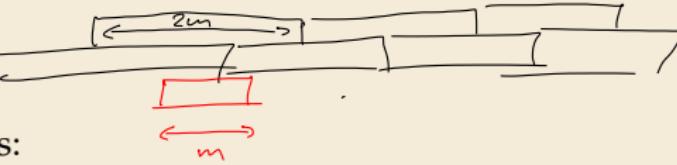
~~ Time: $\Theta(m)$ using sequential checks
 $\Theta(\log m)$ on CREW-PRAM (~~ tutorials)
 $\Theta(1)$ on CRCW-PRAM (~~ tutorials)

Work: $\Theta((n - m)m)$ ~~ not great
... much more than best sequential

Parallel string matching – Blocking



Divide T into **overlapping** blocks of $2m - 1$ characters:
 $T[0..2m - 1], T[m..3m - 1], T[2m..4m - 1], T[3m..5m - 1] \dots$



- ▶ Search all blocks in parallel, each using efficient *sequential* method

```
1 procedure blockingStringMatching( $T[0..n], P[0..m]$ ):  
2     for  $b := 0, \dots, \lceil n/m \rceil$  do in parallel  
3         result := KMP( $T[bm .. (b+1)m - 1], P$ )  
4         if result  $\neq$  NO_MATCH then report match at result  
5     end parallel for
```

Parallel string matching – Blocking



Divide T into **overlapping** blocks of $2m - 1$ characters:
 $T[0..2m - 1], T[m..3m - 1], T[2m..4m - 1], T[3m..5m - 1] \dots$

- ▶ Search all blocks in parallel, each using efficient *sequential* method

```
1 procedure blockingStringMatching( $T[0..n], P[0..m]$ ):  
2     for  $b := 0, \dots, \lceil n/m \rceil$  do in parallel  
3         result := KMP( $T[bm .. (b+1)m - 1], P$ )  
4         if result ≠ NO_MATCH then report match at result |  $\Theta(m)$   
5     end parallel for
```

~~ Time:

- ▶ loop body has text of length $n' = 2m - 1$ and pattern of length m

~~ KPM runtime $\Theta(n' + m) = \Theta(m)$.

~~ Work: $\Theta(\frac{n}{m} \cdot m) = \Theta(n)$ ~~ work efficient!

Clicker Question

Is the string-matching problem “embarrassingly parallel”?



- A** Yes
- B** No
- C** Only for $n \gg m$
- D** Only for $n \approx m$



→ *sli.do/cs566*

Clicker Question

Is the string-matching problem “embarrassingly parallel”?



- A Yes
- B No
- C Only for $n \gg m$ ✓
- D Only for ~~$n \approx m$~~



→ *sli.do/cs566*

Parallel string matching – Discussion

- 👍 very simple methods
- 👍 could even run distributed with access to part of T
- 👎 parallel speedup only for $m \ll n$

Parallel string matching – Discussion

-  very simple methods
-  could even run distributed with access to part of T

-  parallel speedup only for $m \ll n$

► work-efficient methods with better parallel time possible?

- ~~ must genuinely parallelize the matching process! (and the preprocessing of the pattern)
- ~~ needs new ideas (much more complicated, but possible!)

Parallel string matching – Discussion

- ◀ very simple methods
- ◀ could even run distributed with access to part of T
- ▶ parallel speedup only for $m \ll n$

- ▶ work-efficient methods with better parallel time possible?
 - ~~ must genuinely parallelize the matching process! (and the preprocessing of the pattern)
 - ~~ needs new ideas (much more complicated, but possible!)
- ▶ Parallel string matching – State of the art:
 - ▶ $O(\log m)$ time & work-efficient parallel string matching (very complicated)
 - ▶ this is optimal for CREW-PRAM
 - ▶ on CRCW-PRAM: matching part even in $O(1)$ time (easy)
but preprocessing requires $\Theta(\log \log m)$ time (very complicated)