

10

Parallel Algorithms

17 December 2024

Prof. Dr. Sebastian Wild

Learning Outcomes

Unit 10: *Parallel Algorithms*

1. Know and apply *parallelization strategies* for embarrassingly parallel problems.
2. Identify *limits of parallel speedups*.
3. Understand and use the *parallel random-access-machine* model in its different variants.
4. Be able to *analyze* and compare simple shared-memory parallel algorithms by determining *parallel time and work*.
5. Understand efficient parallel *prefix sum* algorithms.
6. Be able to devise high-level description of *parallel quicksort and mergesort* methods.

10 Parallel Algorithms

10.1 Parallel Computation

10.2 Parallel String Matching

10.3 Parallel Primitives

10.4 Parallel Sorting

10.1 Parallel Computation

Clicker Question



Have you ever written a concurrent program (explicit threads, job pools library, or using a framework for distributed computing)?

- ☐ A Yes
- ☐ B No
- ☐ C Concur... what?



→ *sli.do/cs566*

Types of parallel computation

£££ can't buy you more time . . . but more computers!

↪ Challenge: Algorithms for *parallel* computation.

Types of parallel computation

£££ can't buy you more time . . . but more computers!

↪ Challenge: Algorithms for *parallel* computation.

There are two main forms of parallelism:

1. shared-memory parallel computer ← *focus of today*

- ▶ *p processing elements* (PEs, processors) working in parallel
- ▶ **single** big memory, **accessible from every PE**
- ▶ communication via shared memory
- ▶ think: a big server, 128 CPU cores, terabyte of main memory

2. distributed computing

- ▶ *p* PEs working in parallel
- ▶ each PE has **private** memory
- ▶ communication by sending **messages** via a network
- ▶ think: a cluster of individual machines

PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 2)
- ▶ the p PEs are identified by ids $0, \dots, p - 1$
 - ▶ like w (the word size), p is a parameter of the model that can grow with n
 - ▶ $p = \Theta(n)$ is not unusual maaany processors!
- ▶ the PEs all **independently** run the same RAM-style program (they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:
in each time step, every PE executes one instruction

PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 2)
- ▶ the p PEs are identified by ids $0, \dots, p - 1$
 - ▶ like w (the word size), p is a parameter of the model that can grow with n
 - ▶ $p = \Theta(n)$ is not unusual maaany processors!
- ▶ the PEs all **independently** run the same RAM-style program (they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:
in each time step, every PE executes one instruction
- ▶ As for RAM:
 - ▶ assume a basic “operating system”
 - ↪ write algorithms in pseudocode instead of RAM assembly
 - ▶ **NEW:** loops and commands can be run “**in parallel**” (examples coming up)

PRAM – Conflict management



Problem: What if several PEs simultaneously overwrite a memory cell?

- ▶ EREW-PRAM (exclusive read, exclusive write)
any **parallel access** to same memory cell is **forbidden** (crash if happens)
- ▶ CREW-PRAM (concurrent read, exclusive write)
parallel **write** access to same memory cell is **forbidden**, *but reading is fine*
- ▶ CRCW-PRAM (concurrent read, concurrent write)
concurrent access is allowed,
need a rule for write conflicts:
 - ▶ common CRCW-PRAM:
all concurrent writes to same cell must write **same** value
 - ▶ arbitrary CRCW-PRAM:
some unspecified concurrent write wins
 - ▶ (more exist ...)
- ▶ no single model is always adequate, but our default is CREW

PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish assuming sufficiently many PEs!
sometimes called depth or span
- ▶ **work:** total #instructions executed on **all** PEs

PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish assuming sufficiently many PEs!
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

Holy grail of PRAM algorithms:

- ▶ minimal time (=span)
- ▶ work (asymptotically) no worse than running time of best sequential algorithm
 \rightsquigarrow “*work-efficient*” algorithm: work in same Θ -class as best sequential

Clicker Question



Does every computational problem allow a work-efficient algorithm?

☐ A Yes

☐ B No



→ sli.do/cs566

Clicker Question



Does every computational problem allow a work-efficient algorithm?

A Yes ✓

B ~~No~~



→ *sl.i.do/cs566*


The number of processors

Hold on, my computer does not have $\Theta(n)$ processors! Why should I care for span and work!?

Theorem 10.1 (Brent's Theorem)

If an algorithm has span T and work W (for an arbitrarily large number of processors), it can be run on a PRAM with p PEs in time $O(T + \frac{W}{p})$ (and using $O(W)$ work). ◀

Proof: schedule parallel steps in round-robin fashion on the p PEs.



↪ span and work give guideline for *any* number of processors