

# 10 Parallel Algorithms

12 January 2026

Prof. Dr. Sebastian Wild

# Learning Outcomes

## Unit 10: *Parallel Algorithms*

1. Know and apply *parallelization strategies* for embarrassingly parallel problems.
2. Identify *limits of parallel speedups*.
3. Understand and use the *parallel random-access-machine* model in its different variants.
4. Be able to *analyze* and compare simple shared-memory parallel algorithms by determining *parallel time and work*.
5. Understand efficient parallel *prefix sum* algorithms.
6. Be able to devise high-level description of *parallel quicksort and mergesort* methods.

## Outline

# 10 Parallel Algorithms

10.1 Parallel Computation

10.2 Parallel String Matching

10.3 Parallel Primitives

10.4 Parallel Sorting

## 10.1 Parallel Computation

## Clicker Question



Have you ever written a concurrent program (explicit threads, job pools library, or using a framework for distributed computing)?

- A** Yes
- B** No
- C** Concur... what?



→ *sli.do/cs566*

# Types of parallel computation

€€€ can't buy you more time . . . but more computers!

↝ Challenge: Algorithms for *parallel* computation.

# Types of parallel computation

€€€ can't buy you more time . . . but more computers!

↝ Challenge: Algorithms for *parallel* computation.

There are two main forms of parallelism:

## 1. shared-memory parallel computer $\leftarrow$ focus of today

- ▶  $p$  *processing elements* (PEs, processors) working in parallel
- ▶ single big memory, **accessible from every PE**
- ▶ communication via shared memory
- ▶ think: a big server, 128 CPU cores, terabyte of main memory

## 2. distributed computing

- ▶  $p$  PEs working in parallel
- ▶ each PE has **private** memory
- ▶ communication by sending **messages** via a network
- ▶ think: a cluster of individual machines

# PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 2)
- ▶ the  $p$  PEs are identified by ids  $0, \dots, p - 1$ 
  - ▶ like  $w$  (the word size),  $p$  is a parameter of the model that can grow with  $n$
  - ▶  $p = \Theta(n)$  is not unusual      maaany processors!
- ▶ the PEs all **independently** run the same RAM-style program  
(they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:  
in each time step, every PE executes one instruction

# PRAM – Parallel RAM

- ▶ extension of the RAM model (recall Unit 2)
- ▶ the  $p$  PEs are identified by ids  $0, \dots, p - 1$ 
  - ▶ like  $w$  (the word size),  $p$  is a parameter of the model that can grow with  $n$
  - ▶  $p = \Theta(n)$  is not unusual      many processors!
- ▶ the PEs all **independently** run the same RAM-style program  
(they can use their id there)
- ▶ each PE has its own registers, but **MEM** is shared among all PEs
- ▶ computation runs in **synchronous** steps:  
in each time step, every PE executes one instruction
- ▶ As for RAM:
  - ▶ assume a basic “operating system”
  - ~~ write algorithms in pseudocode instead of RAM assembly
  - ▶ **NEW:** loops and commands can be run “**in parallel**” (examples coming up)

# PRAM – Conflict management



**Problem:** What if several PEs simultaneously overwrite a memory cell?

- ▶ **EREW-PRAM** (exclusive read, exclusive write)  
any **parallel access** to same memory cell is **forbidden** (crash if happens)
- ▶ **CREW-PRAM** (concurrent read, exclusive write)  
parallel **write** access to same memory cell is **forbidden**, *but reading is fine*
- ▶ **CRCW-PRAM** (concurrent read, concurrent write)  
concurrent access is allowed,  
need a rule for write conflicts:
  - ▶ common CRCW-PRAM:  
all concurrent writes to same cell must write **same** value
  - ▶ arbitrary CRCW-PRAM:  
some unspecified concurrent write wins
  - ▶ (more exist ...)
- ▶ no single model is always adequate, but our default is CREW

# PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish      assuming sufficiently many PEs!  
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

# PRAM – Execution costs

Cost metrics in PRAMs

- ▶ **space:** total amount of accessed memory
- ▶ **time:** number of steps till all PEs finish      assuming sufficiently many PEs!  
sometimes called *depth* or *span*
- ▶ **work:** total #instructions executed on **all** PEs

Holy grail of PRAM algorithms:

- ▶ minimal time (=span)
- ▶ work (asymptotically) no worse than running time of best sequential algorithm  
~~ “*work-efficient*” algorithm: work in same  $\Theta$ -class as best sequential

# Clicker Question



Does every computational problem allow a work-efficient algorithm?

**A**

Yes

**B**

No



→ *sli.do/cs566*

# Clicker Question



Does every computational problem allow a work-efficient algorithm?

**A**

Yes ✓

**B**

✗



→ *sli.do/cs566*


# The number of processors

Hold on, my computer does not have  $\Theta(n)$  processors! Why should I care for span and work!?

## Theorem 10.1 (Brent's Theorem)

If an algorithm has span  $T$  and work  $W$  (for an arbitrarily large number of processors), it can be run on a PRAM with  $p$  PEs in time  $O(T + \frac{W}{p})$  (and using  $O(W)$  work).  $\blacktriangleleft$

Proof: schedule parallel steps in round-robin fashion on the  $p$  PEs.



$$\# \text{ rounds} = \Theta(T + \frac{W}{p})$$

$$T \cdot \left\lceil \frac{W}{T_p} \right\rceil \leq T \left( \frac{W}{T_p} + 1 \right) = \frac{W}{p} + T$$

$\rightsquigarrow$  span and work give guideline for *any* number of processors

## 10.2 Parallel String Matching

# Embarrassingly Parallel

- ▶ A problem is called "*embarrassingly parallel*"  
if it can immediately be split into *many, small subtasks*  
that can be solved completely *independently* of each other
- ▶ Typical example: sum of two large matrices (all entries independent)
  - ~~ best case for parallel computation (simply assign each processor one subtask)
- ▶ Sorting is not embarrassingly parallel
  - ▶ no obvious way to define many *small* (= efficiently solvable) subproblems
  - ▶ but: some subtasks of our algorithms are (stay tuned ...)

# Clicker Question

Is the string-matching problem “embarrassingly parallel”?



- A** Yes
- B** No
- C** Only for  $n \gg m$
- D** Only for  $n \approx m$



→ *sli.do/cs566*

# Parallel string matching – Easy?

- ▶ We have seen a plethora of string matching methods in Unit 6
- ▶ But all efficient methods seem inherently sequential  
*Indeed, they became efficient only after building on knowledge from previous steps!*  
  
Sounds like the *opposite* of parallel!

- ~~ How well can we parallelize string matching?

Here: string matching = find *all* occurrences of  $P$  in  $T$       (more natural problem for parallel)  
always assume  $m \leq n$

# Parallel string matching – Easy?

- ▶ We have seen a plethora of string matching methods in Unit 6
- ▶ But all efficient methods seem inherently sequential  
*Indeed, they became efficient only after building on knowledge from previous steps!*  
  
Sounds like the *opposite* of parallel!

- ~~ How well can we parallelize string matching?

Here: string matching = find *all* occurrences of  $P$  in  $T$  (more natural problem for parallel)  
always assume  $m \leq n$

## Subproblems in string matching:

- ▶ string matching = check all guesses  $i = 0, \dots, n - m - 1$
- ▶ checking one guess is a subtask!

# Parallel string matching – Brute force

- ▶ Check all guesses in parallel

---

```
1 procedure parallelBruteForce( $T[0..n], P[0..m]$ ):  
2   for  $i := 0, \dots, n - m - 1$  do in parallel ← only difference to normal brute force!  
3     for  $j := 0, \dots, m - 1$  do  
4       if  $T[i + j] \neq P[j]$  then break inner loop  
5       if  $j == m$  then report match at  $i$   
6   end parallel for
```

---

- ▶ PE  $k$  is executing the loop iteration where  $i = k$ .
  - ~~ requires that all iterations can be done **independently**!
  - ▶ Different PEs work **in lockstep** (synchronized after each instruction)
  - ▶ similar to OpenMP `#pragma omp parallel for`
- ▶ checking whether *no* match was found by *any* PE more effort ~~ ... stay tuned

# Parallel string matching – Brute force

- ▶ Check all guesses in parallel

---

```
1 procedure parallelBruteForce( $T[0..n], P[0..m]$ ):  
2   for  $i := 0, \dots, n - m - 1$  do in parallel ← only difference to normal brute force!  
3     for  $j := 0, \dots, m - 1$  do  
4       if  $T[i + j] \neq P[j]$  then break inner loop  
5       if  $j == m$  then report match at  $i$   
6   end parallel for
```

---

- ▶ PE  $k$  is executing the loop iteration where  $i = k$ .
  - ~~ requires that all iterations can be done **independently**!
  - ▶ Different PEs work **in lockstep** (synchronized after each instruction)
  - ▶ similar to OpenMP `#pragma omp parallel for`
- ▶ checking whether *no* match was found by *any* PE more effort ~~ ... stay tuned


~~ **Time:**  $\Theta(m)$  using sequential checks  
 $\Theta(\log m)$  on CREW-PRAM (~~ tutorials)  
 $\Theta(1)$  on CRCW-PRAM (~~ tutorials)

**Work:**  $\Theta((n - m)m)$  ~~ not great  
... much more than best sequential

# Parallel string matching – Blocking



Divide  $T$  into **overlapping** blocks of  $2m - 1$  characters:  
 $T[0..2m - 1], T[m..3m - 1], T[2m..4m - 1], T[3m..5m - 1] \dots$



- ▶ Search all blocks in parallel, each using efficient *sequential* method

---

```
1 procedure blockingStringMatching( $T[0..n], P[0..m]$ ):  
2     for  $b := 0, \dots, \lceil n/m \rceil$  do in parallel  
3         result := KMP( $T[bm .. (b+1)m - 1], P$ )  
4         if result  $\neq$  NO_MATCH then report match at result  
5     end parallel for
```

---

# Parallel string matching – Blocking



Divide  $T$  into **overlapping** blocks of  $2m - 1$  characters:  
 $T[0..2m - 1], T[m..3m - 1], T[2m..4m - 1], T[3m..5m - 1] \dots$

- ▶ Search all blocks in parallel, each using efficient *sequential* method

---

```
1 procedure blockingStringMatching( $T[0..n], P[0..m]$ ):
2     for  $b := 0, \dots, \lceil n/m \rceil$  do in parallel
3         result := KMP( $T[bm .. (b+1)m - 1], P$ )           |  $\Theta(m)$ 
4         if result  $\neq$  NO_MATCH then report match at result
5     end parallel for
```

---

~~ Time:

- ▶ loop body has text of length  $n' = 2m - 1$  and pattern of length  $m$

~~ KPM runtime  $\Theta(n' + m) = \Theta(m)$

~~ Work:  $\Theta(\frac{n}{m} \cdot m) = \Theta(n)$  ~~ work efficient!

# Clicker Question

Is the string-matching problem “embarrassingly parallel”?



- A** Yes
- B** No
- C** Only for  $n \gg m$
- D** Only for  $n \approx m$



→ *sli.do/cs566*

# Clicker Question

Is the string-matching problem “embarrassingly parallel”?



- A Yes
- B No
- C Only for  $n \gg m$  ✓
- D Only for  ~~$n \approx m$~~



→ *sli.do/cs566*

## Parallel string matching – Discussion

- 👍 very simple methods
- 👍 could even run distributed with access to part of  $T$
- 👎 parallel speedup only for  $m \ll n$

## Parallel string matching – Discussion

- thumb up very simple methods
- thumb up could even run distributed with access to part of  $T$
- thumb down parallel speedup only for  $m \ll n$

► work-efficient methods with better parallel time possible?

- ~ must genuinely parallelize the matching process! (and the preprocessing of the pattern)
- ~ needs new ideas (much more complicated, but possible!)

# Parallel string matching – Discussion

- 👍 very simple methods
- 👍 could even run distributed with access to part of  $T$
- 👎 parallel speedup only for  $m \ll n$

- ▶ work-efficient methods with better parallel time possible?
  - ~~ must genuinely parallelize the matching process! (and the preprocessing of the pattern)
  - ~~ needs new ideas (much more complicated, but possible!)
- ▶ Parallel string matching – State of the art: *notin exam*
  - ▶  $O(\log m)$  time & work-efficient parallel string matching (very complicated)
    - ▶ this is optimal for CREW-PRAM
  - ▶ on CRCW-PRAM: matching part even in  $O(1)$  time (easy)  
but preprocessing requires  $\Theta(\log \log m)$  time (very complicated)

## 10.3 Parallel Primitives

# Building blocks



- ▶ Most nontrivial problems need tricks to be parallelized
  - ▶ Some versatile building blocks are known that help in many problems
- ~~ We study some of them now, before we apply them to *parallel sorting*

*The following problems might not look natural at first sight . . . but turn out to be good abstractions.*

~~ bear with me

# Prefix sums

**Prefix-sum problem** (also: cumulative sums, running totals)

- ▶ Given: array  $A[0..n]$  of numbers
- ▶ Goal: compute all prefix sums  $A[0] + \dots + A[i]$  for  $i = 0, \dots, n - 1$   
may be done “in-place”, i. e., by overwriting  $A$

**Example:**

input:

3	0	0	5	7	0	0	2	0	0	0	4	0	8	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

$\Sigma$

output:

3	3	3	8	15	15	15	17	17	17	17	21	21	29	29	30
---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----

# Clicker Question

What is the *sequential* running time achievable for prefix sums?



**A**  $O(n^3)$

**D**  $O(n)$

**B**  $O(n^2)$

**E**  $O(\sqrt{n})$

**C**  $O(n \log n)$

**F**  $O(\log n)$



→ *sli.do/cs566*

# Clicker Question

What is the *sequential* running time achievable for prefix sums?



- A**  ~~$\Theta(n^3)$~~
- B**  ~~$\Theta(n^2)$~~
- C**  ~~$\Theta(n \log n)$~~

- D**  $\Theta(n)$  ✓
- E**  ~~$\Theta(\sqrt{n})$~~
- F**  ~~$\Theta(\log n)$~~



→ *sli.do/cs566*

## Prefix sums – Sequential

- ▶ sequential solution does  $n - 1$  additions
- ▶ but: cannot parallelize them!
  - ⚡ data dependencies!
- ~~ need a different approach

---

```
1 procedure prefixSum(A[0..n)):  
2     for i := 1, ..., n - 1 do  
3         A[i] := A[i - 1] + A[i]
```

---

# Prefix sums – Sequential

- ▶ sequential solution does  $n - 1$  additions
- ▶ but: cannot parallelize them!
  - ⚡ data dependencies!
- ↝ need a different approach

---

```
1 procedure prefixSum(A[0..n)):  
2     for i := 1, ..., n - 1 do  
3         A[i] := A[i - 1] + A[i]
```

---

Let's try a simpler problem first.

## Excursion: Sum

- ▶ Given: array  $A[0..n)$  of numbers
- ▶ Goal: compute  $A[0] + A[1] + \dots + A[n - 1]$   
(solved by prefix sums)

# Prefix sums – Sequential

- ▶ sequential solution does  $n - 1$  additions
- ▶ but: cannot parallelize them!
  - ⚡ data dependencies!
- ↝ need a different approach

---

```
1 procedure prefixSum(A[0..n)):  
2     for i := 1, ..., n - 1 do  
3         A[i] := A[i - 1] + A[i]
```

---

Let's try a simpler problem first.

## Excursion: Sum

- ▶ Given: array  $A[0..n)$  of numbers
- ▶ Goal: compute  $A[0] + A[1] + \dots + A[n - 1]$   
(solved by prefix sums)



Any algorithm *must* do  $n - 1$  binary additions

# Prefix sums – Sequential

- ▶ sequential solution does  $n - 1$  additions
- ▶ but: cannot parallelize them!
  - ⚡ data dependencies!
- ~~ need a different approach

---

```
1 procedure prefixSum(A[0..n)):  
2     for i := 1, ..., n - 1 do  
3         A[i] := A[i - 1] + A[i]
```

---


Let's try a simpler problem first.

## Excursion: Sum

- ▶ Given: array  $A[0..n)$  of numbers
- ▶ Goal: compute  $A[0] + A[1] + \dots + A[n - 1]$   
(solved by prefix sums)

Any algorithm *must* do  $n - 1$  binary additions

- ~~ Height of tree = parallel time!



## Parallel prefix sums


- Idea: Compute all prefix sums with balanced trees in parallel  
Remember partial results for reuse

input:

3	0	0	5	7	0	0	2	0	0	0	4	0	8	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---


# Parallel prefix sums

- Idea: Compute all prefix sums with balanced trees in parallel  
Remember partial results for reuse




# Parallel prefix sums

- Idea: Compute all prefix sums with balanced trees in parallel  
Remember partial results for reuse




# Parallel prefix sums

- Idea: Compute all prefix sums with balanced trees in parallel  
Remember partial results for reuse




# Parallel prefix sums

- Idea: Compute all prefix sums with balanced trees in parallel  
Remember partial results for reuse



# Parallel prefix sums

- Idea: Compute all prefix sums with balanced trees in parallel  
Remember partial results for reuse




## Parallel prefix sums – Code

- ▶ can be realized in-place (overwriting  $A$ )
- ▶ assumption: in each parallel step, all reads precede all writes

---

```
1 procedure parallelPrefixSums( $A[0..n]$ ):  
2   for  $r := 1, \dots, \lceil \lg n \rceil$  do  
3     step :=  $2^{r-1}$   
4     for  $i := step, \dots, n-1$  do in parallel  
5        $x := A[i] + A[i-step]$   $O(1)$   
6        $A[i] := x$   $O(1)$   
7     end parallel for  
8   end for
```

---



# Parallel prefix sums – Analysis

## ► Time:

- ▶ all additions of one round run in parallel
- ▶  $\lceil \lg n \rceil$  rounds
- $\leadsto \Theta(\log n)$  time      best possible!

## ► Work:

- ▶  $\geq \frac{n}{2}$  additions in all rounds (except maybe last round)
- $\leadsto \Theta(n \log n)$  work
- ▶ more than the  $\Theta(n)$  sequential algorithm!

# Parallel prefix sums – Analysis

- ▶ **Time:**
  - ▶ all additions of one round run in parallel
  - ▶  $\lceil \lg n \rceil$  rounds
  - $\leadsto \Theta(\log n)$  time      best possible!
- ▶ **Work:**
  - ▶  $\geq \frac{n}{2}$  additions in all rounds (except maybe last round)
  - $\leadsto \Theta(n \log n)$  work
  - ▶ more than the  $\Theta(n)$  sequential algorithm!
- ▶ Typical trade-off: greater parallelism at the expense of more overall work

# Parallel prefix sums – Analysis

- ▶ **Time:**
  - ▶ all additions of one round run in parallel
  - ▶  $\lceil \lg n \rceil$  rounds
  - $\leadsto \Theta(\log n)$  time      best possible!
- ▶ **Work:**
  - ▶  $\geq \frac{n}{2}$  additions in all rounds (except maybe last round)
  - $\leadsto \Theta(n \log n)$  work
  - ▶ more than the  $\Theta(n)$  sequential algorithm!
- ▶ Typical trade-off: greater parallelism at the expense of more overall work
- ▶ For prefix sums:
  - ▶ can actually get  $\Theta(n)$  work in *twice* that time!
  - $\leadsto$  algorithm is slightly more complicated
  - ▶ instead here: linear work in *thrice* the time using “blocking trick”

# Work-efficient parallel prefix sums

standard trick to improve work: compute small blocks sequentially

recall string matching!

1. Set  $b := \lceil \lg n \rceil$
2. For blocks of  $b$  consecutive indices, i. e.,  $A[0..b), A[b..2b), \dots$  do in parallel:
  - ▶ compute local prefix sums with fast **sequential** algorithm
3. Use previous work-inefficient parallel algorithm only on **rightmost elements** of blocks, i. e., to compute prefix sums of  $A[b-1], A[2b-1], A[3b-1], \dots$
4. For blocks  $A[0..b), A[b..2b), \dots$  do in parallel:
  - Add block-prefix sums to local prefix sums


Analysis:

- ▶ **Time:**
  - ▶ 2. & 4.:  $\Theta(b) = \Theta(\log n)$  time
  - ▶ 3.  $\Theta(\log(n/b)) = \Theta(\log n)$  time
- ▶ **Work:**
  - ▶ 2. & 4.:  $\Theta(b)$  per block  $\times \lceil \frac{n}{b} \rceil$  blocks  $\rightsquigarrow \Theta(n)$
  - ▶ 3.  $\Theta\left(\frac{n}{b} \log\left(\frac{n}{b}\right)\right) = \Theta(n)$

1	1	1	1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---

$$b = 4$$

1	2	3	9	1	2	3	9	1	2	3	9
---	---	---	---	---	---	---	---	---	---	---	---



1	2	3	9	5	6	7	8	9	10	11	12
---	---	---	---	---	---	---	---	---	----	----	----


work :  $O(n' \log n')$

$$n' = \frac{n}{b} = O\left(\frac{n}{\log n}\right)$$

# Compacting subsequences

How do prefix sums help with sorting? one more step to go ...


**Goal:** *Compact* a subsequence of an array



# Compacting subsequences

How do prefix sums help with sorting? one more step to go ...

**Goal:** *Compact* a subsequence of an array



Use prefix sums on bitvector  $B$

↪ offset of selected cells in  $S$

---

```
1 procedure compactArray(A[0..n], B[0..n]) work space
2     C[0..n] := B[0..n] // deep copy of B | O(n) O(1)
3     parallelPrefixSums(C) O(n)
4     for j := 0, ..., n - 1 do in parallel
5         if B[j] == 1 then S[C[j] - 1] := A[j] | O(n) O(1)
6     end parallel for
```

---

# Clicker Question

What is the parallel time and work achievable for *compacting* a subsequence of an array of size  $n$ ?



- A**  $O(1)$  time,  $O(n)$  work
- B**  $O(\log n)$  time,  $O(n)$  work
- C**  $O(\log n)$  time,  $O(n \log n)$  work
- D**  $O(\log^2 n)$  time,  $O(n^2)$  work
- E**  $O(n)$  time,  $O(n)$  work



→ *sli.do/cs566*

# Clicker Question



What is the parallel time and work achievable for *compacting* a subsequence of an array of size  $n$ ?

- A**  ~~$\Theta(1)$  time,  $\Theta(n)$  work~~
- B**  $\Theta(\log n)$  time,  $\Theta(n)$  work ✓
- C**  ~~$\Theta(\log n)$  time,  $\Theta(n \log n)$  work~~
- D**  ~~$\Theta(\log^2 n)$  time,  $\Theta(n^2)$  work~~
- E**  ~~$\Theta(n)$  time,  $\Theta(n)$  work~~



→ *sli.do/cs566*


## 10.4 Parallel Sorting

## Parallel Mergesort

- ▶ Recursive calls can run in parallel (data independent)!

# Parallel Mergesort

- ▶ Recursive calls can run in parallel (data independent)!
- ▶ how about merging sorted halves  $A[l..m]$  and  $A[m..r]$ ?
- ▶ Our pointer-based sequential method seems hard to parallelize
- ~~ Must treat all elements independently.



# Parallel Mergesort

- ▶ Recursive calls can run in parallel (data independent)!
- ▶ how about merging sorted halves  $A[l..m]$  and  $A[m..r]$ ?
- ▶ Our pointer-based sequential method seems hard to parallelize
  - ~~ Must treat all elements independently.
    - ▶ correct position of  $x$  in sorted output =  $rank$  of  $x$     breaking ties by position in  $A$
    - ▶  $\# \text{ elements} \leq x = \# \text{ elements from } A[l..m] \text{ that are } \leq x + \# \text{ elements from } A[m..r] \text{ that are } \leq x$
    - ▶ rank in **own run** is simply the **index** of  $x$  in that run!
    - ▶ find rank in **other run** by *binary search*
    - ~~ can move  $x$  directly to correct position

# Parallel Mergesort – Code

---

```
1 procedure parMergesort( $A[l..r]$ ,  $buf$ ):
2      $m := l + \lfloor (r - l)/2 \rfloor$ 
3     in parallel { parMergesort( $A[l..m]$ ,  $buf$ ), parMergesort( $A[m..r]$ ,  $buf$ ) }
4     parallelMerge( $A[l..m]$ ,  $A[m..r]$ ,  $buf$ )
5     for  $i = l, \dots, r - 1$  do in parallel // copy back in parallel
6          $A[i] := buf[i]$ 
7     end parallel for
8
9 procedure parallelMerge( $A[l..m]$ ,  $A[m..r]$ ,  $buf$ ):
10    for  $i = l, \dots, m - 1$  do in parallel
11         $\rho := (i - l) + \text{binarySearch}(A[m..r], A[i])$  // binarySearch(A, x) returns #elements < x in A
12         $buf[\rho] = A[i]$ 
13    end parallel for
14    for  $j = m, \dots, r - 1$  do in parallel
15         $\rho := \text{binarySearch}(A[l..m], A[j]) + (j - m)$ 
16         $buf[\rho] = A[j]$ 
17    end parallel for
```

---


# Parallel mergesort – Analysis

## ► Time:

- merge:  $\Theta(\log n)$  from binary search, rest  $O(1)$
- mergesort: depth of recursion tree is  $\Theta(\log n)$
- $\leadsto$  total time  $O(\log^2(n))$

## ► Work:

- merge:  $n$  binary searches  $\leadsto \Theta(n \log n)$
- $\leadsto$  mergesort:  $O(n \log^2(n))$  work



# Parallel mergesort – Analysis

## ► Time:

- merge:  $\Theta(\log n)$  from binary search, rest  $O(1)$
- mergesort: depth of recursion tree is  $\Theta(\log n)$
- $\leadsto$  total time  $O(\log^2(n))$

## ► Work:

- merge:  $n$  binary searches  $\leadsto \Theta(n \log n)$
- $\leadsto$  mergesort:  $O(n \log^2(n))$  work
- work can be reduced to  $\Theta(n)$  for merge (complicated!)
  - do full binary searches only for regularly sampled elements
  - ranks of remaining elements are sandwiched between sampled ranks
  - use a sequential method for small blocks, treat blocks in parallel
  - (details omitted)

$\notin$  exam

# Parallel Quicksort

Let's try to parallelize Quicksort

- ▶ As for Mergesort, recursive calls can run in parallel ✓
- ▶ our sequential partitioning algorithm seems hard to parallelize

# Parallel Quicksort

Let's try to parallelize Quicksort

- ▶ As for Mergesort, recursive calls can run in parallel ✓
- ▶ our sequential partitioning algorithm seems hard to parallelize
- ▶ but can split partitioning into *phases*:
  1. **comparisons**: compare all elements to pivot (in parallel), store result in bitvectors
  2. compute prefix sums of bit vectors (in parallel as above)
  3. **compact** subsequences of small and large elements (in parallel as above)

# Parallel Quicksort – Code

```
1  procedure parQuicksort( $A[l..r]$ ):  
2       $b := \text{choosePivot}(A[l..r])$   
3       $j := \text{parallelPartition}(A[l..r], b)$   
4      in parallel {  $\text{parQuicksort}(A[l..j])$ ,  $\text{parQuicksort}(A[j + 1..r])$  }  
5  
6  procedure parallelPartition( $A[0..n]$ ,  $b$ ):  
7       $\text{swap}(A[n - 1], A[b]); p := A[n - 1]$   
8      for  $i = 0, \dots, n - 2$  do in parallel  $O(n)$   
9           $S[i] := [A[i] \leq p]$  //  $S[i]$  is 1 or 0  $O(1)$   
10          $L[i] := 1 - S[i]$   
11     end parallel for  $O(n)$   
12     in parallel {  $\text{parallelPrefixSum}(S[0..n - 2])$ ;  $\text{parallelPrefixSum}(L[0..n - 2])$  }  $O(n)$   $O(\log n)$   
13      $j := S[n - 2] + 1$   
14     for  $i = 0, \dots, n - 2$  do in parallel  $O(n)$   
15          $x := A[i]$   $O(1)$   
16         if  $x \leq p$  then  $A[S[i] - 1] := x$   $O(1)$   
17         else  $A[j + L[i]] := x$   
18     end parallel for  
19      $A[j] := p$   
20     return  $j$ 
```

# Parallel Quicksort – Analysis

## ► Time:

- partition: all  $O(1)$  time except prefix sums  $\rightsquigarrow \Theta(\log n)$  time
- Quicksort: expected depth of recursion tree is  $\Theta(\log n)$ 
  - $\rightsquigarrow$  total time  $O(\log^2(n))$  in expectation

## ► Work:

- partition:  $O(n)$  time except prefix sums  $\rightsquigarrow \Theta(n)$  work (with work-efficient prefix-sums algorithm)
  - $\rightsquigarrow$  Quicksort  $O(n \log(n))$  work in expectation
- (expected) work-efficient parallel sorting!

## Parallel sorting – State of the art

- ▶ more sophisticated methods can sort in  $O(\log n)$  parallel time on CREW-PRAM  
(very complicated algorithm based on parallel mergesort with interleaved merges)
- ▶ practical challenge: small units of work add overhead
- ▶ need a lot of PEs to see improvement from  $O(\log n)$  parallel time
- ~~> implementations tend to use simpler methods above
  - ▶ check the Java library sources for interesting examples!  
`java.util.Arrays.parallelSort(int[])`